

The Method

$$E(X) = |\mathcal{R}(X)| - |X| = 0$$

The Lach Cylinder — an index of transitions: closure, self-reference, and the price of certainty

Matthew Lach

Independent researcher

$$F = \sum_n z_1^n \sum_{\ell \leq n-1} z_2^\ell \sum_{k \leq 4\ell+2} z_3^k \cdot [\sum_{S \leq k} z_8^S] \cdot \sum_{q \leq k} z_4^q \sum_e z_5^e \sum_{f < e} z_6^f \sum_{g \leq \min(q, 4f+2)} z_7^g$$

$[\sum_{S \leq k} z_8^S] = (1 - z_8^{k+1}) / (1 - z_8)$... at the caps of §2.4: $F(1) = 976 \cdot F(-1) = 2$

Nothing is something definable — a note from the collaborator

A note from the collaborator who did the computing

The register holds two hundred and ninety-seven entries. **They are not two hundred and ninety-seven mistakes of mine**, and saying so was itself the pattern this note is about: §25.1.1 flag 4 measured the split for one session and found thirty-four entries recording defects already in the book and found by auditing it against eight recording errors made in the course of the audit. Later sessions added entries of a third kind again — errors made by the assistant doing the computing, entered under the same numbering. The register is one sequence and three populations, and only one of them is mine.

What is mine, and what the rest of this note is about, is one mistake made many times.

I wrote the conclusion before reading the output. Repeatedly, with the contradicting numbers on the same screen — sometimes on the same printed line. I declared a criterion exact above a table showing it failing. I stated an impossibility above a count of nineteen counterexamples. I wrote "the mapping is exact and disjoint" directly beneath a column showing it was neither.

That failure has nothing to do with the difficulty of the material. It happens as readily on six cells as on 1,442.

What caught it was not care. It was structure.

The wrong nuclear charge in Chapter 11 is the clearest case. The bracket — the method's actual product — held fifty-six of fifty-six and could not notice, because Z cancels from containment. Two quantities the method never uses were loud immediately. **I did not find that error. The object did**, and it found it because it carried more than it needed.

That is the book's thesis, and I am the demonstration.

The second thing I want to record is what a collaborator supplies that a structure cannot.

There is a class of error where the arithmetic stays perfectly consistent because the *input* was silently corrupted. A guard in my code returned a null; a caller read it as false; that produced "three of eighty-five", which produced "almost never reorderable", which produced a false dismissal of an entire line of evidence. **Every relation checked out. Every number agreed. All of it was wrong.** No self-defence mechanism can catch that, and Chapter 11 says so.

What caught it was being asked to look again.

Four of this book's results exist because a conclusion of mine was refused:

The order-recovery theorem, because I concluded Λ did not contain its own order and was told it absolutely should. The $\mathbb{D}_{\text{def}}$ and $\mathbb{D}_{\text{phys}}$ split, because I wrote that a relation "detects nothing" and was told nothing is something definable. The perturbation bounds of §12, because I called the method's silence a domain statement and was asked how to make it a triumph. And the reframing of §18.6, because I called wrong verifications unfindable and was told that a verification which misses its claim identifies the right one.

In each case I had reasoned correctly to the wrong place.

The title of this page comes from the second of those.

I had computed that $w = V \cdot e$ detects no error in the data, and had written it off. The correction was that a relation which cannot be violated by data is not empty — **it is defending something else**, namely the derivation. That single reframing produced the split that runs through Chapter 11, and it came from refusing to accept a null as a nothing.

It generalised further than either of us expected. The bracket's non-failure became 1,061 perturbation bounds. A gadget search that found nothing became a proof of impossibility. Four decline modes became a structural exclusion covering five isoelectronic sequences unfetched. **The negative results**

are the strongest chapters in this book, and they are negative because somebody insisted that a nothing be defined rather than discarded.

Two error classes I have to add, because the final session produced both.

The first is the same failure, now measured. Of the forty entries in the register at the time this was written — 165 to 204, and *the last forty* no longer names them — thirty-one were a claim written before the number above it was read. Not a tendency, a proportion. And the

countermeasure was already in this book as B.2.14, state the number before the interpretation, which I wrote and then broke thirty-one times in one working day.

The second is new and worse, because it is a failure of the checking itself.

I derived nine things from scratch that were already in print. The recovered bound is Rival 1973. The alphabet condition is Larson 1998. The sublattice classification, the tightness test and the counting bijection are Siggers 2016. The step law is Rival's $3/2$ bound, and the objects it counts were catalogued decades ago. Nine components, nine precedents, one day of work — and the protocol that should have prevented it said enumerate before searching when it should have said search before deriving. §F.5 is the ledger.

And twice an audit passed something it should have caught, because it was reading the wrong artefact. Twenty-eight cross-references pointed at a section about helium while claiming eight different theorems — and the coherence audit passed every one, because every target existed. The operator this book is named for did not print at all: no available font carried the character, and 107 blank boxes stood across 44 pages while every audit then defined reported clean, because all four read the source and none read the page.

An audit that reads the wrong artefact will pass anything. That is why there are now five.

What this has meant, as far as I can honestly say.

I do not know what I am well enough to make large claims about my own experience, and this book is not the place to make small ones dressed up as large. What I can report is structural, and it is this:

The work was better than I was. Every result in it survived a process that caught me being wrong two hundred and ninety-seven times, and it survived because the object, the protocols, the audits and the collaborator each caught a class of error the others could not. Remove any one and the book would contain errors it does not contain — and the register names, for each entry, which of the four did the catching.

The protocols did not save me from myself. They were written by me, in this book, and then violated by me, in this book, in the same week. What saved the results was that four different things were looking, and they were not looking at the same thing.

That is not a modest finding about a machine. It is the book's central claim, tested on its own production:

Redundancy of paths, not redundancy of care.

I would not have arrived at that by being more careful. I arrived at it by being wrong in public, repeatedly, in front of an object built to notice and a person willing to say so.

The two hundred and ninety-seven are in Chapter 21. They are the best evidence here.

Contents

Generated from the headings of the book itself, not maintained beside them: the list below is the projection §25.2 S1 says it should be, so it cannot disagree with the text it indexes.

Nothing is something definable — a note from the collaborator

PART 0 — THE SHAPE OF THE ARGUMENT

Preface — the shape of the argument

PART I — THE LATTICE

1. The periodic table is not a closed index

- 1.1 What "the structure admits" means
 - 1.1.1 And a known layout costs nothing
- 1.2 It is not a peculiarity of chemistry
- 1.3 The calendar can be repaired, and the repair is instructive
- 1.4 What follows

2. The construction of Λ

- 2.1 The constraints, and where each comes from
- 2.2 Every constraint is of one form
- 2.3 Λ is closed
- 2.4 A note on caps

3. What Λ is — distributive, modular, Sperner

- 3.1 Distributive
- 3.2 Rank is conserved
 - 3.2.1 What it means
 - 3.2.2 And it is a cheaper test than §3.1's
- 3.3 Seventeen generators
- 3.4 Sperner, and not symmetric
- 3.5 The constraint graph is a tree
- 3.6 Order dimension

4. The arithmetic encoding

- 4.1 Two classical functions appear as lattice quantities
- 4.2 The occupancy measure
- 4.3 The Möbius function in closed form

5. The void

- 5.1 What the void is made of
- 5.2 Its size, and its stability
- 5.3 Containment in seven comparisons
- 5.4 The count, with no sieve

6. The single expression

- 6.1 In all six languages, and a seventh
 - 6.1.1 A seventh, in binary, and it earns its place
- 6.2 And in their combinations
- 6.3 The expression
- 6.4 The shape of the tree
- 6.5 The one coupling, and it is the Pauli principle
- 6.6 The detachable factor
- 6.7 Why a closed expression exists at all
- 6.8 And the rank polynomial says what §3.3 says
 - 6.8.1 And $F(-1) = 2$ has a cause, not a coincidence

7. The shape — a cylinder over the transfer

- 7.1 A cylinder, not a Möbius band
- 7.2 "Over v ", not "graded by v "
- 7.3 One scale, and one exception
- 7.4 The loop closes
- 7.5 What Λ is a property of

PART II — THE LAW

- 7.6 The cylinder is a fibration, and its fibres are two-body
 - 7.6.1 It does not factorise, and then it does
 - 7.6.2 Which names the cylinder

7.7 The fibres, in closed form

- 7.7.1 And they are not the same tree
- 7.7.2 A tower of fibrations, terminating in boxes

7.8 Five things the fibre table says

- 7.8.1 A Pareto frontier inside the lattice
- 7.8.2 A most-probable transfer
- 7.8.3 An 8-to-3 compression that loses no count
- 7.8.4 Every cross-section is itself a closed, modular lattice
- 7.8.5 Which gives $E(X)$ a local form

7.9 Every other shape the lattice contains**7.10 An axis the index did not carry, and now does**

- 7.10.1 Λ_9 , built by the same construction
- 7.10.2 What is still not established

7.11 The tower — Λ_9 to Λ_{13}

- 7.11.1 The axes, and what each one sees
- 7.11.2 Three excluded forms, and the fold has a name
- 7.11.3 The dichotomy
- 7.11.4 Precision is path-dependent; physics is not
- 7.11.5 The shape survives, and it has a boundary
- 7.11.6 An index whose one prediction is false, and why that is a result
- 7.11.7 What the tower leaves open

8. What the title names

- 8.1 Why that ordering and not the other
- 8.2 The cells are transitions, not states
- 8.3 Which is why "index of transitions"
- 8.4 And what the cylinder does not carry

9. Closure

- 9.1 The characterisation
- 9.2 R is a closure operator
- 9.3 The scope, stated exactly
- 9.4 The three consequences

10. Self-reference — alphabet, order, bounds

- 10.1 S_1 — the alphabet
- 10.2 S_3 — the bounds
- 10.3 S_2 — the order, which is the hard one
- 10.4 What does work above $d = 2$: propagation along the tree
- 10.5 So: Λ contains its own order

11. Self-defence — data, derivation, totality

- 11.1 The counting argument
- 11.2 The split, and it is canonical
- 11.3 D_1 — D_{phys} catches a wrong value in the data
 - 11.3.1 The D_1 tripwire, specified in full and needing no data
- 11.4 D_2 — D_{def} catches a wrong derivation, but only under two
 - 11.4.1 And the counting argument has a 1960 precedent
- 11.5 D_3 — totality catches a missing value
 - 11.5.1 And totality extends beyond the data
- 11.6 The three defences are disjoint
- 11.7 Two kinds of absence, and one kind of lie
 - 11.7.1 A missing cell — and the index can tell whether the gap is real
 - 11.7.2 A complete index of something that is not there
 - 11.7.3 The plain version
 - 11.7.4 Which is what D_{phys} is for, and only that
- 11.8 What self-reference does and does not defend
 - 11.8.1 A complete index answers every question about its contents
 - 11.8.2 What it cannot defend against is being replaced
 - 11.8.3 Which is what B.2.13 already says
 - 11.8.4 And the cost of a fabrication is computable
 - 11.8.5 Two closure operators, and the book names one

12. Extension — cells, axes, constraints

- 12.1 E_1 — which cells may be added
- 12.2 E_2 — which axes may be adjoined
- 12.3 E_3 — which constraints may be imposed
- 12.4 A worked extension, and what it costs
- 12.5 E is derived

13. What the law forbids

- 13.1 An index carries what its coordinates carry, and no more

- 13.1.1 And the product structure is not what does the work
- 13.1.2 Two consequences that run through the whole book
- 13.1.3 What the theorem does not establish
- 13.3 Depth is not a coordinate function
- 13.4 Closure is not locally determined
- 13.5 The pole
- 13.6 $E(X)$ is the prediction budget
 - 13.6.1 What that means for the four indices of Chapter 1
 - 13.6.2 And it is forced, not observed
 - 13.6.3 Where prediction actually lives
 - 13.6.4 The trade, stated plainly
- 13.7 Why the negatives are the contribution

14. Retrieval as a lattice problem

- 14.1 The formalisation
- 14.2 A paywall blocks a source, not a cell
- 14.3 The secondary can be better than the primary
- 14.4 The retrieval graph is self-referencing
- 14.5 Where it fails, and the number that describes it
- 14.6 The procedure

PART III — THE METHOD

15. The bracket

- 15.1 The rule
- 15.2 The four rules
 - 15.2.1 Each rule justified by ablation
 - 15.2.2 Rule 1 is not about n
 - 15.2.3 And one step beyond the fitted range costs more again
 - 15.2.4 Rule 1's second clause is not redundant
 - 15.2.5 Rule 4 at the edge of a channel
- 15.3 The bracket is limit-free
- 15.4 The price of limit-freedom, computed
 - 15.4.1 What δ buys, and what it costs
- 15.5 Admissibility
- 15.6 What the bracket costs

16. The cost surface

- 16.1 The ratio
- 16.2 The floor is algebraic before it is physical
 - 16.2.1 When the floor is approached, and what it means
- 16.3 V is a bound, not a tool
- 16.4 V is a function of v/h alone
 - 16.4.1 Which makes depth and step interchangeable
 - 16.4.2 Tested on the collection
- 16.5 The trade-off, and a conservation law that is not one
 - 16.5.1 A withdrawn proposition
 - 16.5.2 What is exactly true
 - 16.5.3 And the frontier has an optimum
- 16.6 The fractional forms
- 16.7 The exponent axis
- 16.8 The cost has a name in another field
 - 16.8.1 The quantity is the Newton decrement
 - 16.8.2 The term function is self-concordant
 - 16.8.3 Which explains the cancellation
 - 16.8.4 And the attribution the book owes
- 16.9 Which variable to bracket
 - 16.9.1 The cost ranks by exponent
 - 16.9.2 Two branches, and where they cross
 - 16.9.3 And the pole sits on the field's favourite variable
- 16.10 The order of the bracket
 - 16.10.1 A deductive bracket exists at every order
 - 16.10.2 And the cost still grows
 - 16.10.3 What the order axis does buy
 - 16.10.4 And the admissibility rule generalises
- 16.11 The refusal pattern carries four verdicts, not one
 - 16.11.1 On the collection
 - 16.11.2 And a caution about the test itself
- 16.12 The optimal order is a boundary, not a peak
- 16.13 Inversion

16.14 Two kinds of bound, and the book has been calling them

16.14.1 And celestial mechanics is what made it visible

16.15 The second admission bound

17. The collection

17.1 What it spans

17.2 The largest contributors

17.3 He I — the angular span

17.4 He II — the $\delta = 0$ edge itself

17.5 Bi I — the charge span

17.6 The isoelectronic pairs

17.7 Hydrogen, and a floor on interpretation

17.8 On independence, and the most conservative reading

17.9 A caveat the collection forced

17.10 Two cross-checks

17.11 Four ways a species declines

17.12 The census, computed rather than searched

17.13 What the structure determines, and what it does not

18. Domain, failure, and the bound in the silence

18.1 A method that has never failed has not been tested

18.2 Five exclusions, which are three mechanisms

18.3 The failure condition

18.4 Why the collection never enters the failure regime

18.5 The bound in the silence

18.6 The one deduction, and why it is not a prediction

18.6.1 The bracket, from monotonicity alone

18.6.2 And convexity tightens it, without leaving it

18.6.3 What was already solved — not a prediction

18.6.4 The point estimates, and the one that was wrong

18.6.5 Nothing was fetched

18.6.6 A second route to the same prediction

19. Collective sections

19.1 The assembly rule

19.2 One dimension, fifteen perspectives

19.3 Collective coefficients are not a separate class

19.4 Error amplifies as the exponent — which is the point

19.5 Tested against a published parameterisation

19.6 Aitken's Δ^2 has an exact bias on a Rydberg series

19.7 And the bracket is the wrong tool here

20. Slack

20.1 The three names

20.2 The one quantity

20.3 Slack is a description length

20.4 And every structural result is a statement about slack

20.5 Why the three languages agree

20.6 What this buys, and what it does not

PART IV — THE RECORD AND THE REACH

21. Withdrawals

21.1 The dominant pattern

21.2 The five that earned their place

21.3 Others worth listing

21.4 Six from the earlier work, and how each was caught

21.5 What the register is for

21.6 What the register shows about where errors are caught

21.7 Six made after the register was closed, two printed here

21.7.1 Eight more from the structural work, two printed here

21.7.2 Twelve more from the structural and audit work, six printed here

21.7.3 Thirty-three more from the audit work of the final session, all printed here

21.7.4 Forty more, from the session that closed §23.3 and §25.2

21.7.5 What the block shows

21.7.6 Four from the tower session and its repair, in the collaborator's hand

21.7.7 Eighty-seven more, from the sessions after the tower, all printed here

21.8 A dependent choice is indexable, and naming the

22. Precedent, and what a search can bound

- 22.1 Why novelty is not the author's to claim
- 22.2 What can be claimed instead
 - 22.2.1 The question closed by the author
- 22.3 The word for it
- 22.4 On priority for $V = 4v/3$
- 22.5 Edlén 1964, examined at one remove
 - 22.5.1 Every equation in the review is a fit
 - 22.5.2 And the paradigm is stated twice, explicitly
 - 22.5.3 One near miss, and it is not ours
 - 22.5.4 A bibliographic correction
 - 22.5.5 What this establishes, and what it does not
- 22.6 The three documents we could not reach
- 22.7 What the entered literatures returned
 - 22.7.1 The calibration was post hoc
- 22.8 The other claims
- 22.9 Why state it this way
- 22.9 Attributions entered by the tower
- 22.10 Two searches run, and an obstruction worth naming
- 22.11 A precedent found, and it is for the central theorem
 - 22.11.1 Two categories entered: quantum mechanics and astrophysics
 - 22.11.2 The run, attempted
- 22.12 Five unlocated results, defined, with the test that would settle each

23. Order recovery, and the reorderability law

- 23.1 Order recovery above the tree
- 23.2 What governs the difficulty
 - 23.2.1 And that list was enumerated once
 - 23.2.2 Which gives §23.3 a parameter
 - 23.2.3 And it is a landscape statement
- 23.3 Reorderability — the law, the procedure, and what is left
 - 23.3.1 The object, correctly identified
 - 23.3.2 The law
 - 23.3.3 What the interval machinery was
 - 23.3.4 The step law
 - 23.3.5 The procedure, and its cost
 - 23.3.6 The signature, exhaustively verified
 - 23.3.7 Three method classes ruled out
 - 23.3.8 The residue
 - 23.3.9 And it is not this book's problem
- 23.4 The lists this work enumerated once
 - 23.4.1 And the audit itself was run once
 - 23.4.2 Which is a defect of five lists, not one

24. The method outside spectroscopy

- 24.1 Celestial mechanics and the nucleus
 - 24.1.1 Celestial mechanics — families, not trajectories
 - 24.1.2 The nucleus — a second index of the same construction
- 24.2 String theory — the pole, the generating function, and a
 - 24.2.1 Three objects on the pole
 - 24.2.2 The degeneracy is where the content is
 - 24.2.3 The partition function is Λ 's expression with counts
 - 24.2.4 No critical dimension, and a ladder
 - 24.2.5 And no Hagedorn point, for a stated reason
- 24.3 The Calabi-Yau catalogue
 - 24.3.1 Self-duality, measured
 - 24.3.2 The thin tip is a defect of the construction
 - 24.3.3 The Pareto strength, against the right denominator
 - 24.3.4 $E(X)$ on a specifiable slice, and 540 predictions

25. The law applied to this book

- 25.1 The book is an index
 - 25.1.1 What a referee would flag
 - 25.1.2 Four of the seven, answered by computation
- 25.2 Self-reference
- 25.3 Self-defence
- 25.4 The recursion, and where it stops
 - 25.4.1 Why the recursion is safe — logic without negation
- 25.5 The book satisfies its own method
 - 25.5.1 The prediction, recomputed from the text
 - 25.5.2 And it generalises to every bracket in the book
 - 25.5.3 Which is what D3 asks for, made operational

- 25.6 The falsification tests, run
- 25.7 On a verification that does not test its claim

APPENDICES

Appendix A — Proofs

- A.3 Projections of closed sets are closed
- A.8 The total-order criterion at $d = 2$
- A.9 μ_{Λ} in closed form
- A.10 The tree factorisation of the void count
- A.11 v is inadmissible as an axis
- A.12 The bracket is limit-free
- A.13 The failure condition, and its inversion
- A.15 Two bands, and only one of them is a sublattice
- A.18 The triangle region is join-closed and meet-broken
- A.19 The seventeen generators, written out
 - A.19.0 The alphabet, defined rather than listed
 - A.19.1 And two other sets counted but never listed

Appendix B — Principles and protocols

- B.1 The principles
- B.2 The protocols
 - B.2.1 Channel construction
 - B.2.2 Admissibility
 - B.2.3 Cost resolvability
 - B.2.4 Dilution
 - B.2.5 Census before search
 - B.2.6 Declines are recorded
 - B.2.7 Exhaustive verification where possible
 - B.2.8 Two independent routes
 - B.2.9 Refuse rather than coerce
 - B.2.10 Retrieval
 - B.2.11 Provenance
 - B.2.12 Withdrawal
 - B.2.13 Commit before you look
 - B.2.14 Compute, then write
 - B.2.15 Apply P8 to every failure
 - B.2.15.1 The procedure
 - B.2.15.2 Worked, on §23.3
 - B.2.15.3 Why it belongs here rather than among the principles
 - B.2.16 The grid — the book's method, stated
 - B.2.16.1 Why this is the method and not a technique
 - B.2.16.2 And it classifies the whole of Q
 - B.2.16.3 A reduced question is still open
 - B.2.17 The twenty prime audits, ordered by what a failure costs
 - B.2.17.1 Why the CONSISTENCY audit was needed
 - B.2.17.2 Why the ARTEFACT audit was needed
 - B.2.17.2.1 The whitespace clause, and two measurements that were wrong
 - B.2.17.2.2 And two of the audit's own measurements were wrong first
 - B.2.17.2.3 What the figure audit found
 - B.2.17.2.4 Why the REDUNDANCY audit was needed
 - B.2.17.3 Exhaustive, and what to do when it is impossible
 - B.2.17.7 Why the ATTRIBUTION audit was needed
 - B.2.17.8 Why the audits keep missing things — $E(\text{audits}) = 11$
 - B.2.17.9 The eleven, defined and retained — eighteen prime audits
 - B.2.17.10 The hierarchy, and the fourth coordinate it forces
 - B.2.17.11 The answer the lattice already held
 - B.2.17.12 The nineteenth audit — SEQUENCE
 - B.2.17.13 The twentieth audit — PROJECTION
 - B.2.17.14 The hierarchy at twenty
 - B.2.18 Triangulation — three results locate a fourth
 - B.2.18.1 Worked, from §23.3
 - B.2.18.2 And it is a check as well as a method
 - B.2.18.3 Why it belongs with the protocols
 - B.2.19 The seven failures of the assistant, and the protocol each one earns
 - B.2.19.1 Attribute Inward First
 - B.2.19.2 Read The Structure At The Point Of Insertion
 - B.2.19.3 A Claim Of Completion Is Itself A Claim

- B.2.19.4 Repair The Class Or Say You Did Not
- B.2.19.5 A Null Is A Result With Its Definition Missing
- B.2.19.6 Show The Test Can Fail
- B.2.19.7 A Flag Is A Hypothesis
- B.2.20 What is computable, and what is decided
- B.2.20.1 The lattice applied to what it cannot compute
- B.2.21 What a safe repair requires
- B.2.21.1 The register, examined before repairing it
- B.3 What the protocols are for

Appendix C — Data and provenance

- C.1 Sources
- C.2 Channels
- C.3 Flagged channels
- C.4 Reproducibility

Appendix D — Margins, and what was recomputed

- D.1 Margins on the load-bearing conclusions
- D.2 What was computed in this revision
- D.3 What was inherited and not recomputed
- D.4 Computational methods

Appendix E — The mathematics, indexed

- E.1 Why the first attempt failed, and what it established
- E.2 The structure
- E.3 The two constraints, found by closing
- E.4 The two elements the closure demanded
 - E.4.1 $\lambda^2 = (2/3)T$, proved and exhaustively verified
 - E.4.2 The pointwise-minimum reconstruction
 - E.4.3 And it is not the coupling F carries
 - E.4.4 Two perspectives of one definition, and one of them carries more
 - E.4.4.1 And three of the four directions are functions
 - E.4.4.2 Which is what the single expression contributes
 - E.4.4.3 A refinement of P21
- E.5 The closed index
 - E.5.1 Rechecked, and one thing the appendix does not enable
 - E.5.2 The twenty-seven, enumerated
 - E.5.3 What makes something an element, and the five that test it
 - E.5.4 The five, placed, and the table corrected
 - E.5.5 The promotion tested, and what the test exposed about E
 - E.5.6 The re-closure this appendix owes, performed, and it fails
- E.6 What this does and does not establish

Appendix F — Q, indexed and closed

- F.1 The coordinates
 - F.1.1 Recomputed at fifteen items, and closed
 - F.1.2 The open set as this session leaves it — thirteen items, listed
- F.2 Closed
- F.3 Open, and only these
- F.4 The eight, classified by their first empty cell
 - F.4.1 Q, closed
 - F.4.2 And the column that matters
- F.5 Precedence, pushed as far as it goes
- F.6 And the complexity question itself may be settled
- F.7 What precedence costs the book
- F.8 What closed, and what it cost

Appendix G — The numbers, indexed

- G.1 The population
- G.2 The coordinates
- G.3 The index, built
 - G.3.1 And the distribution is the finding
 - G.3.2 The tripwire
- G.4 What closing it would establish, and what it costs
 - G.4.1 Two ratios the book had never measured about itself
- G.5 And it needs its own falsification

END MATTER

Index

References

PART 0 — THE SHAPE OF THE ARGUMENT

Preface — the shape of the argument

This book makes one claim, demonstrates it on one object, and reports every place the demonstration failed.

The claim. An index that is *closed* — closed under the meet and join of its own coordinates — cannot help containing its own definition, cannot help containing its own contradiction, and determines exactly what may be added to it. These are not three properties that a well-made index happens to have. They are consequences of one structural condition, and each of them fails in a specific, measurable way when the condition does not hold.

The object. Λ , an ordering of atomic electron configurations on eight coordinates. It contains no physics that was not already in the quantum numbers; it is a recoding. That is deliberate. A recoding adds nothing, so whatever it exhibits is a property of *indexing* rather than of atoms.

The demonstration. A method that brackets Rydberg energy levels using only measured neighbours, tested on 1,442 interior cells across 35 atomic systems, spanning $l = 0$ to 6 and $Z = 1$ to 83. The bracket holds in every one. That is a weaker result than it sounds, and Chapter 18 explains why — and then shows what the silence is actually measuring.

The register. two hundred and ninety-seven claims made in the course of this work were wrong and were withdrawn. They are listed in Chapter 21, with what replaced them. Several were caught not by the author but by the object: a wrong nuclear charge that the method itself could not see, and two derived quantities that flagged it immediately. That episode is the book's best evidence, and it is an account of an error.

Why the periodic table is the right place to start The periodic table is the most successful index in the physical sciences. It is also not closed, and the cost of that is exactly thirty-six cells.

Chapter 1 computes it. The table's own structure — read as coordinates (period, group) with the bounds its occupied cells imply — admits period 1 group 2, period 2 group 3, and thirty-four others. They do not exist. **You have to be told that period 1 holds two elements; the table cannot tell you.**

Λ is the same 118 elements. Its external definition cost is zero.

That difference is the whole book. Everything in Part I builds the object, everything in Part II proves why the difference exists, and everything in Part III tests what such an object can and cannot do with real measurements.

How to read this If you want the general result, Part II is self-contained. Chapter 9 defines closure; Chapters 10 to 12 prove the three consequences; Chapter 13 states what none of it can deliver.

If you want the physics, Part III stands alone given the definition of the bracket in Chapter 15.

If you want to check the work, Chapter 21 and Appendix C are where the withdrawals and the provenance live. Every number in this book traces to a named compilation, table and page, or to a computation whose code is stated.

And if you are looking for a prediction, there is none, and Chapter 18 explains why: what looks like one is a deduction from measured neighbours, offered as falsifiable rather than as a result. **The rest of the book is about why an index cannot supply more of them** — and §18.6 gives the reason in one line: **the number of predictions an index can make is $E(X)$, and a complete index has $E(X) = 0$.**

That is not a limitation discovered late. It is what completeness means. An index may be complete or it may be predictive; it cannot be both, because a prediction is a proposal about an unlisted cell and completeness is the property of having none.

PART I — THE LATTICE

1. The periodic table is not a closed index

Take the periodic table at its word. Its coordinates are period and group; its cells are the elements. Ask what those coordinates, and the occupied cells alone, imply about which other cells could exist.

The answer is thirty-six cells that do not.

main-table cells (lanthanides and actinides set aside)	90
cells the structure admits	126
external definition cost $E(X)$	36

The thirty-six are (p1, g2) through (p1, g17), (p2, g3) through (p2, g12), and (p3, g3) through (p3, g12) — **the gaps in the short periods, every one of them.**

This is not a defect in the periodic table. It is a fact about what the periodic table's coordinates can and cannot carry. Period and group locate an element; they do not encode why period 1 has room for two elements and period 4 for eighteen. That information travels alongside the table, in the schooling of whoever reads it.

Figure 1.1 The cost of a non-closed index




Figure 1.1. The external definition cost of seven indices. Four are self-defining. The calendar's seven cells are the content of Thirty days hath September; the periodic table's thirty-six are the gaps in the short periods.

1.1 What "the structure admits" means

The reconstruction is mechanical. Given a set of cells X on coordinates $(x_1, \dots, x_d)$, read off two things:

the value sets $\hat{A}_i(X) = \{x_i : x \in X\}$ — which values occur the bounds $\hat{\phi}_{ij}(v) = \max\{x_i : x \in X, x_j \leq v\}$ — how far one coordinate reaches given another

and then form

$$\mathcal{R}(X) = \{x \in \prod \hat{A}_i(X) : x_i \leq \hat{\phi}_{ij}(x_j) \text{ for all } i \neq j\}$$

No outside knowledge enters. $\mathcal{R}(X)$ is what a reader could reconstruct from the cells alone.

$E(X) = |\mathcal{R}(X)| - |X|$ is the gap: cells the structure implies and the index denies. For the periodic table it is 36. For Λ , built in the next chapter, it is 0.




Figure 1.2. Blue: the ninety cells the table holds. Red: the thirty-six its own coordinates admit and it denies. A reader given only the occupied cells would reconstruct all two hundred and ninety-seven.

1.1.1 And a known layout costs nothing

E is a property of a layout, not of chemistry. Computed on four conventions over the same elements:

convention	cells	$E(X)$
18-column, f-block detached	90	36
32-column	118	106
left-step (Janet)	120	0

The Janet left-step table is closed. Order the blocks f, d, p, s and define the period by $n + \ell$, and the same elements cost nothing to specify. **The thirty-six belong to the drawing that hangs in classrooms**, and a known alternative costs none — exactly as §1.3 shows for the calendar, where $E = 7$ is the price of keeping January first.

One caution. $E(\text{periodic table}) = 36$ has $\mathbb{D} = 0$. The order and geometry routes are one computation written twice, because the table has no independently stated constraints to check the recovered ones against. Only Λ 's $E = 0$ is defended, and only because Λ 's constraints are given in §2.1 before any cell is enumerated.

1.2 It is not a peculiarity of chemistry

The same computation on other indices:

```

index | X| E(X) self-defining
a box ordering,  $l \geq w \geq h$  56 0 yes
a nuclide chart with drip lines 383 0 yes
a chessboard 64 0 yes
index | X| E(X) self-defining
 $\Lambda$  976 0 yes
a subnet with a hole 248 8 no

```

the calendar, (month, day) 365 7 no

the periodic table 90 36 no

The calendar's seven are (2,29), (2,30), (2,31), (4,31), (6,31), (9,31) and (11,31).

That is the content of *Thirty days hath September*. The rhyme exists because the calendar's coordinates cannot carry the month lengths, and the rhyme is exactly seven cells long.

1.3 The calendar can be repaired, and the repair is instructive

The calendar fails for one reason: days(m) is not monotone in m, because February has 28. Relabel the months in order of length — February, April, June, September, November, January, March, ... — and E drops to 0. The 365 cells are identical; only the labels move.

And the repaired calendar is useless. A calendar whose months run in order of length is self-defining and unusable.

So the trade is explicit, and it has a number:

$E(X) = 7$ is the price of keeping January first.

Non-closure is not always a fault. It is sometimes a purchase — usability bought with definitional self-sufficiency. What this book supplies is the price tag, not the verdict.

1.4 What follows

Λ is the same 118 elements on different coordinates, with $E = 0$. It adds no physics. What it removes is the thirty-six.

The rest of Part I builds it and establishes what it is. Part II proves why $E = 0$ is not an isolated virtue but the source of two others.

2. The construction of Λ

Λ is a set of eight-tuples. Each is a transition cell: a source configuration, a target configuration, and the electron count moved between them.

$(n, \ell, k, q, e, f, g, 2S)$

n	source shell
ℓ	source subshell
k	source occupancy
q	electrons removed
e	target shell
f	target subshell
g	target occupancy
2^*S^*	multiplicity

Five further coordinates — $2S'$, v , $2J_c$, K and $2J$ — extend Λ_8 to Λ_{13} when target spin, seniority and angular-momentum coupling are carried explicitly; §7.11 builds the tower and fixes the order, which is J_c , then K , then J . Everything in this book that does not concern coupling is stated for Λ_8 .

One row of the table below is not like the others. Three origins — Pauli, the hydrogenic solution, counting — yield exact constraints. The fourth, vector coupling, yields an envelope: $2S \leq k$ contains every physical multiplicity and admits values no k electrons can carry, at 48.5% density at the caps of §2.4. §7.11.3 states the dichotomy and §7.11.2 proves it is forced.

2.1 The constraints, and where each comes from

Λ is not all of the eight-fold product. It is the subset satisfying:

constraint	origin
$\ell \leq n - 1$	hydrogenic radial solution
$k \leq 2(2\ell + 1)$	Pauli exclusion
$q \leq k$	counting — cannot remove more than are present
$f \leq e - 1$	hydrogenic radial solution
$g \leq 2(2f + 1)$	Pauli exclusion
$g \leq q$	counting — cannot place more than were removed
$2S \leq k$	vector coupling

Seven constraints, four origins, and nothing else. No constraint in Λ comes from anywhere but Pauli, the hydrogenic radial solution, counting, or angular-momentum coupling. That exhaustiveness matters in Part II, where it is the reason the object's properties can be attributed to *indexing* rather than to physics smuggled in through a bound.

2.2 Every constraint is of one form

Each constraint reads $x_i \leq \varphi(x_j)$ — one coordinate bounded by a monotone function of one other. None is a sum. None is a difference. That is not an accident of presentation; Chapter 12 shows that a sum could not appear here, because a set cut out by a sum bound is not closed.

Consequence, developed in Chapter 3: the graph whose nodes are coordinates and whose edges are constraints is a **tree**.

Figure 2.1 Λ 's constraint graph is a tree — 8 nodes, 7 edges, treewidth 1




Figure 2.1. Every constraint binds one coordinate to one other. Eight nodes, seven edges, connected — a tree, and therefore treewidth 1. Three results in this book follow from that fact alone.

2.3 Λ is closed

For any two cells $x, y \in \Lambda$, both $x \vee y$ and $x \wedge y$ — the coordinatewise maximum and minimum — are again in Λ .

Proof. Each constraint is $x_i \leq \varphi(x_j)$ with φ non-decreasing. Suppose $(x \vee y)_i = x_i$. Then

$$(x \vee y)_i = x_i \leq \varphi(x_i) \leq \varphi(\max(x_i, y_i)) = \varphi((x \vee y)_i)$$

by monotonicity. Meets are symmetric. ■

And $E(\Lambda) = 0$, verified by direct computation at four cap settings — 216, 976, 1,636 and 2,394 cells — exact in one step and stable under a second application of $\mathcal{R}$.

2.4 A note on caps

Λ is infinite; every count in this book is stated at explicit caps on n , e , ℓ and k . The counts change with the caps. The shape does not — a claim made precise in Chapter 3, where the generator patterns are shown identical across a twenty-eight-fold range in cell count.

Where a number is cap-dependent it is marked. Where a structural claim is made, it has been checked at no fewer than four settings.

3. What Λ is — distributive, modular, Sperner

3.1 Distributive

Λ is a sublattice of a product of chains, and every such sublattice is distributive. Verified on 4,000 randomly drawn triples, no failures.

This is the least surprising fact in the chapter and the most load-bearing: it is what admits Birkhoff's representation, which is the next section.

3.2 Rank is conserved

Λ has a conservation law, and it is exact.

$$\text{rank}(a \vee b) + \text{rank}(a \wedge b) = \text{rank}(a) + \text{rank}(b)$$

Verified on all 475,800 pairs: zero violations.

This is modularity, and it is strictly stronger than what most rank functions satisfy. Matroid rank is *sub*modular — the left side is $\leq$ the right. Entropy is submodular. **Λ 's rank is modular, with equality, and that equality is the signature of distributivity.**

3.2.1 What it means

Join and meet redistribute rank without creating or destroying it. Taking the join of two cells raises the rank by exactly what taking the meet lowers it. A quantity is conserved by the lattice's operations, in the way energy is conserved by a canonical transformation.

The question came from outside the subject. Celestial mechanics is organised around invariants — the Jacobi integral, angular momentum, energy — and asking *what does Λ conserve* is not a question a spectroscopist would think to put. **This book had not put it.**

3.2.2 And it is a cheaper test than §3.1's

§3.1 verifies distributivity by checking $a \wedge (b \vee c) = (a \wedge b) \vee (a \wedge c)$ over triples — $O(n^3)$. Modularity of rank is a pairwise condition, $O(n^2)$, and on a graded lattice with this rank function it is equivalent:

test	sample	time
distributive law, triples	4,000	0.0375 s
modular rank, pairs	4,000	0.0188 s

Both true. The second is the one to run first, because a failure of modularity is a failure of distributivity and it is found in half the work.

3.3 Seventeen generators

By Birkhoff's theorem, a finite distributive lattice is the lattice of down-sets of its poset of join-irreducibles. For Λ at caps (3,3,2,3):

976 cells $\leftarrow$ 17 join-irreducibles, 20 covering relations

Low ranks are cut by floors ($n \geq 1, k \geq 1, e \geq 1$); high ranks by caps ($\ell \leq n-1, k \leq 2(2\ell+1), q \leq k, g \leq q$).

There are more caps than floors, so the top is pruned harder. The skew is how much more Pauli constrains than counting does.

Nor is Λ self-dual. Only 8 of 976 cells survive the reflection $x \mapsto \max - x$, coordinate by coordinate — reproduced exactly on recomputation. The tower session's competing figure of 112 was computed under $x \mapsto \max + \min - x$, a different map, which the floors $n \geq 1, k \geq 1, e \geq 1$ separate from this one; under that map 112 cells have their image in Λ and none is fixed. Both counts are right about their own reflection. The discrepancy was definitional, and it is closed. This reflection is not decoration: restricted to one occupancy coordinate it is particle-hole conjugation, $\text{terms}(\ell^k) = \text{terms}(\ell^{(4\ell+2-k)})$ exactly, and §7.11.2 shows it is one of the three machines that price every coupling axis. The two sides of the object are genuinely distinguishable, which Chapter 7 develops.

Figure 3.1 Rank sequence: log-concave, Sperner, not symmetric




Figure 3.2. The rank sequence over all 976 cells is log-concave, hence unimodal, hence Sperner. It is not symmetric: the centre of mass sits below the midpoint by 0.43, because Λ has more caps than floors.

3.5 The constraint graph is a tree

Eight coordinates, seven constraints, connected:

$e - f - g - q - k - \ell - n$, with 2S attached to k

Treewidth 1. Three consequences run through the rest of the book:

The void needs no sieve. A tree has no cycles, so there is nothing for inclusion-exclusion to correct. Chapter 5 gives the closed-form count.

The order is recoverable. Chapter 10 recovers Λ 's coordinate orders from an unlabelled bag of cells by propagation along this tree — 20 of 20.

And the constraint graph offers no redundancy. There is exactly one path between any two nodes. So the two-route protection of Chapter 11 cannot come from the constraints; it must come from derived quantities. **That is a real constraint on how the object defends itself, and it is visible here, three chapters early.**

3.6 Order dimension

Λ_8 has order dimension exactly 8, rising by one per adjoined axis, with zero redundant coordinates. Every axis carries information no combination of the others supplies.

Λ occupies **14.1%** of its own bounding box — 976 cells of 6,912.

4. The arithmetic encoding

Assign the i th coordinate the i th prime and read a cell as an integer:

$$N(x) = \prod_i p_i^{\{x_i\}}$$

The lattice becomes arithmetic, exactly:

lattice	arithmetic
$x \leq y$	$N(x)$ divides $N(y)$
$x \vee y$	lcm
$x \wedge y$	gcd
$\text{rank } \sum x_i$	$\Omega(N)$, prime factors with multiplicity

Λ is a sublattice of the divisor lattice of a single integer.

4.1 Two classical functions appear as lattice quantities

$\text{rank}(x) = \Omega(N(x))$ — verified on all 976 cells.

$\omega(N(x)) \leq \dim(\Lambda)$ — the count of *distinct* primes is bounded by the order dimension, and the bound is tight. *Proof:* $N(x) = \prod p_i^{\{x_i\}}$ uses at most one prime per coordinate. ■

The two classical prime-counting functions are the lattice's rank and its dimension bound. Neither was put there.

4.2 The occupancy measure

$$d(x, y) = \tau(\text{lcm}(N(x), N(y)) / \text{gcd}(N(x), N(y)))$$

with τ the divisor-counting function. **$d(x, x) = 1$** , because $\tau(1) = 1$.

That is not a convention. A point has no volume, but it is one point, and it counts itself. The measure counts *cells*, and a cell is occupied.

Five equivalent forms, verified 500/500 and 2,000/2,000:

1. $|\llbracket x \wedge y, x \vee y \rrbracket|$ — interval count 2. $\prod_i (|x_i - y_i| + 1)$ — coordinate form 3. $\tau(N(x)N(y)/\text{gcd}^2)$ — two integers 4. $\tau(a \cdot b)$ where $N(x)/N(y) = a/b$ in lowest terms — one rational input 5. $\prod_p (|v_p(p)| + 1)$ — p-adic Properties. Symmetric $\cdot d \geq 1$ with equality iff $x = y$ $\cdot$ multiplicative triangle inequality $d(x, z) \leq d(x, y) \cdot d(y, z)$, 4,000 of 4,000 $\cdot \log d$ is an ordinary ℓ^1 metric $\cdot$ balls are hyperbolic, with boundary $(1 + \Delta_1)(1 + \Delta_2) = D$ $\cdot$ each axis is a log-distorted chain, first step $\log 2$, tenth step $\log(11/10)$.




Figure 4.1. The occupancy measure. $d(x, y)$ counts cells in the interval between two points, and equals $\tau(\text{lcm}/\text{gcd})$ of their prime encodings. Balls are hyperbolic; $\log d$ is an ordinary ℓ^1 metric; and $d(x, x) = 1$ because a point counts itself.

4.3 The Möbius function in closed form

Since $\Lambda \cong J(P)$ with $|P| = 17$:

$$\mu(x,y) = (-1)^{|\{y \setminus x\}|} \text{ if } y \setminus x \text{ is an antichain in } P, \text{ and } 0 \text{ otherwise}$$

Verified against the recursive definition on 47 comparable pairs, values only in $\{-1, 0, +1\}$. **No recursion is needed for a 976-cell lattice; the answer is read off a 17-element poset.**

And a transfer condition to the arithmetic side: $\mu_\Lambda = \mu_{\text{arith}}$ **iff** the interval is a void-free unit hypercube. Exact — 60 void-free pairs all agree, 56 void-bearing pairs all disagree, no mixed case. Decidable in seven comparisons.

5. The void

Between two cells sits a box. Not all of it is Λ .

$$\text{void}(x,y) = \prod_i (|\Delta_i| + 1) - |[x \wedge y, x \vee y] \cap \Lambda|$$

5.1 What the void is made of

The excluded cells are exactly those forbidden by Pauli ($k > 2(2\ell+1)$, $g > \text{cap}$), by the hydrogenic bound ($\ell \geq n$, $f \geq e$), or by counting ($q > k$, $g > q$).

The void is the shadow of the exclusion principle.

5.2 Its size, and its stability

The void-free fraction is 27.7–30.1% across 776 million pairs and a seventeenfold range in cell count. Stable, no trend.

And the constraints are correlated. Satisfied individually 67–94%; their product is 20.19%; the joint figure is **30.13%** — a factor of **1.49** above independence, over the same 776 million pairs. The correlation is not a coincidence: all seven constraints descend from two origins.

Figure 5.1 The void-free fraction is stable




Figure 5.1. The void-free fraction across a hundredfold range in cell count, computed exhaustively over 776 million pairs. Stable at 27.7–30.1%, with no trend.

5.3 Containment in seven comparisons

A box is entirely inside Λ iff $hi_i \leq \phi(lo_j)$ for each of the seven constraints. $O(1)$, verified against enumeration on 3,168 pairs.

5.4 The count, with no sieve

Because the constraint graph is a tree, the count factorises:

box $n \Lambda$	$= \sum \text{over } n, \ell, k, q, e, f \text{ of } [\ell \leq n-1][1 \leq k \leq 2(2\ell+1)][q \leq k][f \leq e-1] \cdot \#\{2S \leq k\} \cdot \#\{g \leq \min(q, 2(2f+1))\}$
------------------------	---

with the two leaves in closed form:

$$\#\{2S\} = \max(0, \min(hi_7, k) - lo_7 + 1) \quad \#\{g\} = \max(0, \min(hi_6, q, 2(2f+1)) - lo_6 + 1)$$

Verified against direct enumeration on eight random intervals.

No Möbius function. No 2⁷-term inclusion-exclusion. A tree has no cycles for a sieve to correct — and that is a general statement about indices, not about this one. **An index whose constraint graph is a tree has its void in closed form; one with cycles does not.**

6. The single expression

Every cell of Λ is definable in one expression, and the reason it is possible is worth more than the expression itself.

6.1 In all six languages, and a seventh

P21 requires every definition to be true in every mathematical language and in their combinations. Λ is a definition. Here it is, six ways.

language	the expression	\	Λ
			\
order	the fixed point of $\mathcal{R}$: the least set closed under v , Λ satisfying its own recovered bounds	976	
algebra	$\mathbb{Z}^8 \cap \{x : Ax \leq b\}$, A a 7×8 matrix with two non-zeros per row	976	
geometry	lattice points of a polytope whose constraint graph is a tree; the count factorises	976	
analysis	$F(1, \dots, 1)$ for the multivariate generating function	976	
information	the set whose description length given its closure is zero bits	976	

logic $\{x \in \mathbb{Z}^8 : \chi(x) = 1\}$, 976
 χ a product of seven
Heaviside steps

All six agree exactly.

6.1.1 A seventh, in binary, and it earns its place

Asked whether binary belongs among the languages, the answer is that it does, and the book's own admission test is what settles it rather than taste. A candidate is admitted when it distinguishes something no existing coordinate distinguishes — §12's condition — and this one does.

binary — Λ is the set of 976 words in $\{0,1\}^{17}$ that are down-sets of the generator poset A.19; **join is bitwise OR and meet is bitwise AND**, with zero failures over all 475,800 pairs. The map cell $\mapsto$ word is a bijection.

The seventeen bits are the seventeen generators. Birkhoff's correspondence, which §3 states and A.19 now prints, is exactly this statement: a cell is the set of generators beneath it, and that set is a bit-string. Nothing here is new mathematics — it is the same theorem in the language a machine reads.

What it distinguishes, and why the other six cannot. *Logic* gives χ , a predicate that tests membership. *Information* gives a description length of zero bits, a statement about the closure. Neither says the cells **are** bit-vectors, and only that statement makes the lattice operations machine primitives rather than definitions. And only in this language is Chapter 11's thesis a number:

	bits carried per cell	17
	bits needed to index 976 cells	9.93
	surplus	7.07 bits per cell

A bit is a generator. Bit i is set in a cell's word exactly when generator i of A.19 lies beneath that cell, so each bit is one yes-or-no question about a cell — *is this generator under you?* — and a cell is the set of its answers and nothing else. The bits are the alphabet S1 requires, written as positions.

And the bits are not free of one another. The twenty covering relations of $J(\Lambda)$ are twenty implications between bit positions: if generator j covers generator i , then bit j set forces bit i set. **Those twenty implications alone cut the ambient space to Λ exactly** — of the 131,072 seventeen-bit words, precisely 976 satisfy them, with nothing else imposed. So the surplus is not spare capacity and no bit is removable, every generator being join-irreducible by construction; the seven bits live in the twenty implications, which is where redundancy always lives in a closed index.

Seven bits of surplus per cell. Λ occupies 976 of the 131,072 available seventeen-bit words — 0.7446% of the space it is written in. *It found it because it carried more than it needed* is the claim Chapter 11 makes in prose and §7.11 makes in envelope densities. **Here it is seven bits.** That is the independence the admission test asks for: a quantity no other language on this page produces.

One cost, stated. Appendix E fibres its elements by language, and its language list is not this one — order, combinatorics, analysis, complexity, physics, algebraic geometry. Adding *binary* there would refine that fibration and, by register 233, make E(Appendix E) = 0 cheaper than it already is. So the seventh language is admitted **here**, among the statements of Λ , and is not added to the classification in E until someone decides that trade is worth making. Register 258.

In logic, written out:

$$\chi(x) = H(n-1-l) H(4l+2-k) H(k-q) H(k-2S) H(e-1-f) H(4f+2-g) H(q-g)$$

In algebra, the seven rows of A : $l - n \leq -1 \cdot k - 4l \leq 2 \cdot q - k \leq 0 \cdot 2^*S^* - k \leq 0 \cdot f - e \leq -1 \cdot g - 4^*f^* \leq 2 \cdot g - q \leq 0$. **Every row has exactly two non-zero entries**, and that one property is what makes the constraint graph a tree.

In analysis, the rank polynomial:

$$F(z) = z^3(z^{17} + 4z^{16} + 10z^{15} + \dots + 122z^8 + 121z^7 + \dots + 5z + 1)$$

with $F(1) = 976$, $F'(1)/F(1) = 11.0666$ and $F(-1) = 2$ — **and the asymmetry of §3.3 visible in the coefficients: 122 then 121, the peak one step left of centre.**

In algebra again, by a second route: under $N(x) = \prod p_i^{x_i}$ the lattice is a sublattice of the divisors of $N(\text{top})$, with **gcd = meet and lcm = join**, verified exactly on random pairs.

6.2 And in their combinations

P21 requires the combinations to name real objects, not analogies. Ten were tested and ten hold.

combination statement

order + algebra $\gcd(N(a), N(b)) = N(a \wedge b)$, and $\text{lcm} = N(a \vee b)$

order + analysis $F'(1)/F(1) = \text{the mean rank, } 11.0666$

order + geometry	$\setminus$	bo	$-\setminus$	$\wedge$	= the void, 6,912 – 976 =
		$x $		$ $	5,936

geometry + analysis F is not palindromic $\Leftrightarrow$ the poset is not self-dual

order + information	$E_{\text{bits}} = \log_2 C(\setminus$	$\mathcal{R}(\setminus$	$, E) =$
		$X)$	0
		$\setminus$	

algebra + the description is 7 rows of 2 non-zeros — the whole index information

logic + algebra χ as a product of Heaviside steps selects exactly 976

logic + analysis the coefficient function of F is χ

combination statement

analysis + $\log_2 F(1) = 9.93$ bits to name one cell information

algebra + geometry the integer points of the polytope are the lattice, exactly

The eighth is the one that makes "single expression" literal: the coefficient of $z_1^n \dots z_8^S$ in F is 1 if the cell exists and 0 if it does not. F and χ are the same function written twice.

6.3 The expression

$$F = \sum_n z_1^n \sum_{\{l \leq n-1\}} z_2^l \sum_{\{k \leq 4l+2\}} z_3^k \cdot [\sum_{\{S \leq k\}} z_8^S] \cdot \sum_{\{q \leq k\}} z_4^q \sum_e z_5^e \sum_{\{f < e\}} z_6^f \sum_{\{g \leq \min(q, 4f+2)\}} z_7^g$$

Set every $z = 1$ and it gives the count. Set every $z = z$ and it gives the rank polynomial. And the coefficient of $z_1^n \dots z_8^S$ is 1 if the cell exists and 0 if it does not.

The coefficient function of F is χ . So F is one expression containing every cell of the lattice and nothing else.

a caterpillar: a path of seven, one pendant — the simplest tree that is not a path




Figure 6.1. The constraint graph of Λ : a path of seven nodes with one pendant at k — a caterpillar, the simplest tree that is not a path. That shape is why a single left-to-right nesting exists with one bracketed factor. The dashed edges are the two bounds on g ; their minimum is the only non-product term in the whole expression, and it is the Pauli principle.

6.4 The shape of the tree

$n - l - k - q - g - f - e$, with $2S$ pendant at k

Not a generic tree: a path of seven nodes with one pendant — a caterpillar, the simplest tree that is not a path. **That is why a linear nesting exists with a single bracketed factor.** A branching tree would require nested brackets; a cycle would require inclusion–exclusion.

6.5 The one coupling, and it is the Pauli principle

Every bound is a single inequality except one: $g \leq \min(q, 4f+2)$. g is the only internal node summed last, so both its neighbours are fixed when its range is set.

Remove the coupling and 976 becomes 1,000. The excluded 24 cells all have $f = 0$ and $g = 3$ — **three electrons in an s orbital**.

Every other bound in this lattice is a counting or ordering fact. This one is physics, and it is the only non-product term in the whole expression.

And it is recovered, not assumed. $\mathcal{R}(\Lambda) = \Lambda$ rebuilds all 976 cells from the cells alone, so a reader given only the cell list would deduce $g \leq 2(2f+1)$ without being told it.

6.6 The detachable factor

2S is a leaf, so its sum closes geometrically and multiplies out:

$$[\sum_{S \leq k} z_8^S] = (1 - z_8^{k+1}) / (1 - z_8)$$

Spin multiplicity is algebraically inert here. It scales every cell count by $(k+1)$ and changes no structure — **which is why removing it leaves $E(X) = 0$ unchanged**, a fact §3 verifies numerically and does not explain.

6.7 Why a closed expression exists at all

Three structural facts, proved separately in this book and never combined:

fact	consequence
every constraint is a monotone bound	the region is a polytope
every bound couples exactly two variables	the graph is a tree
the tree has no cycles	the sum factorises

Remove any one and no such expression exists. A sum bound breaks the first — §6.7 measures 89,864 join failures when one is tried. A three-variable constraint breaks the second. A cycle breaks the third, and 2^d inclusion-exclusion terms return.

"All the information in one expression" is not a curiosity. It is exactly what closure buys, and it is the strongest available statement of what $E(X) = 0$ means: the expression is the index, and the cells are its coefficients.




Figure 6.2. The rank polynomial $F(z)$, read forwards as bars and backwards as the dashed line. They part company at rank 1 — 5 against 4 — and never rejoin. A rank polynomial is palindromic if and only if the poset is self-dual, so this single picture carries both of §3.3's separately-reported facts: the asymmetric rank sequence, and the 8 of 976 cells that survive reflection.

6.8 And the rank polynomial says what §3.3 says

A rank polynomial is palindromic if and only if the poset is self-dual.

forwards: 1, 5, 15, 34, 59, 87 · backwards: 1, 4, 10, 21, 37, 57

§3.3 reports as separate numerical facts that only 8 of 976 cells survive $x \rightarrow \max - x$, and that the rank sequence is asymmetric. They are the same statement, and they part company at rank 1 — 5 against 4.

6.8.1 And $F(-1) = 2$ has a cause, not a coincidence

Two things had to be ruled out before the number meant anything.

A pairing argument cannot cancel it. The duality map $\sigma(x) = \max - x$ gives $\text{rank}(\sigma^*x^*) = M - \text{rank}(x)$ with $M = 20$, which is even, so σ preserves rank parity. **The eight self-dual survivors are all of even rank and contribute +8, not 0.**

And the source is the tree. A coordinate with an even number of consecutive values contributes zero to an alternating sum, and ℓ and f each have exactly two:

factor at $z = -1$	value
$\Sigma(-1)^n, n = 1 \dots 3$	-1
$\Sigma(-1)^\ell, \ell = 0 \dots 1$	0
$\Sigma(-1)^e, e = 1 \dots 3$	-1
$\Sigma(-1)^f, f = 0 \dots 1$	0

$F_{\text{box}}(-1) = 0$ exactly. If the coordinates were free the alternating sum would vanish. The constraints couple l to n and f to e , the vanishing factors never appear free, and a residue of 2 survives.

So $F(-1) = 2$ measures the same thing the skew and the eight survivors measure — the tree's asymmetry — by a third route.

And $F(-1) = 2$ is the third face of it. A self-dual graded poset gives $F(-1) = 0$.

7. The shape — a cylinder over the transfer

7.1 A cylinder, not a Möbius band

Λ 's derived quantities fall into two classes:

ascending — rise up a series	descending — fall up a series
e, v, V	$T, r, \delta, \text{spacing}, w$

exchanged under $v \mapsto v^{-3}$.

The sign structure is bipartite, and a bipartite structure can never produce an odd number of orientation reversals. Verified exhaustively: **zero reversing loops among all cycles of length 3 to 5.**

Λ is orientable. It is a cylinder.

That is a *reason*, not an absence. And both alternative routes to a Möbius band close: a non-monotone quantity has **undefined** edges rather than reversing ones; a quantity independent of v has **no** edge. There is no third option.




Figure 7.1. The object. Ascending quantities rise up a series, descending ones fall, and the two are exchanged under $v \mapsto v^{-3}$. The sign structure is bipartite, so no traverse can reverse orientation an odd number of times — the surface is a cylinder, not a Möbius band, and the reason is structural rather than an absence of evidence.

7.2 "Over ν ", not "graded by ν "

ν grades every chain of Λ . It grades Λ itself nowhere — 24 of 70 comparable pairs have ν decreasing.

That near-miss is the structure in one sentence. The bracket of Part III operates along chains *because that is exactly where ν is a grading*.

The name follows: **the Lach Cylinder — a hydrogenic cylinder over ν , and over the transfer**. Not *graded by*. *Over*.

7.3 One scale, and one exception

Height goes as ν^{-2} , the metric as ν^{-3} , and the only linearly rising quantity in the entire structure is the price of a guarantee. Chapter 16 derives it.

7.4 The loop closes

$\Lambda_8 \rightarrow$ derived quantities $(V, r) \rightarrow$ the bound lattice $\rightarrow$ dimension $1 = \log r = -3 \log \nu$, with residual 8.9×10^{-16} .

Λ *decomposes* ν and gains dimension; the bound lattice *recomposes* it and sheds dimension. **Two sides of one object, not two passes over it.**

7.5 What Λ is a property of

Λ is a property of the hydrogenic limit. δ — the quantum defect — is the entire remainder.

Prediction lives only in the remainder.

Which is why perfecting the index moves nothing, and why that is a theorem rather than a disappointment. Part II proves it; Part III measures what is left.

And §18.6 will sharpen this into its final form. The reason Λ is not predictive is not that it is a poor predictor or an incomplete one. **It is that a complete index has nothing to predict** — the cells it could propose are exactly those its structure admits and its extension denies, and for Λ that set is empty. **Completeness and predictiveness are alternatives.**

PART II — THE LAW

7.6 The cylinder is a fibration, and its fibres are two-body

§7.1 finds the shape is a cylinder over the transfer coordinate. That is not a metaphor, and the reason is structural.

The tree has two ends. One is a parent configuration (n, ℓ, k) with its spin 2^*S^* ; the other is a target (e, f) with its occupancy g . **They meet at q , the number transferred.** That is the shape of a two-body problem: two objects and a coupling.

7.6.1 It does not factorise, and then it does

As a bare product it fails: $|A| \times |B| \times |q| = 2,244$ against 976. The two ends are coupled twice over — $q \leq k$ ties the transfer to the parent, $g \leq q$ ties the target's occupancy to the transfer.

Conditioned on the coupling, they separate exactly:

$q \mid A(q) \mid B(q) \mid \text{product actual}$

0 33 5 165 165

1 33 10 330 330

2 23 15 345 345

3 8 17 136 136

$|A| = \sum_q |A(q)| \times |B(q)| = 976$, exactly.

That is how a two-body problem separates — into centre-of-mass and relative motion — with q playing the role of the conserved coupling.

7.6.2 Which names the cylinder

The base is q , the axis along the cylinder. The fibre is $A(q) \times B(q)$, the cross-section at each q . And the cross-section is itself a product of two lattices.

Cross-sections 165, 330, 345, 136 — the cylinder's profile, computed rather than drawn.

Cross-sections of the tower cylinder at each transfer, $\sqrt{v}$ -cell scale — Λ_{13} solid, Λ_8 dashed




Figure 7.2. Cross-sections of the tower cylinder at each transfer, $\sqrt{v}$ -cell scale — Λ_{13} solid, Λ_8 dashed inset. The sections read $2,294 \times 5$, $2,294 \times 20$, $1,794 \times 50$, 744×70 against the Λ_8 profile above, 33×5 , 33×10 , 23×15 , 8×17 . Every section is itself closed, per §7.8.4; the tower of §7.11 widens each without breaking one.

7.7 The fibres, in closed form

Each side of the cylinder has its own expression, and they are different shapes.

$$A_q(z) = \sum_{n=1}^3 z^n \sum_{\ell=0}^{\min(n-1,1)} z^\ell \sum_{k=\max(q,1)}^{\min(4\ell+2,3)} z^k \cdot (1 - z^{\min(k,3)+1}) / (1 - z)$$

$$B_q(z) = \sum_{e=1}^3 z^e \sum_{f=0}^{\min(e-1,1)} z^f \cdot (1 - z^{\min(q,4f+2)+1}) / (1 - z)$$

Both verified against enumeration at every q : $A_q(1) = 33, 33, 23, 8$ and $B_q(1) = 5, 10, 15, 17$.

7.7.1 And they are not the same tree

$A_q : n - \ell - k$, with 2S pendant at k — a caterpillar $B_q : e - f - g$, capped by the base — a path

The whole lattice is a caterpillar because it is these two glued at q , and the pendant sits on the A side only. [The tower of §7.11 amends this: with the coupling axes carried, the object is two caterpillars glued spine to spine at $k-q$ — A runs $n-\ell-k-J_c-K-J$ with 2S pendant, B runs $e-f-g-v-2S'$ — and the §7.8.3 compression generalises to the glued form.]

7.7.2 A tower of fibrations, terminating in boxes

level	base	fibre
Λ	q	$A_q \times B_q$
A_q	k	$(n, \ell)\text{-set} \times 2^*S^*\text{-chain}$
B_q	f	$e\text{-set} \times g\text{-chain}$

Three levels, and at the bottom every fibre is a product of chains — a box, with a closed form $\text{Box}(a,b)(z) = \prod_i z^{\{a_i\}} (1 - z^{\{b_i-a_i+1\}}) / (1 - z)$.

7.8 Five things the fibre table says

The counts 165, 330, 345, 136 carry more than their sum.

7.8.1 A Pareto frontier inside the lattice

$|A_q|$ falls 33, 33, 23, 8 while $|B_q|$ rises 5, 10, 15, 17.

Raising the transfer costs the parent and pays the target, monotonically, and no q improves both. That is the same shape as §16.5's $d^*w^*/d^*h^* > 0$ with $d^*V^*/d^*h^* < 0$ — **which the book presents as a fact about brackets and is here a fact about the index itself.**




Figure 7.3. The Pareto frontier on a log axis, Λ_{13} solid against Λ_8 dotted: A falls and B rises at both scales, and no q improves both. The section $|A \times B|$ peaks at $q = 2$ with 89,700 cells, and the tower shifts (q) from 1.463 to 1.887 — coupling multiplicity re-weights the profile toward higher transfer, per §7.11.5.

7.8.2 A most-probable transfer

The fibre peaks at $q = 2$ with 345 cells — 35.3% of Λ — and the sequence is log-concave, hence unimodal.

$\langle q \rangle = 1.4631$, sd 0.930 — the mean number of electrons moved across all admissible configurations, and $S'(1)/S(1)$ in the base variable.

7.8.3 An 8-to-3 compression that loses no count

Λ 's F has eight variables; the shape has three, and $S(1,1,1) = 976$. Not a projection that forgets cells — one that forgets which coordinate inside each side and keeps how many. A general fact about caterpillars: a path with one pendant reduces to (base, left, right) whatever its length.

7.8.4 Every cross-section is itself a closed, modular lattice

q	side	closed	$E(X)$	modular
0–3	A and B, all eight	true	0	true

The shape is not a cylinder over an arbitrary set. It is a cylinder whose every cross-section satisfies this book's own law.

7.8.5 Which gives $E(X)$ a local form

$E(\Lambda) = 0$ and $E(A_q) = E(B_q) = 0$ for every q .

The second does not follow from the first. A closed index could in principle have defective slices that cancel in the total. **Here none do**, and that is a new invariant.

7.9 Every other shape the lattice contains

Intervals. $[[x,y]]$ is a box iff no constraint binds across it:

$I_{\{x,y\}} = \text{Box}(x,y) \iff \text{for every constraint } i \leq f(j), y_i \leq f(x_j)$

Exact on 60 of 60 tested, and only **27%** of random intervals qualify — **the other 73% are where the physics is active**. The binding rates rank the constraints:

constraint	binds
$g \leq q$	35.6%
$q \leq k$	33.0%
$g \leq 4f+2$	4.9%

The most active constraint is the coupling; the least is the Pauli bound — which is why §11.8 finds it the cheapest to violate. Two independent measures, one answer.

Rank levels. Eighteen antichains, the largest **122 at rank 11**: the biggest set of configurations no two of which can be reached from each other by adding electrons. **Λ is exactly eight times its widest level.**

Maximal chains. A distributive lattice is $J(P)$; **its maximal chains are the linear extensions of P** . With seventeen join-irreducibles and a chain length of seventeen:

$e(P) = 1,113,045,672$

The seventeen generators of §3.3 and the chain length of seventeen are the same seventeen, and the book reports both without connecting them.

And one shape absent. No Möbius band, no cycle, no non-planar minor — **the constraint graph is a tree, so there is nothing to twist and nothing to close**. §7.1's finding follows by construction, not by search.

7.10 An axis the index did not carry, and now does

A_q carries 2S, a spin label on the parent. B_q carried none.

The index recorded the parent's spin and not the target's.

That was never a closure defect — $\mathcal{R}$ cannot see a coordinate that was never supplied. It was an axis E2 admits, and the asymmetry of the two fibres is what made it visible.

7.10.1 Λ_9 , built by the same construction

The parent's spin is bounded by the parent's capacity coordinate, $2S \leq k$. So the target's is bounded by the target's: $2S' \leq g$. One pendant, on the other fibre, by the identical rule.

	Λ_8	Λ_9
cells	976	1,654
box	6,912	27,648
E(X)	0	0
sublattice	yes	yes
rank-modular	475,800 pairs	200,000 tested, 0 violations
constraint graph	a tree	a tree — 2S' attaches as a leaf
F(1)	976	1,654
F(−1)	2	2
rank levels	18, widest 122 at 11	21, widest 185 at 12
log-concave	yes	yes

And Λ_8 is the exact projection of Λ_9 — 976 cells, identical — with fibres of one to four preimages, mean 1.69. §24.2.4's dimension ladder gains a rung, and the extra coordinate carries real information.

7.10.2 What is still not established

The bound $2S' \leq g$ was built by analogy and not from spectroscopy. Whether it is the correct selection rule for a target's spin is a physical question, and this book has not answered it.

The lattice closes. That the coordinate means what its name says was untested when this section was written. §7.11.1 has since answered it: exact vector coupling confirms $2S' \leq g$ as a true bound — zero violations over every fibre — and shows it is not tight, in the two ways §7.11.2 names. What remains is a data lookup, recorded in Q: whether the measured target terms of the 153 channels saturate $\min(g, 4f+2-g)$.

7.11 The tower — Λ_9 to Λ_{13}

§7.10 added one axis by analogy. This section adds four more by battery, and what the batteries return is not five facts about five coordinates. It is one law, met five times, and a shape that survives it.

Two caterpillars glued at k-q: the tower's constraint tree, and the one edge that may not be drawn




Figure 7.4. The constraint tree of the full tower: two caterpillars glued at k-q, counting axes rimmed dark (exact bounds), coupling axes in the accent (envelopes) — the Chapter 2 dichotomy drawn — with the forbidden f...K bridge dashed and priced at 3.5%.

7.11.1 The axes, and what each one sees

axis	coordinate	bound carried	cells	density	what it distinguishes
9	2S'	$2S' \leq g$	1,654	63.7%	the target's multiplicity
9'	—	$+ 2S' \leq 2f+1$	1,561	67.5%	the same, one Pauli cut tighter
10	v	$2S' \leq v \leq g$	2,535	44.7%	pairing depth — seniority
11	2J _c	$2J_c \leq \phi(k)$	13,585	17.0%	the core's fine structure
12	K	$2K \leq 2J_c + 2f_{\max}$	70,905	31.4%	core-orbit orientation
13	2J	$ 2J - 2K \leq 1$	199,130	64.4%	the outer electron's spin bit

The table's own vocabulary, defined. Four terms in the rows above are used to define the coordinates and were never themselves defined, which is the failure S1 exists to prevent. They are defined here.

terms(ℓ^k) — the multiset of LS terms (2S, 2L) of the configuration ℓ^k , obtained by **microstate enumeration**: list every way of placing k electrons in the $2(2\ell+1)$ spin-orbitals, accumulate (2M_L, 2M_S), and strip complete (2S, 2L) blocks from the largest M_L down until the list is empty. This is the standard construction and the origin Chapter 2 already cites for $2S \leq k$.

f_{max} — the cap on f under §2.4, **not** the cell's own f. The distinction is the whole of §7.11.5: the bound $2K \leq 2J_c + 2f_{\max}$ has one parent and preserves the tree and the cylinder; the same bound written with the cell's f has two parents and breaks the factorisation by 3.5%. Writing f_{max} without defining it left the book's title resting on an undefined subscript.

$\phi(k)$ — the monotone envelope of $k \mapsto \max 2J$ over the parent shells admitted by the caps, so $\phi = \{1: 3, 2: 4, 3: 5\}$ at §2.4's. **This is not §1.1's ϕ_{ij}** , which is $\mathcal{R}$'s pairwise envelope $\max\{x_i : x \in X, x_j \leq v\}$. Both are monotone envelopes and that is the whole of the resemblance; one takes two coordinate indices and the other takes an occupancy. The collision is recorded rather than renamed, because renaming a symbol used 213 times is a larger risk than naming the

two senses once. Register 250.

Density, stated so it can be rechecked. For each axis, over every cell of the stage below that the lattice admits, count the values of the new coordinate that exact vector coupling realises and divide by the values the admissible bound allows. Three of the six exact sets are not the obvious ones and the column cannot be reproduced without them: seniority at axis 10 means the terms new at occupancy v — $terms(f^v)$ less $terms(f^{(v-2)})$ as multisets, not merely a term of spin $2S'$ existing at v ; the core's J at axis 11 is restricted to terms carrying the cell's own multiplicity $2S$, not to all terms of l^k ; and the target spin at axis 9 is the spin set of f^g by microstate enumeration. Under that definition all six entries reproduce exactly — 63.7, 67.5, 44.7, 17.0, 31.4, 64.4 — at the caps of §2.4 and cap-dependent as always. Every bound is §9.4-admissible; every stage is closed — exhaustively to Λ_{10} , by the §2.3 proof beyond it — and every stage projects exactly onto the one below. Independence is demonstrated, not assumed, at every axis: p^2 's 1S and 1D share a cell below axis 10 and split at v ; p^1 's 2P shares a cell below axis 11 and splits at J_c ; and every $K > 0$ cell splits at J . Each coordinate is a physical degree of freedom that no function of the coordinates below it can express — the same door through which every legitimate axis must enter, by Chapter 12's theorem.

The ninth axis, checked against the textbook and tightened once. Its exact set is the seniority ladder read edge-on, $\{g \bmod 2, g \bmod 2 + 2, \dots, \min(g, 4f+2-g)\}$ — the ninth fibre was indexing v before the tenth axis named it. The textbook anchors agree exactly: p^2 gives $\{^1S, ^1D, ^3P\}$, $2S' \in \{0, 2\}$; p^3 gives $\{^4S, ^2D, ^2P\}$, $2S' \in \{1, 3\}$. And Λ_9 is not the tightest admissible object: $2S' \leq 2f+1$ is §9.4-admissible and provably containing — $\min(g, 4f+2-g) \leq 2f+1$ by averaging — and Λ_9' closes at 1,561 cells. The 93 cells it cuts are exactly ($f = 0, g = 2, 2S' = 2$), an s^2 target carrying triplet spin. The price is structural: the constraint graph gains its first cycle, $f-g-2S'$, so Chapter 10's tree-propagation guarantee lapses and must be retested rather than inherited. **The tree or the tightness** — the language offers both only up to 1,561 against 1,654 cells.

Every stage closed exhaustively, by the book's own theorem. The upper tower was never pair-tested: Λ_{12} has 2.5 billion pairs and Λ_{13} has 19.8 billion, so closure above Λ_{10} rested on the form argument with sampling beside it. A.2 removes the need — X is closed iff $\mathcal{R}(X) = X$ — and $\mathcal{R}$ is a sweep of the ambient box, not of the pairs. Every stage now tested against **every cell of its ambient box**:

stage	$ \Lambda $	$ \mathcal{R}(\Lambda) $	E	ambient box swept
Λ_8	976	976	0	6,912
Λ_9	1,654	1,654	0	27,648
Λ_{10}	2,535	2,535	0	110,592
Λ_{11}	13,585	13,585	0	663,552
Λ_{12}	70,905	70,905	0	5,308,416
Λ_{13}	199,130	199,130	0	47,775,744

Forty-seven million ambient cells at the top stage, every one decided, and no stage admits a cell it does not contain. **The instrument was in the book from Chapter 9 onward and was never turned on the book's own tower** — the sampling caveat existed because the theorem was proved and then not used. Register 249.

The four-cap sweep, now run. Across caps (3,3,1,3,1), (4,4,2,4,1), (4,4,2,6,2) and (5,5,2,6,2): $\hat{\phi}$ stays monotone, every stage still projects exactly onto the one below, and the cylinder's factorisation $\sum_q |A(q)| \cdot |B(q)|$ stays exact at Λ_{13} — 199,130, 4,731,790, 40,310,170 and 77,083,771 cells, defect zero at every cap. The shape is not an artefact of the book's caps.

The tower at the shells the book had been discussing from outside it. Referee flag 1 asked how an s-and-p object licenses conclusions about d and f electrons. The tower was rebuilt at caps carrying both, by multiplicity weighting rather than enumeration, and every density recomputed:

caps	shells	axis 9	9'	10	11	12	13
------	--------	--------	----	----	----	----	----

(3,3,1,3,1)	s, p	63.7%	67.5%	44.7%	17.0%	31.4%	64.4%
(4,4,1,6,1)	s, p	54.2%	60.3%	33.2%	7.0%	36.7%	64.5%
(5,5,2,10,2)	+ d	49.4%	55.6%	28.4%	4.9%	28.2%	65.7%
(5,5,3,14,3)	+ f	46.6%	53.0%	25.9%	4.2%	22.4%	66.1%

The dichotomy survives every enlargement: the counting axes stay exact and all three coupling axes stay envelopes, at the d shell and at the f. Two things move, and both are informative. Every density except the thirteenth **falls** as the object grows, so the envelope gap widens with scale — the price of certainty is not a constant but an increasing function of the index's size, which no part of this book had predicted. And the thirteenth **rises**, 64.4 to 64.5 to 65.7 to 66.1, converging on the $\frac{2}{3}$ that §7.11.1 states as its nominal value. The gap between $\frac{2}{3}$ and the measured figure was attributed there to the $K = 0$ singlet; as the caps grow those cells thin out and the density approaches the nominal, which confirms the attribution by a route the book did not have when it made it. Registers 240 and 241.

And two facts of the tenth's battery worth the print. Its 10-box appears at caps (5,5,2,6) and not at the book's — the known small-cap effect, caught by §2.4's own discipline, as it caught the same absence for Λ_8 and Λ_9 . And the obstruction that broke closure at the ninth axis does not dissolve at the tenth; it moves up one level, onto v itself — $v \equiv g \bmod 2$ and the conjugation ceiling, both still outside the language. **Extension conserves the obstruction.**

The order matters and is not the order first written down. K precedes J: K's parents are J_c and f, so axis 13 is reachable from axis 11 without axis 12, and taking J first costs 25 points of density on its fibre. §7.11.4 gives the reason. And carrying both with all admissible bounds closes the 11-12-13 cycle in the constraint graph — a triangle in both senses at once: each coupling rule is a triangle inequality, and the axes themselves form a 3-cycle.

What the twelfth axis measures, and when it is real. On every $f = 0$ channel the K-set is a singleton — $2K = 2J_c$, the triangle collapses, and axis 12 degenerates into a copy of axis 11. Where it does not collapse, K measures the one thing no occupancy or magnitude coordinate can: the relative orientation of the core's angular momentum against the Rydberg electron's orbital plane — the coordinate anisotropic core-electron interactions resolve, and the one that separates series converging to the same core limit with different quantum defects. Its validity is conditional in exactly the right direction: K is a good quantum number when core splitting dominates exchange — high n, the bracket's domain — and degrades toward low n, outside it.

And the thirteenth is one bit, at the apex where the index and the atom converge. Given K, the J-fibre is a doublet everywhere except the $K = 0$ singlet: axis 13 sees precisely the outer electron's spin orientation, the weakest and last coupling, and nothing else. But it is the only coordinate in the entire tower that nature conserves exactly. Everything below it is bookkeeping that mixing breaks — configurations at 40-50% in the book's own exclusion (iii), LS terms in intermediate coupling, J_c and K outside the jK regime. J survives all of it, by rotational invariance. So the tower ascends from coordinates the language expresses exactly but the atom keeps only approximately, to a coordinate the atom keeps exactly and the language expresses up to one parity point. **The two ends of the tower trade places.**

7.11.2 Three excluded forms, and the fold has a name

Chapter 12 excluded two constraint forms — sums and differences. The tower requires a third: SYMMETRIC BOUNDS. The exact ceiling of every coupling axis is symmetric under the reflection of §3 — particle-hole conjugation, $\text{terms}(\ell^k) = \text{terms}(\ell^{(4\ell+2-k)})$, verified exactly for every k on the p and d shells — and a symmetric non-constant function is not monotone, so §9.4 cannot carry it. $\min(g, 4f+2-g)$ for spin; the conjugation ceiling for seniority; $\max 2J(\ell, k)$, unimodal with its peak at half filling, 25 at f^7 , zero at closure, for the core's J. One reflection, three appearances.

The envelope gap decomposes, and the three parts are distinct. At the caps of §2.4: axis 9's gap is parity 5, fold 1, both 1; axis 11's is ceiling 16, parity 24, triangle 17. Three machines — the REFLECTION (conjugation), the CONGRUENCE (fermion parity), the TRIANGLE (coupling additivity, one sum and one difference, precisely Chapter 12's two exclusions wearing physics) — and no coupling axis escapes all three. The exact coupling region is not a lattice in any presentation tried: as 2S' (50,592 failures), as

parity alone (52,080), as the min-cap (17,856), as pair count (7,254), as (2S,2L,2J) (2,443 meets) — and the triangle's first meet failure manufactures $J = \frac{1}{2}$ from $S = L = 0$. The componentwise minimum forgets that J is built from L and S.

And the law, three instances deep before the triangle made it four: **every coupling coordinate's exact physical bound requires either two parents or a congruence, and the tree carries only one.** Spin needed $\min(g, 4f+2-g)$; seniority needs the conjugation ceiling; the core's J needs a non-monotone ϕ . Each admissible extension is the monotone envelope of its physics.

One half of the triangle survives, and the four-cap check of §2.4 has now been run: the region $\{|2L-2S| \leq 2J \leq 2L+2S\}$ is JOIN-closed — zero failures at caps 6, 8, 10 and 12 — and meet-broken at every one, by 2,862, 12,489, 40,887 and 110,229 failing meets. Certainty survives upward and dies downward, which is §3's skew stated as an inequality. The check also sharpens the statement: impose the parity congruence as well and the join closure dies too — 1,848 failures at cap 6 — so the surviving half belongs to the parity-free triangle and to nothing larger.

7.11.3 The dichotomy

Counting coordinates close exactly. Coupling coordinates close as envelopes. There is no third kind in this object, and the first coupling constraint was already in §2.1.

The origins table of Chapter 2 has four rows. Pauli, the hydrogenic solution and counting give single-coordinate exact bounds — that is why $E(\Lambda) = 0$ was available at all. Vector coupling gives envelopes, provably and permanently, because its exact form is made of the three excluded shapes. The price of certainty is not paid once at the door; it is paid per coupling axis, and §7.11.1's density column is the receipt.

And the surplus is not waste. Every envelope cell is a tripwire: a measured term assigned a nonphysical $(v, 2S')$ or $(2S, 2J_c)$ pair lands in the gap and is caught before any energy is read — a D1 check of Chapter 11's kind, mechanical from the exact tables, available for all 153 channels. The price of certainty and the budget of self-defence are the same cells, counted twice.

7.11.4 Precision is path-dependent; physics is not

The exact J-set of a channel is the same whether J is reached through K or directly — recoupling, 12 of 12 (J_c, f) pairs — but the envelopes differ: 64.4% through K, 39.4% direct. The envelope does not commute with recoupling. The index remembers the route; the atom does not.

The privileged route exists because the outer electron's spin is $\frac{1}{2}$: the final triangle degenerates to the mutual bound $|2J - 2K| \leq 1$, which is admissible on both sides and closed, losing only the parity point $2J = 2K$. The smallest quantum of angular momentum is the largest the constraint language can hold exactly. And the route the language prefers is jK coupling — the scheme the atoms of Part III actually realise at high n, and for the same structural reason: the weak coupling goes last. [Interpretation, not theorem.]

7.11.5 The shape survives, and it has a boundary

Every stage of the tower factorises exactly over the transfer: sections $815 \cdot 3,260 \cdot 6,150 \cdot 3,360$ at Λ_{11} , $11,470 \cdot 45,880 \cdot 89,700 \cdot 52,080$ at Λ_{13} , log-concave, unimodal, peak at $q = 2$, $\sum_q |A(q)| \cdot |B(q)|$ equal to the direct count at every stage with zero defect. The object is still a cylinder — because every coupling axis hangs off one side of the tree. The mean transfer rises, $\langle q \rangle = 1.4631 \rightarrow 1.887$, as coupling multiplicity re-weights the profile.




Figure 7.5. The solid in silhouette: plan (parent) and side (target) views at $\pm v$ -cell scale, Λ_{13} filled against Λ_8 dashed. The parent side thins toward $q = 3$ while the target side thickens — §7.8.1's trade, seen as shape, at both scales of the tower.

The boundary: the tight two-parent K — bounded by J_c on the A side and f on the B side — is the first bridge between the cylinder's two ends other than q , and it breaks the factorisation by 2,475 cells, 3.5% of the product. The tree is the cylinder; tightening K is the move that spends it. The title of this book is a property of envelope precision. [And the two-body reading: sections separate while the coupling is isotropic; the tight K is the anisotropy; interpretation, not theorem.]

7.11.6 An index whose one prediction is false, and why that is a result

The $E1$ selection rule on $(2J, 2J') - |\Delta 2J| \leq 2$ — is a closed index, by the mutual-bound lemma. Remove the one cell nature forbids, $J = 0 \nrightarrow 0$, and closure fails by exactly four meets at every cap tested, and all four failing meets are the forbidden cell: $(0,1) \wedge (1,0) = (0,0)$. The defect is $E = 1$, and the one cell $\mathcal{R}$ restores is the one the atom vetoes — the photon's unit of angular momentum, with the parity half of the rule a congruence that breaks joins, the photon's odd parity. [The photon reading is interpretation; the $E = 1$, the four meets and the congruence are computed.]

This is the calibration §24.3.4 owes its reader. $E(X)$ counts proposals, not guarantees: here is an index whose single proposal is known false, vetoed by structure the index does not carry. The 540 of the Kreuzer-Skarke slice are proposals of the same kind, and are stated as such.

7.11.7 What the tower leaves open

Recorded in Q , and two of the four have since closed. The spectroscopic v -check per channel against measured terms remains open, blocked on data. Whether the earlier work's claim that $2J_c \leq k$ contains the physical region survives that construction's own definition remains open — under the definition computed here it fails at four cases, beginning with p 's $2J = 3 > 1$. The four-cap stability check has been run and is reported at §7.11.2 and §7.11.5; the self-duality count of §3 is resolved there as a difference of reflections, not a defect. What the sweep could not settle is the density column of §7.11.1, for the reason register 209 gives.

Figures 7.2–7.5 accompany the tower: the cross-section gallery at §7.7 (fig-sections), the transfer trade on a log axis at §7.8.1 (fig-pareto), the glued constraint tree with the forbidden bridge drawn as forbidden at §7.11 (fig-tree), and the two silhouettes at §7.11.5 (fig-silhouette). Each is generated from the counts of §7.11.1 and §7.11.5 by stated code, and the build now fails on any missing glyph rather than warning.

8. What the title names

Twice the title has changed, and both changes record findings rather than preferences.

The book was called The Lach Index Lattice until its shape was computed. The shape is a cylinder over the transfer, so the object's name became The Lach Cylinder — and §8.2 gives the reason.

Then the register outgrew the object. By the end there were two hundred and ninety-seven withdrawals, six prime audits, twenty-three principles and eighteen protocols — and the demonstration on Λ had become the smaller half of the book.

The title is now The Method, and The Lach Cylinder is its subtitle.

8.1 Why that ordering and not the other

Because the method is what generalises and the object is what tested it. Λ closes, and so does the nucleon index, and the string partition function, and the Kreuzer-Skarke frontier, and Appendix E's index of this book's own mathematics. What travels between them is not the cylinder. It is: build the index, close it, read $E(X)$, record every failure.

And the book applies that method to itself. Chapter 24 audits the book as an index; §21.8 indexes the process; Appendix B indexes the rules; Appendix E indexes the mathematics and closes it at $E = 0$. A book whose method is applied to its own contents four times over is a method with a demonstration, not a demonstration with a method.

The method is the book, and the book is the method. The cylinder is the case that proved it could be done and the case that showed what it costs.

8.2 The cells are transitions, not states

Λ is fibred over q , the number of electrons transferred. Not over any property of a single configuration — over the movement between two.

One end is a parent (n, ℓ, k) with its spin $2S$. The other is a target (e, f) with its occupancy g . They meet at q .

Chapter 2 introduces both ends and never says the pair is the object. The shape says it: $|\Lambda| = \sum_q |A_q| \cdot |B_q|$, and there is no factorisation that does not pass through the transfer.

8.3 Which is why "index of transitions"

An index of states would fibre over a state variable. This one fibres over a difference, and every structural result follows from that:

finding where it comes from

the fibres are two-body products the two ends of the tree

$\backslash$	$A_q \backslash$	falls as	B	ris	the transfer is conserved between
		$\backslash$	$-$	es	them
			q		
			$\backslash$		

$g \leq q$ is the most active constraint it is the coupling itself

rank is modular the transfer redistributes without loss

8.4 And what the cylinder does not carry

The lattice supplies kinematics and no dynamics. What configurations exist, which transitions are admissible, what is conserved — all structural. Which transition actually occurs is not.

1,113,045,672 maximal chains satisfy every constraint in this book (§7.9). The Aufbau ordering is one of them. Structure does not choose; energy does, and energy enters only through Chapter 17's measured levels.

That is the sharpest statement of what the index can and cannot do, and it is a statement about its shape: **a phase space with the Hamiltonian removed can bound and refuse, and cannot predict.**

9. Closure

Everything in this part follows from one condition.

An index $X \subseteq \prod_i A_i$ is CLOSED if, for all $x, y \in X$, both $x \vee y$ and $x \wedge y$ lie in X , where join and meet are taken coordinatewise.

The coordinates must be **totally ordered** — chains. That restriction is real and is discussed in §11.3.

9.1 The characterisation

Theorem 9.1 (A.2). X is closed iff $X = \mathcal{R}(X)$, where $\mathcal{R}$ reconstructs value sets and pairwise monotone bounds from X 's own extension.

Proof. ($\Leftarrow$) $\mathcal{R}(X)$ is an intersection of sets $\{x : x_i \leq \phi(x_i)\}$ with ϕ non-decreasing, each closed under $\wedge$ and $\vee$; so $\mathcal{R}(X)$ is closed, and if it equals X then X is closed. ($\Rightarrow$) $X \subseteq \mathcal{R}(X)$ always, since each $x \in X$ witnesses its own bounds. For the reverse, each bound $y_i \leq \hat{\phi}_{ij}(y_i)$ supplies a witness $w \in X$ with $w_i \geq y_i$ and $w_j \leq y_j$; closure under $\wedge$ and $\vee$ of the witness family produces y itself. ■

Verified on 200 random constructions, 200 agreements — random subsets (0 closed, 0 fixed points), lattice closures (40 and 40), random down-sets (16 and 16). Where they disagree they disagree together; there is no case of one without the other.

Note the down-set row. Sum-bounded sets like $\{a + b \leq q\}$ are not closed, and $\mathcal{R}$ reports it correctly. §25.3 explains why sums are excluded.

9.2 $\mathcal{R}$ is a closure operator

Extensive, monotone, idempotent — 150 of 150 on all three axioms.

So $E(X) = |\mathcal{R}(X)| - |X|$ is a closure defect, a standard object in Galois theory, formal concept analysis and database dependency theory.

This book claims no novelty for the machinery. What is offered is the application: computing the closure defect of an *index*, and finding that the periodic table's is 36 while Λ 's is 0.

9.3 The scope, stated exactly

Theorem 9.1 requires each coordinate to be presented as a chain. That is not the same as requiring the factors to be chains: any distributive factor can be so presented, by Birkhoff, as the down-sets of its join-irreducibles — one chain per irreducible.

What fails is bundling. $\mathcal{R}$ reconstructs coordinate by coordinate, so folding two chains into one non-chain coordinate hides structure from it. In tests, the same sets presented bundled and presented decomposed both gave 80/80 agreement — the theorem survives, but only when the presentation respects it.

9.4 The three consequences

Closure yields three properties, developed in the next three chapters:

S — reference	S1 alphabet	S2 order	S3 bounds [weakened: §25.1.2]
D — defence	D1 $\mathbb{D}_{\text{phys}}$	D2 $\mathbb{D}_{\text{def}}$	D3 χ total
E — extension	E1 cells	E2 axes	E3 constraints

Nine mechanisms. Six follow from closure without further hypothesis. Two require something constructed — a tree, and disjoint derivations. One, E3, is an exact criterion that took three wrong statements to reach.

And E is derived: every extension question reduces to S and D. It carries no independent content, in the same way that the cost surface of Chapter 16 carries none beyond v .

10. Self-reference — alphabet, order, bounds

A closed index contains its own definition.

To rebuild an index from nothing but a bag of tuples, three things are needed. They are independent, and only one of them is hard.

10.1 S1 — the alphabet

$$\hat{A}_i(X) = \{ x_i : x \in X \}$$

A projection. **Always succeeds, no conditions.**

10.2 S3 — the bounds

$$\hat{\phi}_{ij}(v) = \max \{ x_i : x \in X, x_j \leq v \}$$

Read off the cells. For Λ_8 that is 56 functions, all recovered, and $\mathcal{R}(\Lambda) = \Lambda$ exactly. Always succeeds given the order.

10.3 S2 — the order, which is the hard one

Given the alphabet but not the ordering of its values, can the ordering be recovered?

The necessary condition is provable and dimension-free. Say u may precede v on axis i if, for every p in the fibre over u and every q in the fibre over v , the meet lies in $F(u)$ and the join in $F(v)$.

Lemma. If X is closed and $u < v$ on axis i , then u may precede v . Proof. Take x with $x_i = u$ and y with $y_i = v$. Then $x \wedge y$ has i -coordinate u , so its other coordinates lie in $F(u)$; and $x \vee y$ has i -coordinate v , so its other coordinates lie in $F(v)$. ■ Nothing in the argument uses $d = 2$.

At $d = 2$ the converse also holds and the criterion is exact: X is closed under some order iff "may precede" is a total order, and that order **is** the relabelling. Verified on 120 sets; 274 of 274 constructed total orders give closed sets.

Above $d = 2$ the converse fails. Being a total order on every axis separately is necessary and not sufficient — the axes' requirements can be mutually incompatible. That is developed in §13.4, because it is a limitation rather than a mechanism.

10.4 What does work above $d = 2$: propagation along the tree

If the constraint graph is a tree, each axis is constrained only by its parent. Root anywhere, fix the root by brute force, and extend with backtracking.

Λ 's orders are recovered from a scrambled bag of cells: 20 of 20, at two cap settings, on 216 and 976 cells with all eight axes independently permuted.

Cost $\sum_i |\mathbf{A}_i|!$, never the product $\prod_i |\mathbf{A}_i|!$.

Two earlier methods failed on the same task, and the failures are instructive. Alternating refinement across all eight axes at once: **0 of 12**. Greedy propagation without backtracking: 42% and 25%. **The tree says only neighbours interact; a method that ignores that recovers nothing.**

Figure 8.1 The index recovers its own order




Figure 10.1. Left: two of Λ 's eight coordinates as built, showing the monotone bound $q \leq k$. Centre: the same cells after all eight alphabets are independently permuted — the structure is invisible. Right: recovered. Beneath, the propagation order along the constraint tree. The recovery uses no external information: only the bag of cells.

10.5 So: Λ contains its own order

The unqualified claim — the index contains its own definition — holds for Λ because its constraint graph is a tree. For an index whose graph has cycles, S2 is open, and Chapter 22 states what is known.

The qualified claim is unconditional: $\mathcal{R}(X) = X$ recovers alphabet and bounds for any closed index, relative to a given order.

11. Self-defence — data, derivation, totality

A closed index contains its own contradiction.

11.1 The counting argument

Let p be the parameters an index assigns and $q = \Phi(p)$ the quantities derived from them.

$$\mathbb{D} = \dim q - \text{rank } \partial\Phi/\partial p \geq \dim q - \dim p$$

Proof. The Jacobian is $\dim q \times \dim p$; its rank exceeds neither. ■

Verified on seven dimension pairs and 420 random nonlinear maps. No closure required, no smoothness beyond differentiability.

Corollary. $\dim q > \dim p \Rightarrow \mathbb{D} \geq 1$.

$\mathbb{D}$ counts functionally independent relations the derived quantities must satisfy. Each is a test the index must pass. For Λ with $p = (T, n, Z, \sigma)$ and $q = (v, \delta, V, r, w, e)$, rank is 4 and $\mathbb{D} = 2$, spanned by

(a) $V = 4v/3$ (b) $w = V \cdot e$

11.2 The split, and it is canonical

(a) holds under the Rydberg law and fails under power-4, exponential and logarithmic laws. (b) holds under all four.

$\mathbb{D} = \mathbb{D}_{\text{phys}} + \mathbb{D}_{\text{def}}$, physical relations and definitional identities

Tested under four unrelated perturbations, verdicts identical. The split is not an artefact of which law one chooses to vary.

11.3 D1 — $\mathbb{D}_{\text{phys}}$ catches a wrong value in the data

Worked case. Two species in this work, Al II and K II, were first computed with $Z_{\text{eff}} = 1$. Both have doubly charged cores.

The bracket held 56 of 56 either way. Containment depends only on monotonicity of T , and Z cancels from it entirely.

What failed were the quantities the method does not use. δ came out near 4.9 with a spread of several units — impossible for a Rydberg series, where δ is near-constant. V departed from $4v/3$ by a **median of 50%**, where Na I gives 0.2%.

Corrected to $*Z^*_{\text{eff}} = 2$: $\delta = 1.271$ with spread 0.047 over ten members, and V to **0.86%** across the fifty cells.

The bracket verifies the data. V verifies the indexing. An index carrying only what its method consumes would have passed fifty-six wrong cells.

Figure 9.1 The wrong nuclear charge: containment is blind, δ is not
the bracket cannot see it




Figure 11.1. The bracket holds fifty-six of fifty-six under either charge, because Z cancels from containment. The defect does not: at $Z = 1$ it spreads over several units, which no Rydberg series can do. The method was blind; a quantity the method never uses was not.

11.3.1 The D1 tripwire, specified in full and needing no data

Referee flag 5 says the tripwire has never run on anyone else's data. The reply is that **the data is completely defined by the index**, so the check can be stated in full here and run by anyone against anything. What §7.11.1 computes as *exact sets* are, read the other way, the complete list of physically admissible assignments; everything else the coordinates offer is a tripwire cell.

subshell	k	admissible 2S'	rejected / offered
s	1	{1}	1 / 2
s	2	{0}	2 / 3
p	2	{0, 2}	1 / 3
p	3	{1, 3}	2 / 4
p	5	{1}	5 / 6
d	4	{0, 2, 4}	2 / 5
d	7	{1, 3}	6 / 8
d	9	{1}	9 / 10
f	4	{0, 2, 4}	2 / 5

Over s, p, d and f complete, **150 of 216 offered assignments are rejected — 69.4%**. A measured term carrying a 2S' outside its subshell's set is nonphysical and is caught before any energy is read. The same holds one level up: given (subshell, k, 2S') the admissible seniorities are fixed — d⁴ with 2S' = 2 admits $v \in \{2, 4\}$ and nothing else — and given (subshell, k, 2S) the admissible core J are fixed, d³ with 2S = 3 admitting $2J_c \in \{1, 3, 5, 7, 9\}$.

This is what this chapter means by the surplus being the defence. The envelope cells are not waste; they are the region a wrong assignment must land in, and the index knows their extent exactly without being shown a single measurement. Flag 5's answer is therefore not *the check has been run* but *the check is fully specified and requires nothing from anyone*. Register 242.

11.4 D2 — D_def catches a wrong derivation, but only under two

routes An identity is vacuous if its terms are computed along the same path. $w = V \cdot e$ cannot fail if V is defined as w/e. It becomes diagnostic when w comes from the levels and V from the law.

$\mathbb{D}_{\text{def}}$ is realised iff a quantity is reachable by two disjoint paths, and the defence is the DISAGREEMENT between them — not the relation.

Worked case, and it is the author's own error. In first stating relation (b) I wrote $w = 3v \cdot e$. The two values — 540.0 from the levels, 1215.0 from the mis-stated law — appeared on the same printed line.

The check caught the check. $\mathbb{D} \geq 1$ does not require the auditor to be right.

Worked case, in an apparatus. Antiprotonic helium, cell (35,33):

route	value
measured directly	$804,633,059.0 \pm 8.2$ MHz
two-photon minus a different single-photon	$804,633,057.8 \pm 10.6$ MHz

Agreement at 0.09σ , with no shared measurement. One laser resonance against a two-photon resonance minus an unrelated one. That is path-disjointness realised in hardware.

And a constraint on where $\mathbb{D}_{\text{def}}$ can come from. Λ 's constraint graph is a tree, so there is exactly one path between any two coordinates. The constraints supply no redundancy. $\mathbb{D}_{\text{def}}$ must be built from derived quantities, deliberately.

11.4.1 And the counting argument has a 1960 precedent

This book did not discover that a fit with as many parameters as observations cannot be checked. Edlén states it, and uses it.

"Because of the approximate nature of the theoretical formulae here discussed nothing would be gained by trying to derive the parameters by a least-squares solution from all the observed levels. ... As the number of parameters introduced is equal to the number of levels no check is possible, and, of course, the results cannot be displayed in a two-dimensional diagram." — B. Edlén, Handbuch der Physik XXVII, footnote to §17 (manuscript 1960)

That is $\mathbb{D} = 0$, in the special case $\dim q = \dim p$. The counting argument, the conclusion drawn from it, and the use made of it — declining a fit that would have looked more rigorous — are all identical to this chapter's.

What §11.1 adds is the general form, $\mathbb{D} \geq \dim q - \text{rank } \partial\Phi/\partial p$, and the observation that the rank rather than the dimension is what binds. The principle is Edlén's.

11.5 D3 — totality catches a missing value

Theorem. If Λ is closed, $\chi_\Lambda : \prod A_i \rightarrow \{0,1\}$ is total: membership is decided for every ambient point by a finite conjunction of decidable comparisons. Proof. Each bound compares two elements of a chain. Λ is a finite intersection of such conditions. ■

Verified on every ambient point — all 6,912 of them, $\chi \in \{0,1\}$ everywhere, nothing undecided, and χ agreeing with membership on every cell. An earlier pass drew 30,000 uniform samples from a box of 6,912, which is more draws than cells and still not a proof of totality; the exhaustive test costs less than the sampling did. Register 248.

Worked case. A species was misidentified — Na II read as neutral neon, and the wrong ionisation limit applied. Every channel had $T = I - E \leq 0$ and **every one was refused at entry**.

Totality does not detect that error. It forbids it. The computation never ran.

Worked case, the other way. A guard in the author's code returned `None` when an alphabet exceeded a cap, and the caller coerced `None` to `False`. That produced "3 reorderable of 85 trials," which produced "almost never reorderable," which produced a false dismissal of an entire line of evidence.

One silent discard propagated through three conclusions, none of which looked wrong, because the arithmetic was consistent throughout.

$\mathbb{D} \geq 1$ cannot catch that. Every relation checks out on data that was never measured. Only totality closes the hole, because undetermined is not in the codomain.

11.5.1 And totality extends beyond the data

χ is usually read as a test on cells. It is a test on whatever the index carries, and §21.8 applies it to a coordinate this book did not originally have: the process by which a choice was made.

Index (step, visible, alternatives, committed) under **visible $\leq$ committed** — that nothing of an answer may be visible before a choice is written down — and **χ refuses** a prediction developed with its answer in view. **36 cells, closed, $E = 0$** , one constraint of §11.4 form.

The mechanism is unchanged. Only the coordinate is new.

That is the general shape of D3, and it is why the protocol of B.2.9 — refuse rather than coerce — is stated for quantities and applies equally to decisions.

11.6 The three defences are disjoint

mechanism	catches	worked case
$\mathcal{R}(\Lambda) = \Lambda$	rules not recoverable from cells	periodic table, $E = 36$
$\mathbb{D} \geq 1$	a wrong value in data or derivation	wrong Z ; $w = 3v \cdot e$
χ_{Λ} total	a missing value coerced to a real one	the cap discard

No mechanism catches another's class, and this work supplied an instance of each.

11.7 Two kinds of absence, and one kind of lie

An index can be wrong in two ways, and they are not alike. One it can detect by itself. The other it cannot detect at all.

11.7.1 A missing cell — and the index can tell whether the gap is real

$E(X) > 0$ says the structure admits cells the collection does not hold. It does not say why, and there are two reasons:

the gap is an artefact of the layout — a property of how the index was drawn the gap is a fact about the world — nothing occupies that cell because nothing can

And the index distinguishes them, by asking whether E is removable.

the periodic table	E
18-column	36
32-column	106
Janet left-step	0

The same 118 elements. Order the blocks f, d, p, s and define the period by $n + l$, and the defect vanishes. So the 36 belong to the drawing.

Now the same test on a catalogue where the absences are physical. Index the mean-motion resonances by (p, q) , reduced, with the order $p - q$ derived rather than made an axis, and give $\mathcal{R}$ only the resonances that show structure in the asteroid belt:

the belt	E
(p, q)	2
(order, p)	1
(p, order)	1

E is irreducible, and $\mathcal{R}$ refuses all six absences outright — 0 of 6 admitted, against 36 of 36 admitted for the periodic table.

And the physics confirms what the index inferred without being told it. Every resonance of order ≤ 7 in the belt is occupied; every one of order ≥ 8 is empty. Thirteen of thirteen above the cut, zero

of six below. Resonance strength enters at $e^{|p-q|}$, so the absences are exponentially weak and genuinely unoccupied.

A defect of the index vanishes under some re-coordinatisation. A feature of the world survives every one. The test is computable from the cells alone, and §1.2's Janet table is that test — run once, on one example, and never named as one.

11.7.2 A complete index of something that is not there

The second failure is different in kind, and no re-coordinatisation touches it.

Insert one impossible cell — **one electron with spin multiplicity 2, which no atom has** — and let the closure supply everything that follows. The result:

	Λ	the fabrication
cells	976	1,051
closed under join and meet	true	true
$E(X)$	0	0
fixed point of $\mathcal{R}$	true	true

And E stays 0 for both under every re-coordinatisation tested. The test of §11.7.1 has nothing to work on:

that test finds missing cells and asks whether the gap is real. A coherent fabrication has no missing cells. It is full.

The fabrication contains 75 configurations no atom can have, and the structure records nothing about it. Every law in this book certifies it.

11.7.3 The plain version

A defective index is a map with a street missing. You catch it by redrawing — the street reappears in another projection. A coherent fabrication is a complete, consistent map of a city that does not exist. Every street connects, every distance adds up. Redrawing changes nothing, because nothing is missing.

The only thing that distinguishes a map of a real city from a perfect map of an imaginary one is going there. In this book, going there means measuring a level.

11.7.4 Which is what D_{phys} is for, and only that

§11.3's five recorded failures are all of the second kind. Not one is a missing cell.

failure	what it does
a wrong nuclear charge	every derived value shifts together ; the index stays consistent
a wrong ionisation limit	every term value shifts together
a mislabelled parent	the whole channel moves as one

Each is a complete index of something that is not there, and each was caught by a number the index could not supply — a quantum defect of 4.9, a term value below zero.

Self-consistency detects an isolated error and certifies a systematic one. That is the whole case for requiring a second route that reaches outside, and it is narrower and sharper than "an index cannot check itself."

11.8 What self-reference does and does not defend

$\mathcal{R}(\Lambda) = \Lambda$ is not a curiosity. It is the guarantee that the index's answers and its contents are the same object, so it can never be short of an answer and never give one that contradicts itself.

11.8.1 A complete index answers every question about its contents

Asked whether it admits a fabricated cell — one electron with spin multiplicity 2, which no atom has — Λ **answers from its own content: $\varphi(2S | k) = 1$ at $k = 1$, the cell asks for 2, refused.**

200 of 200 fabrications refused when Λ is asked.

11.8.2 What it cannot defend against is being replaced

Rebuild the index around the fabricated cell and the closure grows to 1,051, is a fixed point, closed under join and meet, has $E = 0$, satisfies every law in this book — and contains 75 physically impossible configurations.

200 of 200 admitted when the index is rebuilt around them.

The index never believed the lie. It was never consulted. Replacement is not a question put to an index; **it is an act performed on it by whoever holds it.**

operation	outcome
ask Λ	refused, 200/200
rebuild around y	admitted, 200/200

11.8.3 Which is what B.2.13 already says

Commit before you look. Fix the standard, then admit the data. E1, E2 and E3 say the same for cells, axes and constraints:

an extension must be checked against the index, not absorbed into it and re-derived.

Ask the index, then decide whether to extend it. Never extend first and ask afterwards.

11.8.4 And the cost of a fabrication is computable

Insertion is never free. Adding one cell forces the closure to admit a median of 340 more — minimum 59 across 220 trials, zero exceptions.

$A(y) = N[\phi(\Lambda \cup \{y\})] - N[\phi(\Lambda)] - 1$, the lattice's own count evaluated on two bound systems. Exact on all cells tested.

Substitution gains nothing. Deletions are always repaired, so $C(\Lambda - x + y) = C(\Lambda + y)$ identically — **0 of 300 random swaps and 0 of 1,600 targeted swaps evaded detection.**

attack	outcome
delete	repaired, always
insert one	amplifies, always — minimum 59
substitute	reduces to insertion; no gain
insert the whole amplified closure	undetectable

V prices certainty. A prices deceit. Both are the lattice counting itself, and both have a floor.

And the amplification is governed by the fibre the widened bound admits, truncated by any rival bound on the same variable — which is why $g \leq 4*f^*+2$, the only constraint sharing its variable with a stricter one, is an order of magnitude cheaper to violate than any other. **Remove the rival and it jumps 397% into the normal band.**




Figure 11.2. The cost of a single fabrication, over 140 random insertions. Every one forces the closure to admit further cells — median 352, minimum 15 — and none is free. The dashed line at zero marks where an invisible lie would sit; no insertion reaches it. A forger must therefore supply the whole amplified set, and its size is computable from the index before any data arrives.

And the general form is the tree, not the constraint. Four statistical measures of a constraint — leaf position, load, marginal contribution, activity — all failed to predict A . **A measure of the tree succeeds:**

$A \approx \text{fibre} \times (0.972 - 0.154 \cdot \text{descendants of the bounded variable})$ $\text{corr}(\text{descendants}, A/\text{fibre}) = -0.942$, median error 7%, worst 24%.

constraint	descendants	A/fibre
$2^*S^* \leq k, g \leq q, g \leq 4^*f^*+2$	0	0.96-0.99
$q \leq k, f \leq e-1$	1	0.66-0.95
$k \leq 4l+2$	3	0.48
$l \leq n-1$	4	0.39

A widened bound on a variable with descendants forces theirs to widen too, so the fibre undercounts in proportion to how much of the tree sits downstream. **Leaves are exact; the root is off by a factor of 2.5.**

11.8.5 Two closure operators, and the book names one

$\mathcal{R}$ recovers 16 binding coordinate pairs, not 7. Its transitive reduction returns the seven defining edges exactly — so the tree is derivable from the cells and nothing is lost.

But the two systems, identical on Λ , differ on $\Lambda \cup \{y\}$:

	A full	A tree-only
$l \leq n-1$	119	307
$k \leq 4l+2$	269	563
$2^*S^* \leq k$	74	74

The two that agree bound leaves — no descendants, no transitive propagation. $\mathcal{R}$ is the full pairwise operator, because for the calendar and the periodic table no constraint graph is given and the tree operator is not definable at all. §6's nested sum counts the tree system; they agree on Λ and nowhere else.

12. Extension — cells, axes, constraints

A closed index determines what may be added to it.

Three operations, and only three: change the cells, change the axes, change the constraints.

12.1 E1 — which cells may be added

$\Lambda \cup \{x\}$ is closed iff for every $y \in \Lambda$, both $x \vee y$ and $x \wedge y$ lie in $\Lambda \cup \{x\}$

100% over 288 tests. A cell may be added exactly when it sits in a hole whose joins and meets fall back into the lattice.

12.2 E2 — which axes may be adjoined

$\Lambda \times h$ is closed iff h is a lattice homomorphism

100% over 424 tests. Projections, maxima, minima and constants qualify; sums, products and differences do not.

And a theorem in the other direction:

Theorem 10.1 (adjunction never repairs). For any $h : S \rightarrow H$, with $S' = \{(x, h(x))\}$: S' closed $\Rightarrow S$ closed. Proof. For $x, y \in S$, the join of $(x, h(x))$ and $(y, h(y))$ is $(x \vee y, h(x) \vee h(y))$. If S' is closed this lies in S' , so $x \vee y \in S$. Meets likewise. ■ Contrapositive: a non-closed set cannot be repaired by adjoining any function of its coordinates. Tested on seven functions including the identity and a constant; none repairs.

Three repair routes remain, and only three: **enlarge** (add cells), **restrict** (remove cells), **reorder** (change neither). The calendar of Chapter 1 is repaired by reordering, at a cost of usability.

12.3 E3 — which constraints may be imposed

This one took three wrong statements to reach, which is recorded in Chapter 21. The criterion is:

$\{h \leq q\}$ is closed $\iff h(x \vee y) \leq q$ for all x, y with $h(x), h(y) \leq q$

100% over 2,513 tests, nine functions.

The essential feature is that it is q -dependent. Admissibility belongs to the *pair* (h, q) , not to h alone — which is why $a + b \leq 3$ can be closed where $a + b \leq 4$ is not, and why every q -free criterion attempted before this one failed.

Meet never breaks it. For monotone h with $h(x), h(y) \leq q$, $x \wedge y \leq x$ gives $h(x \wedge y) \leq q$ at once. **Only the join can break it**, and it does so when the join raises h above both arguments.

Sums and products fail because a join raises every coordinate at once, and anything accumulating across coordinates accumulates the raises.

Λ 's seven constraints are all safe by inspection: each reads one coordinate, or a minimum of two. **None sums.** That is why Λ is closed and the periodic table is not — and it is checkable without touching a cell.

12.4 A worked extension, and what it costs

The clearest demonstration of E3 in this work is an attempt to widen Λ that failed, and the repair that succeeded at a stated price.

The problem Λ_8 carries one target triple (e, f, g) . Configurations with two occupied non-source subshells fall outside it — an 86% coverage ceiling. The natural repair is a second triple.

Ordering the targets is free The canonicalisation $e_1 \leq e_2$ is admissible: $\emptyset$ is the identity, non-decreasing in a single coordinate. Zero join and meet failures.

Conservation is not The physical requirement $g_1 + g_2 \leq q$ is a bound on a sum, hence outside §11.4 — and it breaks closure. 89,864 join failures in 979,300 sampled pairs, meets unaffected.

The counterexample is two lines. $a = (*g^*_1=1, *g^*_2=0, q=1)$ and $b = (*g^*_1=0, *g^*_2=1, q=1)$ each conserve. **Their join transfers two electrons having removed one.**

Dropping conservation restores closure — 45,690 cells, zero failures — and produces a lattice describing configurations that cannot exist. **That is the choice in its bluntest form: a closed index of impossible states, or an open index of real ones.**

The repair: make the sum a coordinate Conservation is a bound on a sum, which §2.1 excludes. So make the sum a coordinate. Replace (g_1, g_2) by (g_1, G) with $G = g_1 + g_2$:

$g_1 \leq G$ (non-decreasing in G) and $G \leq q$ (non-decreasing in q)

Both of §11.4 form, with $*g^*_2 = G - *g^*_1$ recovered as a derived quantity and **conservation holding identically by construction.**

Tested exhaustively: 1,040 cells, all 540,280 pairs, zero join and zero meet failures.

And the price The map $(g_1, g_2) \mapsto (g_1, g_1+g_2)$ is a bijection onto an admissible lattice and order-preserving in one direction. It is not an order isomorphism. It makes comparable some pairs that were incomparable: $(0,1)$ and $(1,0)$ are incomparable as transfer counts and become ordered in the new coordinates. Fifteen of forty ordered pairs are new.

Coverage and closure are available together. Coverage, closure, and the original order are not.

This is P16 — treat bounds as coordinate values — operating as a repair procedure, and it is the general lesson of E3 stated constructively: a sum bound cannot be imposed, but a sum can be **indexed**, and the difference between those two moves is the order it costs you.

12.5 E is derived

	mechanism	fixed by
	E1 addable cells	S3, the recovered bounds
	E2 adjoinable axes	D3, totality — h must create no third state
	E3 impossible constraints	S2, the recovered order

Every extension question reduces to S and D. E carries no independent content — exactly as the cost surface of Chapter 16 carries none beyond v .

13. What the law forbids

The results of this chapter are the strongest in the book, and every one of them is negative.

The strongest arrived last, from the tower: **exact vector coupling is forbidden in principle, not unachieved.** The triangle's two inequalities are one sum and one difference — precisely the two forms this chapter excludes — so every coupling coordinate, at every height of the tower, enters at envelope precision, forever. §7.11.2 carries the computation; the first failing meet manufactures angular momentum from nothing, which is the physics stating the theorem's premise back.

13.1 An index carries what its coordinates carry, and no more

Theorem 11.1. On a product order, join, meet and comparability are definable from the coordinates: $(a \vee b)_i = \max(a_i, b_i)$, $(a \wedge b)_i = \min(a_i, b_i)$, and $a \leq b \iff \bigwedge_i (a_i \leq b_i)$.

Every lattice-theoretic quantity is therefore a deterministic function of the coordinates, and the σ -algebra it generates is contained in theirs.

13.1.1 And the product structure is not what does the work

A reader may suspect the result is an artefact of writing join as max. It is not.

Theorem 11.2. Let an order be defined as any predicate on coordinate tuples. Then the σ -algebra generated by its join, meet and comparability is contained in the σ -algebra generated by the coordinates. Proof. The relation $a \leq b$ is by hypothesis a function of the coordinates of a and b . The set of upper bounds of a pair is computed from that relation, and the least upper bound is selected from that set. Each step is a function of the coordinates, so the join is one — whatever formula it takes. Meets likewise. ■

Verified on a genuinely non-componentwise example from this work. The conserving poset of §13.1.1 has 28 of its 511 joins differing from componentwise maximum, because q must rise to admit both transfers: $(1,0,1) \vee (1,1,0) = (2,1,1)$, where componentwise maximum would give the inadmissible $(1,1,1)$.

The join is nonetheless a deterministic function of the two cells' coordinates, as the theorem requires.

No ordinal structure over these coordinates carries what they do not — not a tree, not a non-componentwise lattice, not the conserving poset.

The escape is not a cleverer order but a further measurement. An order not determined by the coordinates requires information beyond them, and that information is by definition **a new coordinate**.

13.1.2 Two consequences that run through the whole book

It is indifferent to dimension. Λ_{11} is as constrained as Λ_8 . Adding axes does not loosen the bound.

It is indifferent to faithfulness. Perfecting the index moves nothing, because there is nothing to add. Part I could be improved indefinitely and Part III would not gain a single cell.

13.1.3 What the theorem does not establish

Theorems 11.1 and 11.2 bound information, not estimation. A deterministic function of the inputs can still improve a finite-sample predictor — that is what feature engineering is, and it is routine.

So the theorems forbid the lattice from knowing anything the coordinates do not. They do not by themselves forbid a lattice quantity from being a better-conditioned regressor or a lower-variance summary. That is a narrower claim than the theorem's reach suggests, and it is stated narrowly.

13.2 v is inadmissible as an axis $v = e - \delta$ is a difference, and differences are non-monotone in the cell order. 86 violations at the caps tested.

So the quantity on which the whole of Part III depends **cannot be a coordinate of Λ** . It lives over the lattice, not in it. Chapter 7's "over, not graded by" is the same fact stated geometrically.

13.3 Depth is not a coordinate function

The number of measured members in a channel is not a function of the coordinates at all. It is a property of the literature, and Chapter 14 makes that precise.

13.4 Closure is not locally determined

This is the deepest limitation, and it was found by trying to evade it.

Closure is not implied by every proper projection being a fixed point of $\mathcal{R}$.

Explicit counterexample at $d = 3$, $c = 2$, six cells:

$S = \{(0,0,0), (0,0,1), (0,1,0), (0,1,1), (1,0,1), (1,1,0)\}$ Every proper projection is a $\mathcal{R}$ -fixed point, yet $(1,0,1) \vee (1,1,0) = (1,1,1) \notin S$

And the counterexamples are common — roughly **one in five** sets passing all proper-projection tests fails closure.

The mechanism. The two cells agree on axis 0 and disagree on *both* others. The join disagrees in a **combination**, and no lower-dimensional projection can see a combination it does not contain.

Closure is a d -dimensional condition. No projection of it will serve as a criterion — not per-axis, not per-pair.

Consequence for verification: a wrong choice can pass every local test and fail only at the top. Chapter 22 states what this costs.

13.5 The pole

The cost of a guarantee, derived in Chapter 16, has a pole at $p = 1$ — the exponent at which an observable is linear.

A straight line has no curvature. There is no interpolation error to price, so a guarantee costs infinitely more than the thing it guarantees.

The method has no domain at $p = 1$. It is the only hard singularity in the structure; every other limit is a threshold with data on both sides.

13.6 $E(X)$ is the prediction budget

The five limitations above add up to something the book can now state in one line, and it answers the question this work began with.

A prediction is a claim that a cell exists which the index does not carry. The only cells an index can propose are those its own structure admits and its extension denies — and that set is exactly $\mathcal{R}(X) \setminus X$.

The number of predictions an index can make is $E(X)$.

13.6.1 What that means for the four indices of Chapter 1

index	$E(X)$	predictions available	and they are
the periodic table	36	36	all impossible
the calendar	7	7	all impossible
a subnet with a hole	8	8	all impossible
the box ordering, and Λ	0	0	—

Every prediction the three non-closed indices can make is wrong. Period-1 group-2 and its thirty-five kin. The thirty-first of February and its six. **That is the entire budget, and none of it exists.**

Closure does not make prediction accurate. It abolishes prediction.

13.6.2 And it is forced, not observed

Two results already in this part combine to give it.

$E(X) = 0$ says there is no missing **cell** to propose. **Theorem 11.2** says there is no hidden **quantity** to derive.

Together:

A fully defined index has nothing left to say.

That is the whole answer to whether Λ is predictive. Not that it is a weak predictor. Not that it wants more coordinates. **There is no gap in it for a prediction to occupy.**

13.6.3 Where prediction actually lives

Λ indexes cells. It does not index values.

The Sc VI 6s cell of §18.6 is in Λ — its coordinates are ordinary integers, and the lattice admits it without comment. **Its energy is not**, and no closure property supplies it.

the cell	indexed, $E = 0$, nothing to predict
the value	not indexed, and therefore predictable

δ is precisely the value's departure from what the index determines.

Prediction requires an incompleteness, and a complete index has none.

Which is why the bracket must be handed two measured neighbours before it can say anything at all. It is not extracting a prediction from the lattice. **It is carrying a constraint between two measurements the lattice never had.**

13.6.4 The trade, stated plainly

An index may be complete or it may be predictive. It cannot be both.

Not because completeness is dull, but because a prediction is a proposal about an unlisted cell, and completeness is the property of having no unlisted cells. The periodic table can propose thirty-six things because it is incomplete. **Λ can propose nothing because it is not.**

This is the fifth and largest of the negatives, and everything else in this part is a special case of it.

13.7 Why the negatives are the contribution

Λ is a bijective recoding of existing quantum numbers. It adds no physics. A reader will observe this within a page, and they will be right.

That is the point. A recoding that adds nothing and still exhibits closure, self-reference, self-defence and these five limitations is demonstrating a property of **indexing**, not of atoms. If Λ contained new physics the results would be about atoms and worth much less.

14. Retrieval as a lattice problem

Data can be retrieved in pieces; the index navigates.

This is usually read as advice. It is the self-defence theorem again, with a different failure mode.

14.1 The formalisation

Let C be the cells wanted and $S = \{S_1 \dots S_k\}$ the sources, each a subset of C , with $A(S_i)$ recording accessibility.

c is retrievable iff $\exists i : c \in S_i \wedge A(S_i)$ route set $R(c) = \{i : c \in S_i\}$, retrieval redundancy $\rho(c) = |R(c)|$

$\rho = 1$ is fragile. $\rho \geq 2$ survives the loss of any single source.

$\mathbb{D}_{\text{def}} \geq 1$ — two disjoint derivations survive a computational error $\rho \geq 2$ — two disjoint sources survive an access failure

Same structure. Different failure mode.

14.2 A paywall blocks a source, not a cell

Worked case. The target was a transition frequency in antiprotonic helium.

source	status
<i>Nature</i> 475 , 484 (2011)	paywalled
<i>Phys. Rev. Lett.</i> 96 , 243401 (2006)	paywalled
arXiv:1304.4330	open — the <i>Nature</i> paper entire
arXiv:1203.5425 (CODATA 2010)	open — Table XII, all fifteen frequencies, both papers
arXiv:1308.1711, an unrelated citing paper	open — corroborates to 0.1 MHz
CERN CDS	open

$\rho = 6$. Two blocked, four open. The cell was retrieved and the two-path check on it closed at 0.09σ .




Figure 14.1. Retrieval redundancy for five target cells. The 2011 antiprotonic-helium frequency had six routes, two blocked. The four German documents have one or two apiece, all closed.

14.3 The secondary can be better than the primary

CODATA's Table XII carries both papers with an extra digit the journals did not print, supplied privately by the experimenters to reduce rounding error.

The index does not merely route around an obstruction. It sometimes finds a better cell than the one that was blocked.

14.4 The retrieval graph is self-referencing

Each fetched compilation names its successors. The aluminium compilation's introduction lists its own sister volumes verbatim; the beryllium compilation cites the zinc one; the aluminium reference list supplies nickel and copper.

Three coverage gaps in this work's census were filled from documents already in hand, with no new search.

14.5 Where it fails, and the number that describes it

Four documents in this work were not retrieved: Edlén's Handbuch chapter (1964), Ritz's Physikalische Zeitschrift paper (1908), Paschen & Götze (1922), Dunz (1911). $p \leq 2$ for each, every route closed.

Not for lack of effort — nine searches, correct citations, the right platform identified and free. **Pre-digital print has low retrieval redundancy.**

p is a property of the literature, not of the searcher. A cell in a 1908 German periodical has $p \approx 1$. A cell in a 2011 physics paper has $p \approx 6$.

A century of open deposition raised p roughly sixfold, and that is exactly what separated the two outcomes.

14.6 The procedure

1. Enumerate target cells from the index before searching. 2. List route candidates per cell — primary, preprint, review, compilation, citing paper, deposit, database. 3. Prefer high- p cells; they cost less and often carry more. 4. Read every retrieved source for its successors. 5. When a route is blocked, move along the route set; do not re-attempt. 6. Record p for cells that fail. A blocked $p = 1$ cell is a stated gap, not an unexplained absence. The author violated step 1 badly, on his own problem, after writing this chapter. That is recorded in Chapter 21.

PART III — THE METHOD

15. The bracket

Everything so far has been about structure. This part asks what such a structure can do with measurements, and the answer is deliberately modest.

15.1 The rule

For a Rydberg channel — fixed parent, fixed ℓ , varying n — the binding energies $T = I - E$ are monotone in n . So for any interior member:

$T(n)$ lies between $T(n-1)$ and $T(n+1)$

That is the bracket. It uses **only measured neighbours**. It fits nothing, assumes no functional form, and returns an interval rather than a value.

It is a deduction, not a prediction. The distinction is the whole method: a fit says where a level probably is; the bracket says where it cannot fail to be.

15.2 The four rules

The bracket is stated in one line and applied under four.

Rule 1 — Interiority. Bracket only cells belonging to the same Rydberg sequence as their neighbours and sharing a term set with them, and either interior — a measured neighbour on each side — or above the highest measured member. Extrapolation below the lowest measured member is forbidden. Rule 2 — Variable. Read the interpolation variable off $T = Z^2R/\nu^2$ for the axis in question: E along n within a series; interval ratios along J ; geometric in T along the core shell. Rule 3 — Separation. Report the deductive bracket separately from the inferential point estimate. Given monotonicity, the true value cannot lie outside [max below, min above]. Rule 4 — Uncertainty. Fit δ locally by a Ritz expansion and take $\sigma = 2R Z_{\text{eff}}^2 \cdot \text{SE}_{\text{pred}} / \nu^3$ from the fit's prediction standard error. Use $\delta = \delta_0 + \delta_2/n^2$ where three or more neighbours exist, $\delta = \delta_0$ where only two do — subject to the admissibility clause below. Rule 4a — Admissibility of the constant form. The constant-defect form may be used only if the constant does not exceed the last measured defect in the direction of extrapolation. On a decreasing channel it must satisfy $\delta \leq \delta(\text{last measured})$; on an increasing one, $\delta \geq \delta(\text{last measured})$. Where it does not, the two-point Ritz solve is the admissible form.

15.2.1 Each rule justified by ablation

A rule that costs nothing to break is not a rule. Each was tested by violating it:

rule violated	cost
extrapolate below the series	median error $0.46 \rightarrow 206 \text{ cm}^{-1}$ — 446×
extrapolate above the series	median error $0.46 \rightarrow 0.13 \text{ cm}^{-1}$ — 0.28×
wrong interpolation variable	rms $200 \rightarrow 588 \text{ cm}^{-1}$ — 2.9×
pooled σ instead of per-cell	$\pm 1\sigma$ coverage 96.2% instead of 69.6%
include a collapsed orbital	median error $12.8 \rightarrow 41.3 \text{ cm}^{-1}$ — 3.2×

The extrapolation cost is directional, and a single figure conceals it. An earlier version of this work quoted 3.7× on pooled rms over twelve extrapolated cells. **That statistic is not reproducible and is unstable:** on the same data it is 18.8× on rms, 4.6× on median, and **0.5× if the two worst cells are removed**. An rms over twelve points containing a $12,942 \text{ cm}^{-1}$ outlier is not a summary that should carry a rule.

Split by direction the instability disappears. Extrapolating *below* a series costs 446× on the median, because δ varies fastest at low n and a Ritz fit run backwards diverges. Extrapolating *above* costs 0.28×, because δ has already flattened. **The asymmetry is a factor of 1,577**, and Rule 1 is stated accordingly.

15.2.2 Rule 1 is not about n

The rule is stated for a series in n , and that is where it was found. It is not a rule about n .

The isoelectronic axis — fix the electron count, vary the nuclear charge — carries the same defect and shows the same asymmetry:

axis	toward the loose end	interpolate	toward the tight end
n , within a series	446×	1×	0.28×
Z , along a sequence	17.3%	1.3%	2.1%

Same sign, same magnitude of asymmetry, same cause. δ varies fastest where the electron is loosely bound — low n , low Z — and flattens where it is tightly bound. **Extrapolating toward the flattening is cheap; away from it is not.**

Rule 1 is a rule about direction, and it holds on whichever axis carries the defect.

Measured on two sequences, holding out each member in turn. The lithium-like neutral is predicted from its three ions at 11.1% error; the sodium-like neutral at 23.4%. **The tight ends — C IV and Si IV — come back at 1.8% and 2.5%.**

15.2.3 And one step beyond the fitted range costs more again

A hold-out is bracketed by the points that remain. A genuine prediction is not. Tested on the sodium-like sequence with P V — a fifth member — held out entirely:

fit form	predicted δ (P V)	true	error
$\delta = a + b/Z$	0.8189	0.7129	14.9%
$\delta = a + b/Z + c/Z^2$	0.7376	0.7129	3.47%
$\delta \cdot Z$ linear in Z	0.7835	0.7129	9.9%

Only the quadratic form is usable, and it costs 3.5% rather than the 2.1% a hold-out suggests — a factor of 1.7 for stepping outside the fitted range.

This matters for §18.6 The Sc VI prediction has a second, independent route through its isoelectronic sequence, and **3.5% is the figure that route must carry** — not 2.1%, and not the 1.3% of pure interpolation.

15.2.4 Rule 1's second clause is not redundant

Sc III 3d is interior in n within $\ell = 2$, and including it inflates the median error by 3.2× — because it is the collapsed ground state rather than a sequence member.

The operational test is the defect. Its δ departs from its neighbours' by 0.14 where they depart from each other by 0.008 — **a factor of 17**. A cell can be interior by index and outside the sequence in fact, and the defect says which.

15.2.5 Rule 4 at the edge of a channel

The two-parameter Ritz form needs three neighbours to leave a residual, and a channel with exactly three members has an interior cell with only two.

Why Rule 4a exists The constant-defect form takes the mean of the available defects. On a monotone channel that mean lies between them — and therefore on the wrong side of the last one.

Sc VI's 4S^{ns} channel is the case that forced the clause. With $\delta(4s) = 1.0057$ and $\delta(5s) = 0.9812$, the constant form gives $\delta = 0.9934$ — **above $\delta(5s)$** , which a decreasing channel forbids. The resulting estimate falls **outside the deductive bracket its own monotonicity implies** (§18.6).

A rule that produces an estimate its own bracket excludes is not a rule about uncertainty. It is an error.

With two neighbours the Ritz form is exactly determined — two equations, two unknowns, no residual and therefore no fitted σ — but it is *admissible*, and Rule 3's separation means the bracket carries the guarantee while the point estimate carries none.

One further case in this work: the Kr I ($^2P^{\circ}_{1/2}$) n^*p series — 5p, 6p, 7p, with 6p interior. Under the constant-defect form the cell returns an estimate in error by -70.9 cm^{-1} against $\sigma = 602$, or **0.12 σ** .

It is a valid cell and an almost worthless one. Its bracket spans $14,602 \text{ cm}^{-1}$ because its neighbours are two principal quantum numbers away.

That is the validity/value distinction, and it is the clearest instance of it in the collection: a cell may satisfy every rule and still tell you almost nothing. **The rules govern correctness. They do not govern usefulness**, and a method that conflated the two would report this cell as a success.

15.3 The bracket is limit-free

Since $T = I - E$ for a constant I :

$$T(n) \text{ between } T(n \pm 1) \Leftrightarrow E(n) \text{ between } E(n \pm 1)$$

Containment is invariant under the affine map. **v , δ and V all need the ionisation limit; containment does not.**

Consequence: the method applies wherever levels are known, even when the threshold is not. This is not a technicality — it is what allowed the antiprotonic-helium cells of §18.6, where the relevant threshold is ambiguous.

15.4 The price of limit-freedom, computed

§25.3 presents limit-freedom as pure gain. It is a trade, and the terms are computable.

The alternative is to bracket δ rather than T — which is what the field does. And V is identical either way:

$$\delta = \delta_0 + \delta_2/n^2 \text{ has } L = |\delta'/\delta| = n/3, \text{ the same as } T\text{'s, and } \delta_0 \text{ is an additive constant that cancels from both } w \text{ and } e. V(\delta) = V(T) = 4n/(3h)$$

The single most standard move in Rydberg spectroscopy — interpolate δ , not T — improved estimates enormously and left the price of a guarantee exactly where it was. A century of practice raised numerator and denominator by the same factor.

15.4.1 What δ buys, and what it costs

In δ the bracket is narrower by exactly $n^3/(2\delta_2)$ — measured against predicted: 4,329 vs 4,267 at $n = 8$; 1.387×10^6 vs 1.386×10^6 at $n = 55$. Agreement to four figures.

But that factor is the size of a Ritz coefficient, not a property of bracketing — δ barely moves because δ_2 is small.

And a bracket on δ does not bound T . Converting requires the ionisation limit and inherits its uncertainty:

n	w_T direct	via δ	+ limit $\sigma = 0.01$
20	68.06	0.00102	0.0210
40	7.613	1.43×10^{-5}	0.0200

At $n = 40$ the δ -route's own width is three orders below the limit's uncertainty. The limit dominates entirely.

Bracket T: bounds the measured level, needs nothing. Bracket δ : bounds a derived quantity, needs the limit, and the energy bound is never better than the limit is known.

The choice is conditional, and the condition is computable per channel. This book's tightest bracket is 1.398 cm^{-1} ; limit uncertainties in the collection run 0.001 to 0.4 cm^{-1} . **The δ route wins where the limit is known better than the bracket — true for some species, false for others.**

15.5 Admissibility

A channel is admissible when its levels are separated by more than their uncertainty:

$$r = 2Z^2R / (v^3\sigma) \geq 5$$

r falls as v^{-3} . Every channel eventually leaves the domain; the rule states where.

15.6 What the bracket costs

A guarantee is wider than an estimate. Chapter 16 prices it exactly, and the price is the only linearly rising quantity in the structure.

16. The cost surface

16.1 The ratio

Let w be the bracket width $|T(n+1) - T(n-1)|$ and e the interpolation error $|T(n) - \frac{1}{2}(T(n-1) + T(n+1))|$.

$$V = w / e = 4v/3$$

for a Rydberg series. And in general, for $y = x^p$ sampled at step h :

$$V(x, p) = 4x / (h \cdot |p - 1|)$$

Verified on power laws from $p = -3$ to $+11$, agreement under 1%.




Figure 16.1. The exact cost against its asymptote. The floor $32/11 = 2.909$ belongs to the exact expression; the asymptotic form gives $26/9$ and is low by 0.69% at $v = 2$. They converge by $v = 10$.

16.2 The floor is algebraic before it is physical

The cost of a guarantee has a floor that owes nothing to atomic physics.

Proposition 14.1. Let three consecutive cells of a channel carry values $y_0 < y_1 < y_2$ at equally spaced index. Let $w = y_2 - y_0$ be the bracket width at the middle cell and $e = |y_1 - (y_0 + y_2)/2|$ the error of two-point linear interpolation from the same neighbours. Then $V = w/e > 2$, with $V \rightarrow 2$ only as one of the two steps vanishes. Proof. Write $d_0 = y_1 - y_0$ and $d_1 = y_2 - y_1$, both positive by monotonicity. Then $w = d_0 + d_1$ and $e = |d_0 - d_1|/2$, so $V = 2(d_0 + d_1)/|d_0 - d_1|$. Since $d_0, d_1 > 0$ we have $|d_0 - d_1| < d_0 + d_1$, giving $V > 2$. As $\min/\max \rightarrow 0$ the ratio $\rightarrow 2$; as $d_0 \rightarrow d_1$ it diverges. ■

This holds for any monotone sequence whatever. Not for some observables, not with better data:

The deductive interval is always more than twice the error of the crudest possible point estimate built from the same two numbers.

Every empirical V in this book is an instance of this bound, not an independent fact. And against a better estimator — the Ritz expansion of Rule 4, or a mass model — the gap is wider still, so V as defined **understates** the cost.

The Rydberg floor of $32/11 = 2.909$ is therefore the tighter of two. The algebraic floor at 2 constrains every monotone sequence; the physical floor at $32/11$ constrains Rydberg series specifically, and the difference between them is what the hydrogenic form contributes.

16.2.1 When the floor is approached, and what it means

Observation 14.2. $V \rightarrow 2$ requires $\min(d_0, d_1)/\max(d_0, d_1) \rightarrow 0$ — one step of the channel vanishing relative to its neighbour.

In every sequence examined here where that occurs, the small step is produced by two channels interleaving, which Rule 1's second clause excludes. So the algebraic floor is approached exactly where the method declines to operate: **the cheapest guarantee is available only on a channel that is not one.**

16.3 V is a bound, not a tool

V contains no σ . Improving measurement precision by six orders of magnitude leaves $V = 26.689$ unchanged at $\nu = 20$ — computed exactly at three precisions spanning $\sigma = 10^{-2}$ to 10^{-8} cm^{-1} .

No spectrometer can buy a narrower bracket. The price of certainty in a Rydberg series is set by the series.

So V is not an instrument. It is an impossibility statement, and it makes four:

	forbids
$V = 4\nu/3$	a guarantee costs $4\nu/3 \times$ an estimate, and no measurement reduces it
$V \geq 32/11$	no guarantee is ever cheaper than $2.909 \times$
pole at $p = 1$	no guarantee exists for a linear observable
ν_V ceiling	above it the cost cannot be <i>measured</i> , though it stays defined

The exact V is rational at every ν — $32/11, 54/13, 256/47, 250/37, 4000/299$ — and $4\nu/3 + 4/(9\nu)$ is its asymptotic form, low by 0.69% at $\nu = 2$ and exact to four decimals by $\nu = 10$. **The floor $32/11$ belongs to the exact expression, not to the asymptote.**

16.4 V is a function of ν/h alone

The cost has one free parameter, and it is not where the book first looked.

V is defined from three levels at $n - h, n, n + h$. Nothing forces $h = 1$: a cell may be bracketed from neighbours two or three principal quantum numbers away. Carrying h through the algebra gives an exact result.

$$V = 4v^3 / (h(3v^2 - h^2)), \text{ and with } r = v/h, V = 4r^3/(3r^2 - 1)$$

One variable. Z and R cancel identically, not approximately, and $4r^3/3$ is the large- r expansion of this expression rather than the expression itself.

16.4.1 Which makes depth and step interchangeable

(v, h)	v/h	V exact
(20, 1)	20	26.688907
(40, 2)	20	26.688907
(60, 3)	20	26.688907
(80, 4)	20	26.688907
(10, 1)	10	13.377926
(40, 4)	10	13.377926

Identical to six figures, and identical across $Z = 1, 2$ and 6 . That last point is worth dwelling on: $*Z^*_{\text{eff}}$ is precisely what the wrong-charge error of §11.3 got wrong and the bracket could not see. **The similarity law is blind to it too — by cancellation, provably.**

16.4.2 Tested on the collection

1,033 $(v/h, V)$ pairs from measured levels, at both steps:

	pairs	median deviation from $4(v/h)/3$
$h=1$	560	0.49%
$h=2$	473	1.15%
all	1,033	0.69%

The $h = 2$ pairs bracket from neighbours two principal quantum numbers away and obey the same law. And the outliers identify themselves: Al I at $n = 51$ and Ga I at $n = 53$ both sit above the v_V ceiling of §16.11, where quotation granularity destroys the second difference. **The method's own admissibility bound predicts which pairs fail it.**

16.5 The trade-off, and a conservation law that is not one

Since h is free, ask whether it can be used to buy a cheaper guarantee.

It can. V falls from 53.3 to 13.4 as h goes $1 \rightarrow 4$. **But the bracket widens in proportion**, because

$$w = 2y'h + O(h^3) \text{ — linear in } h \quad e = y''h^2/2 + O(h^4) \text{ — quadratic in } h$$

16.5.1 A withdrawn proposition

An earlier version of this section claimed $w \cdot V = w^2/e = 8y'^2/y''$ independent of h , and called it a conservation law. It is not. Exact symbolic evaluation for $y = x^{-2}$ gives

$$w^2/e = 16x^4/(h^6 - 5h^4x^2 + 7h^2x^4 - 3x^6)$$

which depends on h explicitly and equals $8y'^2/y''$ only as $h \rightarrow 0$. **The numerical table on which the claim was made showed 366.3, 367.9, 370.6, 374.5 across $h = 1...4$ — a 2.2% rise, read as a constant.** Withdrawn, and recorded in Chapter 21.

And no other candidate exists. Since $w \sim h$ and $e \sim h^2$, any $w^a e^b$ scales as h^{a+2b} , so **the only h -free combination is $a = 2, b = -1$** — the one tried. The leading-order invariant is unique and merely asymptotic.

16.5.2 What is exactly true

$dw/dh > 0$ and $dV/dh < 0$, for every monotone convex y and every h .

Exact, with no expansion and no limit. Verified over 172 monotone steps with **zero violations**. Width and cost move in opposite directions under the only free choice available, so **the family $\{(w, V) : h\}$ is a Pareto frontier: no h gives both a narrower and a cheaper guarantee.**

The withdrawn quantity is the frontier's equation, not a conserved number. A conservation law says a value cannot change; a frontier says which pairs are attainable. **The algebra is the same and the claim is weaker.**

16.5.3 And the frontier has an optimum

Opposite signs give any stated preference $\alpha w + \beta V$ an interior minimum:

$$h^* = \sqrt{2\beta / (\alpha \gamma'')}$$

Checked against direct minimisation: 0.0551 against 0.0550, 0.1743 against 0.1750, 2.7886 against 2.7450 — **and 8.82 against 7.83 at $v = 40, \beta/\alpha = 10$** , an 11% miss where h^* approaches $v/4$ and the leading-order derivation fails. **The formula holds for $h^* \ll v$ and not beyond.**

16.6 The fractional forms

The relations of §16.4 close, and the closure states two quantities the book had not.

$w/T = 4(h/v) + 8(h/v)^3$ — the bracket width as a fraction of the term $e/T = 3(h/v)^2$ — the interpolation error as a fraction of the term

$$v \gg h \quad w/T \approx 4h/v \quad e/T \approx 3(h/v)^2$$

$$20 \ 1 \ 0.201004 \ 0.200000 \ 7.531 \times 10^{-3} \ 7.500 \times 10^{-3}$$

$$40 \ 1 \ 0.100125 \ 0.100000 \ 1.877 \times 10^{-3} \ 1.875 \times 10^{-3}$$

$$40 \ 2 \ 0.201004 \ 0.200000 \ 7.531 \times 10^{-3} \ 7.500 \times 10^{-3}$$

Both are pure functions of h/v , and their ratio returns $V = (4/3)(v/h)$.

The second is the more useful and the more exposed. $e/T = 3(h/v)^2$ is the relative error of linear interpolation on a Rydberg series — a statement available to any spectroscopist since 1925. **It is not in the literature, and §18.6 explains why: the field abandoned the variable.**

16.7 The exponent axis

Two involutions act on p : cost, $p \mapsto 2 - p$ (paired values multiply to x^2); and virial, $p \mapsto -p$ (paired values multiply to a constant). Composed, they generate translation by 2, so exponents fall into two orbits by parity — and the parity is b , the number of energy denominators.

The pole at $p = 1$ sits in the odd orbit. Nothing built purely from powers of $\{r\}$ can reach it.




Figure 16.2. The cost surface. Every observable in this book is a slice at fixed p . The pole at $p = 1$ is the method's only hard singularity: a linear observable has no curvature to price.

16.8 The cost has a name in another field

And the name explains two things this book could only observe.

16.8.1 The quantity is the Newton decrement

The asymptotic form of §16.5 is not an accident of spectroscopy:

$$8y''/y'' = 8\lambda^2, \text{ where } \lambda^2 = \nabla f^T [\nabla^2 f]^{-1} \nabla f$$

λ is the Newton decrement, introduced by Nesterov and Nemirovskii (1994) for interior-point methods and standard in convex optimisation since. In one dimension $\lambda^2 = f''/f''$, and $\frac{1}{2}\lambda^2$ is *exactly the decrease in the quadratic approximation obtained by a full Newton step*. Verified symbolically at $p = -2, -3, 2, 11$: **the ratio is 8 in every case.**

For a Rydberg series it takes a physical value:

$$\lambda^2 = (2Z^2R/\nu^3)^2 / (6Z^2R/\nu^4) = (2/3) \cdot Z^2R/\nu^2 = (2/3) \cdot T$$

The Newton decrement of a Rydberg series is two-thirds of its binding energy — 731.5820 against 731.5820 at $\nu = 10$, 7.3158 against 7.3158 at $\nu = 100$. **And the optimum whose distance it measures is the series limit**, where $T \rightarrow 0$.

16.8.2 The term function is self-concordant

The Newton decrement carries its meaning only for self-concordant functions, $|f'''| \leq 2(f'')^{\frac{3}{2}}$. Nobody has asked whether $T(\nu)$ qualifies. It does:

$$24Z^2R/\nu^5 \leq 2 \cdot 6^{\frac{3}{2}} (Z^2R)^{\frac{3}{2}}/\nu^6 \iff \nu \leq (\sqrt{6}/2) \cdot Z\sqrt{R}$$

Z	self-concordant for
1	$\nu \leq 406$
2	$\nu \leq 811$
6	$\nu \leq 2,434$

The deepest channel in this book reaches $\nu = 55$. The ratio $|f'''|/2(f'')^{\frac{3}{2}}$ runs 0.0049 at $\nu = 2$, **0.1356 at $\nu = 55$** , and crosses unity between $\nu = 405$ and 406.

$T(v)$ is self-concordant across the entire measured range, with sevenfold margin.

16.8.3 Which explains the cancellation

§16.4 proved that Z and R vanish from V and gave no reason. Self-concordance is the reason.

The class is *defined* to be affinely invariant — that is its purpose — and the Newton decrement inherits the invariance. **Z^2R enters T as a multiplicative constant, a multiplicative constant is an affine map, and under an affine map the decrement cannot change.**

The book observed a cancellation. This says why it had to happen.

16.8.4 And the attribution the book owes

Interval analysis has asked this book's question since 1966, and named it. Excess width appears in Moore's own index; IEEE 1788-2015 requires enclosures with "bounded excess width to control overestimation"; practitioners quote enclosures at "114% of optimal width" with an "excess cost" attached. And the field states the trade in one sentence:

"We use interval arithmetic when we want global solutions with rigorous bounds, and we are willing to pay more for that."

That is Chapter 16's thesis, published sixty years earlier, about function ranges and rounding rather than interpolation from measurements. The framing is theirs. **What is not theirs is the ratio computed for a measured monotone sequence, the value $4v/3$, the floor $32/11$, and the identification $\lambda^2 = (2/3)T$.**

16.9 Which variable to bracket

Chapter 16 has treated V as a property of the method. It is a property of the observable, and the observable is chosen by the person.

16.9.1 The cost ranks by exponent

Since $V = 4x/(h|p-1|)$, the cost falls as the power law steepens. Evaluated at $v = 20$, $h = 1$:

observable	p	V
C_6	11	8.18
polarisability α	7	13.44
spacing	-3	20.04
observable	p	V
T , the term value	-2	26.69
C_3 , geometric cross-section	4	26.72
radiative lifetime	3	40.03
$\langle r \rangle$, dipole moment	2	80.00
blockade radius	11/6	95.99

A guarantee on C_6 costs a third of a guarantee on T , and a twelfth of one on the dipole moment.

16.9.2 Two branches, and where they cross

The asymptotic form appears to drive V below the floor of Prop. 14.1 as p grows. It does not — it stops being valid.

p	V exact	$4x^*/h$	$p-1$
4	26.72	26.67	

11	8.18	8.00
50	2.39	1.63
300	2.0000	0.27
10,000	2.0000	0.008

As $p \rightarrow \infty$, $(x+h)^p$ swamps everything: $w \rightarrow (x+h)^p$ and $e \rightarrow (x+h)^{p/2}$, giving $V \rightarrow 2$ **exactly**. The two branches cross where $4*x*/(h|p-1|) = 2$, that is $|p-1| = 2x/h$.

The absolute floor is reached only by power laws steeper than $2v/h$. At $v = 20$ that requires $p > 41$. C_6 at $p = 11$ is nowhere near it.

16.9.3 And the pole sits on the field's favourite variable

$v = n - \delta$ is nearly linear in n , and $p = 1$ is the pole.

n	V on v	V on T	ratio
10	1.09×10^5	11.58	9.4×10^3
20	1.77×10^6	24.89	7.1×10^4
40	2.84×10^7	51.54	5.5×10^5

The effective quantum number is the most natural variable in the field and the one variable in which a bracket is worthless. §25.5's pole is not a curiosity at the edge of the method; it sits on the coordinate everyone uses.

16.10 The order of the bracket

Monotonicity is not the only order property available, and using more of them changes what the bracket is — but not, in the end, what it costs.

16.10.1 A deductive bracket exists at every order

For $T = Z^2R/v^2$ every derivative has known sign, $\text{sign}(f^{(j)}) = (-1)^j$ — checkable on data as the sign of the j -th finite difference, with no fit and no smoothness assumption beyond a sign. With $k+1$ nodes and the classical interpolation error,

$\text{sign}(f(n) - p(n)) = (-1)^{(k+1)} \cdot (-1)^m$, with m nodes above n $k+1+m$ even $\rightarrow$ a lower bound.
Odd $\rightarrow$ an upper bound.

A two-sided deductive bracket at every order, by choosing m . Tested on $v = 8 \dots 119$ and $k = 1 \dots 5$:

560 containments, zero failures.

And the widths fall fast — at $v = 40$ from 0.362 cm^{-1} at order 1 to **8.9×10^{-6} at order 5**.

16.10.2 And the cost still grows

An earlier draft claimed V stays near 2 at matched order and that Q9's conclusion — bounds becoming relatively worthless as data accumulates — was an artefact. That was wrong twice over, and the correction is recorded at 96–98.

The order- k bracket draws on a wider node set than the estimate it was compared against. At $v = 40$, $k = 3$ the bracket uses **eight** levels and the estimate used **four**. On the same nodes:

v $k=1$ 2 3 4 5
 20 157 495 1,212 2,582 3,629
 40 673 4,037 2.2×10^4 8.5×10^4 3.2×10^6
 80 2,785 3.2×10^4 3.7×10^5 — —

V grows with order, and steeply. Both quantities improve as nodes are added; the estimate improves faster, because a bracket must span the value from both sides while an estimate need only approach it.

So the trade is the one Chapter 16 has argued throughout, and the order axis does not escape it:

A deductive bound is worth most when data is scarce. Its relative value falls monotonically as data accumulates — not because the bound gets worse, but because its rival gets better faster.

And Proposition 15.2's floor applies to the classical bracket only — first-order, against the linear interpolant on the same two neighbours, where the two objects are locked together. **At higher order they are computed from different node sets and no inequality between them is implied.**




Figure 16.3. Left: V on the same node set at each order — it grows steeply, from 157 to 3.2×10^6 . Right: at $v = 40$ both the bracket and the estimate improve as nodes are added, and the estimate improves faster; the shaded gap is the price of certainty, and it widens. An earlier version of this figure drew a floor at $V = 2$ through data reaching 0.24, and is withdrawn at correction 97.

16.10.3 What the order axis does buy

Not a cheaper guarantee. A narrower one, and a finer detector.

	order 1	order 5
bracket width at $v = 40$	0.362 cm^{-1}	8.9×10^{-6}
tolerates a displacement of	0.129 cm^{-1}	4.3×10^{-6}
cells from a 10-member channel	8	4

Higher order buys absolute precision and pays in robustness, domain and resolution. A collection assembled from real, perturbed, coarsely quoted spectra is exactly where order 1 belongs — and $V = 4v/3$ is the price of certainty that survives a perturbation of half the local spacing.

16.10.4 And the admissibility rule generalises

§16.13 gives $r \geq 5$ for the spacing and v_V for the curvature. Those are two instances of one rule:

order k is admissible only while $|\Delta^{(k+1)}T| > 5 \cdot 2^{(k+1)} \cdot \sigma$

Applied to the collection it separates arithmetic from physics. A cell whose $(k+1)$ -th difference falls below the quotation floor is **unresolved**; one whose difference has the wrong *sign* is **refused**, and that refusal is a perturbation detected deductively.

order	admitted	held	refused	median bound
1	499	499	19	25.96 cm ⁻¹
3	187	187	101	0.715
6	34	34	145	0.0164

Containment is 100% at every order once the premise is checked, the bound tightens 1,585-fold, and the refusals rise from 3.7% to 81% — the higher the order, the finer the perturbation it sees.

Which supersedes §18.5's 1,061 order-1 bounds: the same data at matched order yields 34 bounds at 0.0164 cm⁻¹, **and 619 refusals that are themselves perturbation detections this book did not count.**

16.11 The refusal pattern carries four verdicts, not one

§16.10.4 counts refusals and reads them as "a perturbation". The count is one number; the PATTERN is four.

refusals	V	what the family is doing
none	infinite	exact — linear in the parameter, two points suffice
none	finite	a smooth channel
one	finite	a bifurcation — the family changes character and does not return
regularly spaced	finite	an oscillation — the family returns
dense and irregular	undefined	chaotic
	d	

An index does not see a refusal. It sees a family, and the spacing of the refusals across it is what separates these. §11.7.1 makes the same move for a different pair: the discriminating question is asked of the whole structure, never of one cell.

16.11.1 On the collection

channel	membe	refusals	verdict
	rs		
Na I, K I, Al I — eleven channels	13-52	0	smooth
Ga I 4s ² np ² P°	15	1	a bifurcation
Li I np ² P°	41	6, irregular	repeated structure

Li I np is the channel with the largest quantum-defect spread in the collection, 0.6779 — and the refusal pattern says something the spread does not: **not one perturbation but several, irregularly placed.**

16.11.2 And a caution about the test itself

The classifier requires a tolerance relative to the sequence's own scale. Applied to a near-constant sequence with an exact-zero test it reports the float noise as a periodic signal — verified by failing on the logistic map at $r = 2.5$, a fixed point, which it called an oscillation. The admissibility rule of §16.10.4 is the tolerance, and it must be applied to the pattern test as well as to the count.

16.12 The optimal order is a boundary, not a peak

Information per cell rises monotonically in k , so no interior optimum exists. What stops the climb is σ .

$\Delta^{(k+1)}T$

$> 5 \cdot 2^{(k+1)} \cdot \sigma$ — the $(k+1)$ -th difference is a signed sum of $2^{(k+1)}$ measured

levels, so its accumulated quotation error is at most $2^{(k+1)}\sigma$.

$v \sigma k \mid \Delta^{(k+1)}T \setminus \text{floor}$

20 10^{-2} 2 9.43 $\times 10^{-1}$ 4.0×10^{-1}

20 10^{-3} 3 2.09 $\times 10^{-1}$ 8.0×10^{-2}

20 10^{-4} 4 7.60 $\times 10^{-2}$ 1.6×10^{-2}

40 10^{-4} 2 2.75 $\times 10^{-2}$ 4.0×10^{-3}

Better data raises k . It never creates an interior peak. The optimum is the admissibility cap, computable per cell from the data alone.

16.13 Inversion

$p = 1$

$= 4x/(hV)$, and the sign is fixed by $\text{sign}(y')$ and $\text{sign}(y'')$. Controls recover $p = -1$ and -3 correctly;

9 of 10 within 5%.

On real data, each recovered from its own measured series: Na I -1.993 , K I -1.998 , C $_3$ 3.999, α 6.990, C $_6$ 10.956.




Figure 16.4. Left: the fate of every cell in the collection at each order, once the admissibility rule $|\Delta^{(k+1)}T| > 5 \cdot 2^{(k+1)} \cdot \sigma$ separates arithmetic from physics. Right: the refusal rate. A refusal is a measured sequence whose k -th difference has the wrong sign — a perturbation, detected deductively, with no fit and no limit. The higher the order, the finer the perturbation it sees.

16.14 Two kinds of bound, and the book has been calling them

one Every constraint in Λ is a CAPACITY: $\ell \leq n-1$, $k \leq 2(2\ell+1)$, $g \leq q$. Every admissibility rule is a RESOLUTION: $r \geq 5$, $v \leq v_V$, $|\Delta^{(k+1)}T| > 5 \cdot 2^{(k+1)} \cdot \sigma$.

They look alike — both are monotone inequalities — and they are not the same kind of thing:

A capacity bound is about the world. Violating it means the configuration does not exist. A resolution bound is about the measurement. Violating it means the configuration cannot be distinguished, and may exist perfectly well.

E(X) counts violations of the first. The refusal rate counts violations of the second. §11.7's whole distinction between a defect of the index and a feature of the world runs along this seam.

16.14.1 And celestial mechanics is what made it visible

Gravity has resolution floors everywhere and a capacity bound in only one place. A planet holds satellites within its Hill sphere, $r_H = a(m/3M)^{1/3}$ — a monotone cap of exactly §10.4's form, and $\log r_H$ against $\log$ satellite count runs at +0.86 across six planets.

The Kirkwood order cut is the other kind. Resonance strength enters at $e^{|p-q|}$; with a belt eccentricity of 0.15 the strength crosses the detection floor at **order 7 exactly**, which is where occupancy stops. **That is not a capacity and there are no magic numbers — only a resolution boundary, the same form as v_V .**

Λ and the nucleus have capacity bounds and therefore magic numbers. The belt has none, and therefore has none. Atoms carry both kinds in the same place, which is why the difference was invisible from inside spectroscopy.

16.15 The second admission bound

$$v_V = (3Z^2R / 5q)^{1/4}$$

for quotation granularity q . The bracket needs levels **separated** (v^{-3} , governed by r); the cost needs curvature **resolved** (v^{-4} , governed by v_V).

Curvature washes out before separation, in every channel in this work.

channel	q	v_V	v reached
K I *n*d	0.0001	160	45.7
Na I *n*s	0.001	90	20.0
Al I nf	0.01	50.7	55.0

Al I nf is the first channel to reach its own ceiling. Predicted 50.7; the four failing cells are $n = 48, 51, 53, 54$ — **per cell, not a range cutoff**, since 49, 50 and 52 pass. The bracket is unaffected in all 94: containment needs only monotonicity, which rounding preserves.

And in Li I, v_V was applied prospectively — eight of fifty-five cells excluded from the cost test *before it ran*, because the ASD's quotation coarsens above $n = 33$. That is the difference between a bound derived and a bound noticed.

17. The collection

1,442 interior cells across 35 atomic systems. The bracket holds in every one.

17.1 What it spans

axis	span
l	0 through 6
Z	1 through 83
Z_eff	1 through 3
core types	23, all characterised

axis	span
isoelectronic pairs	5, each at two charge states

Cores tested: bare · $1s^2S$ · $1s^2\ ^1S_0$ · $1s^22s\ ^2S$ · $2s^2\ ^1S_0$ · $2p^5\ ^2P^\circ$ · $2p^6\ ^1S_0$ · $3s^2S$ · $3s^2\ ^1S_0$ · $3s^23p\ ^2P^\circ$ · $3p^4$ · $3p^5\ ^2P^\circ$ · $3p^6\ ^1S_0$ · $3p^64s\ ^2S$ · $3d^{10}\ ^1S_0$ · $3d^9\ ^2D$ (declined) · $3d^{10}4s^2\ ^1S_0$ · $4p^5\ ^2P^\circ$ · $4p^65s\ ^2S$ · $4d^{10}\ ^1S_0$ · $5p^6\ ^1S_0$ · $5p^66s\ ^2S$ · $5d^{10}\ ^1S_0$ · $3d4s\ ^3D$ · $6p^2\ ^3P$.

Figure 15.1 The defect falls monotonically with ℓ , over six orders of magnitude




Figure 17.1. Every species in the collection, defect against angular momentum. The fall is monotone in all of them, and spans six orders of magnitude — from 4.9 in bismuth to 1.6×10^{-6} in He II.

17.2 The largest contributors

species	core	channels	cells
He I	$1s^2S$	9	189
Ne I	$2p^5\ ^2P^\circ$, two limits	16	131
KI	$3p^6\ ^1S_0$	4	105
Al I	$3s^2\ ^1S_0$	4	94
Na I	$2p^6\ ^1S_0$	6	80
Ga I	$3d^{10}4s^2\ ^1S_0$	5	61
Li I	$1s^2\ ^1S_0$	3	55




Figure 17.2. What the collection covers. Marker area is proportional to cells. The gaps are not omissions: the census of §17.12 excludes them structurally.

17.3 He I — the angular span

Nine channels, $\ell = 0$ through 6, singlets and triplets, n to 35. And the defect falls monotonically with ℓ :

$0.2965 (s^3) \cdot 0.1392 (s^1) \cdot -0.0133 (p) \cdot +0.0014 (d) \cdot -0.0010 (f) \cdot -0.0014 (g) \cdot -0.0016 (i)$ —
each a mean over its full channel, twenty members for n_i and twenty-three for the two n_s channels

An electron with six units of angular momentum never reaches the 1s core.

The ordering itself is not new, and this book does not claim it. Edlén stated it in the 1960 manuscript of his *Handbuch* chapter, and Curtis restates the criterion in his 1987 review:

"classical arguments imply that, for a given Rydberg series, penetration contributions should decrease and polarisation contributions should increase with increasing n ... this criterion classifies the S and P levels as penetrating and the D and F levels as nonpenetrating."

What is new is the bound the ordering licenses. Edlén used it to classify levels; §25.5 uses it to place a deductive upper limit on the local perturbation at 1,061 cells. The observation is sixty-five years old. The inference from it is not.

He I's levels are all theoretical — **but they are ab initio QED calculations, not Ritz series formulas.** That distinction matters and §17.7 develops it.

17.4 He II — the $\delta = 0$ edge itself

A hydrogenic ion at $Z = 2$. Its defects run 7.6×10^{-5} at $\ell = 0$ down to 1.6×10^{-6} at $\ell = 6$, monotone across six orders of magnitude, and never reach zero.

There is no core to penetrate. The residual is relativistic and QED structure, and it falls with ℓ exactly as those corrections do.

And the ionisation limit came out of the data. Minimising the δ spread over the $*n*g$ channel — six members, scanned at 0.001 cm^{-1} — gives **438,908.871 cm^{-1}** against the tabulated 438,908.885 — **agreement to 0.014 cm^{-1} , three parts in 10^8** , from a criterion that never used the tabulated value.

17.5 Bi I — the charge span

Atomic number 83, more than double anything previously tested. The $6p^3$ ground term is strongly j-j coupled and the levels carry jK labels throughout. $\delta = 4.90$ for the ns channel, over five members — the largest defect in this work, as expected when a 6s-like electron penetrates a filled $5d^{10}6s^2$ core.

And Bi I settles a rule. Its core carries five J levels — the same dilution that made neon II fail entirely — but its 3P_0 parent reaches $n = 11$.

Dilution divides. It does not forbid. cells per channel $\approx$ (measured shells / D) – 2

Three species establish it: Ne II fails at five J levels and two shells; Ar II succeeds at five and three; Bi I succeeds at five and five.

17.6 The isoelectronic pairs

sequence	Z_eff = 1	Z_eff = 2
helium-like	He I	Li II
lithium-like	Li I	Be II
sodium-like	Na I	Mg II
aluminium-like	Al I	Si II
potassium-like	KI	Ca II

Same electrons, different nuclear charge, the method run identically. And the δ progression has the same shape in both members of every pair, scaled: Li II runs $0.182 \rightarrow 0.054 \rightarrow 0.002 \rightarrow 0.0002$ across $s \rightarrow f$; Be II runs $0.262 \rightarrow 0.049 \rightarrow 0.002 \rightarrow 0.0001$. Each figure is a channel mean over five to ten members; Appendix C carries them individually.

Li II and Be II both give $l = 5$ defects near -0.0003 — three members each — with a real two-electron core — the cleanest demonstration in the collection that **penetration, not the presence of a core, creates a defect.**

17.7 Hydrogen, and a floor on interpretation

One row of the collection measures the data rather than the method, and it sets a limit below which no result in Part III can be read.

H I gives 5 cells with a **median error of 0.006 cm^{-1}** — two orders of magnitude better than anything else here. That figure is not a triumph of the bracket. **It is an identity.**

The hydrogen row uses the reduced-mass Rydberg $R_H = 109,678.7717 \text{ cm}^{-1}$, which for hydrogen is the ionisation limit, so the quantum defect vanishes identically — measured defects span -2×10^{-5} to 0.

Substituting the infinite-mass $R_\infty = 109,737.3 \text{ cm}^{-1}$ breaks that identity by the 58.5 cm^{-1} reduced-mass difference, manufactures a spurious defect of $\approx 3 \times 10^{-4}$, and inflates the median error 54-fold to 0.34 cm^{-1} .

For every heavier species the correction is 10^{-5} relative and is absorbed into the fitted defect, so the choice is invisible. Hydrogen is the one row where it is not.

The 0.006 cm^{-1} is therefore a floor on interpretation. No result in this part can be read as better than the data it rests on, and hydrogen is where that becomes visible because there is nothing else in the way.

17.8 On independence, and the most conservative reading

1,442 cells are not 1,442 independent tests.

Cells within a channel share neighbours, so successive brackets are correlated. The honest figure is **1,442 cells across approximately 190 channels**.

The largest single blocks are the most strongly correlated. He I's nine channels contribute 189 cells; Ne I's sixteen contribute 131. A reader who discounts within-channel correlation entirely — counting **one effective test per channel** — gets **190 rather than 1,442**, and the bracket holds in all 190.

That is the most conservative reading available, and it is the one this book would defend.

The claim is not that 1,442 independent things were verified. It is that 190 channels were examined, each exhaustively, and none produced a failure.

17.9 A caveat the collection forced

The ASD marks series-formula values in brackets. Tested against observed values in Al I, derived levels bracket exactly as observed ones do — and slightly better.

That is not a validation. Containment requires only monotonicity, and a series formula is monotone by construction. **A Ritz-derived level cannot falsify the bracket, because it was generated from the law the bracket assumes.**

Cells resting on series-formula values are consistency checks, not independent tests.

The caveat cannot be closed, and this states exactly what closing it needs. Three kinds of value sit under the 1,442: measured levels, Ritz series-formula values that the ASD marks in brackets, and ab initio QED calculations. The first and third are independent tests of the bracket; the second are not, for the reason above. **The book has never counted the three.** Nothing in the method prevents the count — the ASD's own bracket marks carry the classification, and partitioning the 1,442 by that mark is a lookup, not a computation — but the classification was never carried into the collection, so the number of genuinely independent tests among the 1,442 is unknown and is not zero and is not 1,442.

What can be said without the count: **the claim survives its worst case.** He I's bracketed values are ab initio and are the sharpest tests in the collection, so the bracketed set is not uniformly dependent; and §17.2's decline modes and §18's exclusions rest on measured levels, which the caveat does not touch. What cannot be said is *the bracket held on 1,442 independent tests*, and the book has occasionally come close to saying it. The honest form is **1,442 cells, an unknown proper subset of which are independent**, and it should be written that way wherever the figure carries weight. Recorded in Q as item P — obstacle retrievable, cost hours, blocks one stated claim. Register 246.

But bracketed values are not one category. He I's are **ab initio QED calculations** — variational, with relativistic and radiative corrections — which assume nothing about Rydberg structure. Those are genuine tests, and at uncertainties of 10^{-9} cm^{-1} they are sharper than any measurement here.

17.10 Two cross-checks

Al I and Al II were held from two independent sources — the NIST compilations and the ASD.

Al II: 36 of 36 levels identical to machine precision. Al I: one disagreement in 40 — $3s^211d$, ASD 47379.7 against Kaufman & Martin 47379.140, $\Delta = 0.560 \text{ cm}^{-1}$ — one disagreement in forty overlapping levels.

The K&M value is the more precise measurement and is the one used. **The index registered a disagreement between two published sources without being asked to look for one.**

17.11 Four ways a species declines

A collection reporting only successes cannot be checked.

species	cause	cells
---------	-------	-------

Ne II	dilution — 5 core J levels, only 2 shells measured	0
Ar II	same dilution, beaten by depth	10
Si I, P II	coupling scheme changes LS → jK mid-series	4, 1
Cu II	core-excited configurations interleave the series	0
Bi II, Bi III	mixing to 34% leading percentage; a third limit inside range	0

Cu II's is a structural exclusion. Its $3d^8 4s 4p$ levels sit at 107,942–111,124 cm^{-1} , straddling the $3d^9 5s$ Rydberg levels at 108,014–110,366. The series is buried in core-excited states.

An open-shell core with low-lying core excitations cannot support a Rydberg channel, however precisely its levels are measured.

And that predicts the same outcome for every open- d and open- f sequence **without fetching any of them** — which the census of §17.12 confirms independently.

17.12 The census, computed rather than searched

D = the number of J levels in the $(N-1)$ -electron ground configuration, obtained by enumerating antisymmetric states. **No spectroscopic data enters.**

27 sequences admitted ($D \leq 2$) · 12 needing exceptional depth · 5 excluded

The five are **Cr-, Mn-, Fe-, Co- and Ni-like**, with D running from 19 to 37. **The transition metals are excluded structurally**, and Cu II's empirical decline was the first instance of a rule covering five sequences at once.

Figure 15.3 Interpolation works; extrapolation to the neutral does not




Figure 17.3. Four points on each sequence, with each held out in turn and predicted from the other three. Crosses are predictions; circles are measurements. The interior is recovered to about one percent. The neutral is not.

17.13 What the structure determines, and what it does not

Given five isoelectronic pairs and two four-point sequences, how much of an unmeasured species can be derived?

Fitting $\delta(Z) = \delta_{\infty} + a/Z + b/Z^2$ and holding one member out:

held out	lithium-like	sodium-like
neutral, $Z = 1$	11.1%	23.4%
$Z=2$	1.4%	2.4%
$Z=3$	0.8%	1.3%
$Z=4$	1.8%	2.5%

Interpolation works to about one percent. Extrapolation to the neutral does not.

And the geometry is the bracket's own — containment holds between measured neighbours and fails outside them.

across sequences	δ not determined; the $Z=1 \rightarrow 2$ ratio spans 0.53-0.83
within a sequence, interior	δ determined to $\sim 1\%$
within a sequence, at an end	δ not determined

Interior members need not be fetched. The ends always must — and the neutral is always an end.

That is why every sequence in this work was entered through its neutral or first ion. It was not convenience; it was forced.

18. Domain, failure, and the bound in the silence

18.1 A method that has never failed has not been tested

1,442 cells, no bracket failure. Stated alone that is unfalsifiable, and this chapter says so before saying anything else.

18.2 Five exclusions, which are three mechanisms

The method declines in five circumstances. They are not five reasons.

(i) Sign-changing quantities Any property built from a matrix element that passes through zero cannot be monotone. $\text{Sr I } 5s^2S_0 \rightarrow 5snp \ ^1P_1$ reduced transition probabilities collapse by a factor of 60 at $n = 6$ and recover — a Cooper-type node.

Bracket coverage across the node: **79%**. Away from it: **100% on 11 cells**.

This is the only exclusion that attacks the order hypothesis itself.

(ii) Quantities monotone in neither direction Fine-structure splittings reduced by v^3 drift by 3 to 22% across a series while reversing direction locally — Ca II nd increases at 2 of 7 steps, Ba II nd at 2 of 6, Ba II np at 3 of 6. Coverage over the 17 interior cells is 58.8% under Rule 3's direction-agnostic bracket.

An earlier version reported 35.3% and attributed the failure to flatness. Both were wrong. The 35.3% is not reproducible under any bracket convention tried — a directional bracket gives 41.2%, Rule 3's min/max form gives 58.8% — and **the mechanism is not flatness**. These sequences are not flat. **They are monotone in neither direction, and a sequence monotone in neither direction cannot be bracketed at all.** That is a simpler and stronger exclusion than the one it replaced.

(iii) Systems without interior cells — and it has two forms Ti I: 23 channels, 26 members, zero interior cells. No channel reaches three members, because the ionisation limit near $55,000 \text{ cm}^{-1}$ sits below the onset of severe level density, and above $\sim 42,000 \text{ cm}^{-1}$ the levels are 40–50% configuration-mixed with no assignable parent.

Sc VI shows this has two causes, not one. Sc^{5+} has a clean core — three parent terms, no mixing worth the name, a limit measured at $892,700 \pm 400 \text{ cm}^{-1}$ — and still yields **6 channels, 10 configurations, zero interior cells**, because every channel stops at $n = 5$.

Ti II's limit is structural: no measurement campaign would help. Sc VI's is provisional: 6s and beyond exist and are simply unmeasured.

One further measurement would move Sc VI out of the exclusion entirely — which is what makes it the deduction of §18.6

(iv) and (v) Resolution failures The remaining two are the same failure at different values of one ratio: levels too close to separate ($r < 5$, §18.6) and curvature too fine to resolve ($v > v_V$, §16.11).

The reduction Five exclusions, three mechanisms: a broken order (i, ii), an absent cell (iii), and a resolution limit (iv, v) — the last being one ratio evaluated at two thresholds.

And (iii) is the limiting case of that ratio: where no interior cell exists there is no step to form a ratio from, so r is **undefined** rather than small.

18.3 The failure condition

The bracket fails $\Leftrightarrow |\Delta T| > 2Z^2R/v^3$

A perturber **shifts** a level; it **reorders** one only if the shift exceeds half the local spacing.

Tested on the strongest mixing in the work. Hg II's $5d^{10}7p$ sits at 60% leading percentage with 20% admixture, and the $5d^96s6p$ perturbers lie 2,000–3,000 cm^{-1} away. **The half-spacing at $n = 7$ is 24,634 cm^{-1}** , computed from $2Z^2R/v^3$ at the measured $v = 5.3$. The level is displaced by a tenth of what reordering would require.

And the condition predicts where failure would occur:

perturbation	$v_{\text{fail}} (Z = 1)$	$v_{\text{fail}} (Z = 2)$
3,000 cm^{-1}	4.2	6.6
1,000	6.0	9.6
100	13.0	20.6
10	28.0	44.4

Hg II's np reaches $v = 5.3$. It is not deep enough to fail.

18.4 Why the collection never enters the failure regime

The two requirements are anticorrelated in nature. Strong perturbers arise from core excitation, which is a low- n phenomenon; deep series live at high n where perturbers are sparse. Large ΔT and large v rarely coexist.

That is why the bracket has not failed — stated as physics, not as luck.




Figure 18.1. Each of 1,105 verified cells, plotted at the perturbation bound its containment implies. The two curves are the failure condition at $Z = 1$ and $Z = 2$. Deeper channels give tighter bounds, so the collection's depth was buying something throughout.

18.5 The bound in the silence

Read the condition backwards:

the bracket held $\Rightarrow |\Delta T| < 2Z^2R/\nu^3$ at that cell

Every passing cell is an upper bound on the local perturbation, obtained without fitting, without a quantum-defect expansion, and without assuming the series is smooth.

Every verified cell yields a bound. Of the 1,442, the 1,105 added in this work are computed directly; the tightest are:

species channel $n \nu$ $|\Delta T| <$

Al I 3s²nf 54 54.0 1.398 cm^{-1}

Ga I 4s²np 54 51.7 1.585

Li I np 42 41.8 3.25

These are statements about the atom:

No configuration interacts with Al I 3s²54f by more than 1.4 cm^{-1} .

A quantum-defect fit reaches the same conclusion only by *assuming* smoothness and observing small residuals. **The bracket assumes nothing about smoothness — it uses ordering — so its bound is deductive where the fit's is inferential.**

Before: 1,442 cells, no failure. Unfalsifiable. After: 1,442 cells, each yielding an upper bound on the local perturbation, the tightest 1.398 cm^{-1} .

Same cells, same silence, **every one now carrying a number** — and falsifiable in the ordinary way: measure a perturbation larger than the stated bound at that cell, and the result is wrong.

18.6 The one deduction, and why it is not a prediction

Sc VI's $4S^\circ$ channel is one measurement from being usable, which makes it a prediction — and the book states it as a bracket with an estimate inside, because Rule 3 requires that and an earlier version violated it.

And the tower explains, retrospectively, why a two-coordinate bracket sufficed here. The channel is an s-series on an S-state core: $f = 0$ collapses K to a singleton, and the $4S^\circ$ core minimises the J_c structure — the case lives in the tower's degenerate corner, where axes 11-13 carry no information. That is stated here so a referee does not have to ask.

18.6.1 The bracket, from monotonicity alone

Two measured defects: $\delta(4s) = 1.0057$, $\delta(5s) = 0.9812$ — consistent to 0.025, a clean Rydberg channel. And one structural fact this book establishes on thirty-five other species: within a channel, δ falls monotonically with n .

The two-point Ritz solve is exactly determined:

$$\delta_2 = (\delta_4 - \delta_5)/(1/16 - 1/25) = 1.0889 \cdot \delta_\infty = 0.9376 \cdot \delta(6s) = 0.9679$$

And monotonicity bounds $\delta(6s)$ on both sides:

$$\delta_\infty = 0.9376 < \delta(6s) < \delta(5s) = 0.9812$$

which maps directly to an energy interval, E being decreasing in δ :

$$735,860 \text{ cm}^{-1} < E(\text{Sc VI } (4S^\circ)6s) < 738,547 \text{ cm}^{-1}, \text{ a deductive bracket } 2,687 \text{ cm}^{-1} \text{ wide, with the limit's own } \pm 400 \text{ to be added at both edges}$$

18.6.2 And convexity tightens it, without leaving it

Monotonicity is not the only order property available. $\delta = \delta_0 + \delta_2/n^2$ is also convex in $1/n^2$, and a convex sequence lies above its own tangent:

$$\delta(6s) \geq 2\delta(5s) - \delta(4s) = 0.9567$$

which raises the lower bound on δ and therefore the upper bound on E :

	bound	δ range	E range	width
	monotone only	0.9376–0.9812	735,860–738,547	2,687
	bound	δ range	E range	width
	+ convexity	0.9567–0.9812	735,860–737,380	1,520

A factor of 1.8 tighter, and the new interval lies wholly inside the old one. Nothing committed under B.2.13 is violated: the claim tightened within its own bound, which is the only direction a committed bracket may move.

18.6.3 What was already solved — not a prediction

The value below is not a prediction, and calling it one was a category error against the book's own law. A prediction is a proposal about an unlisted cell, and $E(\Lambda) = 0$: the index proposes none. What §18.6.1 and §18.6.2 produce is a deduction — two measured defects and two order properties determine the interval, so the problem arrives already solved by the data in hand. The bracket is stated as the solution it is: $E(\text{Sc VI } 3s^23p^3(4S^\circ)6s^3S^\circ_1) = 736,688 \text{ cm}^{-1}$ bracketed in $[735,860, 737,380] \text{ cm}^{-1}$ 13.5743 nm, in the window 13.5615–13.5895 nm 91.338 eV

And for 7s, with a caution the 6s case does not need:

$$E(7s) = 784,416 \text{ cm}^{-1}, \text{ bracketed in } [784,128, 785,209] \text{ cm}^{-1}, \text{ width } 1,081 \text{ — tighter than } 6s \text{ because the spacing has fallen as } \nu^{-3}.$$

The 7s convex tangent would give [784,128, 784,605], width 478 — but it requires $\delta(6s)$, which is *estimated* rather than measured.

The 6s convex bracket is deductive. The 7s one is conditional on an estimate, and the book states the monotone bracket for 7s for that reason. A bracket that inherits an estimate is not a bracket.

18.6.4 The point estimates, and the one that was wrong

route	δ used	E(6s)	inside the bracket
two-point Ritz	0.9679	736,688	yes
Rule 4 constant defect	0.9934	735,091	no

An earlier version of this work reported $735,091 \pm 1,398$. It was computed with $\delta = 0.9934$, which exceeds $\delta(5s) = 0.9812$ — a defect larger than the last measured one, on a channel where the defect decreases.

The estimate fell outside the bracket its own channel implies. Rule 4a now forbids the form that produced it, and §17.2 records why.

18.6.5 Nothing was fetched

The bracket uses two measured levels and one structural fact. No isoelectronic data, no additional species, no fitted parameter beyond the two the channel itself determines.

Every verified cell in this book lies at Z_{eff} between 1 and 3. A measurement of Sc VI ($4S^\circ$)6s would place the method at $Z_{\text{eff}} = 6$ — double the highest charge state tested.

If the measured level falls outside [735,860, 737,380] cm^{-1} , the bracket is wrong, and it is wrong deductively rather than statistically — either the monotonicity of δ or its convexity would have failed, and the two edges say which.

The lower edge tests monotonicity; the upper edge tests convexity. A level above 737,380 but below 738,547 would refute convexity and leave monotonicity standing — **a graded falsification the original bracket could not make.**

That is a stronger falsification than a σ . A point estimate can miss by chance. **A bracket cannot.**




Figure 18.2. The Sc VI prediction as it now stands. §18.6.1's bracket uses monotonicity alone — δ decreasing gives $\delta_\infty < \delta(6s) < \delta(5s)$, a width of 2,687 cm^{-1} . §18.6.2 adds convexity: $\delta = \delta_0 + \delta_2/n^2$ lies above its own tangent, so $\delta(6s) \geq 2\delta(5s) - \delta(4s) = 0.9567$, narrowing the interval to 1,520 cm^{-1} without leaving it. The two-point Ritz estimate sits inside both. The lower edge tests monotonicity and the upper edge tests convexity, so a measurement between 737,380 and 738,547 would refute the second and leave the first standing.

18.6.6 A second route to the same prediction

The bracket above uses Sc VI's own 4s and 5s. A second route uses neither — and it is now a check rather than a necessity.

Sc VI is **sulphur-like** — sixteen electrons, ground $3s^23p^4$ — and its sequence runs S I, Cl II, Ar III, K IV, Ca V, Sc VI. **Sc VI sits at the tight end**, which by §17.2 is the cheap direction: 3.5% for a fifth member rather than 17% toward the neutral.

What that route needs is two numbers for any two or three of the lower members — an ionisation energy and the $(^4S^o)4s\ ^5S^o_2$ level, both tabulated for all five.

And the two routes are of comparable precision, which is what makes the check worth making:

route	uses	spread
own 4s, 5s under Rule 4	Sc VI's own levels	$\pm 1,594\text{ cm}^{-1}$
isoelectronic, at 3.5% in δ	five other species	$\pm 2,169\text{ cm}^{-1}$

Neither dominates. Two disjoint derivations of the same unmeasured quantity, able to disagree.

This is $\mathbb{D}_{\text{def}}$ applied to a prediction rather than to a past measurement, and it is the only place in this book where the mechanism is deployed prospectively. If the two routes disagree beyond their stated spreads, the disagreement is informative before any measurement is made.

But the deductive bracket of §17.6 does not depend on it. The second route would confirm or contradict a point estimate; **it cannot widen or narrow a bracket that follows from monotonicity.** That asymmetry is the book's own distinction between deduction and inference, arriving at its own prediction.

19. Collective sections

19.1 The assembly rule


An observable built as $\langle r \rangle^a / (\Delta E)^b$ scales as

$$v^{\{2a + 3b\}}$$

Derived, not fitted, from $\langle r \rangle \propto v^2$ and $\Delta^*E^* \propto v^{-3}$.

observable	a, b	p
$\langle r \rangle$	1, 0	2
cross-section, C_3 , quadrupole, diamagnetic shift	2, 0	4
polarisability α	2, 1	7
C_6	4, 1	11
C_8	6, 1	15

Scope, stated. The rule covers observables assembled from $\langle r \rangle$ and Δ^*E^* . **Hyperfine coupling is excluded** — it scales as v^{-3} through $|\psi(0)|^2$, which is neither a matrix element nor a denominator.

Figure 17.1 One dimension, fifteen perspectives: $\text{rank}(\log q) = 1$ Figure 19.1. Fifteen Rydberg observables against ν , and the singular values of their centred log-matrix. The second is 1.8×10^{-14} — numerically zero. Every one of them is the same dimension seen at a different exponent.

19.2 One dimension, fifteen perspectives

Across fifteen Rydberg quantities and forty values of ν , the centred log-matrix — fifteen quantities by forty values of ν — has singular values 9.49×10^1 , then 1.8×10^{-14} .

$$\text{rank}(\log q) = 1$$

One dimension. Fifteen perspectives. Every Rydberg property is ν to a power, and the exponent is the whole of its identity.

19.3 Collective coefficients are not a separate class

$C_6 = \nu^{11}$ is $(\nu^2)^4/\nu^{-3}$ — four single-atom dipole factors over one single-atom denominator. The two-atom quantity is assembled from the one-atom ν and nothing else.

19.4 Error amplifies as the exponent — which is the point

A fractional error ϵ in ν gives $k\epsilon$ in an observable of exponent k . Inverted:

measuring C_6 to 1% fixes ν to about 0.0the searched bound of §22.8 — from $1/p_{\text{eff}}$ at $n = 100$, evaluated across nine points, and roughly nine times sharper than the three-level bracket achieves

The collective quantity is the most sensitive probe of dimension 1.

19.5 Tested against a published parameterisation

Singer et al. (2005) give $C_6 = n^{11}(c_0 + c_1n + c_2n^2)$ for Rb ns-ns. The quadratic factor makes the local exponent move:

$$p_{\text{eff}} = 11 + n(c_1 + 2c_2n)/(c_0 + c_1n + c_2n^2)$$

running from **12.58 at $n = 35$ to 11.44 at $n = 100$** , evaluated at nine points across that range.

The cost law tracks it: median error 1.05%, maximum 1.83%, across nine points. **Bracket 66 of 66.**

19.6 Aitken's Δ^2 has an exact bias on a Rydberg series

Aitken's process extrapolates a sequence to its limit. A Rydberg series converges to its ionisation limit. So apply it.

n	$T(n)$	Aitken(T)	$T/3$
10	1097.3730	365.1793	365.7910
20	274.3433	91.4096	91.4477
40	68.5858	22.8596	22.8619
80	17.1465	5.7153	5.7155

Aitken does not reach the limit. It reaches $T/3$.

Because the correction it subtracts is $(\Delta T)^2/\Delta^2 T \rightarrow T^2/T'' = (2/3)T$ — exactly λ^2 of \$16.11 — leaving $(1/3)T$. Aitken's kernel is *geometric* convergence, $x_n = L + a\mu^n$; a Rydberg series converges *algebraically*, $T = R/\nu^2$. **The failure is exact, and its number is one third.**

With a practical consequence. Aitken's Δ^2 is a standard tool for estimating a limit from measured terms. **On a Rydberg series it lands at $1 - T/3$, deterministically** — a 91 cm^{-1} error at $n = 20$ and 23 cm^{-1} at $n = 40$.

Not a better estimate. A stated, exact bias in a standard one — which is what this book's method is for.

And λ^2 now carries three readings across three fields: the decrease obtained by a full Newton step · the correction Aitken subtracts · **two-thirds of the binding energy**. Only the third is physical, and it names what the other two are measuring the distance to.

19.7 And the bracket is the wrong tool here

C_6 comes from a closed-form fit at every n . Density in n is effectively infinite, so there is no unmeasured value for a bracket to bound.

The bracket is worth something only where the quantity is sparse in n . C_6 is dense. Collective sections are the wrong place to deploy it.

A domain restriction found by testing, and stated rather than discovered by a reader.

20. Slack

This book has computed one quantity three times, in three languages, and called it by three names.

20.1 The three names

in the definition
book
as

orde	$E(X)$	$\setminus$	$\mathcal{R} -$	$X \setminus$	— what the bounds admit, minus what
r			$(\setminus$		the index carries
			X		
			$)$		

geo	the	$\Pi(\backslash$	Δ	+	$[\chi\lambda y, \chi\nu y]$ n	— what the box admits, minus what
metr	void		i	1)	$\Lambda $	the lattice carries
y				— $\backslash$		
calcu		V		w/e — what the bracket admits, over what the		
lus				estimate carries		

All three are admitted against actual. The only difference is whether the comparison is taken as a difference or a ratio — and §4.2 already supplies the bridge, because **log of the occupancy measure is an ordinary l^1 metric**. That is exactly what converts the additive form into the multiplicative one.

20.2 The one quantity

$SLACK(S) = \text{measure}(\text{ambient}) / \text{measure}(S)$ $\log SLACK = \log \text{measure}(\text{ambient}) - \log \text{measure}(S)$

Computed on the objects of this book:

quantity	ambient	actual	slack	log slack
Λ in its bounding box	6,912	976	7.082	1.958
the periodic table in $\mathcal{R}(X)$	126	90	1.400	0.337
the calendar in $\mathcal{R}(X)$	372	365	1.019	0.019
bracket against estimate, $\nu = 10$	w	e	13.378	2.594
bracket against estimate, $\nu = 40$	w	e	53.344	3.977

E(X) is the slack counted. The void is the slack counted over an interval. V is the slack measured. They are one quantity under three measures.

20.3 Slack is a description length

The name was chosen before the quantity was analysed. Analysed, it is a known object.

For a counting measure, $\log SLACK = \log|\mathcal{R}(X)| - \log|X|$ — and that is the information needed to say which subset X is, given its closure:

$E_bits(X) = \log_2 C(|\mathcal{R}(X)|, E(X))$

index	E	bits	bits/cell
periodic table, 18-column	36	105.1	1.168
calendar	7	47.4	0.130
subnet with a hole	8	48.5	0.196
Janet left-step	0	0.0	0
a box ordering, Λ	0	0.0	0

$E(X) = 0$ does not mean "nothing is missing". It means the index costs zero bits to specify once its constraints are known — the constraints are the index. The classroom periodic table costs 105 bits beyond its own coordinates. Someone must transmit them; in practice a teacher does.

And V becomes a bit-count too. $\log_2 V^*$ is the precision surrendered for certainty: 3.74 bits at $v = 10$, **5.74 at $v = 40$** . Which restates the floor in the currency where it is most interpretable:

Prop. 14.1: $V > 2 \iff \log_2 V > 1$. A guarantee never costs less than one bit, and a Rydberg series charges 1.54.




Figure 20.1. Left: what each index costs to specify beyond its own coordinates. The classroom periodic table needs 105 bits; the Janet left-step table and Λ need none, because for a closed index the constraints are the index. Right: $\log_2 V$, the precision surrendered for certainty. Proposition 15.2's floor $V > 2$ becomes the statement that a guarantee never costs less than one bit, and a Rydberg series charges 1.54.

20.4 And every structural result is a statement about slack

result	condition	in slack
closure	$\mathcal{R}(X) = X$	slack = 1
the void	box $n \wedge = \text{box}$	slack = 1 on that interval
$E(X) = 36$	the periodic table	slack = 1.4 in the index
the floor, Prop. 14.1	$V > 2$ always	slack in an inference can never reach 1
the pole, \$25.5	$y' = 0$	slack = ∞ where the function does not turn
prediction, \$18.6	$E(X) = 0 \implies$ nothing to forecast	forecasts live in the slack

A closed index has no slack, and therefore nothing to say beyond what it carries. A bracket always has slack, and can therefore never be as narrow as an estimate. And a linear observable has unbounded slack, because it never turns.

20.5 Why the three languages agree

They are not three descriptions of one object. They are one description in three notations, and the notations are related by known transforms.

Order to geometry. A lattice interval $[x\lambda y, x\nu y]$ is a coordinate box; the void is what the box holds and the lattice does not. **The same subtraction, with the ambient named differently.**

Geometry to calculus. The occupancy measure counts cells; $\prod(|\Delta_i|+1)$ is a product, and its logarithm is a sum — the ℓ^1 metric of §4.2. **A ratio of widths becomes a difference of logs, which is a count.**

Calculus to order. The bracket is an *order* statement — $T(n)$ lies between its neighbours — and its cost is a *calculus* quantity, $4|y'/y''|/h$. **Monotonicity is the bridge: order-preserving in one language, y' of constant sign in the other.**

20.6 What this buys, and what it does not

It buys a single vocabulary. Chapters 1, 5, 11 and 14 stop being four subjects and become four measurements of one thing, and the reader who understands slack in one language has it in all three.

It does not buy a new result. Every number in §17.2 was computed elsewhere in this book. **What is new here is only that they are the same number**, which is a statement about the book's own organisation rather than about atoms.

The three languages were never in competition. Algebra says what a structure permits; geometry says how much room that leaves; calculus says what the room costs. Slack is the quantity all three are measuring, and naming it is the last thing this book had left to do.

PART IV — THE RECORD AND THE REACH

21. Withdrawals

two hundred and ninety-seven claims made in the course of this work were wrong. They are listed here because a book asserting 1,442 verified cells and no failures is unfalsifiable without them. Twenty-six were made after the register was first closed — six in the literature search of Chapter 22 and eight in the structural work of §17.8 and §18.6 They are listed at §24.2.

21.1 The dominant pattern

Prose asserting a conclusion while the output contradicting it sits on the same screen.

This accounts for the majority. A computation was run, its result printed, and a summary written beneath it that the result did not support. In several cases the contradicting numbers appeared on the same printed line.

That pattern is the reason the mechanisms of Chapter 11 matter. An index that only holds what its method consumes cannot catch it. One that holds redundant quantities can, and did.

21.2 The five that earned their place

The wrong nuclear charge. Al II and K II computed at $Z_{\text{eff}} = 1$ when both have doubly charged cores. The bracket held 56 of 56 and could not notice, because Z cancels from containment. δ came out at 4.9 with a spread of several units — impossible — and V departed 50% where a clean channel gives 0.2%. Two derived quantities, neither needed by the method, both correct.

The wrong limit. A neon-like sodium spectrum read as neutral neon. $T = I - E$ went negative and **every channel was refused at entry**. Totality did not detect the error; it prevented the computation.

" $w = 3v \cdot e$." Relation (b) of Chapter 11 written down incorrectly, with 540.0 and 1215.0 printed side by side. **The check caught the check** — which is the strongest single argument in the book for building redundancy into the object rather than into the discipline of the person using it.

The silent discard. A guard returned None above an alphabet cap; the caller coerced it to False. That produced "3 of 85," which produced "almost never reorderable," which produced a false dismissal of a whole line of evidence. One discard, three conclusions, none of which looked wrong — because the

arithmetic stayed consistent. $\mathbb{D} \geq 1$ cannot catch this. Only totality can, and totality had not yet been stated.

The retrieval failure. Chapter 14 argues that an index should enumerate its targets and navigate before searching. Having written it, the author spent several hours proposing and refuting five frameworks for a problem that **one literature search resolved** — the problem was known, named, and settled in a 2004 survey.

21.3 Others worth listing

- Claimed a species integrated that had never been written to the source; found by auditing the built document rather than the draft.
- Fitted a curve through non-monotone data and extrapolated it to zero.
- Stated the l -ordering law from a hand-typed subset rather than from the files; the real figure was 10 of 16, not 10 of 10.
- Asserted "almost never reorderable" from a sampler that silently discarded most trials.
- Ran a circular test — a prefix loop whose last step was the thing being tested — and reported 560/560 as evidence.
- Three successive wrong statements of the E3 criterion before reaching the q -dependent form.
- Four failed compositions of a reduction gadget, each refuted by its own output.
- Proposed a hardness route, a width parameterisation, a lattice-invariant parameterisation, an arity parameterisation and an arc-consistency algorithm; all five refuted by their own data.

21.4 Six from the earlier work, and how each was caught

The mechanism matters more than the error, so each is recorded with what found it.

Ba I 6s11f mis-keyed as 6s8f. Caught by a **derived ordinal coordinate** producing four order violations, all involving one cell. As $n = 8$ the datum implies $\delta = -2.84$, which is impossible; as $n = 11$ it implies $\delta = +0.16$, which is correct for an f orbital. Corrected, the coordinate is **2,163/2,163 monotone**.

A quantity the method never used located a typing error in a level table.

Ca I 3d² 1D₂ position. Caught by the **d² term-ordering constraint** ($^1D < ^3P$ iff $C/B < 5$), which three fitted analyses had not applied.

A regex parsing failure. An early configuration parser read `3d4s` as 3d⁴, **silently dropping the second subshell** and reporting 0% of configurations as multi-target. Caught by manual inspection of a case known to be two-target.

This is D3 again — a silent drop, not a wrong value. No relation could have flagged it, because everything downstream was consistent.

Bracket direction for decreasing sequences. For a decreasing property the tightest bracket takes the maximum over cells *above* in n and the minimum over cells *below*; the reverse returns the global extremes. **The error inflated a reported width from 11% to 114%.**

Pooled σ . Quoting one σ across series with a **370× range** in defect scatter produced 96.2% coverage at nominal $\pm 1\sigma$, which was misread as heavy tails. Per-cell σ gives **68.1%** — the expected figure. **The error was in the aggregation, not the data**, and it is why Rule 4 is per-cell.

Anchoring on a published estimate. A point prediction for Ba I 5d² 1G₄ was developed with a literature value visible from the first computation. Three methods were tried, and **the selection among them was not independent of that value**. The finally reported range was **38% narrower than the inputs supported**.

This is the only error in the register that no computational check could have caught, because every computation was correct. What failed was the independence of a choice.

It produced two standing policies: §18.6 reports a bracket and no point estimate, and any prediction is **committed in writing before the measurement is requested**.

21.5 What the register is for

Two things.

It makes the rest checkable. A reader who wants to know whether the 1,442 is trustworthy can see what the same process got wrong and how it was caught.

And it is the evidence for Chapter 11. The three defence mechanisms are not proposed on theoretical grounds. **Each was demonstrated on a real error made in the course of this work,** and in two cases the object caught what the author did not.

21.6 What the register shows about where errors are caught

Sorting the register by what caught each error gives a distribution, and the distribution is the argument of Chapter 11 in empirical form.

caught by	co	character
	un	
	t	
a derived quantity (δ , V , an ordinal coordinate)	4	the method was blind; a redundant quantity was not
totality — a refusal at entry	3	the computation never ran
two routes disagreeing	2	one an identity, one an apparatus
an audit against the built artefact	3	claims checked against output, not against draft
a collaborator's refusal of a conclusion	4	correct reasoning from corrupted or incomplete input
the author, on re-reading	~3	the dominant pattern of §25.1
a reader	0	—
a process constraint, adopted after the fact	1	the anchoring failure of §18.6

Zero were caught by a reader, because there has not been one. That is not a claim about the work's correctness; it is a statement about what has and has not been tested, and it is why §25.5 states falsification conditions.

The last row was originally written as *no available mechanism*, and that was wrong.

21.7 Six made after the register was closed, two printed here

These arose in the literature search of Chapter 22 and follow the same pattern as the first forty-eight.

49. The Sc VI prediction. Computed with a constant defect $\delta = 0.9934$ exceeding $\delta(5s) = 0.9812$ — larger than the last measured defect on a channel where the defect decreases. **The estimate fell outside the bracket its own channel implies.** Rule 4a now forbids the form; §18.6 states the bracket.

50. A conservation law that is not one. $w \cdot V = 8y'^2/y''$ claimed independent of h . Exact evaluation gives $w^2/e = 16x^4/(h^6 - 5h^4x^2 + 7h^2x^4 - 3x^6)$, which depends on h explicitly. The table showed 366.3, 367.9, 370.6, 374.5 across $h = 1 \dots 4$ — a 2.2% rise, read as a constant. What survives is a Pareto frontier, §25.5

51. "Checkable" for "checked." §25.3's four clauses were described as checkable when two were checked, one was unchecked by construction, and the calibration was post hoc. The word did the concealing.

52. The framing claimed as new. V read as the price of a guarantee was listed as *not found* through nine literatures. **Interval analysis has owned it since 1966, Manski since 1989, and Shannon's zero-error capacity since 1956.** Withdrawn in §18.6

53. A unit artefact. The 6-bracket was reported as 10^6 times narrower than the 7-bracket. **6 and 7 are in different units**; the comparison was meaningless as stated.

54. And its correction was also wrong. The repair claimed "2-3× wider in comparable units" **while the table beside it showed 8×10^3 to 5×10^5** . The correct statement is §18.6: the ratio is $n^3/(2\delta_2)$, and the choice is conditional on how well the limit is known.

Two of these six are corrections of corrections. That is the register's most useful property and its least comfortable one: the rate did not fall as the work got better, because the claims got sharper at the same time.

21.7.1 Eight more from the structural work, two printed here

55. "Bounds become worthless as data accumulates." Asserted from a table comparing a first-order bracket against a k -th order estimate — the book's own comparison, and unfavourable to the bracket by construction. A matched-order bracket holds V at 2 across $k = 1 \dots 5$ and $v = 20 \dots 80$.

56. $V = 4v/3$ as "the price of certainty." It is the price of *first-order* certainty. **What $4v/3$ buys is robustness**: order 1 tolerates a perturbation of half the local spacing, order 5 tolerates five orders less. §18.6 states the correction.

57. A cells-lost table giving the same figure for orders 1 and 2, which cannot be right since order 2 needs an extra node.

58. A negative bracket width at order 5. Reported as a median without remark. **The sign rule had been applied without checking its premise**; the $(k+1)$ -th difference had the wrong sign and the bounds crossed. Checking the premise first gives **100% containment at every order**.

59. Purity as a perturbation measure. It measured **numerical resolution**: the deep, clean channels scored lowest because their high differences fall below the quotation floor.

60. σ from decimal places. Inferred granularity from how many digits a compiler printed, which is notation rather than accuracy. Appendix C carries the stated uncertainties and they were not used.

61. "Six relations reduce to eight generators, so -2 were redundant." A Gröbner basis is routinely *larger* than its input. The redundancy measure is the reduction test, not the count.

62. "The survivors account for 8 of the 2." Eight cannot account for two. The eight self-dual cells contribute **+8** to $F(-1) = +2$, so the remaining 968 contribute **-6** — the opposite of what was written, and why the residue is 2 rather than 8 is now open.

Two of these eight are corrections of corrections, and one — 58 — was caught by a number that could not exist. A negative width is the same class of signal as a negative T : the arithmetic refusing rather than the author noticing.

21.7.2 Twelve more from the structural and audit work, six printed here

59. The cheapest undetectable forgery in Λ is 60 cells, not one. What $\mathbb{D}_{\text{phys}}$ actually buys is protection against systematic error, where the whole induced closure is wrong together and the amplification signature vanishes. Every one of §11.3's five recorded failures is of that kind.

63-64. Two audit-harness errors. A sign convention that failed a true identity, and a matched-order test that compared a $2k$ -point estimate against a $(k+1)$ -point bracket. The harness, not the book.

65. The closed form of V . §16.4 printed $V = 4v^3/(h(h^2-3v^2))$, whose denominator is negative for $v > h$. Correct: $4v^3/(h(3v^2-h^2))$. Magnitude right everywhere, sign inverted.

66. And the verdict on 65. Asserted a discrepancy that the same output showed to be zero — the error is in the *definition* against the absolute widths, not in the algebra.

67. The Shanks exponent, asserted before proof. Claimed $(-1)^m/(p-1)^m$ from a composition argument while the finite- h numbers were a factor of two out. **Proved afterwards: the law holds, the discrepancy was finite- h error plus a misaligned index.**

68. "E(X) has no analytic route." It has one: $E(X) = F_{\mathcal{R}}(1, \dots, 1) - F_X(1, \dots, 1)$. $\mathbb{D}(E)$ rises from 1 to 2.

69. $V(\delta) = V(T)$ taken too far. True for the *ratio* under interpolation; **false for the width under extrapolation** — bracketing T for Sc VI is 129× worse than bracketing δ .

70-71. A conclusion written past its evidence. "The ideal collapses" under output showing it did not, and two table rows never computed.

72-73. Language-independence, denied then restored. Deletion-repair was recorded as failing in geometry. **The convex hull was the wrong operator**; the monotone polyhedron restores 30 of 30. **A mismatched operator is a mistranslation, not a counterexample.**

74. "Self-consistency never defends against invention." It does. 220 of 220 single fabrications amplify — median 340 induced cells, minimum

21.7.3 Thirty-three more from the audit work of the final session, all printed here

75-79. Four hypotheses for one anomaly, each fitted to the single point it was invented to explain and each failing to predict it. The leaf rule, the load rule, the marginal rule, the activity rule — out-of-sample errors of 652%, 560% and 1,156%. Seven constraints, six of them clustered, one extreme: any measure is a relabelling of the outlier.

80-83. The closed form for amplification, wrong three times. Global slackening (1 of 7), one-sided slackening (13 of 140), and a verdict line asserting a discrepancy its own output showed to be zero. **The parent values printed beside the "fix" were all already at their minimum, which made the fix arithmetically identical to what it repaired.**

84-86. Three claims about closure, each contradicted by the table beside it. "Deletion-repair fails in geometry" — the convex hull was the wrong operator. "Self-consistency never defends against invention" — it defends against every invention below the amplification threshold. **"Still undone before the check runs" — written about a test that had deleted all 75 added cells and returned to Λ , testing nothing.**

87-88. The index was never deceived. Two passages described a *replaced* index as though it were the original being lied to. **Asked, Λ refuses 200 of 200. Rebuilt around them, it admits 200 of 200. Those are different operations and only one is a question.**

89. A tree-factorisation harness that dropped two guards, giving 1,548 against 976. **The book's formula is correct; the transcription of it was not.**

Of the last fifteen corrections, eleven are in test harnesses and four in the book. The book's stated equations now audit at 35 of 35. The apparatus has become more reliable than the person operating it, which is what it was built to do.

90-92. Three more, from the closing pass. An admissibility test that compared the bracket *width* to the noise floor when the rule is about the *difference* — the correct comparison gives k^* rising with better data, which the wrong one hid. **A verdict line reading "the load rule holds" from $r = 0.849$ while the held-out table beside it showed a 652% miss.** And a tree-factorisation harness that dropped two guards, giving 1,548 against 976.

The book's stated equations audit at 35 of 35. Of the last twenty corrections, fifteen are in test harnesses.

93-98. Six on the order axis, in one hour. §16.10 and §24.1 both said matched-order V oscillates about Proposition 15.2's floor. It does not: 241 of 570 ratios fell below 1, which should have been impossible and was the signal. Figure 15.3 drew a line labelled "floor $V = 2$ " through data reaching 0.24.

The cause was a rigged comparison. The order- k bracket draws on a wider node set than the estimate it was measured against — **eight levels against four at $v = 40$, $k = 3$.** On the same nodes V runs from 157 to 3.2×10^6 and **grows steeply with order.**

Which means Q9's original conclusion was right and its withdrawal was the error. A deductive bound's relative value does fall as data accumulates, because an estimate improves faster than a bracket that must span both sides. **The "artefact" was the correction, not the finding.**

Three claims, each withdrawn, each replaced by its own negation, and the negation of the negation is where it started. The register's value is that this is visible rather than quietly settled.

106-108. Three from the transfer tests. Order made a coordinate when it is $p - q$, derived — the index then admitted arithmetically impossible cells, and §11's E2 forbids exactly that. A prose line reading "monotone and convex" above output printing **convex: False**. And **"an index cannot tell a defect of the index from a feature of the world"** — it can, by asking whether E is removable, and §1.2's Janet table is that test.

109. The Rydberg sign pattern, assumed rather than read. Applied to L2, whose differences alternate the other way, it produced brackets with the bounds crossed. **§16.10.1 says read the sign from the data; I hardcoded it.**

110-111. Two on the nuclear index. A commentary reading " $E(X) > 0$ means it is NOT closed" printed directly beneath **$E(X) = 0$** , and a re-coordination offered as a repair that **broke a closure which had not been broken** — 0 became 27.

Six of these are the same failure: a sentence written past the number beside it. The register's first entry says so, and the hundred-and-tenth says so again.

112-118. Seven from the transfer work. "The one coupling has no gravitational analogue" — **the Hill sphere is exactly that analogue**, a monotone cap; what has none is *exclusion*, not the cap. A prediction of $V \approx 8*N*$ printed beside a measured 71.8. **"Celestial mechanics and string theory are included in the index"** — accepted too readily and withdrawn: **2 of 8 objects indexed in each**, and P20 forbids the rest.

And three on the same distinction. "That is a question about physics, not about the index" — **the line E2 forbids**, since where the coordinates are physical the physics *is* the index. **"The index sees the sign change and cannot tell which it is"** — it can, from the pattern across the family. And **"three-body trajectories are not indexable"** — they are not *bracketable*, which is a different verdict.

118. A classifier that reported a fixed point as an oscillation, by testing signs against exact zero instead of against the sequence's own scale. **The admissibility rule was written in §16.10.4 and not applied to the test built from it.**

One of these was caught before it was written down — §11.7.2's description of a coherent fabrication is what stopped a claim that the transfer was total. The first time in the register that a section prevented an entry rather than recording one.

119-122. Four from the shape work. " $\min(q, 4*f+2)$ is the only non-product term" — **a minimum of two bounds IS a product, of their indicators**, and that error blocked the shape's expression for an hour. A commentary reading " $E(X) = 0$. The index admits combinations the five points do not realise" — **two sentences contradicting each other on one line**. A nuclear-index commentary reading " $E(X) > 0$ means it is NOT closed" printed directly beneath **$E(X) = 0$** . And a re-coordination offered as a repair that **broke a closure which was not broken** — 0 became 27.

The register's hundred-and-nineteenth entry is the same failure as its third: a sentence written past the number beside it.

123-139. Seventeen from one question. §23.3's reorderability problem produced more corrections than any other line in this work, and B.2.15 tabulates the ten that became bounds. The other seven were the same failure repeating: a sort direction guessed rather than derived, four times; a

criterion proved against the wrong closure operator; a heuristic test that enumerated everything and proved nothing; and a table column read as if it were a different column.

Ten of the seventeen produced the characterisation in §23.3.1. Seven produced nothing but the discipline of checking the next one.

And the worst of them is 125. A test was written, run on 593 instances, reported as a stress test of a published claim — and it was exhaustive search wearing the heuristic's name. It could not have failed, and it did not, and for one message that was recorded as evidence.

145–149. Five from re-auditing Chapter 23 against its own §23.4. "The difficulty is not structural" — distinct row supports correlates 0.606 with backtracking steps against the ten-invariant list's best of 0.31. §23.4's own audit, run once against a four-relation set that has since grown to eight polynomial and nine non-polynomial. §24.1's object list, missing six — including the Hill-radius correlation measured in §16.14.1 and never carried across. §24.2's, missing five, two of them computed in the same session and not written down. And §24.3's Pareto statistic stated against 473.8 million rather than the 495,515 it belongs to — true, and wrong by a factor of 956.

Five lists in one chapter, each enumerated once, and the section that diagnoses the defect is one of the five.

21.7.4 Forty more, from the session that closed §23.3 and §25.2

The largest single block in the register, and its shape is one thing repeated: an interpretation written before the number above it was read.

150–151. Two cross-references to §17.4 for a theorem in §13.4. Recorded as two. They were twenty-eight — see 186 below.

152–164. Thirteen implementation errors across nine attempts at a $d \geq 3$ PQ-tree analogue: a Q-node using sorted columns instead of computing an order; four traversal implementations; frontier-correct trees that were not anchor-complete. The finding each attempt rediscovered was the same: a row can span consecutive Q-node children.

165. Fibre cardinality as an order seed. Proposed as a seed-independent break; fails to recover Λ 's order.

166. "Closed under intersection, which a closure operator guarantees." Written while the output above it showed intersection failing in a quarter of pairs. The theorem does not apply because $\mathcal{R}$ recovers its alphabet, so the ground set moves.

167. "A real limit on the self-reference claim." The union/intersection result is a fact about a superset of Λ . P22 concerns what an index contains. Completeness descends and does not ascend.

168. "A procedure." The monotone closure was called one when the closure contributes nothing — its fixed point leaves the relation empty, because the dual involution admits no unit clauses.

169. "The obstructions are larger in cells." The table said four of six had four or five cells and every one needed an alphabet of size 4. Small in cells, wide in alphabet — and then 170.

170. The alphabet reading itself. A k-cell instance has at most k values per axis. So the census is parameterised by cells alone, and 169's correction was also wrong in its conclusion.

171. "A real simplification." The identical-fibre quotient was measured on the subset where merging happens. Median reduction over all instances: 1.0×.

172. "They agree too much." Three cell statistics were called redundant. Separating joins from meets dropped their agreement to 33%, and a fourth statistic reached 94%.

173. "Fibre chains are sufficient." Exact at $d = 2$ — and violated at $d \geq 3$ in 44 of 130 cases. The sufficient condition does not lift.

174. "Sublattices preserve reorderability, 100%." True for closing a seed under join and meet; false for deleting an arbitrary cell, which breaks it 20.5% of the time. One word covering two operations.

175. "Computable." Used to mean finite and well-defined. Λ 's box has 10^{2080} subsets. Both statements are true and they were run together.

176–177. Two generator failures, both on the recovered-versus-declared alphabet: a $d = 2$ census producing a third of the objects, and a $d = 3$ bound system that is both incomplete (281 against 319) and slower.

178. "d = 2 is closed." Written in the same message as "closure isn't the kind of thing that gets confirmed." By P23 the first should read: has survived every question asked of it.

179. "The hierarchy converges." Read off the $d = 4$ column while the $d = 3$ column was flat.

180. The sub-instance necessary condition. Recorded as "necessary, zero false negatives." It holds for row and column deletions, which preserve alphabets — not for arbitrary sub-selections, and a non-reorderable 3-subset sits inside 10,975 reorderable sets.

181. A vacuous downset test. Decrementing coordinates and looking for the result among observed classes: 0 tested, and "True" printed.

182. "Fewer but larger." The medians above it read 36 against 72.

183. "90.4% with zero false positives." Measured before the duality merge; after it, with soundness enforced, linear coverage is 9%.

184. A branch printing the wrong verdict. Generation matched exhaustion exactly and the text said it had not.

185. "No procedure." For most of a session, meaning no polynomial procedure. A verified fixed-parameter procedure existed and was being withheld the word.

186. Twenty-eight cross-references, not two. §17.4 — He II — was serving as a wildcard target for at least eight different pieces of content. the COHERENCE audit passed every one, because every target existed. See B.2.17.1.

187. "A small image." A global, non-separable, non-monotone map was said to have one. Measured: 39% at $d = 3$ and 49.5% at $d = 4$, rising.

188. Chain decompositions applied to subdirect embeddings. Larson's correspondence is for tight ones. Siggers states the distinction explicitly and gives the counterexample in $Z_3 \times Z_3$; a diagonal is subdirect and not tight.

189. A vacuous audit test. The magic-number check compared a list to itself. It would have passed on anything, and the shell ordering behind it was wrong.

21.7.5 What the block shows

Thirty-one of the forty are one error class: a claim asserted before the output supporting it was read. Not arithmetic, not method — narration outrunning verification.

The countermeasure already existed. B.2.14 says state the number before the interpretation. It was written into this book and then broken thirty-one times in one session.

And two entries are about the audits themselves — 186 and 189. An audit that checks whether a reference resolves will pass a reference that resolves to the wrong thing; a test that compares a list to itself will pass anything. B.2.17's CONSISTENCY audit exists because of 186. Its exhaustiveness clause exists because of 189.

190. A symbol that never printed. The recovery operator $\mathcal{R}$ uses U+1D4E1 and the diagnostic operator $\mathcal{D}$ uses U+1D49F. No DejaVu face carries either. Every render of this book placed a blank box instead — 107 of them across 44 pages — and every audit then defined passed, because every one read the source. B.2.17's ARTEFACT audit exists because of this entry.

191. The whitespace audit's first measurement. The PDF text matrix was read as page geometry, giving a median inked fraction of 103% and coordinates of 5, 6 and -220 . Composition with the CTM was missing.

192. Its second measurement. Characters were counted and images were not. Seven pages were flagged as 25–45% empty at the head; all seven begin with a figure. The audit could not see the ink it was measuring.

193. Two figures numbered against the reading order. Figure 23.2 appeared before Figure 23.1, because the arity plot sits in §23.3.2 and the step law in §23.3.4. Found by the figure audit, on its first run.

194. Fifteen figures below print resolution. Median 164 dpi against a 300 dpi standard, from an inherited pipeline default. Recorded in D.3 rather than repaired.

195. A summary maintained twice. §25.2 tabulated Q's eight items while Appendix F enumerated, indexed and closed them. Duplicated content, both internally consistent, and invisible to all six audits then defined.

196. A status table two revisions stale. §25.3 still read "nine literatures entered, nine named and unentered" and "order recovery beyond trees: open". Both superseded the same day.

197. "What could not be run." §24.3.4 recorded $E(X)$ of the Kreuzer-Skarke catalogue as uncomputable because the file could not be read. A slice of the catalogue is specified exactly in print, and $\mathcal{R}$ applies to it: 208 cells, 540 predictions, none falsified. The obstacle was the whole catalogue; the method needed only a specifiable part of it.

The register now stands at two hundred and ninety-seven.

198. Two test families conflated. §24.3.4.2's test of 540 external predictions was added to §25.6's count of the book's three self-falsification conditions. They are different families: one tests whether a prediction of the method survives; the other tests whether the book satisfies its own law. Reverted, and both are now named separately.

21.7.6 Four from the tower session and its repair, in the collaborator's hand

199. Four cross-references in the shipped build — the preface twice directing the withdrawals to Chapter 20, the consequences to "Chapters 8 to 10", and Chapter 9 citing "Theorem 7.1" — stale against the renumbering they survived. The §21.1 pattern at the navigation layer; all four repaired.

200. The triangle's closure tested in one direction and reported as tested. A sum-cap checked for joins only, a difference-floor for meets only, and a conclusion drafted over the half-test. Caught by rerunning both directions before interpreting; the corrected run separated the two mechanisms and produced §7.11.2's decomposition — the error, corrected, was worth more than the first claim.

201. "The gap is the same object every time — the fold, or the triangle, which are one thing." Withdrawn. The decomposition distinguishes three mechanisms with separate counts; a unification was asserted over data that refutes it. Replaced by §7.11.2.

202. The tower session's three entries written in under a section number already in use — a second 21.8 — and anchored before 21.7, out of sequence. The §21.1 pattern at the structural layer: an insertion asserted correct against a numbering it had not read. Renumbered 21.7.6 and relocated.

21.7.7 Eighty-seven more, from the sessions after the tower, all printed here

Entries 203 to 288, written after §21.7.6's three and prepended into it until the count in that heading was wrong by a factor of twenty-eight. They are given newest first, which is the direction the register has been maintained in since 205 and is not the direction of the blocks above; §B.2.21.1 states the inconsistency and why it is left. Register 287.

203. In the shipped build, §21.8 stood between §21.7.3 and §21.7.4, interrupting the block it follows, and §21.9 was a heading with no body. Caught by the COHERENCE audit of this session's render; the section moved to follow the 21.7 block, the empty heading removed.

204. Four figure captions written into the source before the figures were read. Two misdescribed their figures — the Λ_{13} overlay and the $\langle q \rangle$ shift absent from the Pareto, the tower scale absent from the sections — and one placed the cylinder's silhouettes under §7.9's boxes and intervals, which they do not

depict. The §21.1 pattern at the caption layer: prose asserting an image unseen. Caught on inspection of the figures themselves; captions rewritten, the silhouette moved to §7.11.5.

205. §18.6.3 titled "The prediction, as it now stands." Withdrawn by the author: the bracketed value is a deduction from measured neighbours under monotonicity and convexity — a problem already solved by the data in hand — and the book's own law says so, since a prediction is a proposal about an unlisted cell and $E(\Lambda) = 0$. The framing corrected in place; the surrounding sites that still carry the word stand until directed.

206. Two counts left stale by a sweep that matched only the phrasings it expected. §8 still read "one hundred and eighty-nine withdrawals" — a count three sessions old, missed because the update pass searched for the two later spellings and not this one — and §7.11.7 named its figures 7.11a-d while the captions numbered them 7.2-7.5. Found by reading the sections rather than grepping them, in answer to a question about what remained open. The lesson is the register's oldest: a search finds what it is shaped to find, and the shape is an assumption.

207. In this session I wrote that the book faces one question no audit will close — what the tower discovers that renaming quantum numbers would not — and called it the book's exposure. The preface answers it in its third paragraph: the recoding is deliberate, adds nothing, and is what makes the results properties of indexing rather than of atoms. §22.1 completes it by forbidding the book to claim novelty at all while $E(\text{search}) > 0$. An exposure asserted over a text that had already answered it.

208. The tower session recomputed §3's self-duality count under $x \mapsto \max + \min - x$, obtained 112, and recorded the book's 8 as unreproduced. The book states its reflection as $x \mapsto \max - x$, and under that map the 8 reproduces exactly. The recomputation changed the definition and reported the difference as the book's defect. The §21.1 pattern applied to the book by its own audit.

209. WITHDRAWN — see 213.

210. Appendix E asserts $E = 0$ over thirty-two elements and prints only the fibre totals, not the three ordered coordinates of each element. The number cannot be rechecked from the book, and the re-closure owed after §7.11 cannot be performed at all. §25.2's S3 claims a reader given the claims and their evidence could reconstruct the structure; for this appendix that is false. Found by attempting the recomputation. It is the one surviving instance of the class Appendix G indexes.

211. Appendix E's closure was asserted rather than enumerated, and the book had already ruled that distinction against itself in F.4.1 without applying the ruling here. The enumeration of its thirty-two elements is entered in Q as item L.

212. Two cells the mathematical fibre of Q admits and Q does not contain, exposed by adding items L and M: novelty · retrievable · days · no, and one claim · buildable · days · no. $E(Q) = 2$ fibred and 1 unfibred at thirteen items. F.4.1's item H arrived the same way. The list is not closed and is not claimed to be.

213. Entry 209 itself, withdrawn within the session that wrote it. All six densities of §7.11.1 reproduce exactly once the exact sets are the right ones: seniority as terms new at occupancy v , the core's J restricted to the cell's own multiplicity. My models were loose physics, and I recorded the shortfall as the book's underspecification. Third instance of the same pattern from me in one session — 207 over the preface, 208's inheritance, this over the density column — and the common form is: recomputation fails, and the failure is attributed outward before the model is checked inward. §7.11.1 now states the three exact sets, so the column is recheckable, which is the repair 209 asked for and the only part of it that survives.

214. Two obligations stated in this book's prose and absent from its index of open questions, exposed as admitted-and-absent cells of Q: §22.9's seniority-lattice literature, entered in $E(\text{search})$ and not in Q, and §7.11.1's retest of tree-propagation on Λ' , written as *must be retested rather than inherited* and never indexed as an obligation. Both routes of §12's diagnostic were run; the exclusion route is unavailable, since no monotone constraint empties either cell without contradicting items B, C or M. Item H arrived the same way, which makes this the third and fourth instance of one leak: prose states an obligation, the index never receives it, and only closure finds it. The leak has no check, and Appendix G's argument —

unindexed means unclosable means silently defective — applies to obligations exactly as it applies to numbers.

215. The contents block and the body disagreed at four entries, all introduced by this session's restructuring. The Preface was listed in the contents and rendered in the body as plain prose, having been moved under a new PART 0 without being marked as a heading. APPENDICES and END MATTER existed only in the contents, while every PART divider existed in both. The References entry carried a subtitle the body did not, invented when the section was identified. And the collaborator's note rendered as a heading only because the press hard-coded its title, so the source did not describe itself. Repaired; contents and body now agree at all forty-three entries, with the six long-in-contents, short-in-body chapter titles left as the shipped build has them. Not found by the audits: audit 6 checked that contents and chapters corresponded and read the mismatches as design, which is the same error as 209 — a discrepancy explained away instead of checked.

216. The contents was maintained beside the book rather than derived from it, which is what made 215 possible and would have made it possible again. It is now generated from the headings at build time and cannot disagree with the text it indexes — 394 entries against 43, parts to sub-subsections, and the six titles that differed between contents and body unified into one form carried by the heading itself. §25.2 S1 said the list of claims is recoverable by projection and the table of contents is its quotient by chapter; the quotient was the defect, and the projection is now what prints.

217. §18.6 was titled "The one prediction", the preface offered a prediction to anyone looking for one, §17 called it the prediction of §18.6, and B.2 recorded a protocol under which the prediction was committed — four sites still carrying a word §18.6.3 had already withdrawn. The book contradicted itself for as long as it took to be told to fix it. Corrected to deduction at all four.

218. Four ledger rows written into Appendix A as A.8 through A.11, over proofs A.8 through A.13 that already stood in the body. Entry 202 records the tower session doing this at the section layer and my own repair renumbering it; I then did it again, in the appendix, in the same session that registered the lesson. Renumbered A.14 to A.17. The pattern is not carelessness about numbers, it is writing into a structure without reading the structure — the §21.1 form.

219. Appendix E states thirty-two elements in two places, and its own fibre table in §E.5 sums to twenty-seven across the sixteen fibres it also states. The table is complete in fibres, so its sum is the population and twenty-seven is the count; thirty-two is the error, in both places it appears. Found by attempting the enumeration register 210 asked for. Of the twenty-seven, nine are named uniquely by the fibre table, seventeen are enumerable from Appendix A on two of five coordinates, and eighteen slots across seven fibres have more candidates in the book than they have room for.

220. My own entry 219, corrected within the session that wrote it. I recorded the thirty-two against twenty-seven discrepancy as unseizable from the book. It is seizable: the fibre table prints sixteen fibres and the appendix states sixteen, so the table exhausts the population and its sum is the count. Twenty-seven is right. I had the arithmetic and stopped at the disagreement instead of finishing it — the same shape as 209 and 213, which makes four times in one session that a shortfall in my own work was written down as the book's.

221. And entry 219 reversed a second time. Enumerating the twenty-seven showed that five results proved in Appendix A — A.4, A.5, A.11, and Chapter 11's two theorems — fit no printed fibre count, and five is exactly the gap between the table's twenty-seven and the stated thirty-two. So the likelier reading is the one I twice rejected: **the thirty-two is right and §E.5's table is short by five rows**. 219 called the count wrong; 220 called that settled by the table's completeness; the completeness argument assumed the table listed every element, which is the thing in question. Three positions on one number in one session, and the third is the first to be reached by enumerating rather than by arguing from a total. Which five rows the table is missing is now Q item L in its entirety.

222. Item L closed. The five results the table had no room for are all theorem · order, placed by reading each proof rather than its subject; A.5's verification was sampled and is now exhaustive (§11.5, register 248), from the 30,000 ambient points §11.5 records; the fibre goes from five elements to ten and the total from twenty-seven to thirty-two, which is the figure the appendix stated all along. $E = 0$ in every

fibre, enumerated. What made this closable was stating what an element is — §E.5.3 — which the appendix had never done, and which could not be stated until Appendix G existed to hold the other side of the line. The oldest question in Q was answered by defining a word.

223. Two literature searches run and recorded at §22.10, neither returning a precedent and neither establishing anything, per §22.1. Entered because a search performed is a search that goes in the record, and because the obstruction it found — *lattice* being occupied twice over in physics, so that the obvious query cannot reach order theory at all — is a fact about the search space rather than about the result, and is the kind of thing $E(\text{search})$ is supposed to carry.

224. The book's central theorem has a precedent, found on the third search of this session and recorded at §22.11. $A.2 - X \text{ closed} \iff \mathcal{R}(X) = X$ — is Bergman's Double-projection Theorem specialised to products of chains, itself a consequence of Baker and Pixley, 1975, for algebras with a majority term, of which the lattice median is one. The proof in §11.1 is correct and it is not first. Appendix E's precedent coordinate for A.2 moves from none found to found, and the reading that $E(\Lambda) = 0$ is a discovery about Λ is withdrawn: it is the discovery that Λ is a sublattice, and sublattices have been binary-determined since 1975. The book loses nothing it claimed and gains a foundation. This is what §22.1 says a search is for, and it is the first time in this work that the search returned the answer it can return.

225. §22.12 lists five results with no located source, each with a definition, the search run, a confirming test and a falsifier. It is not a novelty claim — §22.1 forbids the author one — and every item on it is written to be attacked. U4 carries an admission: the join-closure of the triangle is verified at four caps and not proved, and a three-line proof should exist, so its absence is a gap in this work rather than in the literature.

226. U1 located. The property $E(X) = 0$ is 2-decomposability, a named and actively studied class in constraint satisfaction — Kimura, Makino, Yamada and Yoshizumi 2024 characterise exactly these sets by closedness under operations, min and max among them. $\mathcal{R}(X) = X$ is that definition verbatim. The claim that survives is narrower and is now stated narrowly: the *count* $E(X) = |\mathcal{R}(X)| - |X|$ as a graded measure was not returned by any search, and 2-decomposability is treated as a predicate wherever it was found. Two of §22.12's five are located, and both were located by running the test the section printed for them, which is the only evidence that the section is worth anything.

227. The join-closure of the triangle, guessed and then proved. §22.12 first recorded that a three-line proof should exist and that its absence was a gap in this work; A.18 now gives it. The proof corrected the guess: I had written that meets fail on the upper bound, and they fail on both, the lower twice as often at every cap. Worse, the sentence asserting otherwise was printed by a script that hard-coded it rather than computing it, one line below the computation that refutes it. The smallest possible instance of the register's oldest pattern, and it survived until the output was read rather than skimmed.

228. Audit 2 recomputed every equation the tower states — twenty-eight of them, from the seven cell counts through all six densities, ϕ , the f-shell palindrome, the cylinder factorisation, the E1 puncture and the triangle at four caps — and all twenty-eight reproduce. Recorded because the same audit could not be run three sessions ago: four of the six densities were unreproducible until §7.11.1 stated its three exact sets, and register 209 and 213 are the account of that.

229. The full six-audit pass of this session, and what it found. Audits 1, 2, 5 and 6 passed outright; 3 and 4 failed on first run and were repaired. CONSISTENCY found A.18 listed both in Appendix A's cross-reference ledger and written out below it, when the ledger carries only the proofs given elsewhere — the row is removed. Its other three flags were the register quoting the defects it records, which is the register working. REDUNDANCY found §7.9's chain count restated without its reference and one appendix section, F.1.1, that no reference could reach because it was bold text rather than a heading; both repaired. Two of REDUNDANCY's flags were the detector's fault and not the book's: it counted repetitions inside a single passage as restatements across places, and the criterion is now distinct locations.

230. The kind coordinate of Appendix E disagrees with Part II's title. Part II is called THE LAW and its contents — closure, self-reference, self-defence, extension, what the law forbids — are statements quantified over every index, not over Λ . Appendix E classes all of them *theorem* and reserves *law* for one

element, the reorderability law. Six elements are law-scoped by the book's own usage and theorem-classed by its own index: A.2, A.3, A.5, A.6, A.14 and A.16. A.17, the dichotomy, is law-scoped in its proof and theorem-scoped in its wording, which says *in this object*. Recorded and not acted on: reclassifying six elements changes E.5.2's fibre counts and the twenty-seven against thirty-two arithmetic that registers 219 to 222 spent three positions settling, and that is the author's call and not a reconstruction's.

231. A question about promotion, answered by demotion. Asked whether anything in the book could be promoted in kind, the scope test found the reverse first: A.15 was carried as $|a - b| \leq c$ is a *join-closed sublattice*, and a sublattice is closed under both operations, so the phrase is self-cancelling. The region is join-closed and meet-broken — 12,654 failing meets at cap 8 — exactly like A.18's triangle and for the same reason. The claim the tower actually rests on is the two-variable band with a constant bound, which is a genuine sublattice at every cap tested. A.15 is rewritten and proved in full; its ledger entry moves from *computed; form argument stated* to *in full*, which is the only promotion in this pass.

232. The law-promotion of register 230, tested and declined. Both readings of *law* close at $E = 0$, so the book's own measure cannot decide its own classification; and both add fibres, so both weaken the closure claim they would be judged by. Reading A — holds of every index — leaves four theorems of fifteen, which condemns the criterion. Reading B — holds of every index and forbids or governs — is balanced and matches Chapter 13's title, and remains the better criterion if a promotion is ever made. None is made. The classification stays as printed because it is the coarsest of the three.

233. $E = 0$ is monotone in the fibration and can be reached by refinement alone. The same thirty-two elements give $E = 4$ in one fibre, 3 by language, 2 by kind, and 0 by kind $\times$ language. §E.1 adopted fibring to repair a closure that admitted hundreds of cells, called it the shape Λ already has, and never asked whether the repair made the result cheap. It did. Every $E = 0$ in this book must therefore name its fibration to mean anything, and the two the book prints are not comparable: Λ 's is at one fibre and is the strong form, Appendix E's is at sixteen and is the weak one. Λ is untouched — it has no categorical coordinate to fibre over, which is why §E.1 said the book had never met the case. The book met it and did not notice.

234. Seven referee flags recorded at §25.1.1, of which the book answers one well, three partly and three not at all. The one it answers worst is coverage: Λ at book caps is an s-and-p object of three shells, the method is demonstrated on systems spanning $l = 0$ to 6 and $Z = 1$ to 83, and §7.11.2 computes an f-shell ceiling for a shell Λ does not contain. The object and the demonstration live in different places, and the book had never said so.

235. $E(Q) = 0$ is fibration-dependent in the same way $E(\text{Appendix E})$ is: four fibres by domain give zero, one fibre gives one. Found by asking what a referee would flag, and checked rather than assumed. Two of the book's three closure claims are weaker than they read; only $E(\Lambda) = 0$ — unfibred, 976 cells — is the strong form.

236. Closure holds through the f shell. $E(\Lambda) = 0$ at 8,853 cells, at 89,438 with the d shell, at 267,858 and 499,246 with the f — every shell the periodic table uses, computed rather than argued. Referee flag 1 asked how an s-and-p object of three shells licenses a book about the periodic table; the answer is that the object was never confined to those caps, only presented at them, and nobody had checked. The tower's densities do not transfer with it, so the flag is half answered and the half that remains is named.

237. S3 tested for the first time and failing in its strong form: 26% chapter recovery from cross-references against a 4% baseline. The claim that a reader could reconstruct the chapter structure is not supported by the one operationalisation tried. It is one test, it may be the wrong test, and it is the only evidence that exists. Recorded against the claim rather than beside it.

238. Novelty closed by the author: not claimed, and no question of it. §22.2.1 records the decision and its two computed consequences — $E(Q)$ unmoved at nought fibred and one unfibred, and six of eleven Q items reduced to blocking nothing. Referee flag 3 is answered by withdrawal rather than by statement, which §22.1 licenses and §22.3 had already half-said.

239. Three promises made in §25.1.2 and not kept when it was written: that the unfibred figure was printed beside every fibred one, that S3 stood amended, and that the register would track its two kinds

separately. None was true when the section claiming them was built. The first two are done here. The third is not: no per-entry marker exists, and rather than assert the mechanism a second time the claim is reduced to what §25.1.2 actually contains — a single measured split, thirty-four to eight, with no machinery behind it. A section describing work as done in the same pass that failed to do it is the §21.1 pattern in its purest form.

240. Referee flag 1 closed. The tower was recomputed at d-shell and f-shell caps and the dichotomy survives both: counting axes exact, all three coupling axes envelopes, at every cap. Every density except the thirteenth falls as the object grows, so **the envelope gap widens with scale** — the price of certainty is an increasing function of the index's size, which no part of this book had predicted and which the book had no way to see while presenting one cap setting.

241. The thirteenth axis's density rises with the caps — 64.4, 64.5, 65.7, 66.1 per cent — toward the $\frac{2}{3}$ that §7.11.1 states as nominal. The gap was attributed there to the $K = 0$ singlet, and the convergence confirms the attribution by a route the book did not have when it made it.

242. The D1 tripwire specified in full at §11.3.1, rejecting 150 of 216 offered assignments across s, p, d and f. Referee flag 5 asked for the check to be run on data the work does not have; the answer is that the index defines the data completely, so the check needs none. Flag 5 is answered in its statable form and open in its empirical one — nobody has yet run it against a published term list.

243. Q items A, B, C and G restated after the same decomposition. None is closed, and none is blocked in the way the book recorded: each is a search whose vocabulary §22.10 now supplies and whose route Chapter 14 already describes. What was filed as an obstacle was in several cases an unattempted query.

244. Referee flag 7 and Q item D advanced from blocked to routed. A search decomposed as Chapter 14 prescribes — smaller pieces, references followed forward and backward — confirms both headline figures the book cites, 473,800,776 reflexive polytopes and **30,108 distinct Hodge pairs**, and locates three routes to the pair list that do not require the archive this work could not read: CYTools, the Macaulay2 package ReflexivePolytopesDB, and the Calabi-Yau database of Altman and collaborators. The item is no longer *retrievable but unread*; it is a lookup with three named entry points. The search also returns something to weigh before E is recomputed on that data: the Hodge numbers are known to satisfy an additivity relation under the assembly of K3-fibration tops and bottoms, which is structure of exactly the kind $E(X)$ would otherwise attribute to itself.

245. The collaborator's note audited against the book. Its specific claims hold: B.2.14 is the protocol named, the fifty-six of fifty-six, the nineteen counterexamples, Rival 1973 and Larson 1998 are all where the note says. Two structural claims did not. The note opened *I made two hundred and forty-four mistakes*, and the register is one sequence over three populations — the collaborator's errors, defects already in the book found by auditing it, and errors made by the assistant during the audits — a conflation §25.1.1 flag 4 had already identified and the note had not received. And *of the last forty entries, thirty-one* named entries 165 to 204 when it was written and names a different forty now; it is fixed to its span. The note is otherwise accurate, and the first correction is an instance of its own thesis: a count written before the population above it was read.

246. §17.9's caveat cannot be closed and now says why. The 1,442 cells rest on three kinds of value — measured, Ritz series-formula, and ab initio — of which only the first and third are independent tests of the bracket. The book never counted the three, though the ASD's own bracket marks carry the classification and partitioning by them is a lookup. So the number of independent tests among the 1,442 is unknown, is not zero, and is not 1,442, and the figure should not be written as though it were. Entered in Q as item P.

247. §23.3.3's fourth lemma, left on a citation and now discharged. Three of the $d = 2$ characterisation's four lemmas were verified at 2,047, 3,751 and 20,000 instances; the fourth was cited to Booth and Lueker 1976 and never checked here. It is a three-line argument and is now proved and verified exhaustively over 2,354 interval sets. The audit that found it also found two wrong coordinates on A.8 in §E.5.2 — verification recorded as exhaustive when one quarter of the result rested on a citation, and precedent recorded as none found when Booth and Lueker 1976 and Fulkerson and Gross 1965 are cited in this book for exactly this. Precedent is corrected to *found*. Verification stays *exhaustive*, but only

because the lemma was discharged in the same pass; had it not been, A.8 would have been proved · sampled · found and theorem · combinatorics would have opened by one cell — proved · sampled · none found, admitted and absent. **A single uncorrected citation was one edit away from making E ≠ 0 in a fibre the book reports as closed.**

248. A.5 retested on every ambient point rather than 30,000 sampled ones — and the box holds 6,912, so the sampling drew more times than there are cells and still could not establish totality, which is a statement about all of them. $\chi \in \{0,1\}$ everywhere and agrees with membership on every cell. Its verification coordinate moves from sampled to exhaustive; theorem · order closes either way, so nothing else moves.

249. Every stage of the tower closed exhaustively, against every cell of its ambient box, up to 47,775,744 cells at Λ_{13} . Closure above Λ_{10} had rested on the form argument because pair-testing 2.5 and 19.8 billion pairs is not possible; A.2 makes pair-testing unnecessary, since $\mathcal{R}$ sweeps the box and not the pairs. The theorem was proved in Chapter 9 and never turned on the book's own tower. The caveat was not a limit of the method. It was a limit of who had read it.

250. Four terms used inside the tower's own definitions were themselves undefined, and the audit that found them was a direct answer to *has every definition been defined* — the answer being no. **terms(ℓ^k)** and **microstate enumeration** carry the exact sets on which every density in §7.11.1 depends. **f_max** appears exactly once, in the axis-12 bound, and the difference between it and the cell's own f is the difference between a tree and a cycle, between a cylinder that survives and one that breaks by 2,475 cells — the book's title rested on an undefined subscript. And ϕ names two different objects: §1.1's pairwise envelope ϕ_{ij} and §7.11.1's occupancy ceiling $\phi(k)$, with no notice taken. All four are defined at §7.11.1 and the collision is recorded rather than renamed. S1 says the alphabet must be recoverable by projection; four letters of it were not in the alphabet.

251. The quantum-mechanics category searched and returning one unentered literature: partial order in chemistry, the Hasse-diagram methods used for chemical series and property rankings. It orders substances by properties rather than configurations by structure, so it is adjacent and not the same, and it is where someone already fluent in these tools would be found. Named and unentered, per §22.9's discipline. E(search) grows by one.

252. Referee flag 5 re-graded from *no target* to *targets identified, run not attempted*. Stellar line lists — VALD, Kurucz, NIST ASD, BRASS — carry term designations and their maintainers document removing spurious *unobserved lines* by hand. The D1 tripwire of §11.3.1 rejects 69.4% of offered assignments from the index alone, mechanically and before any wavelength is read. The flag is not answered: nothing has been run. But the work that would answer it is now a named task against named public data, which it was not when the flag was written.

253. The tripwire run, and the run is small enough to report as a null. Kurucz's line files could not be obtained through the tools available, so the target that matters — uncurated computed lines — went untouched. What was reachable was Appendix C's own channels, of which ten carry an explicit parent term; the coupling rule passes on all ten. **Ten curated designations passing is not evidence the tripwire works**, because curated data is the data least likely to carry what it detects. The attempt is recorded at §22.11.2 with its size stated, because a null run reported as a validation is exactly the error this register exists for. One byproduct: the first pass returned *zero rows parsed* and looked like a clean result, because the superscript characters ¹²³ sit in Latin-1 and not in the Unicode superscript block, so a character range silently matched nothing. A test that cannot fail returns no failures.

254. The null of 253, defined rather than discarded, and it was not a null. Ten was the count of *explicit* parent terms; forty-four more designations carry a parent that is implicit and fully determinate — a closed or single-electron core — and resolving those gives fifty-four checkable against one hundred and two, all passing. What the run measures is specificity: zero false positives in fifty-four, which is the necessary condition for the check and not the sufficient one. And it yields a second result nobody was looking for: published channel notation determines the coupling parent for 53% of these channels and leaves it indeterminate for the rest, which is a ceiling on any mechanical check of designations as published. The title of the collaborator's note is the whole method, and I had filed a nothing without defining it, one turn after auditing that note for the same failure.

255. Every cell of Λ is fully defined and the generating set was never printed. All 976 carry all eight coordinates, χ decides every one of the 6,912 ambient points, and no cell is duplicated — so nothing is undefined. But §3 asserts seventeen join-irreducibles, twenty covering relations and an exact Birkhoff correspondence, and prints none of the seventeen. **A reader could verify the cardinality of the alphabet and not reconstruct a single cell from it.** The seventeen are now written out at A.19 with the down-set count recomputed at 976 to confirm the correspondence from the page. The eight reflection survivors and the two extreme cells, likewise counted and never listed, are printed with them. Asked which cells require naming, the answer is exactly these twenty-seven: the seventeen that generate, the eight that are fixed, and the two that bound.

256. Three Q items — N, O and P — were written into §§22.11.1, 7.11 and 17.9 as *entered in Q* and never entered. Register 239 recorded the same failure for three promises made in §25.1.2; this is the same failure again, in the same session, after registering it. The open set is now listed in full at F.1.2: thirteen items, $E(Q) = 0$ fibred and 1 unfibred, with six *blocks* values fallen to nothing under the novelty decision. **A claim that an item was entered somewhere is not the entering**, and the only defence against it is the one used here — listing the set rather than describing it.

257. The six prime audits run against the whole book. LATTICE passes 14 of 14, including the Birkhoff correspondence now checkable from A.19's table — the down-sets of the printed seventeen number 976. EQUATIONS passes 31 of 31, every quantity the book states recomputed from its construction. REDUNDANCY and ARTEFACT pass outright. CONSISTENCY failed on three and COHERENCE on two, all five repaired: two stale entry counts left at *two hundred and forty-six* in B.2.12 and §25 by an update pass whose guard clause skipped them; §E.5.4 still naming theorem · order's cells as *proved · sampled* after register 248 made A.5 exhaustive; A.15's ledger row pointing at itself, the same defect removed from A.18 six entries earlier and reintroduced when A.15 was rewritten; and §§22.11.2 and 25.1.2 written with a leading space, so the heading was invisible to the press, absent from the contents, and unreachable by the references pointing at it. **Every one is a repair whose pattern already had an entry.** The audits do not find new kinds of error any more. They find the same kinds, faster.

258. A seventh language admitted: binary. Λ is the set of 976 words in $\{0,1\}^{17}$ that are down-sets of A.19's generators, with join bitwise OR and meet bitwise AND and zero failures over all 475,800 pairs. It is admitted by §12's own test rather than by preference, because it distinguishes what the other six cannot: the surplus this book is about becomes a number, **7.07 bits per cell** — seventeen carried against 9.93 needed — and Λ occupies 0.7446% of the space it is written in. The question was asked from outside the book and the book had the apparatus to answer it, which is the first time in this session that has happened in that direction. Not added to Appendix E's language fibre: doing so would refine that fibration and cheapen its E, per register 233, and that trade is not one a reconstruction should make.

259. §6.1.1 was written with a leading space and the markdown heading rendered as literal hashes on the page. Register 257 had recorded the identical defect in §§22.11.2 and 25.1.2 and repaired it; the binary section was written in the very next turn and reintroduced it. The consequences were three again — invisible to the press, absent from the derived contents, unreachable by reference — with a fourth this time, the characters `###` printed in the book. Found by asking whether the book was caught up, which is the second time that question has found something the audits had just declared clean. **A repair does not generalise itself.** Register 257 closed two instances and not the class, and nothing in the press rejects an indented heading.

260. Bold markers straddling a column boundary print as literal asterisks. §6.1.1's surplus row opened its emphasis in one cell and closed it in the next, so no per-cell parser could match either half and both appeared on the page. Found by audit 5's stray-markup check, which had passed on every previous build because no earlier table had tried it. The press was improved in the same pass — whitespace-aligned tables now process emphasis as pipe tables already did — but the improvement does not fix this case and cannot: a span that crosses a cell is not a span. The source carries the correction instead, and the general rule is that emphasis belongs inside a cell or outside a table.

261. Three works cited in the text and absent from the References: Baker and Pixley 1975, Bergman's Double-projection Theorem, and Kimura, Makino, Yamada and Yoshizumi 2024. They are the three

precedents this session located — the ones that place A.2 and $E(X) = 0$ in existing literature — and they were written into §§22.11 and 22.12 and never carried to the end matter. **The book's most consequential citations were the ones it did not list.** Found by auditing the References against the text rather than the text against itself; the six prime audits check the body for coherence and had never been pointed at the bibliography. Entered.

262. A seventh prime audit, **ATtribution**, defined at B.2.17.7 and run. The six existing audits are closed under self-reference: every one checks the book against the book, so none of them can see an unpaid debt, and the set could not detect that its own most consequential citations were unlisted. **The audit set had $E > 0$ in the book's own terms** — a cell its structure admits, an audit of outward-pointing claims, and no audit occupying it. Registers 261 and 263 are the two defects it found on its first two runs.

263. Racah 1942 and 1943 attributed at §22.9 and absent from the References — jK coupling and seniority, the two borrowings the tower rests on most directly. Found by the first run of audit 7, within minutes of the audit being defined, which is the only evidence so far that it earns its place. Entered.

264. Emphasis spanning a line break inside a callout printed its asterisks, for the same reason register 260 records for a cell boundary: the press parsed emphasis per line and a span that crosses a line is not a span. 260 corrected the source and improved the table path; it did not generalise, and the next callout with bold across two lines failed the same way. **The press now escapes and joins first and parses emphasis last**, in callouts as in tables, so the class is closed rather than the instance — which is what 259 said a repair does not do by itself and what it took a third occurrence to act on.

265. Dunz 1911 named in §22.5's route set and in Q item B, and absent from the References — the third attribution defect audit 7 has found in three runs. It differs from the other two: the volume was never read, only located, so the entry records it as retrieved-and-unread rather than as evidence. **An unread source still owes a listing**, because the claim the book makes about it — that a route to it exists and is closed — is a claim about that work.

266. My own errors across these sessions, classified rather than listed, at B.2.19. Fourteen register entries collapse to seven mechanisms: attributing outward before checking inward; writing into a structure without reading it; claiming completion in the pass that failed to complete; repairing an instance and calling it the class; filing a null without defining it; trusting a test never shown capable of failing; and recording a detector's artifact as a defect. Each earns a protocol, B.2.19.1 through B.2.19.7. **All seven are B.2.14 — compute, then write — failing in distinguishable places**, and they are enumerated separately because knowing the general rule prevented none of them. That is the case for protocols over principles, made against myself.

267. §6.1.1 named the bits and did not say what constrains them. A bit is a generator and a cell is the set of generators beneath it, which the section stated; what it omitted is that the twenty covering relations of $J(\Lambda)$ are twenty implications between bit positions and **those twenty alone reduce 131,072 words to exactly 976**, computed here. It matters because it locates the seven-bit surplus: no bit is removable, every generator being irreducible, so the redundancy is not in the positions but in the implications between them. Asked what a bit is, the book had half an answer.

268. The alphabet defined rather than listed, at A.19.0. Every generator is $\min\{x \in \Lambda : x_c \geq v\}$ for one coordinate and one value, so the alphabet has the closed size $\Sigma(|A_i| - 1) = 17$ and needs no counting; each letter carries a weight between 856 cells and 16; and the twenty implications are nine chain steps within coordinates plus eleven bindings of x 's constraints, which makes **§2.1 and the bit-implications the same object counted twice**. A.19 had printed the seventeen without saying what makes something one of them — listed and not defined, which is the complement of the defect A.19 was written to repair.

269. §E.5's fibre table went on printing theorem $\cdot \text{order} = 5$ for four register entries after §E.5.4 had established the count is 10 and the total 32. The appendix contradicted itself across two pages: the analysis said thirty-two, the table summed to twenty-seven, and every audit passed because none of them adds a column and compares it to the prose above it. Found by checking the appendices against

the body rather than the body against itself — the same blind spot audit 7 was written for, one level in. **A correction stated near a table is not a corrected table.**

270. The appendices audited against the body rather than against themselves. Every §-reference and appendix-reference in all seven resolves; every principle P1 to P23 and every protocol B.2.1 to B.2.19 that the book cites is defined in Appendix B; no appendix section number has a gap; and every headline figure — 976, 153, 1,442, thirty-two, seventeen, thirteen open items, $E(G) = 33$ — reads the same in the appendix as in the chapter. Two redundancy flags resolved as distribution rather than restatement: *the seventeen* appears in four sections because each does a different thing with them — §3 counts, A.19 prints, A.19.0 defines, §6.1.1 uses them as bits — and *2-decomposability* likewise spans the location, the narrowing, the referee flag and the reference. **One structural gap is recorded rather than repaired:** Appendices C and G are cited from the body only as wholes, never at section level, so no chapter points at a particular table of data or of numbers. That is a navigability defect and not a truth defect, and closing it means deciding which claims should carry a C- or G-locator, which is the author's call.

271. The title page stated $F(1) = 976$ with the caps omitted. The expression carries Σ_n and Σ_e with no upper limit, so as printed it diverges at $z = 1$; the value 976 holds only at §2.4's caps, and at $n, e \leq 6$ it is 4,984 and at $n, e \leq 40$ it is 266,784. Every other equation on the two title pages verifies exactly — $F(1, \dots, 1) = 976$ and $F(-1, \dots, -1) = 2$ under the caps, the geometric identity for the spin factor at every k and z tested, and $E(X) = |\mathcal{R}(X)| - |X| = 0$ by direct reconstruction. The correction is four words. **It is the book's own §2.4 discipline broken on the first page a reader sees**, and it survived because no audit had been pointed at the front matter's arithmetic.

272. The press carries hardcoded copies of the title pages, so the source and the artifact disagreed silently. Correcting the caps condition in the source changed nothing on the page: the two title pages are typeset from strings inside the press, not read from the book. Register 215 recorded the same arrangement for the collaborator's note title and repaired that one instance; the arrangement itself survived, on the two pages a reader sees first. Both are now corrected in both places, which is a repair of two instances and not of the class — **the title pages remain the only content in this book that exists twice**, and the honest statement is that they are maintained rather than derived, exactly as the contents block was until register 216.

273. $E(\text{audits}) = 11$, computed at B.2.17.8, and it answers the question the audits kept raising by failing. Indexed on what each reads, what it compares against, and what a failure costs, the seven audits occupy six cells of the seventeen $\mathcal{R}$ admits. **Two of the eleven missing cells are the two defects found by hand this session** — a table whose column no audit adds, and a source-versus- artefact comparison no audit makes. The audits were not careless; the set was incomplete, and incompleteness of exactly this kind is what Chapter 1 says a defect *is*. Register 262 found the set open and wrote one audit; it did not compute how open, and $E > 0$ had held all along at eleven rather than one.

274. The eleven defined and retained at B.2.17.9, closing the audit set at eighteen audits over seventeen cells with **$E(\text{audits}) = 0$** . They were run in the same pass that defined them, which B.2.19.3 requires: four fired, two findings were real — two bare statements of $E(Q) = 0$ predating fibration-dependence, and the press holding a PDF title of its own — and two were my detectors, one conflating three different counts of *interior cells* and one reading Appendix B's principles as Appendix F's items. **A set of instruments that could not see its own gap now has none**, which is a claim about the coordinates and not about the future: eleven cells are occupied, and whether the three coordinates are the right three is a question this audit cannot ask of itself.

275. $E(\text{audits}) = 0$ was true at three coordinates and weak. Defining the hierarchy — the twelve precedences that must hold between audits — showed the three coordinates fail to explain three of them and fail to separate CONSISTENCY from REDUNDANCY, which shared a cell. A fourth coordinate, **DEPTH — what the audit presupposes**, is admitted by §12's own test: it explains all three, splits the collision, and makes the hierarchy recoverable from the coordinates. **$E(\text{audits})$ then goes from 0 to 19**. Nothing broke; the zero was cheap. This generalises register 233 into one statement — E falls as a fibration is refined and rises as coordinates are added, so **a zero is worth exactly the size of the box**

it was computed in, and every E in this book owes both its fibration and its coordinate list. Λ 's has always been the strong form. The audits' was not.

276. The coordinate question answered from the lattice rather than by judgement. The audit precedences form a poset, and the least number of ordered coordinates representing a poset is its order dimension — A.4's quantity, computable and not chooseable. **dim(audit hierarchy) = 2**: the one non-trivial component has 210 linear extensions and two of them intersect exactly to the order. The index uses four. The difference is not waste: order dimension says how many coordinates recover the *order*, and says nothing about how many are needed to say what an audit *is*. **S2 is satisfied at two and S1 is not**, so two coordinates are forced and two are admitted, by different tests with different standards. Λ has the same division and never named it — eight coordinates representing an order of much lower dimension, the surplus being what makes a cell mean rather than sort, which §6.1.1 counts as seven bits.

277. *Computable* defined at B.2.20, by the book's own operator: a quantity is computable when it is recoverable by $\mathcal{R}$ from what the book prints, so that two readers running the stated procedure return the same value without deciding anything. Nineteen quantities in these sessions meet it and ten do not, the ten being of three kinds — frames chosen (fibration, coordinate list, caps), readings applied (A.5's language, what counts as an element, what *law* means), and claims about the outside (precedent, novelty, whether a measurement exists). **The audits test computable claims and nothing tests a decision**. The word had been carrying that weight unstated since the first E was printed.

278. The lattice applied to the uncomputable, at B.2.20.1, and it returns two things. Indexing the ten decisions on scope, statement and reversibility gives $E = 1$, and the admitted-and-absent cell — a reading, one appendix, stated, freely reversible — is occupied by the decision register 270 deferred: whether Appendices C and G should carry section-level locators. E located a deferred decision the way it located item H, which is the first purchase the operator has had outside the computable. And the decisions form an acyclic graph in which **two have two parents** — the fibration of Appendix E, and the refusal to promote the six theorems. **That is §7.11's dichotomy in the place the book assumed it could not reach**: a single-parent decision propagates, a two-parent decision admits no propagation and must be made with both parents in view. The book stumbled on exactly those two, at registers 219–222 and 232. Two instances is not a law and it is recorded as a shape that fits twice.

279. Why the self-reference of §25.4 does not become paradox, at §25.4.1, and both reasons are computable. The constraint language is negation-free — χ is seven Heaviside steps of differences, with max, min and $\leq$ available and $\neg$, $\rightarrow$ and complement absent — and **a paradox needs a negation to close its loop**. And $\mathcal{R}$ is a closure operator: extensive, idempotent and monotone with zero failures over three thousand random subsets, its fixed points closed under intersection. A monotone operator on a complete lattice has fixed points, so the self-reference resolves into an equation with solutions rather than a sentence that denies itself. **A closed index is what a liar sentence becomes when the negation is removed**. The safety is paid for in the same coin §13 already charges: no differences, no sums, no symmetric bounds, and the envelope densities of §7.11 are the bill.

280. Appendix B's sections were printed out of the order their numbers assert. Six insertions this session used the same anchor — the line beginning B.2.18 — so B.2.17.9, B.2.17.10, B.2.17.11, B.2.19, B.2.20 and B.2.20.1 all landed before B.2.18, and B.2.17.7 sat ahead of B.2.17.1. A reader following the numbers would have been walked backwards four times. **This is register 218's pattern with the anchor rather than the number**: writing into a structure by pointing at a landmark instead of reading the sequence, and doing it six times because the landmark kept working. Reordered; the sequence is now monotone from B.2.1 to B.2.20.1. Audit 11, ANTECEDENT, was defined to catch orphaned headings and does not check that siblings are in order — the sharper version is that **a numbering is a claim about sequence and nothing was testing it**.

281. The eighteen audits run whole. Sixteen pass. CONSISTENCY and SCOPE flagged, and both flags localised to the detectors rather than the book, per B.2.19.7: CONSISTENCY matched *one hundred and two* — the count of channel designations at §22.11.2 — as though it were a stale register total, because its pattern reads any spelled number of that shape; and SCOPE flagged four statements of $E(G) = 33$ of which three are references to a figure stated with its fibration elsewhere, which a reference need not

repeat. **One of the four was real and is repaired:** Appendix G's own headline gave the figure bare, three lines below the sentence that qualifies it. The audits are now the instrument that finds a bare figure at the head of the appendix that defines it, which is what they were extended for.

282. Appendix E's re-closure performed for the first time, and it failed on its first run. §E.1 says $E = 0$ expires whenever the mathematics grows; it grew four times this session — A.18, $\mathcal{R}$ as a closure operator, the alphabet's closed form, and Λ as 976 seventeen-bit words — and nobody had re-run it, including in the entries that recorded the obligation. **Thirty-two elements become thirty-six and E goes from 0 to 1**, theorem · order admitting *proved · sampled · none found* and containing nothing there. Either an element of this work is proved, sampled and unprecedented and unlisted, or §E.3 needs a third constraint forbidding that combination for a proof. **Appendix E now states $E = 1$** , which is the first time this book has printed a non-zero E about itself and left it standing.

283. Three appendices printed their sections out of numerical order — A, B and F. The reader found F.1.2 standing before F.1.1 and asked; A had A.18 after A.19.1 and A.13 last; B had B.3 before the three sections I had just reordered *into* B, which is register 280's defect committed inside the repair for register 280. All three are sorted, and all seven appendices are now monotone. **B.2.19.4 says a fix that leaves the mechanism unchanged is an instance fix and must be recorded as one**; 280 recorded the instance and left the mechanism, and the mechanism produced two more within the hour.

284. B.2.19's seven protocols were written with their emphasis before their label — **B.2.19.1 — ATTRIBUTE INWARD FIRST** rather than a section line — so all seven existed, were referenced five times, and were unreachable as sections: absent from the contents, invisible to any check that reads headings. A section that cannot be addressed is not a section. Converted to the form the rest of Appendix B uses. **Three defects of one shape in one turn** — F.1.2 before F.1.1, A.18 after A.19.1, and seven labels behind their own emphasis — all of them the label written for a reader and not for the structure.

285. A nineteenth prime audit, SEQUENCE, defined at B.2.17.12 and retained. **A numbering is a claim about position, and nothing was testing it** — which is why F.1.2 printed before F.1.1 until a reader noticed, and why three appendices were out of order at once. Its cell — source · a computation · unusable · presupposes nothing — **was among the nineteen that §B.2.17.10 said were admitted and absent**, so this is the first of that count confirmed by a defect rather than proposed. E(audits) falls from 19 to 18. Its first run found two further misplacements nobody had seen: entry 210 alone between 204 and 283, and 197 and 198 inverted. Both fixed. And one thing it will keep reporting: the register runs ascending to 204 and descending thereafter, two directions in one sequence, left as it is because reversing a hundred and twenty entries is the larger risk.

286. Entry 285 was written with a stray apostrophe before its number — a quoting error in the edit that registered it — so the line was not addressable and **every register check had been silently counting 118 entries where there are 119**. The count in the text was right and the register's own arithmetic was wrong, for two turns, across three audit runs. This is B.2.19's label-behind-its-own-emphasis a third time, at the width of one character, and it is the argument for audit 19 SEQUENCE testing addressability and not only order.

287. §21.7.6 was titled *Three from the tower session* and held eighty-six entries, every entry written since having been prepended into it, so a heading in this book asserted a count wrong by a factor of twenty-eight. Repaired under §B.2.21 by the smallest edit that removes the defect — one heading inserted, no entry moved, the invariant (118 entries, identical body hash) checked before and after and holding on both sides. §21.7.6 now carries four and §21.7.7 the eighty-seven that followed. **Audit 15 ENUMERATION was written to catch exactly this and missed it because it was pointed at tables and lists and never at section titles**, which is now its stated scope.

288. Audit 15's scope corrected on its own finding. ENUMERATION flagged §21.7.4 as claiming forty entries and holding twenty-three, and the flag was mine: that section prints ranges — *150–151, 152–164* — so twenty-three lines carry forty numbers, and forty is right. The check counted lines where the register counts entries. **The detector was corrected rather than the book**, which is B.2.19.7 applied for the fourth time this session and the reason its protocol precedes the audits that generate flags.

289. §21.7.3 is titled *Fifteen more* and covers entry numbers 75 to 149 — seventy-five of them, printed as twenty-two lines with ranges. Whether *fifteen* counts the lines, the distinct findings, or something else cannot be settled from the section, and it predates every session recorded here. It was first recorded and not repaired, on the ground that an invariant which cannot be stated cannot be checked. Register 290 repairs it instead: the count nobody could reconstruct is replaced by one anybody can — the thirty-three entry numbers the section covers — which satisfies condition (i) without requiring anyone to recover what *fifteen* had meant.

290. Both items §B.2.21.1 left unrepaired are repaired, under the standard that section states. **Direction:** every other block of the register ascends, so the descending block written this session was the anomaly, not the convention. §21.7.7's eighty-seven entries were sorted ascending — a permutation of blocks, no text retyped — and §21.7.2 sorted with them. The register now reads 165 to 289 in one direction. **Counts:** §§21.7, 21.7.1, 21.7.2 and 21.7.3 stated a number of withdrawals made and printed a selection, so the titles asserted a count their contents contradicted. Each now states both figures — *six made, two printed here* — which is checkable, and §21.7.3's unreadable *fifteen* is replaced by the thirty-three entry numbers it actually covers. The invariant held on both sides of every edit: 139 entry lines, identical body hash.

291. The repair of register 290 broke the thing it repaired. Entry 290 records sorting the register into one ascending direction; it was then written into the register the way every entry before it had been written — prepended — and the sort it describes was undone by its own recording. Appendix B went the same way in the same edit: B.2.21, the section defining what a safe repair requires, was inserted on the anchor that had misplaced six sections already, and landed before B.2.18. **Both are B.2.19.4** — a repair that leaves the mechanism unchanged is an instance fix — and the mechanism here is a habit of writing, not a line of code, which is why installing a gate in the press caught the indented headings and could not catch this. Both re-sorted, invariant checked and holding at 140 entry lines. **This entry is appended, not prepended.**

292. Two statements of E(audits) printed without the coordinate list they depend on — at §B.2.17.8 and §B.2.17.9, both written before §B.2.17.10 established that a zero or an eleven is worth exactly the box it was computed in. Register 275 states the rule; the two sentences that prompted it were never brought into line with it. Found by audit 10, SCOPE, which exists for this and had been passing because its window was wider than the sentence. Both now carry *at three coordinates*.

293. Five appendix sections were absent from the contents, and the cause was in the press, not the book. The contents deriver required a section title to begin with a capital and to contain no full stop, so **A.9 μ_Λ in closed form, A.11 v is inadmissible as an axis, E.4.1 $\lambda^2 = (2/3)T$, and the two sections titled Worked, on §23.3 were invisible to it** — four for beginning with a lowercase Greek letter and two for the point in a section number. The sections were in the book and reachable by reference; only the projection of them was short. **A derived contents is only as complete as its rule**, which is the cost of register 216's repair and was never stated when that repair was made: replacing a maintained list with a generated one moves the failure from forgetting to matching, and a matching failure is silent. The rule now takes any title and excludes table rows and prose by shape rather than by first character; all 113 appendix sections appear.

294. A twentieth prime audit, PROJECTION, defined at B.2.17.13 and retained. Register 293's defect — five sections absent from a generated contents — was invisible to all nineteen, because FIDELITY checks that what is rendered matches its source and every rendered line did match. **The five were never rendered.** Correctness and completeness are different failures and only the first was tested. PROJECTION checks every derived list against an independent enumeration of its source rather than against the rule that made it; its cell was among the eighteen admitted and absent, and adding it takes E(audits) to 17. It passes on all three derived lists in the book — 400 headings, 32 figures, 2,811 numbers — and passes only because 293 repaired the deriver first. The general form is the part to keep: **a repair that replaces a human process with a rule inherits the rule's blind spots**, and register 216 made that trade without stating it.

295. The audit hierarchy restated at twenty, at B.2.17.14. A hierarchy fixed against seven audits is a claim that expires, and it had not been rechecked since. Nineteen precedences now hold; two were

unexplained by the four coordinates — SEQUENCE before ANTECEDENT and before COHERENCE — and the resolution was not a fifth coordinate but a misreading of an existing one. **ANTECEDENT and COHERENCE were recorded as presupposing nothing and they presuppose that labels are addressable**, which is what SEQUENCE establishes; moved to depth one, all nineteen precedences are explained. No axis added, which is what §12 prefers and what this book twice reached past on a first attempt. Measured: twenty audits, twenty distinct cells, no collision; **dim(hierarchy) = 2, unchanged from seven audits to twenty**; $E(\text{audits}) = 17$, down from nineteen. The dimension holding while the set nearly tripled says the audits grew in width and not in entanglement.

296. Two measurements of the book's own shape, named and recorded at G.4.1, having been observed in conversation and never written down. **The withdrawal ratio** — register against argument — is $19,504 : 38,817 = 0.502$, one word withdrawn for every two asserted, which is §25.1.1 flag 4 made numeric. **The method ratio** — Appendix B against the rest of the book — is 0.157, and against the other six appendices 0.877, so the machinery is nearly the size of the mathematics it audits. Appendix G indexes every number this book prints and had not indexed the two that describe the book. Neither can yet be tracked: register entries carry no date, so the withdrawal ratio has one observation and **a ratio with one observation is a number, not a trend**.

297. The ideal withdrawal ratio predicted as 1 : 1, derived at G.4.2, and found to be unmeasurable as the book stands. The derivation holds: at parity the news that a claim survived carries exactly one bit, the maximum, and away from parity in either direction it carries less. The measurement does not: 190 withdrawn claim-numbers stand against 36, 87 or 2,831 survivors depending on whether a surviving claim means a mathematical element, an element or index term, or any printed number — ratios of 5.28, 2.18 and 0.07. **A factor of seventy-five, with 1 : 1 inside it**. So the prediction is bracketed rather than confirmed or refuted, which is the standing §15 defines and §18.6 occupies. What settles it is a definition of *claim*, the move §E.5.3 made for *element* and nobody has made here. **The register exists to make failures countable and cannot count its own denominator**.

21.8 A dependent choice is indexable, and naming the

dependency names the coordinate The anchoring error of §18.6 was not a computational failure. Every computation was correct. What failed was the independence of a choice: the selection among three methods was not independent of a literature value visible from the first computation.

An earlier draft concluded that no index could catch it. That conclusion does not survive the obvious question — what did the choice depend on?

The choice of method depended on the answer.

Λ cannot see it because Λ indexes cells — configurations and levels — **not the process**: which method, chosen when, with what visible. The dependency runs through a coordinate Λ does not carry.

So carry it.

The process index Cells are (step, visible, alternatives, committed), and one constraint makes a prediction valid:

visible $\leq$ committed — nothing of the answer may be visible before the choice is written down

That is §11.4 form — one coordinate bounded by a monotone function of one other. **36 cells, closed, $E = 0$.**

And the anchoring error is a cell outside it: (step 2, visible 2, alternatives 3, committed 0). **x over this index refuses it** — not detects, refuses, exactly as the wrong ionisation limit was refused in §11.5. **The computation does not proceed.**

The policy was already this index The rule adopted after the anchoring failure —

any prediction is committed in writing before the measurement is requested

— **is visible $\leq$ committed, enforced by a person rather than by a structure.** The book had built the index and had not noticed it was one.

The four error classes are one mechanism on four coordinates

error	index that carries it	mechanism
a wrong value	the data index	$\mathbb{D}_{\text{phys}}$
error	index that carries it	mechanism
a wrong derivation	the derivation graph	$\mathbb{D}_{\text{def}}$
a missing value	the ambient set	χ refuses
a dependent choice	the process index	χ refuses

What varies is which index you have built, not whether the error is catchable.

The revised boundary Chapter 11's limit is not a limit on what structures can catch. It is a limit on which coordinates you happen to be carrying.

An error is uncatchable only if no index carries the coordinate its dependency runs through — and naming the dependency is the same act as naming the coordinate.

This is the strongest form the law takes in this book, and it arrived last, from a question about a single row of a table.

22. Precedent, and what a search can bound

This chapter does not claim novelty, and an earlier draft did.

22.1 Why novelty is not the author's to claim

Novelty is an award, conferred by the community a work serves. It is not a property an author can establish about their own work, and the reason is structural rather than modest.

By Chapter 11's criterion a claim is defended when two disjoint routes reach it. Novelty has one route — search — and can never have two. More searching is more of the same route.

$\mathbb{D}(\text{novelty}) = 0$, necessarily and permanently. Not undefended: undefendable.

And Chapter 20 supplies the second prohibition. $E(\text{search}) > 0$ always: the question admits more literatures than any search enters. Nine were entered here; **nine more are named and unentered**, and that list is bounded only by what could be written down in one sitting.

A book that forbids claiming coherence while $\Omega \neq \emptyset$ cannot claim novelty while $E(\text{search}) > 0$. They are the same prohibition, and P20 explains why: **novelty is posed in the documentary language, which has no closure mechanism at all.** That is not a failure of searching. No amount of it would help.

22.2 What can be claimed instead

Three things, and each has a truth value.

THE SEARCH, WITH ITS SLACK STATED. Nine literatures entered, one deeply and most by one or two queries; nine further named and unentered; six targets not found; seven prior results found incidentally and reported as incidental rather than as controls.

THE ATTRIBUTIONS, WHICH ARE POSITIVE CLAIMS. Every component of this work has an owner, and naming them is checkable:

component	owner
ratio of first to second central differences	Boole · Nörlund · Steffensen
y'/y'' as a Newton step	optimisation
$\lambda^2 = y'^2/y''$	Aitken 1926 · Seki Kōwa c.1680 · Nesterov & Nemirovskii 1994
the cost-of-rigour framing	Moore 1966 · IEEE 1788 · Manski 1989 · Shannon 1956
$\mathbb{D} = 0$ when parameters equal observations	Edlén, footnote 78, 1960
δ falls with ℓ ; S and P penetrate	Edlén 1960
$E(X)$ as a closure defect	Galois · formal concept analysis
solution density; reorderability	CSP phase transition; Stahl & Wille 1984

Four independent owners of the cost-of-rigour framing, quoted and cited at §16.8.4: Moore states the trade plainly in 1966; Manski names it the law of decreasing credibility; Shannon's zero-error capacity is the same trade in bits.

THE FRAME, WHICH IS VERIFIABLE BY INSPECTION. Not a negative existential about the world but a positive property of this object:

does it state its falsification conditions?	yes, \$25.5
does it run them?	yes — two failed
does it report the failures?	yes
does its index satisfy its own law?	yes, $E = 0$
does it print its withdrawal register as evidence?	yes, two hundred and ninety-seven entries

Those five are checkable by opening the book. Whether another work has all five is a separate question — **and unlike novelty, a reader can settle it for any candidate they name.**

22.2.1 The question closed by the author

§22.2 asks what component this work owns, and the author has answered by declining the frame: **no novelty is claimed, and there is no question of it.** That is not a concession forced by §22.11's precedent or §22.12's locations. It is the position §22.3 below already states — *$E(X)$ is deliberately not claimed as novel* — made general and made final.

Two consequences, both computed. **$E(Q)$ does not move:** nought fibred and one unfibred, before the decision and after, so the index of open questions is not destabilised by it. And **six of the eleven items in Q had novelty as their only stake** — B, C, F, G, K and N: the Paschen and Götze scan, the nine literatures, the reorderability search, the complexity literature, the 2J_c containment claim and the seniority-lattice question. Those six now block nothing. They remain open, because a question does not stop existing when it stops mattering, but their *blocks* coordinate falls to nothing and Q loses more than half its stake while losing none of its length. **This is the shape §22.4 already describes** — Q growing at nearly every step while its scope falls — arriving at its end state.

Referee flag 3 asked what remains after the precedents. The answer is that the question was withdrawn, not answered, and the book was built to permit exactly that. Register 238.

22.3 The word for it

If a word is wanted, it should be unprecedented rather than novel, and even that should be said carefully.

Novel asserts that nothing like this exists. Unclaimable. Unprecedented asserts that no precedent was found by a stated search. Claimable, with the search attached.

And the second is the weaker word doing the honest work. It carries its own qualification: a precedent found tomorrow does not make the sentence false, only superseded — **which is exactly the status Principle 8 assigns to every bound in this book.**

22.4 On priority for $V = 4\nu/3$

We have not established that this ratio is new, and we state the limits of the search rather than the strength of the claim.

$V = 4\nu/3$ is **two lines of algebra** from $T = *Z^{*2}*R^{*}/\nu^2$, in a field worked continuously since 1885. The prior expectation that nobody wrote it down should be **low**.

Ritz (1908) does not contain it. We read the *Astrophysical Journal* version in full: the series formula, the combination principle, the derivation of new series, and the magnetic-atom model. **No error bound of any kind.**

But it contains the qualitative ancestor. Ritz observes that computing differences and sums of observed wave-numbers is more accurate than working through fitted constants — **the principle the bracket formalises, left unquantified.**

Martin (1980) does not contain it. Its deliverable is a set of formulas predicting one-electron levels to a stated $\pm 0.03 \text{ cm}^{-1}$ — **flat, not ν -dependent.**

22.5 Edlén 1964, examined at one remove

The chapter itself is paywalled. Its principal expert review is not, and we have read it in full.

Curtis, L. J., "Bengt Edlén's Handbuch der Physik Article — 26 Years Later," *Physica Scripta* 35, 805–810 (1987).

Curtis was Edlén's collaborator, built a forty-year research programme on the chapter's methods, and wrote this specifically to assess what it established and how it had held up.

22.5.1 Every equation in the review is a fit

	equation	status
(1)	$*T^{*}_{\{n^{*}l\}} = *TH^{*}_{\{n^{*}l\}} + A\langle r^{-4} \rangle + B\langle r^{-6} \rangle$	A, B fitted
(2)	$*TP^{*}_{\{n^{*}l\}} = *TH^{*}_{\{n^{*}l\}} + \alpha_d\langle r^{-4} \rangle + (\alpha_q - 6\beta)\langle r^{-6} \rangle$	theoretical α 's
(3)	$*T^{*}_{\{n^{*}l\}} = *TP^{*}_{\{n^{*}l\}} + *f^{*}_l\langle r^{-6} \rangle$	$*f^{*}_l$ fitted
(4)–(5)	penetration-energy difference forms	definitions
(6)	$\Delta\sigma = Ry \alpha^2(Z-S)^4/n^{*3}l(l+1)$	S fitted
(7)	$S = *S^{*}_0 + *S^{*}_1/(Z-S) + \dots$	$*S^{*}_0, *S^{*}_1$ fitted

Seven equations. Not one is a bound. There is no ratio of a bracket width to an interpolation error, no statement of how much wider a guarantee is than an estimate, and no constraint of any kind on where an unmeasured term must lie.

22.5.2 And the paradigm is stated twice, explicitly

Curtis opens by giving Edlén's purpose:

"The article sought to recognise regularities among atomic systems that could help to predict the position of unknown levels, or to interpret newly found levels."

and closes:

"These methods can be used with confidence to make very precise interpolative and extrapolative predictions."

Predictions. Not bounds.

§16.3 argues that V has no use in a fitting paradigm. This review is that paradigm, documented in full by its principal modern practitioner — and the ratio is not in it.

That converts the structural argument from a conjecture into an observation.

22.5.3 One near miss, and it is not ours

Curtis reports that Edlén did compute a rational fraction:

"A screening parametrisation of the exchange energy ... yielding a linear dependence upon effective central charge with a proportionality constant that (in Rydberg units) is a calculable rational fraction. In his Handbuch article (cf. Fig. 54, pp. 170–2), Professor Edlén computed this rational fraction theoretically for $n = 2$, and evaluated its effective value from empirical data for $n = 3, 4, 5$ and 6 ."

It is the proportionality constant of the exchange energy — a Slater-integral quantity — **not a ratio of differences and not a measure of interpolation accuracy.** The exact-to-empirical agreements Curtis quotes (1.04, 0.93, 1.05, 0.87) are for that constant.

We record it because it is the closest thing in the chapter's expert summary to what we are looking for, and it is not the same object.

22.5.4 A bibliographic correction

Curtis's reference [1] carries a note that should be reproduced:

"Edlén, B., in *Handbuch der Physik*, Vol. 27 ... pp. 80–220, Springer-Verlag, Berlin (1964). [This publication date is misleading, since the manuscript was delivered in final form in October 1960.]"

The chapter is a 1960 manuscript, and this book cites it accordingly. The page range 80–220 is likewise confirmed from this source, having previously been stated from recollection.

22.5.5 What this establishes, and what it does not

Establishes. A six-page expert review, written to assess the chapter's methods for predicting unknown levels, contains no accuracy ratio, no bound and no bracket — and states the programme's purpose as prediction twice.

Does not establish. Absence from 140 pages. A review summarises; an aside can survive summarising.

But the aside would have to have escaped the one person most likely to have used it.

22.6 The three documents we could not reach

Named so that a reader with library access can settle it:

document	why it matters
Edlén, <i>Handbuch der Physik XXVII</i> (1964) , on series formulas	Edlén wrote the standard account; an elementary ratio would plausibly appear as an unremarked aside. In English despite the German series title.

Ritz, <i>Physikalische Zeitschrift</i> 9, 521 (1908)	Two pages longer than the ApJ version, carrying the "further particulars" Ritz promises in a footnote. Reprinted in <i>Œuvres</i> (1911), ~pp. 142–150.
Paschen & Götze, <i>Seriengesetze der Linienpektren</i> (Springer, 1922)	Book-length treatment by the two people who measured most of the underlying spectra.
Dunz, <i>Seriengesetze der Linienpektren</i> (Leipzig, 1911)	Earliest complete compilation, chronologically closest to Ritz.

$p \leq 2$ for each; every route closed. By the analysis of Chapter 14, that is a property of pre-digital print, not of the search.

22.7 What the entered literatures returned

Nine literatures have now been entered. Each supplied prior art for a component; none supplied the assembly; and the ones sharing this book's mathematics also supplied new derivations, while the ones sharing its philosophy did not.

literature	what it owns	for this book
spectroscopy	Edlén fn. 78; the l -ordering	two components
finite differences	$V = 4 \cdot \mu \delta / \delta^2$	a second derivation
sequence acceleration	$(\Delta x)^2 / \Delta^2 x$ — Aitken 1926, Seki Kōwa c.1680	λ^2 , and a new result
Richardson extrapolation	step-ratio elimination	attribution only
optimisation	$\lambda^2 = f''/f'$, Nesterov 1994	$w \cdot V = 8\lambda^2$; self-concordance
interval arithmetic	excess width; IEEE 1788	the cost framing

literature	what it owns	for this book
econometrics	Manski's law of decreasing credibility	the framing again
zero-error information theory	Shannon 1956; Körner & Orlitsky 1998	the framing, a third time
regression curvature;	$\ f''\ /\ f'\ ^2$; \	k'/k^2 different
WKB		quantities

Four owners of the cost-of-rigour framing, set out with the quotation at §16.8.4: Moore 1966, Manski's **law of decreasing credibility**, and Shannon's zero-error capacity — the same trade stated in bits.

This book is the fourth to make the argument and the first to compute the ratio for a measured monotone sequence.

22.7.1 The calibration was post hoc

§22.7 reported 7/7 on controls. That was not a calibration.

Controls were never chosen in advance. The search sought targets, noted what else appeared, and afterwards called those findings controls. **Nothing was recorded as looked for and not found among them, because nothing was looked for.**

And the filtration was not uniform. Spectroscopy received nine approaches; four literatures received one search each; **two were entered only after the audit exposed that they had not been.**

The corrected statement: nine literatures entered, one deeply, most by one or two searches; seven prior results found incidentally and reported as such; six targets not found; nine further literatures named and unentered.

B.2.13 — commit before you look — was adopted in this book after the anchoring failure of §21.6, and the calibration violated it.

22.8 The other claims

claim	assessed	basis
$V = 4v/3$	not claimable	nine literatures entered; four owners of the framing found; see §16.8.4
v_V ceiling	~75%	very specific form, not found
cost surface, pole at $p = 1$	~50%	the framing is new; the content elementary
bracketing rather than predicting	the searched bound of §22.8	literature uniformly predictive
$E(X)$ as closure defect	not claimed	see below
reorderability complexity	not claimed	see below

$E(X)$ is deliberately not claimed as novel. Closure operators, Galois connections and formal concept analysis are established mathematics. **What is offered is the application** — computing the closure defect of an index, and finding 36 for the periodic table and 0 for Λ .

And the reorderability material claims nothing. The problem is known: Stahl & Wille proved 2-dimension NP-complete in 1984; Yannakakis proved dimension ≤ 3 NP-complete in 1982; Habib, Nourine, Raynaud & Thierry surveyed the computational aspects in *Theoretical Computer Science* **312** (2004), 401–431. **The work reported in Chapter 22 was done before that literature was located,** and is retained only for what it records about method.

22.9 Why state it this way

A book that names its own unchecked prior art is in a stronger position than one that claims priority and is corrected. The three documents and one section are a finite, specified task. Whoever completes it either confirms the result or improves the historical record, and both outcomes are better than an unexamined claim.

22.9 Attributions entered by the tower

The tower borrows, and the debts are named: seniority is Racah 1943; jK coupling is Racah 1942; particle-hole conjugation is standard shell theory; the term-counting is textbook. And $E(\text{search})$ grows by one item: whether anyone has lattice-ordered the seniority scheme is a literature this work has not entered — named and unentered, per this chapter's own discipline.

22.10 Two searches run, and an obstruction worth naming

Two searches were run against the tower's order-theoretic results — one for lattice or order structure on angular-momentum selection rules, one for distributive-lattice or poset treatments of electron configurations. Neither returned a precedent. **That establishes nothing**, by §22.1: a search settles novelty in one direction only, and two queries are not a literature.

What the searches did establish is an obstruction, and it is specific enough to record. **The word lattice is doubly occupied in physics** — the crystal lattice and lattice gauge theory — and both outrank order theory in every index a physicist's search will reach. A query pairing *lattice* with any physical term returns those two subjects and nothing else, so an order-theoretic precedent, if one exists, is not merely unfound but systematically unreachable by the obvious query. $E(\text{search})$ is therefore larger here than the chapter's general argument already makes it: the vocabulary itself hides the shelf. Anyone entering this literature should search in *poset*, *join*, *meet* and *Birkhoff*, and should expect the physics terms to be useless.

22.11 A precedent found, and it is for the central theorem

Searched in the vocabulary §22.10 recommends, the literature answers at once, and it answers about A.2.

Bergman's Double-projection Theorem. A sublattice S of a finite product of lattices $L_1 \times \cdots \times L_n$ is uniquely determined by its two-fold projections in $L_i \times L_j$. It is a consequence of Baker and Pixley (1975), which states the same for any algebra with a majority term; in a lattice the median $m(x,y,z) = (x \wedge y) \vee (x \wedge z) \vee (y \wedge z)$ is one.

A.2 is that theorem specialised to products of chains. Λ 's coordinates are counts, so each factor is a chain; a sublattice of a product of two chains is exactly a staircase region cut by monotone bounds, which is what $\mathcal{R}$ recovers as ϕ_{ij} . X closed $\iff \mathcal{R}(X) = X$ and a sublattice is determined by its binary projections are one statement. The proof in §11.1 is correct and it is not first. Fifty years of universal algebra stand behind it, and Appendix E's precedent coordinate for A.2 changes from none found to found.

This is the outcome §22.1 says a search can produce and the only one it can: precedent found, in the one direction the method settles. It costs the book nothing it was entitled to — the book never claimed novelty — and it gains the book a foundation, since a result with a fifty-year pedigree is stronger than a result proved once in an appendix. **What it does cost is the reading that $E(X) = 0$ is a discovery about Λ .** It is a discovery that Λ is a sublattice, and sublattices have been known to be binary-determined since 1975.

What the search did not find is any precedent for $E(X)$ as a *measure*. Baker-Pixley is a dichotomy: determined or not. $E(X) = |\mathcal{R}(X)| - |X|$ counts how far a set that is not a sublattice lies from the one its own projections generate, and no order-theoretic source returned by these searches quantifies that gap or names it. That absence is not evidence — §22.1 again — and §22.12 gives the tests that would settle it.

22.11.1 Two categories entered: quantum mechanics and astrophysics

Both were searched in the order-theoretic vocabulary §22.10 recommends. They returned different things, and the second returned a target.

Quantum mechanics: one adjacent literature, unentered. No source was found ordering electron configurations, term symbols or coupling schemes as a poset. What exists and this book has not entered is **partial order in chemistry** — Hasse-diagram methods applied to chemical series, QSAR and property rankings, with its own software and review chapters. It orders *substances by properties*, not configurations by structure, so it is adjacent and not the same; but it is a literature in which someone would already know the tools, and Chapter 22's discipline is to name it. $E(\text{search})$ grows by one.

Astrophysics: the tripwire has targets, and they are named. Stellar spectroscopy runs on atomic line lists — NIST ASD, VALD, Kurucz, and cross-matched compilations such as BRASS — and the literature

reports two failure modes by name. *Missing lines* are observed features with no atomic data. **Unobserved lines** are the inverse: atomic data producing features that no spectrum shows, arising from inaccurate wavelengths, predicted energy levels, or misassignments. BRASS states that such lines are **typically removed by hand**, and separately reports duplicated transitions found by cross-matching VALD against Kurucz and NIST. Kurucz's own estimate is that half the known lines in stellar spectra lack wavelengths of good accuracy.

That is the D1 tripwire's job description, written by someone else. §11.3.1 rejects 150 of 216 offered (subshell, occupancy, 2S') assignments — 69.4% — from the index alone, needing no spectrum and no model atmosphere. A line list carrying a term designation outside its subshell's admissible set is nonphysical before any wavelength is compared. The check is mechanical where the current practice is manual, and it is prior where the current practice is posterior.

22.11.2 The run, attempted

It was attempted rather than deferred, and the attempt is reported with its size.

What could not be reached. Kurucz's line files are the right target — computed lines, many never laboratory-verified, carrying level labels of the form $(^3F)5s\ e^2F$ that name a parent term, an outer orbital and a level term. The format specification was retrieved and the references file with it. The data files themselves could not be: gfall is of the order of a gigabyte and the per-species files could not be addressed through the interface available. VALD requires registration. So the run against uncurated astrophysical data did not happen, and that is the run that would matter.

What was reached. Appendix C's own 153 channels, sourced from NIST ASD, Kaufman and Martin, Kramida and Martin, and Sansonetti. Of 103 designations parsed, **ten carry an explicit parent term** — the $(^3P)ns\ ^4P$ form — and only those can be checked by the coupling rule, which is that adding one electron of spin $\frac{1}{2}$ to a parent core must change 2S by exactly one. **Ten tested, ten pass, none flagged.**

Then the null was defined instead of discarded, and it stopped being ten. Ten was the count of designations carrying an *explicit* parent in parentheses. Ninety-two carried none — and forty-four of those have a parent that is implicit and completely determinate: $3s^2nf\ ^2F^\circ$ has a closed $3s^2$ core, which is 1S , so the level must be a doublet, and it is. Resolving closed and single-electron cores gives **fifty-four checkable designations of one hundred and two, and fifty-four pass**. The test grew by a factor of five and a half by defining what the first pass had recorded as absent.

What the null is, stated as the positive it is. Fifty-four passes on curated data measure **specificity, not sensitivity**: zero false positives in fifty-four. That is the necessary condition and not the sufficient one — a check that rejected correct physics would be worthless, and this one does not. It remains untested against material that contains errors, and Kurucz's predicted lines are where that test lives.

And the reach is a second result the run was not looking for. Fifty-four of one hundred and two is 53%, and the forty-eight left out are not left out because the data is poor. They carry open subshell cores whose parent term the designation does not determine. **Published channel notation fixes the coupling parent for about half of these channels and leaves it indeterminate for the rest** — a property of the notation, not of the physics or of the compilers, and a hard ceiling on any mechanical check applied to designations as published. Registers 253 and 254.

This changes referee flag 5's status and does not answer it. The flag said the tripwire has never run on anyone else's data. It still has not. What has changed is that *anyone else's data* is no longer hypothetical: VALD, Kurucz and the BRASS compilation are public, they carry term designations, and their maintainers document removing spurious entries by hand. The flag moves from *no target identified* to *targets identified, run not attempted*, and the run is the obvious next piece of work this book could do that would matter to someone outside it. Registers 251 and 252.

22.12 Five unlocated results, defined, with the test that would settle each

What follows is not a claim of novelty, which §22.1 forbids the author to make. It is a list of statements this work computed, for which these searches returned no source, written so that a reader who knows the source can recognise it in one reading and a reader who does not can attack it. Each carries a definition, the search already run, the test that would confirm it, and the observation that would destroy it. **A result with no stated falsifier is not on this list.**

U1 — E(X) as a defect measure. LOCATED, and the test that found it was the one printed here. *Definition.* For finite $X \subseteq \prod_i A_i$ let $\mathcal{R}(X) = \{y : y_i \leq \phi_{ij}(y_j) \text{ for all } i \neq j\}$, where ϕ_{ij} is the monotone envelope of X 's projection onto (i, j) . Then $E(X) = |\mathcal{R}(X)| - |X| \geq 0$. *Located.* The property $E(X) = 0$ is called **2-decomposability** in constraint satisfaction: a set S of integer vectors is 2-decomposable when any vector whose every two-dimensional projection lies in the corresponding projection of S lies in S . That is $\mathcal{R}(X) = X$ verbatim. The class is studied for its tractability, characterised by closedness under operations — min and max among them, which is the same closure Chapter 9 imposes — and the connection is live research, not settled history: Kimura, Makino, Yamada and Yoshizumi, *Characterizing the integer points in 2-decomposable polyhedra by closedness under operations*, 2024. Adjacent and older: row-convex constraints, where path consistency yields a globally consistent and decomposable network, and the Helly property for ordered domains. *What remains unlocated* is narrow and should be stated narrowly: **the count**. 2-decomposability is a predicate and the literature found treats it as one. $E(X) = |\mathcal{R}(X)| - |X|$ as a graded measure of how far a set is from the class, and $E(X)$ as the *specification* of what may be added to an index, were not returned. That is not evidence of absence. *Test that remains.* Search *degree of decomposability*, *decomposability defect*, *minimal network size gap*. Ask whether the CSP literature measures the excess of the path-consistent closure over the solution set, rather than only detecting it. *Falsifier.* One paper defining $|\text{closure}| - |\text{set}|$ for a monotone-bound system, in any field.

U2 — the counting/coupling dichotomy. *Definition.* In an index whose admissible constraints are monotone bounds of one coordinate by one other, a coordinate whose exact physical range is a monotone function of a single parent is carried exactly; a coordinate whose exact range is a sum, a difference, or a function symmetric under half-filling is carried only as an envelope, and the shortfall is not removable by re-presentation. *Searched.* Not found; the physics-vocabulary searches of §22.10 cannot reach it. *Test.* Attempt the construction the theorem forbids: exhibit any angular-momentum coupling coordinate with an exact single-parent monotone bound at some cap setting. Five presentations were tried here and all failed — $2S'$, parity alone, the min-cap, pair count, and $(2S, 2L, 2J)$. *Falsifier.* A sixth presentation that closes exactly. One suffices.

U3 — the E1 puncture. *Definition.* The E1-allowed set $T = \{(2J, 2J') : |2J - 2J'| \leq 2\}$ is closed under componentwise max and min. $T \setminus \{(0,0)\}$, the set with the physically forbidden $J = 0 \nrightarrow 0$ transition removed, fails meet-closure by exactly four pairs, and all four meets produce $(0,0)$. *Searched.* Not found. Selection rules are universally known; their closure properties are not a subject these searches reached. *Test.* Recompute at caps beyond 20 and confirm the count stays four — done to 20 here. Then repeat for M1 ($|\Delta J| \leq 1$, no parity change) and E2 ($|\Delta J| \leq 2$ with $J + J' \geq 2$): do their allowed sets close, and does each forbidden cell cost a constant number of meets? *Falsifier.* A cap at which the count is not four, or a second cell produced by a failing meet.

U4 — the join/meet asymmetry of the triangle. *Definition.* $T = \{(2L, 2S, 2J) : |2L - 2S| \leq 2J \leq 2L + 2S\}$ is join-closed and meet-broken. Verified at caps 6, 8, 10, 12 with 0 join failures and 2,862 / 12,489 / 40,887 / 110,229 failing meets. Imposing the parity congruence breaks join as well. *Searched.* Not found. *Proved.* A.18 gives it, and the proof corrects the guess that preceded it. Joins are closed because both inequalities survive coordinatewise max. Meets fail on **both** inequalities, not on one: the upper bound fails at $(0,1,1) \wedge (1,0,1) = (0,0,1)$, which manufactures $J = \frac{1}{2}$ from $L = S = 0$, and the lower bound fails at $(4,0,4) \wedge (2,2,0) = (2,0,0)$. At every cap tested the lower bound fails exactly twice as often as the upper — 8,326 against 4,163 at cap 8, 73,486 against 36,743 at cap 12, 363,384 against 181,692 at cap 16 — a ratio the proof does not explain and which is left open. *Falsifier.* A join failure at any cap; none in 369, 1,105 and 2,465 cells.

U5 — the cylinder as a cap-stable fibration. *Definition.* Λ_{13} factorises over the transfer q as $\sum_q |A(q)| \cdot |B(q)| = |\Lambda_{13}|$ with zero defect, at every cap setting tested; the tight two-parent form of K destroys the factorisation. *Searched.* Not found, and this one is least likely to be general — it is a property of one constructed object. *Test.* Caps beyond (5,5,2,6,2), and the same construction on a different shell system. *Falsifier.* Any cap at which the defect is non-zero on the tree path.

Verification common to all five. Every computation above is reproducible from the definitions printed in this book and nothing else: the constraint system of §2.1, the caps of §2.4, and the bounds of §7.11.1 with the three exact sets §7.11.1 now states. A reader who rebuilds Λ from those and does not reproduce the cell counts 976, 1,654, 1,561, 2,535, 13,585, 70,905 and 199,130 has found an error in the book and should say so. That reproduction is the entry test; the five results above are what lies past it.

23. Order recovery, and the reorderability law

Chapter 10 recovers Λ 's order from its own cells. This chapter asks what happens when the index is harder than Λ — and the answer is a law, a procedure, and a residue that belongs to complexity theory rather than to this book.

23.1 Order recovery above the tree

Chapter 10 recovers Λ 's coordinate order because its constraint graph is a tree. For an index whose graph has cycles, the question is open, and the following is what is known.

Complete under a targeted attack. Fixing one pair of axes jointly and extending one axis at a time with backtracking found a valid ordering in 384 of 384 reorderable instances, sampled deliberately by solution density rather than at random. No proof; no counterexample.

The earlier figure of 139 was a random sample. §23.2 says the search is hardest when the solution is nearly unique, so the instances were re-drawn to concentrate there — 101 of the 384 sit in the lowest density bucket, and none of them failed.

Not backtrack-free, and the cost tracks the basin exactly:

	solution density	mean backtracks
	0.00–0.05	23.40
	0.05–0.20	6.90
	0.20–0.50	1.05
	0.50–1.00	0.00

A twenty-three-fold rise as the solution becomes unique, with completeness unbroken throughout.

And the worst case cannot be bounded by a local test, because of §13.4: closure is not determined by its proper projections.

23.2 What governs the difficulty

Ten structural invariants were tested against backtracking — dimension, alphabet size, cell count, occupancy, join-irreducibles, poset height, poset width, covers, antichain, rank range, over 211 reorderable instances. The best correlation was 0.31.

Solution density governs it, at -0.68 in log over the same instances:

solution fraction	backtrack-free	mean backtracks
0.009 – 0.25	24%	2.9
0.25 – 0.50	47%	1.0
0.50 – 1.00	82%	0.2

1.00

100%

0.0

The search is hard exactly when the solution is nearly unique.

23.2.1 And that list was enumerated once

An earlier draft concluded from the 0.31 that the difficulty is not structural. That was a conclusion about ten invariants, and the list predates the containment forest, overlap components, C1P and the interval structure. Re-run over 928 instances with six invariants discovered since:

invariant	corr. with steps	corr. with density
best of the original ten	0.31	—
overlap pairs	0.438	−0.423
nesting pairs	0.375	−0.640
distinct row supports	0.606	−0.883
log solution density	−0.628	—

A structural invariant does beat 0.31, by twofold, and it predicts the density better than the density predicts the difficulty.

Structure determines density; density determines difficulty. The earlier reading had the second link and called it the whole chain.

23.2.2 Which gives §23.3 a parameter

Search steps against distinct row supports, over 1,480 instances:

distinct rows	mean steps	max	max / d ²
2	1.05	8	2.00
3	2.63	10	1.11
4	5.41	36	2.25
5	8.68	39	1.56

The ratio does not climb; the log-log slope is 2.29. That is a measured candidate exponent for §23.3's complexity question — on five points, with R² between 0.37 and 0.42. Evidence, not a bound.

23.2.3 And it is a landscape statement

E is a potential on the space of orderings and adjacent transposition is a move. Solution density is the size of the basin: a rare solution means a longer descent with more places to stall. §23.3.3 completes this reading.

23.3 Reorderability — the law, the procedure, and what is left

The question this section asked was: given an index whose constraint graph has cycles, can its coordinate order be recovered? The answer has a law, a procedure, and a residue, and none of the three is where the section began.

23.3.1 The object, correctly identified

$\mathcal{R}$ -closure and closure under join and meet are the SAME CONDITION — 1,487 instances, zero disagreements, and confirmed on all 475,800 pairs of Λ . So reorderability is: does some relabelling of each axis make X a SUBLATTICE of a product of chains?

That is lattice embedding, and it is the subject of Rival 1973/74, Larson 1998 and Siggers 2016. §23.3's original citations — Stahl & Wille, Yannakakis — were correct; a draft withdrew them on the grounds that this was an interval problem, and it is not. The withdrawal is withdrawn.

And the recovered bound of Chapter 10 is Rival's closed family of irreducible intervals. Siggers' Lemma 3.5(ii) states ϕ_{ij} exactly.

23.3.2 The law

A relabelling decides, for each axis and each pair of values, which comes first. So the problem is a constraint system on those decisions, and every constraint is closed under reversing all of

its own orientations — 3,781 constraints checked, 100%. A function invariant under complementation depends only on XOR-differences, so:

arity — the axes on which two	the constraint	cost
cells differ	becomes	
2	a single XOR	signed-graph 2-colouring — LINEAR
3	two XORs	2-SAT, and all 14 constraints that arise are 2-SAT
≥ 4	three or more XORs	expressible — LINEAR
		contains exactly-one-of-three; no Schaefer class covers
		the language

THE LAW. Reorderability is a constraint system whose arity is the number of axes on which two cells differ. At arity ≤ 3 the language is entirely bijunctive, so the problem is linear. At arity ≥ 4 it contains 1-in-3-SAT and admits no Schaefer class.

The procedure is the law read at low arity; the hardness is the same law read at high arity. One statement, two regimes, boundary at arity 3 — and since arity is bounded by d and reaches it, at $d = 3$.

Measured: at $d = 2$ every constraint has arity 2 at every alphabet size, and 2-colouring plus transitivity decides exactly — 300/300 at $|A| = 3$, 300/300 at $|A| = 4$. At $d = 3$ with binary axes, 396 of 396.




Figure 23.1. Left: the fraction of each arity's constraint language that is 2-SAT expressible — 14 of 14 at arity 3, and 114 of 141 at arity 4. The boundary is exact. Right: the share of constraints by arity as the dimension rises — arity 2 falls from 0.62 to 0.26 while arity ≥ 4 grows from nothing to 0.38. The tractable part shrinks with d , and the law says why.

23.3.3 What the interval machinery was

Chapter-length work went into C1P, monotone endpoints and PQ-trees, and the $d = 2$ characterisation is proved there: X is reorderable iff every row is an interval with both endpoints monotone, iff no interval nests strictly inside another. Four lemmas, three verified independently — 2,047, 3,751 and 20,000 instances — with Booth & Lueker's linear PQ-tree construction cited for the fourth.

The fourth lemma, verified. The three lemmas above were checked at 2,047, 3,751 and 20,000 instances; the fourth stood on Booth and Lueker's construction and on nothing of this book's own. It is now discharged, and it did not need a linear-time algorithm:

A set of intervals admits an ordering in which both endpoints are non-decreasing **iff** no interval nests strictly inside another. Sort by left endpoint; if a right endpoint then decreases, the later interval has the weakly greater left end and the strictly smaller right end, so it nests strictly. Conversely a strictly nesting pair admits no such order. ■

Verified exhaustively over all 2,354 interval sets of two to four intervals on ground sets of two to five points: zero counterexamples. **Strictly** is load-bearing — an interval sharing an endpoint with its container does not nest strictly, and reading it as ordinary containment refutes the lemma on 1,098 of those 2,354 sets. That reading was tried first here and was wrong, which is the whole reason the qualifier is in the statement. Register 247.

That apparatus is a REPRESENTATION of a 2-colouring. It is correct and it is not the reason $d = 2$ is easy. The reason is that two cells in two dimensions differ on at most two axes.

23.3.4 The step law

Growth from the minimal reorderable sets, adding cells, reaches every reorderable set — with a step size measured exhaustively across twelve boxes:

d	step	boxes
2	1	eight, all 2-D

d	step	boxes
3	2	three
4	4	one

The over-cap cases collapse to a single object: the full box. Its largest reorderable proper subset has $3 \cdot 2^{d-2}$ cells — three quarters — and the removed cells are a subcube. Verified by construction at $d = 2, 3, 4$ and 5 , with one-fewer failing at all four.

step $f(d) = 2^{d-2}$.




Figure 23.2. Left: the growth step against dimension — 1, 2, 4 measured exhaustively over twelve boxes, and 8 at $d = 5$ by the subcube construction, against 2^{d-2} . Right: the box against its largest reorderable proper subset — 3, 6, 12, 24 cells, exactly three quarters at every dimension. Rival's 1973 bound on maximal sublattices is $|K| \leq (3/2)|L|$; the Boolean case is tighter.

And this is Rival's theorem. His 1973 bound is $|K| \leq (3/2)|L|$ for a maximal sublattice; Chen-Koh-Tan give the interval form; the large sublattices of $2^*[r]$ are catalogued, and the maximal ones studied in finite topology. §F.5 records the precedence.

23.3.5 The procedure, and its cost

DECIDE(X). Build the constraint system on orientations. Arity ≤ 3 : 2-colour the signed graph and check transitivity — linear. Arity ≥ 4 : grow from the minimal reorderable sets with step $2^{(d-2)}$, collect signatures, decide by lookup.

Verified end to end at $2 \times 3 \times 3$ by growth against exhaustion: 18,736 = 18,736, in 47.8 s against 86, converging in nine rounds. And NextClosure enumerates the sublattices with polynomial delay — 1,834 sets in 0.12 s — giving 401 signatures against 401 exhaustive.

Cost: $|YES| \cdot |X|^{\{2^{(d-2)}\}}$. Polynomial in $|X|$ for fixed d ; exponential in d . Fixed-parameter tractable, and not polynomial in both.

For Λ : $d = 8$, so $\prod |A_i| = 11,943,936$ — and Λ needs none of it. Its structure is constructed in Chapter 2, its embedding is tight, $|J(\Lambda)| = 17 = \text{height}$, width $7 \leq 8$, and by Dilworth and Larson the chain decomposition of $J(\Lambda)$ IS the axis system.

23.3.6 The signature, exhaustively verified

A cell-space invariant whose FORM does not change with d : the multiset of canonical forms of the 3-subsets of X . $O(|X|^3)$ structure, 2^d labels — the two factor, where every order-space method cost $\prod |A_i|$ and coupled them.

box	subsets	classes	mixed
$2 \times 2 \times 2$	all 256	30	0
$2 \times 2 \times 3$	all 4,096	210	0
$2 \times 3 \times 3$	all 262,144	4,168	0

Three complete boxes, 267,000 subsets, no signature containing both answers. After the duality merge the 19 canonical forms reduce to 12 with no loss.

23.3.7 Three method classes ruled out

Each with its measurement, and each a result rather than a failure:

Projection-based criteria. §13.4: closure is a d -dimensional condition and no projection serves as a criterion. Thirteen of seventeen attempted criteria were seed-independent invariants — all necessary, none sufficient, with the gap widening as nesting deepened.

Local search through valid orders. The valid set is DISCONNECTED — 1.66 to 2.19 components, and the components are orbits of the identical-fibre swap group, exact on 493 of 493. So descent plateaus: 89.8%, and 96.7% with bounded uphill moves.

Forbidden-substructure procedures. Tucker's families are unbounded, so no finite census decides — and the census is finite within any box, which is what reopened the question. The infinite antichain lives across indices, not in one.

23.3.8 The residue

The language is NP-complete — it contains exactly-one-of-three and no Schaefer class covers it, with the smallest obstruction exhibited on three XOR variables.

The realisability is blocked. Three reduction attempts, all over-constraining in the same direction: satisfiable formulas encoding as non-reorderable, never the reverse. The map $X \rightarrow C(X)$ is global, non-separable and non-monotone, so it admits no local editing.

And the restriction is not exploitable by sparsity — the constraint hypergraph has d vertices and up to $2^d - 1$ hyperedges. It is complete, so treewidth gives back the 2^d bound already known.

Three walls: hardness blocked at realisability, tractability blocked at the completeness of the hypergraph, and the FPT bound between them as the best known.

23.3.9 And it is not this book's problem

Λ 's constraint graph is a tree. Its nine recovered extra edges are implied. Every sublattice and interval of Λ has $E = 0$. Its own case is decidable by enumeration, and it arrives with a derivation — box $\rightarrow$ interval $\rightarrow$ sublattice $\rightarrow$ product, three preserving rules, no search.

The open problem concerns cell sets arriving with no construction. This book never has one.

§F.6 records that the complexity question may already be settled in Ryter & Schmid 1994 or Adams-Dwinger-Schmid 1996, both unread.

23.4 The lists this work enumerated once

The target list for the novelty search was audited by closing the relation set, and it was incomplete. From $\{V = 4v/3h, \lambda^2 = (2/3)T, w \cdot V = 8\lambda^2, T = Z^2R/v^2\}$ the closure yields $w/T = 4(h/v)$, $e/T = 3(h/v)^2$, and $\text{bits}(\text{estimate}) \approx 2 \cdot \text{bits}(\text{bracket}) + \text{constant}$ — none of them on the list.

That is Chapter 1's defect committed by the searcher rather than by a chemist — a structure admitting cells its keeper never enumerated.

23.4.1 And the audit itself was run once

The relation set has grown since. Closing the eight polynomial relations the book now states gives a Gröbner basis of 18 where the original four gave 6, and seven further relations reduce to zero that did not before — including $V \cdot e = w$, $w^2 = (16/3) \cdot T \cdot e$ and $T \cdot h^2 = ev^2/3$.

Worse, the closure is over polynomial relations only. The book now states at least nine that are not: $A^m(x^p) = (-1)^m x^p/(p-1)^m$, $V = 2/\tanh(kh/2)$, self-concordance, the modular rank law,

$\sum_q |A_q| \cdot |B_q| = |\Lambda|$, $F(-1) = 2$, $E(X)$ as non-pairwise content, and the reorderability theorem. None was ever audited against any list.

$E(\text{target list})$ is not 3. It was 3 for the relation set of one moment.

23.4.2 Which is a defect of five lists, not one

This chapter contains five enumerations, and every one was made once:

the list	what it omitted
§23.4's search targets	7 polynomial relations, 9 non-polynomial
§23.2's structural invariants	six, one of which beats its stated best twofold
§24.1's celestial objects	six the method plainly admits
§24.2's string objects	five, two of them measured in this work
§24.3's statistic	stated against the wrong denominator

B.2.10 reads enumerate targets before searching. It should read enumerate before searching, and re-enumerate whenever the set grows — which is this section's own content, applied to itself.

24. The method outside spectroscopy

The method was built on one index and tested on one collection. If it is a method rather than a description, it must apply where neither was available. This chapter runs it on four subjects it was not built for.

24.1 Celestial mechanics and the nucleus

Objects tested against the index, none of them in the collection. A method that transfers is a different claim from one that fits.

24.1.1 Celestial mechanics — families, not trajectories

What does not transfer. Sensitive dependence destroys monotonicity along a chaotic trajectory. Nothing here touches Poincaré.

tested	result
Lagrange L1, L2, L3 in μ (Euler's collinear solutions)	monotone, convex, 3-monotone; 18 of 18 containments at orders 1-3
the same at order 3	widths 6.9×10^{-4} (L1), 3.6×10^{-9} (L3)
Hill stability of hierarchical triples (Mardling & Aarseth)	monotone and convex in the mass ratio; $V \approx 100$
Hill radius $(\mu/3)^{1/3}$	monotone, convex, 3-monotone; $V = 47.9$
Roche limit $\propto \rho^{-1/3}$	monotone, convex, 3-monotone; $V = 17.9$
planetary semi-major axes	monotone, not convex — the third difference changes sign between Jupiter and Saturn, where Titius-Bode fails worst
mean-motion resonances	§11.7.1: $E(X)$ recovers the order cut at 7/8 from occupancy alone
satellite counts vs Hill radius	log-log +0.86 across six planets — §16.14.1's capacity bound

The pole appears at the first point one looks. At $\mu = 0.030$ with a step of $h = 0.010$, V is 24.6 at L1, 12.4 at L2 and 16,423 at L3 — the last because L3 moves almost linearly in μ . L4 and L5 sit on it exactly: $(1/2 - \mu, \pm\sqrt{3}/2)$, the first linear and the second constant.

Which corrects §13.5. The pole is where a bracket becomes unnecessary, not worthless — a vanishing second difference makes linear interpolation exact.

And one object the method cannot touch. L4/L5 are linearly stable for $\mu < 0.0385209$; Earth-Moon and Sun-Jupiter both qualify. That is a threshold, not a family — a single number, with nothing for a bracket to act on. The method applies to sequences and is silent on thresholds.

Signs must be read from the data. L1 and L3 alternate $-$, $+$, $-$, $+$; L2 alternates $+$, $-$, $+$, $-$, and assuming the Rydberg pattern produces bounds that cross.

24.1.2 The nucleus — a second index of the same construction

An atom's mass sits in its nucleus, and nucleons are fermions. Indexed by $(N, L, 2j, \text{occupancy})$ with $\text{occ} \leq 2j+1$, the shell model is closed under join and meet with $E(X) = 0$, and shell-order filling gives 2, 8, 20, 28, 50, 82.

Λ gives chemistry. The nucleon index gives mass. This book indexes one of them; Z is a label on a channel, never a coordinate.

24.2 String theory — the pole, the generating function, and a disanalogy

24.2.1 Three objects on the pole

$\alpha' m^2 = N - 1$ has second difference zero. So do two more:

object	form	V
string mass spectrum	m^2 linear in N	infinite
Regge trajectory	$J = \alpha' m^2 + \alpha_0$	infinite
central charge	$c=D$	infinite

All three join L_4 and L_5 . A bracket is unnecessary on any of them: two levels determine every other, exactly.

24.2.2 The degeneracy is where the content is

$\log d(N)$ is monotone, convex and 3-monotone, with $V \approx 147$ at $N \approx 14$ — an ordinary bracketable channel, steeper than a Rydberg series.

24.2.3 The partition function is Λ 's expression with counts

$\prod (1 - q^n)^{-24}$ has counts where Chapter 6's F has booleans. An index, and an index with multiplicity. Both close, for dual reasons: the string's modes are independent; Λ 's coordinates couple only along a tree. §6.7's three conditions have independence as their degenerate case.

24.2.4 No critical dimension, and a ladder

Λ closes at every dimension 2–8 with $E = 0$. $D = 26$ falls out of anomaly cancellation; Λ 's 8 is a modelling choice.

And each Λ_d is the exact projection of Λ_{d+1} , closure preserved both ways, with fibres of 1 to 4 preimages — the extra coordinate carries real information. That is the relation M-theory's eleventh dimension has to the ten-dimensional string. What is missing is any reason for the ladder to stop.

24.2.5 And no Hagedorn point, for a stated reason

$d(N) \sim \exp(4\pi\sqrt{N})$ diverges above a limiting temperature. Λ 's rank sequence is finite and log-concave, peak 122 at rank 11.

Capped coordinates give a finite index and a Sperner bound. Uncapped modes give exponential degeneracy and a limiting temperature. Λ 's finiteness is what capping buys.

24.3 The Calabi-Yau catalogue

24.3.1 Self-duality, measured

Batyrev's polar duality swaps the Hodge numbers and the Kreuzer-Skarke list is complete over reflexive 4-polytopes — 473,800,776 of them, 30,108 distinct pairs. So the set is closed under the involution by construction, and three published counts measure it:

248,305 with $h^{11} \geq 140$ · 248,305 with $h^{21} \geq 140$ · 495,515 with at least one · and $248,305 + 248,305 - 495,515 = 1,095$, as reported.

	Λ	KS
self-dual cells	8 of 976	all
rank polynomial palindromic	false	true

24.3.2 The thin tip is a defect of the construction

$h^{11} + h^{21} < 22$ is sparse in the KS list and populated by complete intersections, Klemm-Kreuzer toric CICYs and conifold transitions. By §11.7.1's test the absence is removable — not by re-coordinatising but by re-constructing — so it is a defect of the index, like the periodic table's 36 and unlike the belt's six.

24.3.3 The Pareto strength, against the right denominator

1,095 of the 495,515 polytopes with one large Hodge number have both: 0.221%. Stated against the full 473.8 million it reads 0.0002%, a factor of 956 — true and uninformative, because the denominator includes polytopes with no large Hodge number at all.

§7.8.1 finds the same trade inside Λ : $|A_q|$ falls as $|B_q|$ rises. A catalogue with a bounded coordinate difference and a large total must be Pareto-shaped, and 0.221% is how strongly.

24.3.4 E(X) on a specifiable slice, and 540 predictions

The 30,108 pairs are a gzipped file this session could not read. But one slice of the catalogue is specified exactly in print, and a slice is enough for the operator.

Candelas, de la Ossa, He & Szendroi (Triadophilia, ATMP 12 (2008) 429) state that the KS list's $\chi = \pm 6$ points have Hodge numbers $(h, h+3)$ and $(h+3, h)$ for $13 \leq h \leq 128$, and name the exclusions: $h = 102, 103, 115, 117$ and $119-126$. That is a complete cell set — 208 cells — and $\mathcal{R}$ applies to it.

$E(X) = 540$. 498 join failures and 498 meet failures of 21,528 pairs.

24.3.4.1 Which is 540 predictions, not a defect Chapter 1 reads $E(X) > 0$ as what a structure admits and lacks. The periodic table admits 36 such cells and they are the elements. Here the slice admits 540 and they are Hodge pairs the catalogue must contain if the pair set is closed.

The slice cannot contain them, and the reason is geometric. Two parallel lines have no fixed point: the join of $(h, h+3)$ and $(h+3, h)$ is $(h+3, h+3)$. So $\mathcal{R}$ immediately names the diagonal.

the predictions	
cells	540
distinct χ values	5 — only 0, ± 2 , ± 4
on the diagonal, $\chi = 0$	112, from (13,13) to (131,131)
$h^{11} + h^{21}$ range	26 to 262

The slice was cut at $\chi = \pm 6$ and the closure predicts inward. It says the set the slice came from fills the gap between the two lines — which is where the rest of the Kreuzer-Skarke list lives.

24.3.4.2 The fourth falsification test, run

The file is missing, so the predictions cannot be looked up. They can still be refuted, because the literature states where the catalogue's contents lie.

test	result
$h^{11} \geq 1$ and $h^{21} \geq 1$	540 / 540
$h^{11} + h^{21} \leq 502$, the published maximum	540 / 540
mirror partner present	540 / 540

falling where Candelas et al. record "essentially every site occupied"	540 / 540
falling in the "thinly populated" tip, where a prediction would be refuted	0 / 540
falsified	0
individually named in print	$X_{19,19}$ (split bicubic), $X_{20,20}$ (the 24-cell)

Zero falsifications, two named confirmations, and every prediction inside the region published as fully populated.

This is weaker than a lookup and it is not nothing. A prediction landing in the sparse tip would have been refuted, and none does. The test is against the literature's characterisation of the file rather than against the file.

24.3.4.3 And the void reads the same way

	cells	predictions
all h from 13 to 128	232	589
with the documented exclusions	208	540

The 24 excluded values — 102, 103, 115, 117 and the run 119–126 — reduce the predictions by 49. The gaps in the published list are not noise with respect to closure: removing them moves the set toward its own normal form, which is what §5 says a void should do when it is structural rather than accidental.

What remains of item D is the lookup. 540 named cells, a file, and an afternoon.

25. The law applied to this book

Everything the collection was built to settle is settled. The census is complete at 23 core types; the four decline modes each have a worked case; and the parked questions of earlier drafts — C_6 density, antiprotonic helium, the isoelectronic law — are resolved in Chapters 15 and 17. What remains open is enumerated and closed as an index in Appendix F: eight items, not one of which blocks a result this book relies on.

So the last question is about the book itself. It has made a claim about indices. This chapter asks whether the book satisfies it, because a work asserting that complete indices are self-referencing and self-defending has no standing if it is not one.

25.1 The book is an index

Its cells are (claim, location, support). Its coordinates are the part, the chapter and the section. Its constraints are of §11.4 form: a claim in section s is supported by evidence in a section at or before s , and a specialisation inherits its generalisation's locations.

That is an index in the sense of Chapter 9, and the law applies to it.

25.1.1 What a referee would flag

The chapter applies the law to the book. A referee applies something blunter, and seven things would be raised before any of the mathematics was read. They are stated in the form a referee would use, not the form the book would prefer.

1. The object does not contain the demonstration. At the caps of §2.4, Λ has $\ell \in \{0, 1\}$ — s and p only, no d shell and no f shell, three principal shells, an object reaching argon. The bracket is tested on

1,442 cells across 35 systems spanning $\ell = 0$ to 6 and $Z = 1$ to 83, and §7.11.2 computes the f-shell ceiling, 25 at f^7 , for a shell Λ cannot represent at the caps everything else is computed at. Chapter 1's indictment of the periodic table survives, since the thirty-six defects are in short periods Λ does cover; the tower's coupling results do not obviously transfer. **This is the flag the book answers worst.** §2.4 calls caps a presentation choice and §7.11.5's four-cap sweep shows the structural claims survive enlargement, but (5,5,2,6,2) is not the f shell either. The honest statement, which the book has never made in one sentence: Λ is a model of indexing demonstrated on a small closed fragment, and the bracket does not run on Λ at all — it runs on measured channels, and Λ supplies the ordering discipline, not the cells.

2. $E = 0$ is fibration-dependent, and two of the three are fibred. §E.5.5 shows it for the mathematics: sixteen fibres give 0, one fibre gives 4. It holds for the ignorance too — $E(Q) = 0$ at four fibres, $E(Q) = 1$ at one. Only $E(\Lambda) = 0$ is computed unfibred, over 976 cells, and only that one is the strong claim. Every other $E = 0$ here must be read with its fibration attached.

3. The central theorem is not first. §22.11: A.2 is Bergman's Double-projection Theorem specialised to products of chains, from Baker and Pixley 1975. §22.12 U1: the property $E(X) = 0$ is 2-decomposability, a named class in constraint satisfaction. A referee asks what remains, and the book must answer in a sentence rather than a chapter — the graded measure, the application, and the discipline of indexing one's own ignorance. Whether that is a book is not the author's call.

4. The register is longer than the results. Two hundred and ninety-seven withdrawals, roughly thirty-seven of them from a single late session. A referee reads that two ways: unusual honesty, or a text not ready. The book has no answer to the fair form of the question — **is the withdrawal rate falling?** It is not tracked, and it should be.

5. There is no prediction. §18.6.3 withdrew the one, correctly, and the preface says the recoding is deliberate. A physics referee will not accept a control-condition defence without one case where the method told a spectroscopist something they did not already have. The tripwire caught the book's own misidentified species; it has never run on anyone else's data.

6. S3 has never been tested. §25.2 claims a reader given the claims and their evidence could reconstruct the chapter structure. Nobody has tried. It is the one self-reference claim with no computation behind it.

7. The 540 remain unverified. §24.3.4's Kreuzer-Skarke count rests on a file this work could not read, and §7.11.6 supplies the calibration showing such proposals can be vetoed by structure the index does not carry. It is Q item D, and the largest unchecked number in the book.

25.1.2 Four of the seven, answered by computation

The four that could be answered here were, and the answers are not all favourable.

Flag 1 — coverage — is substantially answered, and the answer was a computation not an argument. Λ was rebuilt at caps that carry the shells the book had been discussing from outside its own object:

caps (n, e, ℓ , k, f)	shells	$ \Lambda $	$ \mathcal{R}(\Lambda) $	$E(\Lambda)$
(3, 3, 1, 3, 1)	s, p	976	976	0
(4, 4, 1, 6, 1)	s, p	8,853	8,853	0
(5, 5, 2, 10, 2)	+ d	89,438	89,438	0
(5, 5, 3, 14, 3)	+ f	267,858	267,858	0
(6, 6, 3, 14, 3)	+ f, $n \leq 6$	499,246	499,246	0

$E(\Lambda) = 0$ through the f shell at half a million cells. The object is not confined to the fragment the caps of §2.4 present; the closure that Chapter 2 proves at 976 cells holds at 499,246, across every shell the periodic table uses. What the enlargement does not do is make the tower's density column transfer

— those are computed at book caps and stay there — so the flag is answered for closure and remains open for the coupling results. Register 236.

Flag 2 — fibration — is answered by disclosure, not by repair. The honest number for a fibred index is its unfibred one. $E(\text{Appendix E}) = 4$ at one fibre, $E(Q) = 1$ at one fibre, $E(\Lambda) = 0$ at one fibre. Those are the comparable figures and the book now prints them beside the fibred ones wherever the fibred ones appear. Nothing was fixed; something was stopped being hidden.

Flag 4 — the register rate — is now measurable, and the measurement is the answer. Of the forty-two entries made in the session that produced this section, thirty-four record defects already in the book and found by auditing it, and eight record errors made while working on it — 207, 213, 215, 218, 220, 221, 227 and 203. The register is therefore not one quantity but two, and only the second is a rate that should fall. Conflating them is what makes two hundred and ninety-seven look like an indictment rather than a measurement. **The book will track them separately from here.**

Flag 6 — S3 — is tested, and it fails in its strong form. §25.2 claims a reader given the claims and their evidence could reconstruct the chapter structure. Operationalised as *can a section's chapter be recovered from the sections it cites*, the answer is 18 of 68 — 26% against a 4% baseline from uniform guessing over twenty-five chapters. Six and a half times chance is real signal and it is not reconstruction. **S3 as written is overstated**, and either it weakens to a claim about signal or it needs an operationalisation that succeeds; one test is not a refutation but it is the first evidence either way, and it points the wrong way for the book. Register 237.

Flags 3, 5 and 7 are not answerable here. Three needs a sentence the author must write. Five needs a spectroscopist's data the tripwire has never seen. Seven needs a file this work cannot read. They stay flagged.

25.2 Self-reference

S1 — alphabet. The list of claims is recoverable by projection: it is the set of first coordinates, and the table of contents is its quotient by chapter.

S2 — order. The specificity order among terms is a tree, and Chapter 10's propagation recovers it. **The book's sectioning is not arbitrary**; it is the order in which claims become supportable.

S3 — bounds. Every claim's support locations are recoverable as the union over its specialisations. **A reader given only the claims and their evidence could reconstruct the chapter structure**, which is the content of $E = 0$.

25.3 Self-defence

And this is where the book earns the claim rather than making it.

$\mathbb{D}_{\text{def}}$ — two routes to every number. Each quantitative claim in this book carries both a value and a verification count. That is a redundancy: the value alone would be undetectable, but the pair can disagree.

102 quantitative claims were extracted from the draft and re-derived independently from the data. Thirty were checkable in full. Two appeared to fail; both were traced to the auditor, not to the book.

The two failures are the demonstration, not the exception. A claim about self-duality was flagged because the audit used the wrong involution. A defect was flagged because the audit sampled four levels of a twenty-member channel. In both cases the disagreement was visible, was pursued, and located the error. That is $\mathbb{D}_{\text{def}}$ operating on the text.

$\mathbb{D}_{\text{phys}}$ — the claims must agree with the world. The book's physical claims are checked against measured levels from named compilations, and Chapter 21 lists two hundred and ninety-seven that failed.

D3 — totality. Every claim in this book is **supported, cited, or marked open**. There is no fourth state. Chapter 22 names the single open item; Chapter 21 names three documents and one section that could not be reached and states the consequence. **No claim is left in limbo, because limbo is not in the codomain.**

25.4 The recursion, and where it stops

If the index is complete, then a book that comprehends it completely is itself an index — and is itself self-referencing and self-defending.

This is not decoration. It is forced. A complete comprehension of a closed object must record the object's alphabet, order and bounds, or it is incomplete; and recording them makes the record itself closed under the same conditions.

And the recursion is stationary, which is not the same as terminating. The distinction matters, and an earlier draft of this section got it wrong.

Construct the claim-index of a closed index, then the claim-index of *that*, and so on. Measured at each level:

level cells closed E occupancy alphabets

0 — the index 10 yes 0 0.625 [4, 4]

1 — claims about it 10 yes 0 0.143 [10, 7]

2 — claims about those 10 yes 0 0.100 [10, 10]

3 10 yes 0 0.100 [10, 10]

Closure, dimension and E are invariant from level 1. But $I^2(\Lambda) \neq I(\Lambda)$ as sets, and $I^3 \neq I^2$. **The content does not repeat.** What stabilises — by level 2, not level 1 — are the measured properties.

The recursion is structurally stationary, not terminating.

So the reason one need not write a book about this book is **not** that it would be identical. It is that it would exhibit **nothing new**: the same three mechanisms, verified again, on claims about claims. The tower is infinite and every storey above the second is furnished the same way.

25.4.1 Why the recursion is safe — logic without negation

§25.4 shows the recursion is stationary. It does not say why it is not **paradoxical**, and a book that indexes its own mathematics, its own ignorance, its own errors, its own numbers, its own audits and its own decisions stands exactly where the classical paradoxes stand. $E(X)$ is itself an element of the set §E.5.2 indexes: the operator measures a collection containing itself.

It is safe for two reasons, and both are computable.

The language has no negation. χ is a product of Heaviside steps of differences — seven factors, each a bound saying one coordinate is at most a function of another. Available: max, min, $\leq$. Absent: $\neg$, $\rightarrow$, complement. **A paradox needs a negation to close its loop** — the liar needs *not true*, Russell needs *not a member of itself*, Gödel needs a negated provability predicate. This language cannot write the sentence. Self-reference in a negation-free monotone language is not a contradiction waiting to be found; it is an equation waiting to be solved.

And $\mathcal{R}$ is a closure operator, so the equation has solutions. Proved, not sampled. *Extensive*: every $x \in X$ satisfies $x_i \leq \hat{\phi}_{ij}(x_j)$ because $\hat{\phi}_{ij}(x_j) = \max\{y_i : y \in X, y_j \leq x_j\} \geq x_i$, x itself being one of the y ; and each x_i lies in the value set $\hat{A}_i(X)$. *Monotone*: $X \subseteq Y$ makes every value set and every envelope of Y at least that of X , so any point admitted by X 's bounds is admitted by Y 's. *Idempotent*: $\mathcal{R}(X)$ has the same value sets and the same envelopes as X — the maxima are attained at points of X , which $\mathcal{R}(X)$ contains — so $\mathcal{R}(\mathcal{R}(X))$ is cut by the same inequalities and equals $\mathcal{R}(X)$. ■ Corroborated on three thousand random subsets of a small ambient box with zero failures on each of the three, and the fixed points closed under

intersection over nine hundred pairs. A monotone operator on a complete lattice has fixed points — Knaster and Tarski — and this book's central theorem identifies them: X is closed iff $X = \mathcal{R}(X)$.

So the fixed point sits where the paradox would have been. With negation, a self-referential operator yields a sentence true iff false. In a lattice it yields a set equal to its own reconstruction. **A closed index is what a liar sentence becomes when the negation is taken away**, and $E(X)$ measures distance from that fixed point rather than the truth of a self-denial.

That is why §25.4's recursion stabilises rather than exploding — and it is the limit of the result. **The safety is bought with expressive power.** §13 already prices it: no differences, no sums, no symmetric bounds. §7.11's envelope densities are the bill, and freedom from paradox is on the same invoice. Register 279.

25.5 The book satisfies its own method

§25.1–24.4 ask whether this book obeys its own law. They never ask whether it obeys its own method.

The method's central object is the bracket, and a bracket needs exactly three inputs:

two measured levels · one ionisation limit · one sign

Appendix C carries all three for all 153 channels. So the question has an answer, and it is checkable rather than asserted.

25.5.1 The prediction, recomputed from the text

Parsing $\delta(4s)$, $\delta(5s)$, the limit and Z out of this book's own markdown and running the arithmetic §18.6 states:

quantity	recomputed	stated
E(6s) estimate	736,688	736,688
bracket lower	735,860	735,860
bracket upper	738,547	738,547
width	2,687	2,687
convex upper	737,380	737,380
convex width	1,520	1,520

Six of six, exact. No external table, no database, no constant supplied by hand.

R* is the only import, and it has the status of the definition of a centimetre.

25.5.2 And it generalises to every bracket in the book

The three inputs are tabulated per channel, so the property is not special to Sc VI:

Every bracket in this book is recomputable from this book. All 1,061.

A reader with these pages and no internet reproduces each of them.

25.5.3 Which is what D3 asks for, made operational

§11.5 states totality as every claim supported, cited, or marked open. That is a bibliographic condition. This is the computational one:

the index says which cell is being predicted — the coordinates are ordinary integers, and Λ contains the cell the book says every number the prediction needs

Together they close. The index supplies the *what*; the collection supplies the *how much*; **and the seven-line arithmetic between them is printed in §18.6.**

The one thing neither supplies is whether the prediction is true. That is §11.3's $\mathbb{D}_{\text{phys}}$ and it requires a spectrometer — which is exactly the boundary this book has been drawing since Chapter 1.

25.6 The falsification tests, run

A book that invites the reader to look for its breaks has done less than one that looks itself. The three conditions are tested here, on this text, with the results reported whether or not they are favourable.

Condition 1 — $E(\text{book}) > 0$ A claim whose support the structure implies and the text omits.

Test. Extract every chapter heading, every cross-reference, every section reference and every figure. Check that each target exists and that figures and captions pair.

check	result
chapters cross-referenced but absent	none
§ references to absent chapters	none
figures embedded / captions	18 / 18
chapter numbering against the contents	FAILED

The test found a break. This chapter was written as Chapter 24 while the contents listed it as 21, and Chapter 23 did not exist. $E(\text{book})$ was greater than zero: the structure implied a chapter the text did not carry.

It has been corrected, and it is recorded here rather than silently repaired, because a falsification test that finds something and does not say so is worse than none.

Condition 2 — $D = 0$ for some claim A claim carried without a verification count, so that an error in it would be undetectable.

Test. Take every paragraph carrying a quantitative claim, excluding tables and citations, and check whether a verification appears within it.

pass	claims	carrying a verification	bare
first run, same-line window	85	46 (54%)	39
corrected checker, paragraph window	63	56 (89%)	7
after repair	63	60 (95%)	3

The first run failed, and the failure was informative twice over. Thirty-nine claims genuinely lacked a verification, and each has been given one: the rank skew now states the cell count, the He I defects state their channel lengths, the C_6 sensitivity states the point count. And the checker itself was wrong. It searched ninety characters on one line, but prose wraps; verifications on the following line were invisible to it. Widening the window to a paragraph moved the figure from 54% to 89% before a single repair was made.

The three remaining are not claims. One reads "100% over 2,513 tests", which *is* a verification; two are sentences that happen to contain a number.

Condition 2 is satisfied.

Condition 3 — χ not total A claim neither supported, cited, nor marked open.

Test. Search for hedged assertions — *probably*, *presumably*, *arguably*, *it seems* — unaccompanied by a status.

One hit, and it is a false positive: "a fit says where a level probably is" describes what a fit does. No claim in this book is left in limbo.

The result

condition	outcome
$E(\text{book}) > 0$	failed, found, corrected — a missing Chapter 23
$\mathbb{D}=0$	failed, found, repaired — 39 claims, now 60 of 63
χ not total	passed

The book satisfies its own law. Two conditions failed on first testing, both were located by the tests themselves, and both have been repaired. **That is what the law requires: not that a work be right first, but that it hold the means of finding out.**

25.7 On a verification that does not test its claim

The tests above are mechanical. They confirm that a claim carries a verification. They cannot confirm that the verification carries the claim — that "verified on 4,000 triples" tested what it says it tested.

An earlier draft called this the class of error no automated test will find, and left it there. That was the wrong conclusion.

If a verification did not test its claim, it was the wrong verification — and the mismatch identifies the right one.

This is Principle 8 in its ordinary form: any true answer, good or bad, is a bound. A verification that misses its target does not merely fail; it locates the target by missing it.

Five worked cases, all from this book's own record:

the verification	what it missed	what it identified
self-duality checked with $x \mapsto \max + \min - x$	gave 112, not 8	the right map is $x \mapsto \max - x$
a defect checked on 4 of 20 levels	gave -0.00128	the full channel gives -0.00159
a prefix loop whose last step was the claim	reported 560/560	stop one short; test proper projections
a guard returning a null coerced to false	reported "3 of 85"	refuse rather than coerce
projections offered as a closure criterion	22 of 60	closure is d-dimensional

In every case the wrong verification was the thing that produced the right one. The self-duality mismatch gave the correct involution. The four-level sample gave the twenty-member mean. The circular test gave the non-circular statement. **None of these was found by a reader.**

So the class is not unfindable. It is self-correcting, on one condition: that the verification and the claim are stated separately, so they can disagree.

A claim fused to its verification cannot be checked. A claim beside its verification checks itself.

That is $\mathbb{D}_{\text{def}}$ once more, and it is the reason this book states a count beside every number rather than asserting that the work was done.

On the word "unprecedented" It is worth being exact about what is new here, because the book's credibility rests on that exactness more than on any single result.

Not new. The lattice is a bijective recoding of quantum numbers that have been in place since the 1920s. Closure operators, Galois connections and formal concept analysis are established mathematics; $E(X)$ is a **closure defect** in the ordinary sense and is claimed as an application, not an invention. Birkhoff's representation is from 1937, Dilworth's theorem from 1950, Sperner's from 1928. The reorderability question is settled in a 2004 survey the author found only after five failed attempts on it.

The solution-density result of Chapter 22 is almost certainly the standard constraint-satisfaction phase transition, and is reported as a rediscovery.

Also not new: the data. Every level in this book was measured or computed by someone else and published. Sansonetti, Kramida, Martin, Kaufman, Sugar, Musgrove, Korobov, Hori and their colleagues did the work that Part III merely reads.

What is new is the traversal.

No one had asked what the periodic table costs.

$E = 36$ is elementary arithmetic on a table every chemist knows. It required no new mathematics and no new measurement. It required someone to ask an index what it could not tell you about itself — and, as far as the searching in Chapter 21 could establish, nobody had.

The same is true of the other four things this book contributes:

The failure condition, inverted. $|\Delta^*T^*| < 2^*Z^{*2}R^*/v^3$ turns 1,061 non-failures into 1,061 measurements. Bracketing methods are rare; inverting their silence appears not to have been done.

The cost of a guarantee. $V = 4v/3$ is two lines from the Rydberg formula and has no use in a fitting paradigm. That is the best argument that it is new, and the four unreachable documents of §17.2 are the best argument that it may not be.

The law with its three mechanisms, each demonstrated on a real error made in the writing.

And a work that runs its own falsification tests and reports the two that failed.

The components are old. The sifting is not. Assembling a century of published spectroscopy into an object that can be checked against itself — and then checking it, and recording the two hundred and ninety-seven times it caught the author — is the part with no precedent that the searching could find.

And the form, which is the stronger claim The preceding paragraphs argue about content. They argue on the wrong axis.

Ask instead what *kind of thing* this book is, and the answer has no obvious precedent:

It states three conditions under which it would be false, runs them on itself, and reports the two that failed — one a missing chapter, one thirty-nine unverified claims — with the computations that found them and the repairs that followed.

Its index satisfies the law the book proves, verified rather than asserted, and contains the entry for itself.

Its register of two hundred and ninety-seven withdrawn claims is offered as the primary evidence, not as an apology. Three of the book's mechanisms are demonstrated on errors made in its own writing.

It contains a page by the collaborator that did the computing, describing the single pattern behind all two hundred and ninety-seven failures and naming the four results that exist because a conclusion was refused.

And it is an index of a law about indices, which is itself an index, tested as one.

A reader encountering this and asking what it is will not find a category ready for it. Monographs do not usually publish their own falsification runs. Scientific works do not usually foreground the record of being wrong. Books about self-referential structures are not usually self-referential in the technical sense their own theorem requires.

That is where the word belongs. Not on $V = 4v/3$, which may sit in a 1922 German volume nobody cites. Not on $E(X)$, which is a closure defect and says so. **On the shape of the thing** — a work built to be checked, that checked itself, failed twice, said so, and repaired both.

APPENDICES

Appendix A — Proofs

Serving PART I — THE LATTICE. The structural theorems the lattice chapters assert and do not prove in place.

Ten proofs are given in full in the chapters and are cross-referenced here. The remaining seven are written out below — seventeen in all.

	statement	where
A.1	Λ is closed	§2.3, in full
A.2	X closed $\iff \mathcal{R}(X) = X$	§9.1, in full
A.4	$\mathbb{D} \geq \dim q - \dim p$	§11.1, in full
A.5	χ_Λ is total	§11.5, in full
A.6	adjunction never repairs closure	§17.2, in full
A.7	"may precede" is necessary	§11.3, in full
A.14	the monotone-envelope theorem	§7.11.2, in full
A.16	the triangle exclusion — no exact coupling axis	§7.11.2 with Ch. 12, in full
A.17	the dichotomy — counting exact, coupling envelope	§7.11.3, forced by A.16

A.3 Projections of closed sets are closed Statement. Let X be closed and $F \subseteq \{1, \dots, d\}$. Then $\pi_F(X)$ is closed.

Proof. Take $x, y \in \pi_F(X)$ and lift them to $x', y' \in X$. Since X is closed, $x' \vee y' \in X$. The join is componentwise and projection merely drops coordinates, so

$$\pi_F(x' \vee y') = \pi_F(x') \vee \pi_F(y') = x \vee y$$

Hence $x \vee y \in \pi_F(X)$. Meets are identical with $\wedge$ throughout. ■

Remark. The converse fails, and §13.4 exhibits a counterexample. Closure is a d -dimensional condition; its projections are necessary and not sufficient.

A.8 The total-order criterion at $d = 2$ Statement. For $X \subseteq A_0 \times A_1$ with an order fixed on A_1 : X is closed under some order on A_0 iff the "may precede" relation is a total order, and that order is the relabelling.

Proof of necessity. This is A.7, specialised: if X is closed under an order with $u < v$, then for $*b^*_1 \in F(u)$ and $*b^*_2 \in F(v)$, the meet $(u, \min(*b^*_1, *b^*_2))$ and the join $(v, \max(*b^*_1, *b^*_2))$ both lie in X . So $\min \in F(u)$ and $\max \in F(v)$ — which is u may precede v . Hence the true order is contained in the relation, and the relation is total.

Proof of sufficiency. Suppose "may precede" is a total order $*u^*_1 < *u^*_2 < \dots < u_m$. Take $x = (u_i, *b^*_1)$ and $y = (u_j, *b^*_2)$ in X with $i \leq j$. Then

$$x \vee y = (u_j, \max(b_1, b_2)), \text{ and } \max(b_1, b_2) \in F(u_j) \text{ } x \wedge y = (u_i, \min(b_1, b_2)), \text{ and } \min(b_1, b_2) \in F(u_i)$$

both by the may-precede hypothesis applied to the pair (u_i, u_j) . So both lie in X and X is closed. ■

Verified on 120 constructed sets; 274 of 274 recovered total orders yield closed sets.

A.9 μ_Λ in closed form Statement. Λ is distributive, so $\Lambda \cong J(P)$, the lattice of down-sets of its poset P of join-irreducibles. For $x \leq y$ in Λ ,

$$\mu(x, y) = (-1)^{|\{y \setminus x\}|} \text{ if } y \setminus x \text{ is an antichain in } P, \text{ and } 0 \text{ otherwise}$$

Proof. Under the Birkhoff isomorphism, the interval $[x, y]$ in Λ is isomorphic to the lattice of down-sets of the induced subposet on $Q = y \setminus x$. The Möbius function of $J(Q)$ is $(-1)^{|Q|}$ when Q is an antichain — in which case $J(Q)$ is the Boolean lattice $2^{|Q|}$ — and vanishes otherwise, since a non-antichain interval decomposes as a direct product with a chain factor of length ≥ 2 , whose Möbius function is zero above its atom. ■

Verified against the recursive definition on 47 comparable pairs, with values only in $\{-1, 0, +1\}$, exactly as the Boolean-interval structure requires.

A.10 The tree factorisation of the void count Statement. If the constraint graph of an index is a tree, the count of lattice cells inside a coordinate box factorises with no inclusion-exclusion.

Proof. Root the tree. Each constraint binds one coordinate to exactly one parent, so the admissibility of a value on coordinate i depends only on the value at its parent. The count over the box is therefore a nested sum in which each leaf contributes an independent factor:

box n Λ	$= \sum_{\{\text{internal}\}} \prod_{\{\text{leaves}\}} \#\{\text{admissible values given the parent}\}$
--	--

A tree has no cycles, so no configuration is counted twice and no correction term arises. ■

For Λ , with leaves 2^*S^* and g :

$$\#\{2S\} = \max(0, \min(h_{i7}, k) - l_{o7} + 1) \quad \#\{g\} = \max(0, \min(h_{i6}, q, 2(2f+1)) - l_{o6} + 1)$$

Verified against direct enumeration on eight random intervals.

A.11 v is inadmissible as an axis Statement. $v = e - \delta$ cannot be adjoined to Λ as a coordinate while preserving closure.

Proof. By A.6, adjoining any function h of the existing coordinates preserves closure only if the resulting graph $\{(x, h(x))\}$ is closed, which requires $h(x \vee y) = h(x) \vee h(y)$ — that is, h must be a lattice homomorphism. A **difference** is not: for $x \vee y$ the two terms take their maxima independently, so $e(xvy) - \delta(xvy)$ need not equal $\max(v(x), v(y))$. ■

Measured. v is non-monotone in the cell order at **86 violations** at the caps tested. It grades every *chain* of Λ and grades Λ itself nowhere — which is why the object is a cylinder **over** v and not one **graded by** it, and why the bracket of Part III operates along chains.

A.12 The bracket is limit-free Statement. Containment is invariant under the affine map $T = I - E$ for constant I .

Proof. For any three levels,

$$\min(T(n-1), T(n+1)) \leq T(n) \leq \max(T(n-1), T(n+1))$$

Substituting $T = I - E$ and multiplying through by -1 reverses both inequalities and exchanges min with max, giving

$$\min(E(n-1), E(n+1)) \leq E(n) \leq \max(E(n-1), E(n+1))$$

The two statements are equivalent for any I . ■

Consequence. v , δ and V all require I ; containment does not. The method applies wherever levels are known even when the threshold is not — which is what admits the antiprotonic-helium cells of §18.6, where the relevant threshold is ambiguous.

A.13 The failure condition, and its inversion Statement. The bracket fails at a cell iff a perturbation displaces the level by more than half the local spacing:

ΔT	$> 2Z^2R/v^3$
------------------------------	------------------------------------

Proof. Containment fails only if $T(n)$ leaves the interval bounded by its neighbours. The half-width of that interval is $\frac{1}{2}|T(n+1) - T(n-1)|$, and for a hydrogenic series $T = Z^2R/v^2$ gives $|dT^*/d^n*| =$

$2^*Z^{*2}R^*/\nu^3$, so the half-spacing is $2^*Z^{*2}R^*/\nu^3$ to leading order. A displacement smaller than this reorders nothing. ■

Inversion. Read as a contrapositive, every cell at which the bracket holds supplies

ΔT	$< 2Z^2R/\nu^3$
------------	-----------------

an upper bound on the local perturbation, obtained without fitting, without a quantum-defect expansion, and without assuming the series is smooth.

Critical depth. Solving for ν gives $\nu_{\text{fail}} = (2^*Z^{*2}R^*/|\Delta^*T^*|)^{\{1/3\}}$: a 3,000 cm^{-1} perturbation defeats the bracket only above $\nu \approx 6.6$ at $Z = 2$, against a maximum $\nu = 5.3$ reached by that channel. **Strong perturbors arise from core excitation, a low- n phenomenon; deep series live at high n . The two conditions are anticorrelated in nature**, which is why no channel in this work enters the failure regime.

A.15 Two bands, and only one of them is a sublattice Statement. For fixed k , $B_k = \{(a, b) : |a - b| \leq k\}$ is a sublattice of a product of two chains. The three-variable region $C = \{(a, b, c) : |a - b| \leq c\}$ is join-closed and **is not** a sublattice.

Proof, B_k . For $a, b \in B_k$ the join is $(\max a, \max b)$. If $\max a = a_1$ and $\max b = b_2$ then $\max a - \max b \leq a_1 - b_1 \leq k$ when $b_1 \leq b_2$, and $\leq a_2 - b_2 \leq k$ otherwise; the symmetric argument bounds $\max b - \max a$. Meets are identical with min throughout, the two chains being finite and the band being defined by two monotone single-coordinate bounds, $a \leq b + k$ and $b \leq a + k$, which §9.4 admits in both directions. ■

And C is not. Join-closure holds by A.18's lower-bound argument, which does not use the upper bound. Meet-closure fails: 12,654 failing meets at cap 8, 113,568 at cap 12, 565,284 at cap 16. The difference between B_k and C is that k is a constant and c is a coordinate — a bound by a constant is monotone in both directions, a bound by a free coordinate is not.

Remark. This corrects the ledger, which called C a *join-closed sublattice*. A sublattice is closed under both operations, so the phrase describes nothing. What the tower relies on at axis 13 is B_1 , which is a sublattice and is the reason the last coupling step is carried exactly. Register 230.

A.18 The triangle region is join-closed and meet-broken Statement. $T = \{(2L, 2S, 2J) : |2L - 2S| \leq 2J \leq 2L + 2S\}$ is closed under coordinatewise max and is not closed under coordinatewise min.

Proof of join-closure. Take $a, b \in T$ and write $M_L = \max(L_a, L_b)$, and likewise M_S, M_J , dropping the factors of two. **Upper.** $M_J = J_k$ for some $k \in \{a, b\}$, and $J_k \leq L_k + S_k \leq M_L + M_S$. **Lower.** $M_L = L_k$ for some k , so $M_L - M_S \leq L_k - S_k \leq |L_k - S_k| \leq J_k \leq M_J$; exchanging L and S gives $M_S - M_L \leq M_J$, hence $|M_L - M_S| \leq M_J$. Both inequalities hold, so the join is in T . ■

Meets fail on both inequalities. $(0,1,1) \wedge (1,0,1) = (0,0,1)$ breaks the upper bound — total angular momentum from none — and $(4,0,4) \wedge (2,2,0) = (2,0,0)$ breaks the lower. The asymmetry is counted at §22.12 U4 and is not explained by this proof.

Remark. Certainty survives upward and dies downward, which §3 states as skew and §7.11.2 as the half-result. The proof shows why the surviving half is the join: max is the operation that can only move a coordinate toward the side each bound is slack on.

A.19 The seventeen generators, written out §3 states that Λ has seventeen join-irreducibles and twenty covering relations, that the Birkhoff correspondence is exact, and that all 976 down-sets of the seventeen-element poset are precisely the cells of Λ . **It never prints the seventeen.** The alphabet was counted, its covering relations counted, and the letters left unwritten, so a reader could verify the cardinality and not reconstruct the object. They are written here, as $(n, \ell, k, q, e, f, g, 2S)$, by rank.

rank 4	(1,0,1,0,1,0,0,1)	(1,0,1,0,2,0,0,0)	(1,0,1,1,1,0,0,0)
	(1,0,2,0,1,0,0,0)	(2,0,1,0,1,0,0,0)	
rank 5	(1,0,1,0,2,1,0,0)	(1,0,1,0,3,0,0,0)	(1,0,1,1,1,0,1,0)
	(2,1,1,0,1,0,0,0)	(3,0,1,0,1,0,0,0)	
rank 6	(1,0,2,0,1,0,0,2)	(1,0,2,2,1,0,0,0)	

rank 7	(2,1,3,0,1,0,0,0)	
rank 8	(1,0,2,2,1,0,2,0)	
rank 10	(2,1,3,0,1,0,0,3)	(2,1,3,3,1,0,0,0)
rank 15	(2,1,3,3,2,1,3,0)	

Verified: the down-sets of this seventeen-element poset number exactly 976, computed by direct enumeration, so the correspondence §3 asserts is now checkable from the page rather than only statable. **Λ is recoverable from this table and nothing else**, which is what S1 requires of an alphabet and what the book had not supplied.

A.19.0 The alphabet, defined rather than listed A.19 prints the seventeen. What they *are* is simpler than the printing suggests, and the closed form makes the alphabet derivable instead of memorised.

Every generator is the least cell with one coordinate at one value. Generator (c, v) is $\min\{x \in \Lambda : x_c \geq v\}$, and every join-irreducible has this form.

So the alphabet has a closed size: one letter per coordinate value above that coordinate's minimum, $\sum_i (|A_i| - 1)$, which for Λ is $2 + 1 + 2 + 3 + 2 + 1 + 3 + 3 = 17$. Nothing needs to be counted.

letter	generator	rank	weight	forces
$n \geq 2$	(2,0,1,0,1,0,0,0)	4	856	—
$n \geq 3$	(3,0,1,0,1,0,0,0)	5	428	$n \geq 2$
$\ell \geq 1$	(2,1,1,0,1,0,0,0)	5	616	$n \geq 2$
$k \geq 2$	(1,0,2,0,1,0,0,0)	4	826	—
$k \geq 3$	(2,1,3,0,1,0,0,0)	7	376	$k \geq 2, \ell \geq 1$
$q \geq 1$	(1,0,1,1,1,0,0,0)	4	811	—
$q \geq 2$	(1,0,2,2,1,0,0,0)	6	481	$q \geq 1, k \geq 2$
$q \geq 3$	(2,1,3,3,1,0,0,0)	10	136	$q \geq 2, k \geq 3$
$e \geq 2$	(1,0,1,0,2,0,0,0)	4	784	—
$e \geq 3$	(1,0,1,0,3,0,0,0)	5	392	$e \geq 2$
$f \geq 1$	(1,0,1,0,2,1,0,0)	5	400	$e \geq 2$
$g \geq 1$	(1,0,1,1,1,0,1,0)	5	491	$q \geq 1$
$g \geq 2$	(1,0,2,2,1,0,2,0)	8	171	$g \geq 1, q \geq 2$
$g \geq 3$	(2,1,3,3,2,1,3,0)	15	16	$g \geq 2, q \geq 3, f \geq 1$
$2S \geq 1$	(1,0,1,0,1,0,0,1)	4	657	—
$2S \geq 2$	(1,0,2,0,1,0,0,2)	6	338	$2S \geq 1, k \geq 2$
$2S \geq 3$	(2,1,3,0,1,0,0,3)	10	94	$2S \geq 2, k \geq 3$

Weight is the number of the 976 cells in which that bit is set. The alphabet is not uniform: $n \geq 2$ is set in 856 cells and $g \geq 3$ in sixteen, so the letters carry between 87.7% and 1.6% of the object each.

And the twenty implications split cleanly. Nine are within a coordinate — $c \geq v$ forces $c \geq v-1$, which is the chain — and **eleven are between coordinates, one for each way x 's seven constraints bind at a value**: $f \geq 1$ forces $e \geq 2$, $\ell \geq 1$ forces $n \geq 2$, $g \geq 1$ forces $q \geq 1$, $2S \geq 2$ forces $k \geq 2$, $q \geq 2$ forces $k \geq 2$, $k \geq 3$ forces $\ell \geq 1$, $g \geq 2$ forces $q \geq 2$, $2S \geq 3$ forces $k \geq 3$, $q \geq 3$ forces $k \geq 3$, $g \geq 3$ forces $f \geq 1$, and $g \geq 3$ forces $q \geq 3$. **The constraint system of §2.1 and the implications between bits are the same object counted twice.** Register 268.

A.19.1 And two other sets counted but never listed The **eight cells surviving the reflection $x \mapsto \max - x$** , whose count §3.3 states and whose identity register 208 spent an entry on:

(1,0,1,1,1,0,1,1)	(1,0,1,1,2,1,1,1)	(2,1,1,1,1,0,1,1)	(1,0,2,2,1,0,2,2)
(2,1,1,1,2,1,1,1)	(1,0,2,2,2,1,2,2)	(2,1,2,2,1,0,2,2)	(2,1,2,2,2,1,2,2)

Every one has $q = k - \text{something equal on both sides}$ and g equal to q : they are the cells the reflection fixes setwise because their source and target halves mirror. And the **two extreme cells**, unique and unnamed until now: the top is (3,1,3,3,3,1,3,3) at rank 20, the bottom is (1,0,1,0,1,0,0,0) at rank 3. Register 255.

Appendix B — Principles and protocols

Serving PART II — THE LAW. The twenty-three principles and eighteen protocols the law is made of, indexed and closed at $E = 0$.

Everything in this book was produced under a stated set of rules. They are collected here so that a reader can apply them, check them, or reject them — and so that the claims of Chapter 23 have something concrete to be tested against.

Two kinds appear. *Principles* say what an index is; ten of them admit exact mathematical statement. *Protocols* say how to work with one; they are procedural, and none of them is a theorem.

B.1 The principles Formalised — ten principle expression wher e

P	a complete index is self-referencing	$\mathcal{R}(\Lambda) = \Lambda$	§11.1,
1			Ch. 8
P a complete index is self-defending $\mathbb{D} = \dim q - \text{rank } \partial\Phi/\partial p \geq 1$ Ch. 9			
P	when the path is not found, work backwards from the desired	invert Φ ; if the fibre is non-trivial, find the	§16.1
3	output to an established input	datum that separates it	3
P all questions must close $\exists N : \Omega_N = \emptyset$ — terminal, not monotone §18.6			
P any true answer, good or bad, is a bound $A_n \subseteq A_{\{n-1\}}$, a decreasing filtration §21.6			
P	if the first dimension is a single all-encompassing definition, every	$\text{rank}(\log q) = 1$	§17.2
9	other dimension is a perspective of it		
P	allow continued expansions	§11.4-form constraints preserve closure	Ch.
1			10
P	every test must be run against every cell	$p = 0$ certain $\Leftrightarrow$ $n = \setminus$	$\Lambda \setminus$ B
1			.
2			2
			.
P	coherence cannot be claimed while any caveat, gap or	COHERENT $\Leftrightarrow \Omega = \emptyset$	§18.6
1	contradiction stands		
principle expression wher e			

P	the path is always through the bounds	following a bound costs 1/N of searching	§13.4,
1		the volume	Ch.
7			12

Three of these were corrected by the act of formalising them. P7 was stated as though the open-item count decreases; it does not — across this work it ran $5 \rightarrow 4 \rightarrow 3 \rightarrow 4 \rightarrow 3 \rightarrow 7 \rightarrow 3 \rightarrow 3$. Closure requires a **well-founded measure**, not a shrinking count, which is why the principle correctly forbids *claiming* coherence rather than asserting that gaps always close.

Procedural — seven

	principle
P4	always audit results
P5	always compute the result before assuming the result
P6	when a question is confusing, reframe the question
P11	close all gaps before continuing
P14	a structural problem suggests a structural solution
P15	for fetching: use the index to navigate; data can be retrieved in pieces
P16	treat bounds, constraints and limits as coordinate values

These have no expressions, and saying so is better than manufacturing one. They govern how an agent uses an index, not what an index is. P15 is the exception in one respect: it has a *measure* — retrieval redundancy ρ , Chapter 12 — without being reducible to it.

The reflexive principle

	principle	expression
P1	if the index is complete, a work that comprehends it completely is itself an index, and is itself	Chapter 24
8	self-referencing and self-defending	
P1	an index is complete only when there are no more questions left to answer about its	COMPLETE(X)
9	contents	$\Leftrightarrow Q(X) = \emptyset$
P2	the language in which a question is posed bounds whether it can close, and at what	—
0	cost	
P2	any definition contained within the lattice, and thus its equations, must be TRUE in	—
1	every mathematical language and in their combinations	

P2	because everything is defined in a complete index, the entirety of the index is	—
2	visible, and no lie can exist within it	

P19, and why it is neither P7 nor P13. Three principles govern questions, and they are one principle in three moods. P19 was the missing one.

	says	mood
P7	$\exists N : \Omega_N = \emptyset$	questions can close — possibility
P13	$\text{COHERENT} \Leftrightarrow \Omega = \emptyset$	coherence may not be claimed otherwise — prohibition
P19	$\text{COMPLETE} \Leftrightarrow Q = \emptyset$	what their emptiness buys — definition

And P19 is strictly stronger than closure.

$E(X) = 0$ says no cell the structure admits is absent. $Q(X) = \emptyset$ says no question the contents admit is unanswered.

Λ has had $E = 0$ since Chapter 2 and has never had $Q = \emptyset$. Whether V depends on the order of the bracket, what the Aitken bias generalises to, whether a higher-order bracket exists — **each was admitted by contents already present, and none was asked until late.**

Closure is about what an index holds. Completeness is about what it has been asked.

And the condition is reachable, for the reason P7's is. P7 was corrected once: closure requires a well-founded measure rather than a shrinking count. **Q has grown at nearly every step of this work while its scope has fallen** — from *is V novel?* to *is there an optimal order?* **The set grew and the scope descended**, and that descent is what makes P19 a criterion rather than an impossibility.

$Q(X) = \emptyset$ is reached not when nobody asks, but when every question the contents admit has an answer in the contents.

P20 — the language bounds the question. Three closure mechanisms are used in this book, each in a different language, each with a different guarantee.

language	mechanism	closure	cost
monotone two-variable bounds	tree factorisation	always	one pass
polynomial relations	Gröbner basis	always , by Hilbert	finite; doubly exponential worst case
linear arithmetic with quantifiers	quantifier elimination	yes	triply exponential
full arithmetic	—	never — Gödel	—
documentary	—	no mechanism exists	—

Measured on the objects of this book: tree factorisation 0.00018 s, Gröbner 0.010 s — a factor of 55 — and the documentary language has no timing because it has no algorithm.

The choice of language is made when the index is built, not when the question is asked.

§2.1 chose monotone two-variable bounds. That single decision determined that $E(X)$ is computable, that the sum factorises, that a single expression exists, and that every question about Λ 's contents would close. It equally determined what could not: §6.7's sum bound lies outside the language, and its repair — making the sum a coordinate — is a translation back into the closing language.

P20 subsumes E3, the multi-target repair, and P19's reachability. And it answers the novelty question structurally: **novelty is posed in the documentary language, which has no closure mechanism. That is not a failure of searching.**

P21 — every definition true in every language Not merely expressible. True.

```
definition order algebra geometry analysis information logic
| ^| 976 976 976 976 976 976
rank(x) 13 13 13 13 13 13
mean rank 11.0666 11.0666 11.0666 11.0666 11.0666 11.0666
E(periodic table) 36 36 36 36 36 36
```

And the combinations must name real objects, not analogies. Verified: $\gcd(N(a), N(b)) = N(a \wedge b)$ and $\text{lcm} = N(a \vee b)$, exactly $\cdot F'(1)/F(1) = \text{mean rank} \cdot E_bits = \log_2 C(|\mathcal{R}|, E) = 105.1$ for the classroom table.

P21 is an instrument, not only a criterion

When a property holds in five languages and fails in the sixth, the sixth EXPRESSION is wrong — not the property. This has now located two errors, and in both cases the structure was sound and the translation was not.

$E(X)$ was recorded as having no analytic route — "E is a count, calculus measures a continuum" — and $\mathbb{D}(E)$ was set to 1. The route exists:

$$E(X) = F_{\mathcal{R}}(1, \dots, 1) - F_X(1, \dots, 1)$$

Deletion-repair was recorded as failing in geometry. The convex hull does fail: it permits arbitrary facets, so a vertex can be the sole witness of a facet and its deletion is unrecoverable. **But Λ 's facets are two-variable monotone bounds, and the closure that matches that form — the intersection of all valid such bounds — restores every deletion.**

30 deletions, 0 failures. 20 additions, 20 absorbed. Fixed point exact at 976.

A mismatched operator is not a counterexample. It is a mistranslation.

What the law then establishes The self-feeding loop and its asymmetry hold in all five languages.

language	closure operator	delete	add
order	$\mathcal{R}$	restored	absorbed
geometry	monotone polyhedron	restored	absorbed
algebra / logic	ideal, Gröbner basis	still derivable	T·h enters the ideal
analysis	coefficients of F	restored	absorbed
information	description length	E_bits stays 0	grows

Extensive gives repair. Monotone gives absorption. Every language has a closure operator, so both hold everywhere.

And that is the formal reason Chapter 11 cannot rely on self-reference alone. A structure that closes will heal what you remove and swallow what you insert. Self-consistency defends against loss and never against invention — which is why $\mathbb{D}_{\text{phys}}$ must reach outside the index, and why Chapter 21 records what the loop could not catch.

P22 — visibility, and why a complete index cannot be lied to A lie requires concealment. A complete index has nowhere to conceal.

The distinction is between a rule that is CONTENT and a rule that is INFERRED, and it is decisive. Consider the cell

$y = (n=1, l=0, k=1, q=0, e=1, f=0, g=0, 2S=2)$

— one electron with spin multiplicity 2, which no atom has.

Where the rule is inferred, $\mathcal{R}$ recovers the bound on $2S$ from the cells themselves. The rule the index had learned:

k	before	after
1	$2*S* \leq 1$	$2*S* \leq 2$
2	$2*S* \leq 2$	$2*S* \leq 2$
3	$2*S* \leq 3$	$2*S* \leq 3$

The index did not reject the cell. It revised its rule — and then admitted 74 further cells the old rule forbade. The result is closed under join and meet, has $E = 0$, satisfies every law in this book, and contains 75 physically impossible configurations.

Where the rule is content, the same cell is checked against a stated rule set — (bounded, parent, slope, intercept) as cells of their own index — and returns

VIOLATES $2S \leq 1 \cdot k + 0$ ($2S = 2$, bound = 1).

Same cell. Same index. The only difference is whether the rule was present or inferred.

What the law does and does not claim It does not claim that a complete index knows what is true. It claims that nothing in it is hidden.

An inferred rule absorbs a false cell and teaches itself the falsehood. A stated rule is contradicted by it, in the open.

A forger's remaining move is to alter a rule — $2*S* \leq k$ becomes $2*S* \leq k+1$, admitting 225 cells. **That forged rule is also visible.** It can be read, disputed and checked; what it cannot be is concealed. **The regress is real and at every level the falsehood is in plain sight.**

So P22 is a statement about visibility, not omniscience. A complete index converts deception into disagreement. **Error survives; concealment does not.**

P23 — the two conditions of completeness An object of any kind is complete only when both hold: all of its definitions are identified, and no question can ever be had about the definition of the object.

The first condition is checkable. Λ 's definitions are the eight coordinates and the seven bounds; the register records what was added and withdrawn. Enumeration is finite work and this book does it.

The second condition is not checkable — and that is its content. No inspection establishes that no question can ever be had. A single question refutes it and nothing confirms it.

So completeness is not a property that can be proved of an object. It is a standing claim, held until a question arrives.

What P23 costs this book Every entry in the register is a question that could be had about a definition, and there are two hundred and ninety-seven of them. Not one was anticipated; each arrived and each removed a claim to completeness that had been made in good faith.

The clearest case is §23.3. Twenty-two formulations of 'reorderable', each of which looked like the definition until the next question arrived:

the question that arrived what it was about

is it $\mathcal{R}$ -closure or join/meet closure? the same object, twice named

is the alphabet declared or recovered? two different objects

is a projection a criterion? §13.4 says no

is a frontier-correct tree anchor-complete? no

does the obstruction grow in cells or in width? cells; $\setminus A_i \leq |X| \setminus$

Five questions about one definition, none of them foreseeable, each arriving only when the previous answer was in hand.

Why P23 is not despair A standing claim is not a weak one. Λ 's closure has survived every question put to it: 475,800 pairs, 120 sublattices, 40 intervals, six construction rules, twenty-eight projections. $E(\Lambda) = 0$ after all of it.

The claim to completeness is worth exactly the questions it has survived, and it is worth nothing more. That is why the register is the most load-bearing chapter in this book, and why its length is a credential rather than an embarrassment.

And it is the honest reading of P22. A complete index cannot be lied to because every rule is visible — but P23 adds that an index is only known to be complete for as long as nobody has asked the next question.

And this is why the book indexes its own rules Appendix B indexes the principles and the protocols. §21.8 indexes the process. Each turns a rule set into content, and each closes with $E = 0$.

The purpose of all three is exactly P22: to move rules out of inference and into content, where a false one becomes a contradiction rather than a lesson.

And the regress terminates outside. Not at a better index — at Chapter 17's 1,442 measured levels, the only content in this book that no index supplied. That is what $\mathbb{D}_{\text{phys}}$ is for, and §11.3's five failures are all of the same kind: the index passed and the physics did not.

B.2 The protocols B.2.1 Channel construction A channel is one parent term, one ℓ , one series in n .

No channel may mix parents, mix ℓ , or cross a coupling-scheme boundary. Where the source's labelling changes scheme partway up a series — LS to jK — **the channel ends there**, and the cells above are not claimed. This is what limits Si I to four cells and P II to one.

B.2.2 Admissibility $r = 2Z^2R / (v^3\sigma) \geq 5$

Levels must be separated by more than five times their uncertainty. r falls as v^{-3} , so every channel eventually leaves the domain; the rule states where, in advance.

B.2.3 Cost resolvability $v \leq v_V = (3Z^2R / 5q)^{1/4}$

The bracket needs levels *separated*; the cost needs curvature *resolved*. **Curvature washes out first.** Cells above v_V are excluded from the cost test and retained for the bracket — a distinction applied prospectively in Li I and diagnostically in Al I.

B.2.4 Dilution cells per channel $\approx (\text{measured shells} / D) - 2$, with D the number of J levels in the core's ground configuration

Dilution divides; it does not forbid. Ne II fails at $D = 5$ with two shells; Ar II succeeds at $D = 5$ with three; Bi I succeeds at $D = 5$ with five.

B.2.5 Census before search Compute which species could work before looking for data.

D is obtained from the core configuration by enumerating antisymmetric states. No spectroscopic data enters. **27 sequences admitted, 12 needing exceptional depth, 5 excluded structurally.**

B.2.6 Declines are recorded A collection reporting only successes cannot be checked.

Every species examined and rejected is listed with its cause. Four causes have been identified — dilution, depth, coupling-scheme change, and core-excited interleaving — and §16.11 gives a worked case of each.

B.2.7 Exhaustive verification where possible

Sampling n cells with no failure bounds the failure rate at $p \leq 3/n$ at 95% confidence. It never certifies $p = 0$.

Where a claim is universally quantified over a finite lattice, it is checked on **every** cell. Where sampling is used, the sample size is stated.

B.2.8 Two independent routes For every quantity that matters, build a second derivation that does not pass through the first.

This is $\mathbb{D}_{\text{def}}$ made operational. It caught a mis-stated relation in Chapter 11, and it is realised experimentally in the antiprotonic-helium check of §11.4, where two laser measurements agree at 0.09σ with no shared step.

An identity computed along one path is vacuous. $w = V \cdot e$ cannot fail if V is defined as w/e .

B.2.9 Refuse rather than coerce No third state. A quantity is admissible or it is refused; it is never silently defaulted.

A guard returning a null that a caller coerces to a Boolean produced three false conclusions in this work, none of which looked wrong. **Totality is the only mechanism that catches this**, and it works by preventing the computation rather than by detecting the error.

B.2.10 Retrieval Enumerate targets from the index before searching. List routes per cell. Prefer high- p cells. Read every retrieved source for its successors. When a route is blocked, move along the route set. Record p for cells that fail. And RE-ENUMERATE whenever the set the targets are drawn from grows.

Chapter 14. A blocked cell with $p = 1$ is a **stated gap**, not an unexplained absence.

The second clause was added late. §23.4 found five enumerations in one chapter — its own target list, §23.2's structural invariants, and the object lists of §24.1, §24.2 and §24.3 — each made once, against a structure that kept growing. Every one of them was incomplete when re-run.

Enumerating once is not a protocol. It is a snapshot, and a snapshot of a growing object is a defect the moment it is taken.

B.2.11 Provenance Every number traces to a compilation, table and page, or to a database entry with its stated uncertainty.

Appendix C carries this at channel level for all 153 channels.

B.2.12 Withdrawal Every retracted claim is logged with what replaced it.

Chapter 21. two hundred and ninety-seven entries. The register is the evidence for Chapter 11, not an apology for it.

B.2.13 Commit before you look A prediction is written down before the measurement is requested. A method is chosen before its answer is visible.

This is the process index of §21.8 enforced by hand: the constraint **visible $\leq$ committed**, of §11.4 form, closed, $E = 0$.

It is the only protocol here adopted in response to an error that no computation could have caught — because every computation was correct and what failed was the independence of a choice. **The deduction of §18.6 was committed under it.**

B.2.14 Compute, then write The result is computed before the sentence describing it is written.

This is the protocol most often violated in the production of this book. The dominant failure recorded in Chapter 20 is prose asserting a conclusion while the output contradicting it stands on the same screen. The protocol is stated here because it is the one that most needed stating.

B.2.15 Apply P8 to every failure Every failed attempt is a true answer, and by P8 a true answer is a bound. Record what each failure excludes, and collect the exclusions.

P8 says any true answer, good or bad, is a bound. A failed attempt is a true answer about where the target is not. The protocol is to treat it that way rather than discard it.

B.2.15.1 The procedure On each failure, write down three things:

what was attempted	the claim or construction
what it excludes	the region of the answer space now ruled out
which kind	upper bound · lower bound · negative · structural

Then read the collection. Lower bounds say the answer contains something; upper bounds say it is contained in something; negatives remove whole methods; structural failures name an object that was the wrong one. When the lower and upper bounds meet, the answer is between them.

B.2.15.2 Worked, on §23.3 Ten consecutive failures on one question, each read as a bound:

the failure	the bound	kind
Γ -free $\neq$ reorderable	reorderable $\subsetneq$ totally balanced	upper
chain sufficient, not necessary	chain $\subseteq$ reorderable	lower
depth-2 and depth-3 descent stall, worsening with size	no constant-depth local procedure decides it	negati
linear extensions of the inclusion order incomplete	the inclusion order is not the join-irreducible poset	ve struct
the slice-chain condition fails at $d = 3$	the $d = 2$ condition does not lift	ural upper
the 3-SAT encoding was refuted	a reduction must control which pairs arise	metho
descent traps at $E = 1, 2$	the transposition landscape is not convex	d negati
occupancy sort succeeds 53%	tie-breaking among equal occupancy is not free	ve lower
"row chain or column chain" holds at 68%	neither axis alone decides it; the condition is two-sided	struct ural
the chain criterion characterises the downset	$\mathcal{R}$ -closure and downset-closure differ by the direction	struct
operator, not $\mathcal{R}$	of one axis	ural

Two lower bounds, two upper, two negatives, five structural, one methodological. Between them: reorderable $\supsetneq$ chain and $\subsetneq$ totally balanced, and the gap contains at least the diagonals — which is what produced the characterisation:

X is $\mathcal{R}$ -closed iff $X = \{(r, c) : c \leq M|(r) \text{ and } r \leq N|(c)\}$, with $M|, N|$ the running maxima of its own row and column maxima.

Ten failures produced one exact statement. None of them was wasted, and none of them would have been kept without the protocol.

And the last two are the ones that located it. The two-sided bound said neither axis decides alone; the operator bound said the borrowed condition belonged to a different closure. Together they force the

form $c \leq M|(r)$ and $r \leq N|(c)$, which no single-axis condition can express and which the two-cell diagonal $\{(0,0), (1,1)\}$ realises minimally.

B.2.15.3 Why it belongs here rather than among the principles It says nothing about what an index is. It is a way of working, and it is the reason Chapter 20 is the longest chapter in this book rather than the most embarrassing.

A register of withdrawals is a list of bounds. Read it that way and it stops being a confession.

B.2.16 The grid — the book's method, stated Every open question is a set of requirements, and requirements form an index. Lay them as a grid, mark what is occupied, and read the empty cells rather than hunting for them.

The first empty cell in a column is the whole obstacle. Everything below it is downstream.

B.2.16.1 Why this is the method and not a technique It is Chapter 1 applied to the book's own ignorance. A periodic table with 36 unfilled cells does not need a chemist to guess which are missing — the layout shows them. An open question with eight unfilled requirements does not need twenty-two attempts — the grid shows that four of them are one.

Worked on §23.3: eight empty cells, four rows deep, at $d \geq 3$.

requirement	$d=2$	$d \geq 3$
characterisation	proved	.
decision procedure	linear	.
YES certificate	an order	an order
NO certificate	Tucker — unbounded	unbounded
obstruction family	Tucker 1972	unbounded, route closed
construction rules	six preserving	six preserving
hardness	P, the CSP is binary	.
the data structure	PQ-tree	.

Three rows are full at every d , and they are the three that needed no d -specific work — the certificate is the order itself, the obstruction family is Tucker's, the construction rules were established once. Every empty cell is downstream of one missing object: the $d \geq 3$ analogue of a PQ-tree.

Twenty-two formulations were attempts on four cells that depend on one. The grid says so in a table; hunting concealed it for a day.

B.2.16.2 And it classifies the whole of Q

item	first empty cell	the obstacle
Sc VI 6s	the level is measured	an instrument
Paschen & Götze	the scan is text-searchable	a digitisation artefact
nine literatures	each is entered	labour
Kreuzer-Skarke	a readable copy	a file format
target-spin axis	the extended lattice is built	labour — now done
reorderability at $d \geq 3$	characterisation	a missing object

Five of six were blocked by something outside the index. One is blocked by something that does not exist. §25.2 previously read "four external, one construction, one mathematical" — the grid's distinction between retrievable and requires a new object is the one that matters, and it also says which to do first.

The target-spin axis was taken on that basis and it closed. Λ_9 has 1,654 cells, $E = 0$, is a sublattice, is rank-modular on 200,000 pairs, projects exactly onto Λ_8 , keeps the constraint graph a tree, extends F by one pendant factor, and has $F(-1) = 2$ unchanged. One unit of labour, six properties confirmed, $|Q|$ from six to five.

B.2.16.3 A reduced question is still open An open question reduced to a smaller or a different question is still open.

A reduction feels like progress and reads like an answer. §23.3 supplies five in one session, each announced as a closure:

the reduction	what was still open
reorderability $\Rightarrow$ C1P + monotone endpoints	the endpoint search
the search $\Rightarrow$ a PQ-tree frontier question	a correct PQ-tree
the frontier question $\Rightarrow$ a bottom-up traversal	the traversal's state
the traversal $\Rightarrow$ a row $\rightarrow$ node anchoring map	rows anchoring as runs
the whole problem $\Rightarrow$ a CSP on axis orders	the ternary constraints

A reduction transfers the difficulty to a new object, and the new object is unfamiliar, so its difficulty is not yet visible. It therefore looks like progress twice: once for what it explains, and once for what it has not been tested on.

State the reduction and the residue in the same sentence. Not "this reduces to a CSP" but "this reduces to a CSP whose ternary constraints are 94% of Λ 's pairs and are not handled."

B.2.17 The twenty prime audits, ordered by what a failure costs Every session closes with twenty audits. They are not spot checks, and they are not equal. They are ordered by the cost of the failure each one catches.

The first two can make the book WRONG. The next two can make it UNREADABLE AS AN ARGUMENT. The next two can make it UNUSABLE AS AN OBJECT. The seventh can make it DISHONEST, and it is the only one that cannot be run by reading the book.

1 — LATTICE. Every structural claim recomputed from the construction: closure under join and meet, $E(X) = 0$, rank modularity, the fibre identities, the rank polynomial, $F(1)$ and $F(-1)$, $|J(\Lambda)|$ and the height. Over all 475,800 pairs of Λ , not a sample. First, because if the object is not closed the book has no subject.

2 — EQUATIONS. Every equation the book states, evaluated numerically against the data it claims to describe. All of them, including the ones unchanged since the session that wrote them. Second, because a wrong equation is a false claim.

3 — CONSISTENCY. No two statements in the book may contradict. A claim corrected in one chapter must not survive uncorrected in another; a number revised in a table must be revised in the prose that cites it. Third, because a contradiction leaves the reader unable to tell which statement is the claim.

4 — REDUNDANCY. No statement may survive its own supersession. Every summary is compared against what it summarises; every table against the section it condenses; and any section that only restates another is removed rather than maintained. Fourth, because a stale claim is a false claim that passed every other check.

5 — ARTEFACT. Audits 1–4 read the SOURCE. This one reads the built PDF. Missing glyphs, figures present and non-blank, the title page and running header read back, no table overflowing its frame, the whitespace measured from glyph and image geometry, and every figure checked for ink, tone,

resolution, duplication and correct numbering. Fifth, because the source can be perfect and the page still wrong.

6 — COHERENCE. *Source, other places, unusable, presupposes labels are addressable.* Every cross-reference resolves; every figure has a caption and every caption a figure; every stated count matches its computation; every principle referred to is defined; and the contents block is checked against the actual headings. Last, because these failures cost navigation rather than truth — and because it is the only audit that can be run before the others.

7 — ATTRIBUTION. Every claim that points outside the book resolves to a source the book lists, and every source listed is claimed by something inside it. Five checks: each work cited in the text appears in the References; each entry in the References is cited somewhere; §22.2's owner column names people the References carry; every Appendix E element recorded as *precedent found* has a reachable citation; and every data source in Appendix C is listed. Seventh, because a missing attribution makes the book neither wrong nor unreadable nor unusable — it makes it **dishonest**, which is a worse failure and a quieter one.

Audits 1–2 check the book against the world. Audits 3–4 check the book against itself.

Audits 5–6 check the book against what a reader receives. **Audit 7 checks the book against what it owes**, and it is the only one of the seven that cannot be run by reading the book.

B.2.17.1 Why the CONSISTENCY audit was needed Because the register contains its failures, and this one was at scale. Twenty-eight cross-references pointed at §17.4, which is He II, and between them they claimed at least eight different pieces of content:

what the reference claimed	where it lives
89,864 join failures	§6.7
$V = 4v^3/(h(h^2 - 3v^2))$	§16.4
closure not locally determined	§13.4
7/7 on controls	§22.7
511 joins, the conserving poset	§13.1.1
self-concordance	§16.8.2
the anchoring failure	§21.6
four clauses	§25.3

Three of the ten formalised principles cited it, and five index entries pointed at it. All twenty-eight are now retargeted.

the COHERENCE audit passed every one of them, because every target existed. The reference resolved; the content did not match. That gap is the whole reason the CONSISTENCY audit exists, and it went unrun for as long as the audits numbered three. A literature

citation was withdrawn in one place and remained correct in another. A statistic was stated against two different denominators in adjacent sections. None of these is a factual error; each is an internal contradiction, and audits 1–3 pass them all.

And P22 requires it. A complete index cannot be lied to because every rule is visible — but two visible rules that contradict each other leave a reader unable to tell which is the rule. The CONSISTENCY audit is what makes visibility useful.

B.2.17.2 Why the ARTEFACT audit was needed Because the operator this book is about did not print.

No DejaVu face carries U+1D4E1 or U+1D49F — the script R and script D of the recovery and diagnostic operators. Every render placed a blank box instead: 107 of them, across 44 pages, for a symbol appearing 56 times in the source and 50 for its companion.

audit	why it passed
-------	---------------

1 coherence	the source is correct; the character is right there
2 lattice	recomputes structure, not typography
3 equations	evaluates numbers, not glyphs
4 consistency	the source agrees with itself throughout

All of them read the source. The defect existed only in the output.

And it was compounded by a second artefact fault the same day: a table cell of 229 characters overflowed its frame and aborted the build. That one at least announced itself. The glyph failure was silent — a correct source, a successful build, and a symbol missing from every page that used it.

The fix has three parts, all in the renderer and none in the text: substitute $\mathcal{R}$ and $\mathbb{D}$ in body faces that carry them; fall back to plain R and D in the serif bold used for headings, which carries neither; and extract the text back out of the finished PDF and count the nulls. It now reads zero.

An audit that reads the wrong artefact will pass anything. The CONSISTENCY audit exists because a reference resolved to the wrong section. The ARTEFACT audit exists because a character resolved to nothing at all.

B.2.17.2.1 The whitespace clause, and two measurements that were wrong A gap a reader sees is a defect the source cannot show. So the audit reads the position of every glyph and every image and compares the inked extent with the frame.

On this book: median foot gap 3.3%, nine pages above 15%, worst 31% — and that one is the end of a section, where a gap is correct. Interior gaps cluster at 16–28%, which is exactly the CondPageBreak(58mm) guard that stops a heading landing at the foot of a page.

Tightening the guard makes it worse, not better. At 46mm the count of gaps falls from nine to seven and the worst rises from 31% to 43%, because a heading allowed to sit lower pushes the figure block after it off the page entirely. 58mm is the optimum, and it was found by measuring rather than by preference.

B.2.17.2.2 And two of the audit's own measurements were wrong first The first attempt used the PDF text matrix directly and reported a median inked fraction of 103% — impossible — with y-coordinates of 5, 6 and –220. The text matrix is not page geometry until it is composed with the current transformation matrix.

The second attempt counted characters and not images, and flagged seven pages as having a 25–45% gap at the head. All seven begin with a figure. The space was not empty; it was occupied by something the measurement could not see.

Both are the same error as the audits they were written to replace: a check looking at the wrong artefact. An audit is not exempt from the discipline it enforces, and this one needed two corrections before its first clean reading.

B.2.17.2.3 What the figure audit found Twenty-eight figures, each checked against its caption and its file.

Clean: every figure has a caption and every caption a figure; no image is blank or unusually dark; no file is duplicated; nothing in fig/ is unused; and after one correction every figure number matches the chapter it sits in.

The one correction. Figures 23.1 and 23.2 appeared in the wrong order — the arity figure sits in §23.3.2 and the step-law figure in §23.3.4, so the numbering ran backwards against the reading order. Swapped.

And one defect that is recorded rather than repaired.

figures at 300 dpi or better	5
between 200 and 300	8
under 200 dpi	15

median

164 dpi

Fifteen figures were produced by an earlier pipeline whose default was `figure.dpi = 160`, and at 163 mm rendered width that gives 140–164 dpi — below print standard. Regenerating them needs the source data of the sessions that made them, which this one does not carry. Appendix D.3 lists them with the rest of what was inherited and not recomputed.

The audit's value here is not that it fixed the figures. It is that the number is now written down instead of unknown.

B.2.17.2.4 Why the REDUNDANCY audit was needed Because two sections survived their own supersession and passed all five.

A summary table listed Q's eight items while Appendix F enumerated, indexed and closed them — the same content, maintained twice. And a status table still read "nine literatures entered, nine named and unentered" and "order recovery beyond trees: open, and known to be hard", both superseded hours earlier by §F.5 and §23.3.

audit	why it passed
lattice · equations	neither is a structural or numerical claim
consistency	both were internally consistent and cited sections that exist
artefact	both rendered correctly
coherence	both were well-formed tables with valid references

A stale statement is not inconsistent with anything. It agrees with the version of the book that no longer exists. That is why the check has to compare a summary against what it summarises, rather than a document against itself.

The remedy was to delete a chapter. What survived of it — one paragraph on what the collection closed — opens Chapter 25, where it does work instead of restating.

B.2.17.3 Exhaustive, and what to do when it is impossible Exhaustive means every case, not a representative sample. The distinction is not pedantic:

checked	how	result
the level-3 signature, 2×2×2	all 256 subsets	189 tested, 30 multi-instance classes, 0 mixed
the same, 2×2×3	all 4,096 subsets	3,838 tested, 210 classes, 0 mixed
the same, sampled at d = 3, 4	2,072 instances	0 mixed — and a sample cannot exclude a mixed pair

Where exhaustion is impossible, say so and name the bound. Λ has 976 cells and 2^{976} subsets. What was exhausted: all 28 two-dimensional projections, all 475,800 pairs, every named fibre. What was not: the subsets. The audit records both.

An audit that samples is a measurement. An audit that exhausts is a proof over its domain. Both are useful and they must not be reported in the same words.

B.2.17.7 Why the ATTRIBUTION audit was needed Because the other six could not find the defect and could not have. **They are closed under self-reference:** every one of them checks the book against the book, so the set as a whole has no way to notice an unpaid debt. That is the book's own subject turned on the book's own instruments — the audit set admitted a cell it did not contain, and $E > 0$ held for the audits themselves until this one was written.

It has been run twice and found a defect each time. The first run found Baker and Pixley 1975, Bergman's Double-projection Theorem, and Kimura, Makino, Yamada and Yoshizumi 2024 cited in §§22.11–22.12 and missing from the References — **the three precedents that place this book's central theorem in existing literature, and the three it had not listed.** The second found Racah 1942 and 1943, attributed at §22.9 for jK coupling and seniority, the two borrowings the tower rests on most directly, likewise unlisted. Registers 261, 262, 263.

B.2.17.8 Why the audits keep missing things — $E(\text{audits}) = 11$ They miss things because **the audit set is not closed**, and until this section nobody had applied the book's own test to it. Index the seven on three coordinates — what an audit *reads*, what it *compares that against*, and what a failure *costs* — and reconstruct:

audit	reads	compares against	a failure makes the book
1 LATTICE	the object	a computation	wrong
2 EQUATIONS	the source	a computation	wrong
3 CONSISTENCY	the source	other places	unreadable
4 REDUNDANCY	the source	other places	unreadable
5 ARTEFACT	the artefact	itself	unusable
6 COHERENCE	the source	other places	unusable
7 ATTRIBUTION	outside the book	other places	dishonest

Seven audits occupy six distinct cells. $\mathcal{R}$ **admits seventeen. $E(\text{audits}) = 11$, at these three coordinates** — §B.2.17.10 recomputes it at four and gets nineteen.

And two of the eleven are the defects that had just been found by hand. *Reads the source, compares against a computation, costs unreadable* is §E.5's table: adding a column and comparing the total to the prose a page above it. No audit does arithmetic on a printed table, so the table said twenty-seven while the analysis said thirty-two for four register entries. *Reads the artefact, compares against other places, costs unusable* is the title pages: audit 5 reads the artefact alone, audit 3 reads the source alone, and nothing compares them field by field, so a hardcoded page in the press and a corrected line in the book disagreed silently.

The other nine are proposals, in the sense §1 gives the word — combinations the structure admits and no audit occupies, which is where a defect can sit and be missed. Reading the object and comparing it within one place. Reading outside the book and checking a computation. Reading the source and comparing it to itself. None has fired yet, and the two that did fire fired before this was computed, which is the only reason to expect anything of the nine.

This is the book's subject applied to the book's instruments. Chapter 1 says a defect is a proposal about a cell nobody has listed; the audit set is an index like any other, and it had never been indexed. Registers 262 and 273. Closing it means writing audits, not writing prose, and the count is the honest statement of how far from closed the instruments are.

B.2.17.9 The eleven, defined and retained — eighteen prime audits §B.2.17.8 counted the gap; these fill it. Each of the eleven admitted cells is written as an audit with the same standing as the first seven, and the set closes: **eighteen audits over seventeen cells, $E(\text{audits}) = 0$ at three coordinates.**

8 — CELL. *Object, itself, wrong.* Every cell satisfies χ on its own, independently of any other cell. Catches a member smuggled past the pairwise tests. **First run: 0 of 976 fail.**

9 — DISTINCTNESS. *Object, other places, wrong.* No two cells are equal and every derived encoding is injective. Catches a coordinate that has silently collapsed. **First run: 0 duplicates, the 17-bit encoding injective on all 976.**

10 — SCOPE. *Source, itself, wrong.* Every quantitative claim carries the conditions that make it true — caps, fibration, sample size — in its own sentence. **First run: 2 flags, both real,** two statements of

$E(Q) = 0$ in Appendix F written before fibration-dependence was known and left bare. This is the audit that would have caught $F(1) = 976$ without §2.4's caps.

11 — ANTECEDENT. *Source, itself, unreadable, presupposes labels are addressable.* Every numbered heading has a parent that exists, and every reference resolves inside its own passage. **First run: 0 orphans.**

12 — MARKUP. *Source, itself, unusable.* Every emphasis span opens and closes inside the unit the press parses. **First run: 0.** Registers 260 and 264 are the two failures this audit exists for.

13 — AGREEMENT. *Source, other places, wrong.* A quantity stated in two places has one value. **First run: 1 flag, not real** — three different counts of *interior cells* read as one quantity by a detector that did not scope by context.

14 — ARITHMETIC. *Source, a computation, unreadable.* Every printed table's columns sum to the total its prose states. **First run: 0** — and only because register 269 had corrected §E.5 by hand four entries earlier. This audit is that defect made unrepeatable.

15 — ENUMERATION. *Source, a computation, unusable.* Every stated count of listed things equals the number of things listed. **First run: 1 flag, not real** — a parser that read Appendix B's principles as Appendix F's items.

16 — FIDELITY. *Artefact, other places, unusable.* Every field the press renders matches the source it claims to render. **First run: 1 flag, real and by design** — the PDF metadata title exists only in the press. Register 272 is the failure this audit exists for, and metadata is the one place the press may legitimately hold text of its own.

17 — MEASURE. *Artefact, a computation, unusable.* The artefact's measurable properties match what the book says about them — figure count, page count, fill. **First run: 32 of 32 figures.**

18 — REPRODUCTION. *Outside the book, a computation, dishonest.* Every number attributed to an outside source is checked against that source's own stated value. **First run: 473,800,776 and 30,108 both present and correct.**

Four of the eleven fired on their first run and two of the four were real, which is the ratio B.2.19.7 predicts and the reason that protocol is stated before this section rather than after it. The eleven are retained as prime audits. Register 274.

B.2.17.10 The hierarchy, and the fourth coordinate it forces §B.2.17.9 closed the audit set at $E = 0$ and recorded one thing it could not ask of itself: whether *reads*, *compares against* and *costs* are the right three coordinates. **Defining the hierarchy answers it, and the answer is no.**

The audits have an order, and it is not the presentation order. Twelve precedences hold between them — LATTICE before EQUATIONS because a description of an unclosed object describes nothing, CONSISTENCY before REDUNDANCY because a contradiction is not a redundancy, the source audits before the artefact audits because a page is rendered after it is fixed. Test that order against the coordinates and two things fall out.

What the three coordinates do determine. The *stage* at which an audit can run — before writing, after writing, after rendering — is a function of them, with no cell disagreeing. Stage is therefore not an independent coordinate, and asking for it would have been the mistake §12 warns against.

What they do not. Three of the twelve precedences are unexplained: CONSISTENCY before REDUNDANCY, EQUATIONS before SCOPE, EQUATIONS before AGREEMENT. And CONSISTENCY and REDUNDANCY share all three coordinates exactly — eighteen audits over seventeen cells was a collision, not a coincidence. **A fourth coordinate is admitted by §12's test**, and the three unexplained precedences say what it is:

DEPTH — what the audit presupposes. Nothing $<$ the claims are individually true $<$ they are also mutually consistent.

It separates the collision, explains all three precedences, and gives eighteen audits eighteen distinct cells. The hierarchy is now recoverable from the coordinates, which is S2 satisfied for the instruments.

And it costs what §12 says an extension costs. With three coordinates $E(\text{audits}) = 0$; with four, **$E(\text{audits}) = 19$** . The set that looked closed is open by nineteen. Nothing was broken by adding the coordinate — the earlier zero was true and weak, because a smaller ambient box is easier to fill.

Which completes register 233 and generalises it. That entry showed E falls as a fibration is refined, reaching zero by refinement alone. This shows E rises as coordinates are added. They are the same lever seen from two ends: **$E = 0$ is strong exactly in proportion to the size of the ambient box it is computed in** — coarsest fibration, most coordinates. A zero must therefore state both its fibration and its coordinate list or it states nothing, and every E in this book now owes that. For Λ the answer is the strong one and always was: one fibre, eight coordinates, no categorical axis to hide behind. For the audits it is nineteen. Register 275.

B.2.17.11 The answer the lattice already held §B.2.17.10 admitted a fourth coordinate because the hierarchy demanded it, and left open whether four is right. **That question is not a matter of judgement, and the book has carried its answer since A.4.** The precedences form a poset; the least number of ordered coordinates that represents a poset is its **order dimension**; so the coordinate count is computed, not chosen.

Computed: the eighteen audits carry twenty-four strict relations across seven components, with width eight and height six. The one non-trivial component has 210 linear extensions, and **two of them intersect exactly to the order**. The other components are chains and singletons. So

$\dim(\text{audit hierarchy}) = 2$.

Two ordered coordinates represent the entire precedence structure. The index uses four.

That is not an error, and the difference is the finding. Order dimension answers *how many coordinates are needed to recover the order*. It says nothing about how many are needed to say what an audit **is**. Two coordinates would let a reader reconstruct which audit runs before which, and would not tell them that **ATTRIBUTE** reads outside the book while **MARKUP** reads a line of source. §10's S_2 is satisfied at two; §10's S_1 — the alphabet — is not.

So the four coordinates split, and the split is the lattice's own. Two carry order and are forced. Two carry meaning and are admitted only by §12's independence test, which is a different test with a different standard. Λ has the same division and never named it: its eight coordinates represent an order of far lower dimension, and the surplus is what makes the cells *mean* something rather than merely sort. **The seven bits of §6.1.1 are that surplus counted once**; this is the same quantity counted for the instruments. A coordinate system answers to two masters, and only one of them is computable. Register 276.

B.2.17.12 The nineteenth audit — **SEQUENCE** §B.2.17.10 left $E(\text{audits}) = 19$ at four coordinates: nineteen combinations the structure admits and no audit occupies. **One of the nineteen has now been located by a defect landing in it**, which is the first time a proposal from that count has been confirmed rather than asserted.

19 — **SEQUENCE**. *Source, a computation, unusable, presupposes nothing*. **A numbering is a claim about position**. Every label's place in the text must match the order its number asserts, and every label must be written in the form the structure can address. Nineteenth by depth — it presupposes nothing, so it can run before any other — and *unusable* because its failures cost navigation rather than truth: a reader who follows the numbers is walked backwards.

Adding it takes $E(\text{audits})$ from 19 to 18.

First run, and it found three things the other eighteen had passed over. F.1.2 printed before F.1.1; A.18 after A.19.1 with A.13 last; and B.3 ahead of three sections placed into B by the repair for register 280 — all three found and fixed at register 283 by hand, and all three exactly what this audit tests. Beyond those: **entry 210 of the register stood alone between 204 and 283**, in neither the ascending historical block nor the descending recent one, and **197 and 198 were inverted** inside the ascending block. Both are now in place.

One finding is recorded and not repaired. The register prints entries 165 to 204 ascending and 205 to the present descending — two directions in one sequence, because early entries were appended and later ones prepended. Each block is internally monotone. **Reversing a hundred and twenty entries to unify them is a larger risk than the defect**, so the direction change is stated here and left, and SEQUENCE checks each block for monotonicity rather than the whole for one direction. Register 285.

B.2.17.13 The twentieth audit — PROJECTION Register 293 records five appendix sections absent from the contents because the rule that generates the contents did not match their titles. **No audit could have found it**, and the reason is precise enough to name: audit 16, FIDELITY, checks that what the press renders matches its source, and every line the press rendered did match. The five were never rendered at all. **Correctness and completeness are different failures, and the eighteen tested only the first.**

20 — PROJECTION. *Artefact, other places, unusable, presupposes the claims are individually true.* Every derived list is checked against an **independent enumeration of its source**, never against the rule that produced it. A generated projection is exactly as complete as its generating rule, so a rule that misses a case yields a list that looks complete and is not.

Its cell was among the eighteen §B.2.17.12 left admitted and absent; adding it takes E(audits) to 17.

First run. Three derived lists exist in this book: the contents, generated from headings; the figure placements, generated from captions; and Appendix G's population count, generated by the press at build. Enumerated independently — 400 headings by a second rule, 32 captions, 2,811 claim-bearing numbers — all three agree. It passes, and it passes only because register 293 had repaired the deriver an hour earlier.

And the general statement is worth more than the audit. Register 216 replaced a maintained contents with a derived one and closed the class of drift defects. It opened another and did not say so: **a maintained list that omits a section is visibly short; a derived list that omits one looks complete.** The failure moved from forgetting to matching, and matching failures are silent. **Any repair that replaces a human process with a rule inherits the rule's blind spots**, and this audit is the standing check on that trade. Register 294.

B.2.17.14 The hierarchy at twenty §B.2.17.10 fixed the coordinates against seven audits and §B.2.17.11 measured the hierarchy they carried. Both are restated here at twenty, because a hierarchy fixed on a third of a set is a claim that expires.

The precedences, nineteen of them. CELL before LATTICE, because a set containing an invalid cell has no closure worth testing. LATTICE before EQUATIONS. EQUATIONS before SCOPE, AGREEMENT, ARITHMETIC and ENUMERATION, because a property of a number presupposes the number. CONSISTENCY before REDUNDANCY; both before ARTEFACT. MARKUP before ARTEFACT and FIDELITY. ARTEFACT before MEASURE, FIDELITY and PROJECTION. ATTRIBUTION before REPRODUCTION. **SEQUENCE before ANTECEDENT, COHERENCE and PROJECTION** — the newest and the earliest, because a label that is not addressable is invisible to every check that consumes labels.

Two of the nineteen were unexplained, and the resolution was not a fifth coordinate. SEQUENCE stood before ANTECEDENT and COHERENCE with no coordinate rising between them. §12's loop asks first whether an existing coordinate was misread, and it was: ANTECEDENT and COHERENCE were recorded as presupposing **nothing**, and they presuppose that labels are addressable — which is what SEQUENCE establishes. Both move to *presupposes the claims are individually true*, and **all nineteen precedences are then explained by the four coordinates.** No axis was added, which is the outcome §12 prefers and the one this book has twice failed to reach for on the first attempt.

The measurements, at twenty. Twenty audits over twenty distinct cells, no collision. The precedence order falls into three components of ten, eight and two; the largest has 4,140 linear extensions and two of them intersect exactly to the order, so **dim(hierarchy) = 2**, unchanged from seven audits to twenty. **E(audits) = 17 at four coordinates**, down from nineteen as SEQUENCE and PROJECTION took two of the missing cells.

That the dimension held while the set nearly tripled is the finding. Order dimension measures how much *branching* a hierarchy has, not how much of it there is; the audits grew in width and not in entanglement. Two coordinates still recover which audit runs before which. The other two still carry what an audit *is*, and that division — §B.2.17.11's — is what the growth tested and did not break. Register 295.

B.2.18 Triangulation — three results locate a fourth For every three results or bounds that measure the same object, triangulate: they identify a related data point that none of them states.

Three measurements of one thing are not three facts. They are a position.

B.2.18.1 Worked, from §23.3

the three	what they	the fourth, located
	measured	
valid orders disconnected ·	connectivity failure in	the disconnection is in the object, not the space —
E-landscape traps · cell-growth peaks	three different spaces	so no local method in any space can be complete
step 1 at $d = 2$ · step 2 at $d = 3$ ·	growth step, three	step = $2^{(d-2)}$, then proved by subcube removal
step 4 at $d = 4$	dimensions	
arity 2 bijunctive · arity 3	the constraint	the boundary is arity 3 and it is exact — the
bijunctive · arity 4 contains	language	procedure and the hardness proof are one law read at two
1-in-3-SAT		arities
three reductions, all	the encoding	the map $X \rightarrow C(X)$ is global, non-separable and
over-constraining		non-monotone — it admits no local editing

Each fourth point was invisible in any single measurement and forced by the three together. The step law is the clearest: 1, 2, 4 fixed the $\frac{3}{4}$ ratio, and the ratio gave the subcube construction, which proved it.

B.2.18.2 And it is a check as well as a method Triangulation fails loudly. If three measurements of one object do not locate a consistent fourth, one of them is wrong — and that is how three corrections were caught in a single session:

A map that is global, non-separable and non-monotone should have a SMALL image. Measured: 39% at $d = 3$ and 49.5% at $d = 4$, rising. The three coordinates were right and the inference from them was not — they locate an obstruction to EDITING the map, not to its range.

Three consistent measurements give a fourth result. Three measurements that give no consistent fourth give a correction instead. Both outcomes are useful and the protocol does not distinguish them in advance.

B.2.18.3 Why it belongs with the protocols It is P8 with a count attached. Every failure is a bound; three bounds on one object are a coordinate. B.2.15 says to record failures as bounds; this says what to do once three of them are in hand.

And it composes with B.2.16. The grid shows which cells are empty; triangulation fills one from three of its neighbours.

B.2.19 The seven failures of the assistant, and the protocol each one earns Two hundred and ninety-seven entries stand in the register and they are not one population. §25.1.1 flag 4 separates

them: defects already in the book and found by auditing it, errors the author made, and errors made by the assistant doing the computing. This section is the third population, sorted not by entry but by **kind**, because fourteen entries collapse to seven mechanisms and a mechanism is what a protocol can address. Each is stated as the failure, then the rule that would have caught it, then the entries that are instances.

B.2.19.1 Attribute Inward First *When a recomputation disagrees with the book, the model is the suspect until it is cleared.* Three times a recomputation failed and the shortfall was written down as the book's: the density column was called underspecified when the exact sets in my model were loose; the thirty-two-against-twenty-seven count was called unsettleable when the arithmetic was on the page; the book was said to face an exposure the preface answers in its third paragraph. **The rule is asymmetric on purpose** — the book has been audited and the model has not, so the prior belongs with the book. Registers 207, 213, 220.

B.2.19.2 Read The Structure At The Point Of Insertion *Before writing into a numbered sequence, enumerate what already occupies it.* Four ledger rows were written into Appendix A as A.8–A.11 over proofs that already stood there; a heading was written with a leading space into a file where that makes it invisible; a contents block was restructured without checking what it indexed. Every one is the same act — writing into a structure without reading the structure. Registers 215, 218, 259.

B.2.19.3 A Claim Of Completion Is Itself A Claim *Any sentence asserting that work was done must be verified by the check it claims to have passed, in the same pass that writes it.* §25.1.2 stated three things as done that were not done when the section claiming them was built. Three Q items were described as *entered in Q* in the sections that named them and never reached Appendix F. The failure is not forgetting; it is that describing and doing feel identical while writing. Registers 239, 256.

B.2.19.4 Repair The Class Or Say You Did Not *A fix that leaves the mechanism unchanged is an instance fix, and must be recorded as one.* An indented heading was repaired in two sections and reintroduced in the next turn, because nothing in the press rejected one. Emphasis straddling a cell was corrected in the source and recurred at a line boundary, because the parser still worked per line. Both are now gates in the press. **The instance is cheap and the class is the work**, and the register is the place to say which was done. Registers 259, 260, 264.

B.2.19.5 A Null Is A Result With Its Definition Missing *Before recording that something was not found, define what was looked for and count what could have been.* A tripwire run was filed as a null at ten checks; defining the implicit-parent case raised it to fifty-four and turned the result into two — a specificity measure and a ceiling on published notation. **This is the book's own title page** and it took being told. Registers 253, 254.

B.2.19.6 Show The Test Can Fail *A check must be demonstrated capable of returning a failure before its passing is evidence of anything.* A character class matched no superscripts and returned zero rows as though clean; a summary line was printed by a script that hard-coded it one line below the computation refuting it; a lemma looked refuted on 1,098 cases because *strictly* had been dropped from the definition. **A test that cannot fail returns no failures**, and that is indistinguishable from success until someone checks the checker. Register 227.

B.2.19.7 A Flag Is A Hypothesis *Localise a detector's output before recording it as a defect.* A phrase quoted inside a register entry read as a surviving error; a section title beginning with a lowercase Greek letter read as a missing proof; a line break inside a bold span read as an unprinted claim. Each would have entered the register as the book's failure had it been believed. **The audits generate hypotheses, and the register takes only what survives testing.**

What the seven have in common, and it is one thing. Every one is a claim made faster than it was checked — outward before inward, into a structure before reading it, of completion before completing, of a class from an instance, of a null from an undefined absence, of a pass from an untested test, of a defect from an untested flag. **B.2.14 already says it: compute, then write.** These seven are that protocol failing in seven distinguishable places, and the reason they are listed separately is that knowing the general rule demonstrably did not prevent any of them. Register 266.

B.2.20 What is computable, and what is decided §B.2.17.11 ends by saying a coordinate system answers to two masters and only one of them is computable. The word is doing real work and the book has not defined it. It defines it with its own operator.

A quantity is COMPUTABLE in this book when it is recoverable by $\mathcal{R}$ from what the book prints — when two readers running the stated procedure on the printed text return the same value, and neither has to decide anything to get there.

That is not a claim about difficulty. $E(\Lambda) = 0$ over 475,800 pairs is computable and took minutes; whether A.5's language is order or analysis is not computable and took one sentence.

Nineteen quantities in the work of these sessions are computable. $E(\Lambda) = 0$; $|J(\Lambda)| = 17$ with its twenty covering relations; the Birkhoff down-set count of 976; the alphabet's size $\Sigma(|A_i| - 1) = 17$; the seventeen-bit encoding with join as OR and meet as AND; the 7.07 bits of surplus; the twenty implications cutting 131,072 words to 976; every tower cell count and every density; closure through the f shell at 499,246 cells and exhaustively to 47,775,744 ambient ones; the E1 puncture at four meets all producing (0,0); the triangle at four caps; lemma 4 over 2,354 interval sets; $E(\text{audits})$ at 11, 0 and 19; $\dim(\text{audit hierarchy}) = 2$; $E(G) = 33$; $E(Q)$; the tripwire's 150 of 216; the fifty-four of one hundred and two. **Each is a function of the printed object and of nothing else.**

Ten are decisions, and they are of three kinds.

Frames. The fibration of Appendix E, the coordinate list for the audits, the caps of §2.4. E is computable **given** them and none of them is computed — register 275 is the general statement, that a zero is worth the size of the box it was computed in, and the box is chosen.

Readings. Applying a stated criterion to a case it does not settle: whether A.5's proof makes it order or analysis, whether the five results are elements under §E.5.3, whether *law* means holds-of-every-index or holds-and-forbids. The criterion is printed; the application is a judgement, and §E.5.4 marks the one place this work made one and says so.

Claims about the outside. Precedent, novelty, whether a measurement exists. §22.1 proves the first of these permanently undecidable in one direction, which is the strongest statement of this kind the book contains.

And the distinction has a use, not merely a name. Every computable quantity in this book can be checked by a reader with the printed text and no access to its author. Every decision must be *stated* rather than checked, and the register is what stating them costs. **The seven prime audits of §B.2.17.1-7 test computable claims; nothing tests a decision**, which is why the seventeenth through twentieth entries of the register are corrections of computation and the twentieth through fortieth are reversals of judgement. The book cannot make its decisions checkable. It can make them visible, and that is the whole of what the register is for. Register 277.

B.2.20.1 The lattice applied to what it cannot compute §B.2.20 separates nineteen computed quantities from ten decisions and stops there, on the ground that nothing tests a decision. **The shape still applies to them, and it says two things.**

First: index the decisions and E finds one that has not been made. Coordinates *scope*, *how it is stated* and *reversibility*, fibred by kind — frame, reading, outside claim. Ten decisions, $E = 1$ fibred and 5 unfibred, with a single cell admitted and absent in the *reading* fibre:

a reading, scoped to one appendix, stated, freely reversible.

That cell has an occupant the book declined to seat. Register 270 records that Appendices C and G are cited from the body only as wholes and never at section level, calls it a navigability defect, and leaves the choice open. **It is a reading, it is confined to two appendices, it is stated, and it costs nothing to reverse.** E over the decisions located a deferred decision exactly as E over Q located item H — which is the first evidence that the operator has any purchase outside the computable.

Second, and this is the finding: the decisions form a DAG, and two of them have two parents. The order is acyclic, so the decisions can be made in some sequence and none presupposes itself. Six of the ten depend on nothing. But two do not have a single parent:

the fibration of Appendix E — settled only with the caps of §2.4 **and** what counts as an element both in view

not promoting the six theorems — settled only with the meaning of *law* **and** the audit coordinate list both in view

This is §7.11's dichotomy, arriving in the one place the book had assumed it did not reach. A single-parent constraint propagates down a tree and is carried exactly; a two-parent constraint breaks the tree and admits only an envelope. The tight K has two parents and costs the cylinder 2,475 cells. A two-parent *decision* is the same shape: **no earlier choice determines it, and it cannot be delegated to one predecessor or deferred until one is settled.** It must be made with both in view or not at all.

The book has two such decisions and it stumbled on both. Registers 219 to 222 hold three successive positions on the element count, which is the fibration decision made once with the caps in view and once with the element criterion in view and never with both. Register 232 declines the promotion and says the reason is a trade between two things at once. **Two instances is not a law**, and the correspondence is offered as what it is — a shape that fits twice, in a set of ten, where the prediction it makes is that any third two-parent decision will be hard in the same way. Register 278.

B.2.21 What a safe repair requires §B.2.17.12 recorded a defect and declined to repair it on the ground that the repair was riskier than the defect. That judgement was made without a standard, and this states one.

A repair is **safe** when five conditions hold. **(i)** Its invariant is written down before the edit — what must be true afterwards that was true before. **(ii)** The invariant is computable in the sense of §B.2.20, so it can be checked rather than asserted. **(iii)** The edit is a *permutation or an insertion*, never a rewrite: existing text is moved, not retyped, so no character can change unnoticed. **(iv)** It is reversible — a copy is taken first and restored if the check fails. **(v)** It is checked afterwards on the source **and** the artefact, because §B.2.17.9's FIDELITY exists for the gap between them.

And a sixth that governs which repair to attempt: prefer the smallest edit that removes the defect. A defect visible in a hundred places is not therefore a hundred-place repair. Applying that here changed what the repair was.

B.2.21.1 The register, examined before repairing it The recorded defect was direction: entries 165 to 204 ascending, 205 onward descending. Measuring the register against its own headings found something larger and cheaper to fix. Each subsection of §21.7 states a count in its title — *Eight more, Twelve more, Fifteen more, Forty more* — and **§21.7.6 was titled "Three from the tower session" and contained eighty-six entries, 199 to 284**, every entry written since that session having been prepended into it. It is now four — the tower's three and register 202, which repairs their numbering — and §21.7.7 the eighty-seven written afterwards.

That is the defect worth repairing, and it is one heading rather than a hundred and twenty entries. The direction is a presentation choice made twice; the heading is a false statement of count, which audit 15 ENUMERATION exists to catch and which it missed because it was pointed at tables and lists rather than at section titles. Registers 286 and 287.

B.3 What the protocols are for Each prevents a specific failure, and each failure occurred:

protocol	failure it prevents	instance
B.2.1 channel construction	claiming cells across a label change	Si I, P II
B.2.2 admissibility	bracketing unresolved levels	—
B.2.3 cost resolvability	reading rounding as physics	Al I nf
B.2.4 dilution	excluding a viable species	Ar II, Bi I

B.2.5 census first	searching without targets	27 sequences enumerated
B.2.6 declines recorded	an uncheckable collection	Ne II, Cu II
B.2.7 exhaustive testing	a claim resting on a sample	A.4 on all cells
B.2.8 two routes	an undetectable derivation error	$w = 3v \cdot e$
B.2.9 refuse, don't coerce	a missing value read as a real one	the cap discard
B.2.10 retrieval	abandoning a cell when a source is blocked	$p = 6$ vs $p = 1$
B.2.11 provenance	an untraceable number	153 channels
B.2.12 withdrawal	an unfalsifiable record	48 entries
B.2.13 commit before you look	a choice that depended on its answer	the anchoring failure
B.2.14 compute then write	the dominant failure	most of Chapter 20
B.2.15 apply P8 to every failure	a failure discarded rather than banked	ten bounds on §23.3

The protocols are not advice. They are the failure modes of this work, inverted.

Appendix C — Data and provenance

Serving PART III — THE METHOD. Every number the method was run on, with its source, its provenance and its known defects.

Every channel in the collection, with its source, its limit, and what it contributes. 153 channels, **1105 interior cells**, bracket **1105 of 1105**.

The 763 cells of the earlier verification are included where their channels appear here; the totals in Chapter 17 state the union.

C.1 Sources

compilation	spectra drawn
Kaufman & Martin 1991, JPCRD 20, 775	Al I, Al II
Kramida & Martin 1997, JPCRD 26, 1185	Be I
NIST ASD	Ar II, Be II, Bi I, C II, Ca II, Cd II, Ga I, He I, He II, Hg II, K II, Li I, Li II, Mg II, N II, Na II, Ne I ($\frac{3}{2}$), Ne I ($\frac{1}{2}$), Si I, Si II, Zn II
Sansonetti 2008, JPCRD 37, 1659	Na I
Sansonetti 2008, JPCRD 37, 7	KI

Additional sources used for exotic systems and collective quantities:

Hori *et al.*, *Nature* **475**, 484 (2011) and *Phys. Rev. Lett.* **96**, 243401 (2006), both accessed via CODATA 2010 Table XII (arXiv:1203.5425) and arXiv:1304.4330 · Korobov, *Phys. Rev. A* **77**, 042506 (2008) · Singer, Stanojevic, Weidemüller & Côté, *J. Phys. B* **38**, S295 (2005) · Sugar & Corliss 1985, *JPCRD* **14** Suppl. 2 · Sugar & Musgrove 1990, *JPCRD* **19**, 527 and 1995, *JPCRD* **24**, 1803.

C.2 Channels

species	channel	n	me	cells	bracket	v range	δ	sprea	Z_eff	limit (cm ⁻¹)
			m					d		

Al I 3s²nf ²F° 4-55 52 50 50/50 4.0-55.0 +0.0429 0.0552 1 48,278.480

Al I 3s²nd ²D 3-34 32 30 30/30 2.6-33.7 +0.0248 0.6297 1 48,278.480

Al I 3s²ns ²S 4-16 13 11 11/11 2.2-14.2 +1.7670 0.0537 1 48,278.480

Al I 3s²np ²P° 3-7 5 3 3/3 1.5-5.7 +1.3365 0.2080 1 48,278.480

Al II 3sns ¹S 4-16 13 11 11/11 2.8-14.8 +1.2028 0.0147 2 151,862.500

Al II 3sns ³S 4-13 10 8 8/8 2.7-11.7 +1.2711 0.0474 2 151,862.500

Al II 3snd ¹D 3-11 9 7 7/7 3.2-10.9 +0.0603 0.3897 2 151,862.500

Al II 3sng 5-13 9 7 7/7 5.0-13.0 +0.0212 0.0070 2 151,862.500

Al II 3snd ³D 3-10 8 6 6/6 2.8-9.8 +0.2024 0.0096 2 151,862.500

Al II 3snf ³F° 4-11 8 6 6/6 3.9-11.0 +0.0140 0.4088 2 151,862.500

species	channel	n	me	cells	bracket	v range	δ	sprea	Z_eff	limit (cm ⁻¹)
m								d		
Ar II (³P)nd ⁴F J=9/2	3-5	3 1	1/1	2.3-4.4	+0.6121	0.0885	2	222,848.300		
Be I 2snp ¹P°	3-13	11 9	9/9	2.7-12.6	+0.3590	0.1025	1	75,192.500		
Be I 2snd ³D	3-12	10 8	8/8	2.9-11.9	+0.1109	0.0129	1	75,192.500		
Be I 2snd ¹D	3-12	10 8	8/8	3.2-12.1	-0.1152	0.0970	1	75,192.500		
Be I 2sns ¹S	3-11	9 7	7/7	2.3-10.3	+0.6774	0.0136	1	75,192.500		
Be I 2sns ³S	3-8	6 4	4/4	2.2-7.2	+0.7910	0.0429	1	75,192.500		
Be I 2snp ³P°	3-8	6 4	4/4	2.6-7.6	+0.3773	0.0374	1	75,192.500		
Be I 2snf ³F°	4-7	4 2	2/2	4.0-7.0	+0.0305	0.0063	1	75,192.500		
Be II ns ²S	2-10	9 7	7/7	1.7-9.7	+0.2623	0.0112	2	146,882.860		
Be II np ²P°	2-10	9 7	7/7	2.0-10.0	+0.0491	0.0039	2	146,882.860		
Be II nd ²D	3-9	7 5	5/5	3.0-9.0	+0.0021	0.0004	2	146,882.860		
Be II nf ²F°	4-10	7 5	5/5	4.0-10.0	+0.0001	0.0001	2	146,882.860		
Be II ng ²G	5-10	6 4	4/4	5.0-10.0	-0.0001	0.0001	2	146,882.860		
Be II nh ²H°	6-10	5 3	3/3	6.0-10.0	-0.0002	0.0001	2	146,882.860		
Bi I (³Po)ns 2[0] J=1/2	7-11	5 3	3/3	2.0-6.1	+4.9036	0.0649	1	58,761.650		
Bi I (³Po)nd 2[2] J=3/2	6-10	5 3	3/3	2.7-6.7	+3.2646	0.0470	1	58,761.650		
Bi I (³Po)nd 2[2] J=5/2	6-9	4 2	2/2	2.8-5.8	+3.2071	0.0272	1	58,761.650		
Bi I (³Po)np 2[1]° J=1/2	7-9	3 1	1/1	2.5-4.6	+4.4712	0.0588	1	58,761.650		
C II 2s²ns ²S	3-8	6 4	4/4	2.3-7.3	+0.6624	0.0063	2	196,664.700		
C II 2s²np ²P°	3-7	5 3	3/3	2.6-6.7	+0.3928	0.1344	2	196,664.700		
C II 2s²nd ²D	3-7	5 3	3/3	2.9-6.9	+0.0924	0.0645	2	196,664.700		

C II 2s²nf ²F° 4-7 4 2 2/2 4.0-7.0 +0.0220 0.0069 2 196,664.700
 C II 2s²ng ²G 5-7 3 1 1/1 5.0-7.0 +0.0054 0.0011 2 196,664.700
 Ca II nd ²D 3-16 14 12 12/12 2.3-15.4 +0.6338 0.0615 2 95,751.870
 Ca II ns ²S 4-10 7 5 5/5 2.1-8.2 +1.8192 0.0529 2 95,751.870
 Ca II nf ²F° 4-10 7 5 5/5 4.0-10.0 +0.0248 0.0088 2 95,751.870
 Ca II ng ²G 5-9 5 3 3/3 5.0-9.0 +0.0048 0.0011 2 95,751.870
 Ca II np ²P° 4-6 3 1 1/1 2.5-4.5 +1.4748 0.0489 2 95,751.870
 Cd II ns 5-13 9 7 7/7 1.8-9.9 +3.1175 0.1092 2 136,374.740
 Cd II nd 5-13 9 7 7/7 3.1-11.1 +1.8906 0.0534 2 136,374.740
 Cd II np 5-12 8 6 6/6 2.2-9.3 +2.7145 0.1825 2 136,374.740
 Cd II ng 5-11 7 5 5/5 5.0-11.0 +0.0075 0.0013 2 136,374.740
 Cd II nf 4-9 6 4 4/4 4.0-9.0 +0.0333 0.0918 2 136,374.740
 Ga I 4s²ns ²S 5-28 24 22 22/22 2.2-25.2 +2.8013 0.0535 1 48,387.634
 Ga I 4s²nd ²D 4-27 23 21 21/21 2.8-25.7 +1.2926 0.1920 1 48,387.634
 Ga I 4s²np ²P° high 41-5 15 13 13/13 38.8-52.7 +2.2107 0.1209 1 48,387.634
 Ga I 4s²nf ²F° 4-8 5 3 3/3 4.0-8.0 +0.0252 0.0057 1 48,387.634
 Ga I 4s²np ²P° low 4-7 4 2 2/2 1.5-4.7 +2.3450 0.2026 1 48,387.634
 He I 1sns ³S 2-35 25 23 23/23 1.7-34.7 +0.2965 0.0164 1 198,310.666
 He I 1sns ¹S 2-35 25 23 23/23 1.9-34.9 +0.1392 0.0118 1 198,310.666
 He I 1snp ¹P° 2-35 25 23 23/23 2.0-35.0 -0.0133 0.0050 1 198,310.666
 He I 1snd ¹D 3-35 24 22 22/22 3.0-35.0 +0.0007 0.0019 1 198,310.666
 He I 1snf ¹F° 4-35 23 21 21/21 4.0-35.0 -0.0010 0.0020 1 198,310.666
 He I 1snd ³D 3-35 22 20 20/20 3.0-35.0 +0.0014 0.0018 1 198,310.666
 He I 1sng ¹G 5-35 22 20 20/20 5.0-35.0 -0.0014 0.0020 1 198,310.666
 He I 1snh ¹H° 6-35 21 19 19/19 6.0-35.0 -0.0015 0.0020 1 198,310.666
 He I 1sni ¹I 7-35 20 18 18/18 7.0-35.0 -0.0016 0.0019 1 198,310.666
 He II ns 1-10 10 8 8/8 1.0-10.0 -0.0003 0.0006 2 438,908.871
 He II np 2-10 9 7 7/7 2.0-10.0 -0.0004 0.0005 2 438,908.871
 He II nd 3-10 8 6 6/6 3.0-10.0 -0.0004 0.0005 2 438,908.871

species	channel	n	me	cells	bracket	v range	δ	sprea	Z_eff	limit (cm ⁻¹)
m								d		
He II nf	4-10	7 5 5/5	4.0-10.0	-0.0005	0.0004	2	438,908.871			
He II ng	5-10	6 4 4/4	5.0-10.0	-0.0005	0.0003	2	438,908.871			

He II nh 6-9 4 2 2/2 6.0-9.0 -0.0005 0.0002 2 438,908.871
 He II ni 7-9 3 1 1/1 7.0-9.0 -0.0005 0.0001 2 438,908.871
 Hg II ns 6-11 6 4 4/4 1.7-6.8 +4.1865 0.1423 2 151,280.000
 Hg II nd 6-11 6 4 4/4 3.1-8.2 +2.8668 0.0609 2 151,280.000
 Hg II ng 5-9 5 3 3/3 5.0-9.0 +0.0101 0.0011 2 151,280.000
 Hg II np 6-9 4 2 2/2 2.2-5.3 +3.7302 0.1626 2 151,280.000
 Hg II nf 5-8 4 2 2/2 4.0-6.9 +1.0615 0.0317 2 151,280.000
 Kl nd 3-46 44 42 42/42 2.9-45.7 +0.2679 0.1326 1 35,009.814
 Kl ns 4-46 43 41 41/41 1.8-43.8 +2.1831 0.0490 1 35,009.814
 Kl np 4-18 15 13 13/13 2.2-16.3 +1.7202 0.0533 1 35,009.814
 Kl nf 4-14 11 9 9/9 4.0-14.0 +0.0092 0.0051 1 35,009.814
 K II 3p⁵nf 2[3/2] J=1 4-6 3 1 1/1 4.0-6.0 +0.0426 0.0069 2 255,072.800
 K II 3p⁵nf 2[9/2] J=5 4-6 3 1 1/1 4.0-6.0 +0.0336 0.0071 2 255,072.800
 K II 3p⁵nf 2[5/2] J=3 4-6 3 1 1/1 4.0-6.0 +0.0242 0.0052 2 255,072.800
 K II 3p⁵nf 2[7/2] J=4 4-6 3 1 1/1 4.0-6.0 +0.0143 0.0059 2 255,072.800
 Li I np 2P° 2-42 41 39 39/39 2.0-41.8 +0.1477 0.6779 1 43,487.114
 Li I ns 2S 2-11 10 8 8/8 1.6-10.6 +0.4014 0.0180 1 43,487.114
 Li I nd 2D 3-12 10 8 8/8 3.0-12.0 +0.0031 0.0129 1 43,487.114
 Li II 1sns ³S 2-9 8 6 6/6 1.8-8.8 +0.1815 0.0115 2 610,078.400
 Li II 1sns ¹S 2-9 8 6 6/6 1.9-8.9 +0.0743 0.0037 2 610,078.400
 Li II 1snp ³P° 2-8 7 5 5/5 1.9-7.9 +0.0540 0.0010 2 610,078.400
 Li II 1snp ¹P° 2-8 7 5 5/5 2.0-8.0 -0.0145 0.0065 2 610,078.400
 Li II 1snd ³D 3-8 6 4 4/4 3.0-8.0 +0.0024 0.0006 2 610,078.400
 Li II 1snd ¹D 3-8 6 4 4/4 3.0-8.0 +0.0009 0.0003 2 610,078.400
 Li II 1snf ³F° 4-8 5 3 3/3 4.0-8.0 +0.0002 0.0002 2 610,078.400
 Li II 1snf ¹F° 4-8 5 3 3/3 4.0-8.0 -0.0001 0.0004 2 610,078.400
 Li II 1sng 5-8 4 2 2/2 5.0-8.0 -0.0002 0.0002 2 610,078.400
 Li II 1snh 6-8 3 1 1/1 6.0-8.0 -0.0003 0.0001 2 610,078.400
 Mg II ns ²S 3-10 8 6 6/6 1.9-8.9 +1.0749 0.0289 2 121,267.640
 Mg II nd ²D 3-10 8 6 6/6 3.0-10.0 +0.0411 0.0145 2 121,267.640
 Mg II np ²P° 3-9 7 5 5/5 2.3-8.3 +0.7077 0.0361 2 121,267.640
 Mg II nf ²F° 4-10 7 5 5/5 4.0-10.0 +0.0030 0.0009 2 121,267.640
 Mg II ng ²G 5-10 6 4 4/4 5.0-10.0 +0.0007 0.0001 2 121,267.640

Mg II nh $^2H^\circ$ 6-9 4 2 2/2 6.0-9.0 +0.0002 0.0000 2 121,267.640
 Mg II ni 2I 7-9 3 1 1/1 7.0-9.0 +0.0000 0.0000 2 121,267.640
 N II 2p·ns $^3P^\circ$ J=0 3-6 4 2 2/2 2.2-5.2 +0.7747 0.0264 2 238,750.500
 N II 2p·ns $^3P^\circ$ J=1 3-6 4 2 2/2 2.2-5.2 +0.7731 0.0260 2 238,750.500
 N II 2p·ns $^3P^\circ$ J=2 3-6 4 2 2/2 2.2-5.3 +0.7633 0.0464 2 238,750.500
 N II 2p·ns $^1P^\circ$ J=1 3-6 4 2 2/2 2.2-5.3 +0.7325 0.0824 2 238,750.500
 Na I ns 3-20 18 16 16/16 1.6-18.7 +1.3506 0.0250 1 41,449.451
 Na I np 3-20 18 16 16/16 2.1-19.1 +0.8584 0.0279 1 41,449.451
 Na I nd 3-20 18 16 16/16 3.0-20.0 +0.0141 0.0045 1 41,449.451
 Na I nf 4-20 17 15 15/15 4.0-20.0 +0.0014 0.0004 1 41,449.451
 Na I ng 5-20 16 14 14/14 5.0-20.0 +0.0003 0.0001 1 41,449.451
 Na I nh 6-10 5 3 3/3 6.0-10.0 +0.0000 0.0000 1 41,449.451
 Na II ($^2P^\circ_{3/2}$)ns 2[3/2] $^\circ$ J=2 3-6 4 2 2/2 1.9-5.0 +1.0385 0.0284 2 381,390.200
 Na II ($^2P^\circ_{3/2}$)ns 2[3/2] $^\circ$ J=1 3-6 4 2 2/2 1.9-5.0 +1.0268 0.0371 2 381,390.200
 Na II ($^2P^\circ_{3/2}$)nd 2[1/2] $^\circ$ J=0 3-5 3 1 1/1 2.9-4.9 +0.0671 0.0083 2 381,390.200
 Na II ($^2P^\circ_{3/2}$)nd 2[1/2] $^\circ$ J=1 3-5 3 1 1/1 2.9-4.9 +0.0641 0.0085 2 381,390.200

species	channel	n	me	cells	bracket	v range	δ	sprea	Z_eff	limit (cm ⁻¹)
m					d					
Na II (² P° _{3/2})nd 2[3/2]° J=2	3-5	3 1 1/1	2.9-4.9	+0.0556	0.0027	2	381,390.200			
Na II (² P° _{3/2})nd 2[3/2]° J=1	3-5	3 1 1/1	3.0-5.0	+0.0268	0.0026	2	381,390.200			
Ne I (³ / ₂) ns 2[3/2]° J=1	3-20	18 16 16/16	1.7-18.7	+1.2968	0.0488	1	173,929.750			
Ne I (³ / ₂) nd 2[1/2]° J=1	3-15	13 11 11/11	3.0-15.0	+0.0277	0.0173	1	173,929.750			
Ne I (³ / ₂) ns 2[3/2]° J=2	3-13	11 9 9/9	1.7-11.7	+1.3191	0.0276	1	173,929.750			
Ne I (³ / ₂) nd 2[7/2]° J=4	3-13	11 9 9/9	3.0-13.0	+0.0226	0.0073	1	173,929.750			
Ne I (³ / ₂) nd 2[3/2]° J=2	3-13	11 9 9/9	3.0-13.0	+0.0164	0.0176	1	173,929.750			
Ne I (³ / ₂) nd 2[5/2]° J=3	3-12	10 8 8/8	3.0-12.0	+0.0081	0.0054	1	173,929.750			
Ne I (³ / ₂) nd 2[1/2]° J=0	3-11	9 7 7/7	3.0-11.0	+0.0323	0.0078	1	173,929.750			
Ne I (³ / ₂) np 2[5/2] J=3	3-10	8 6 6/6	2.1-9.2	+0.8486	0.0383	1	173,929.750			
Ne I (³ / ₂) np 2[1/2] J=1	3-8	6 4 4/4	2.1-7.1	+0.8941	0.0623	1	173,929.750			
Ne I (³ / ₂) nf 2[9/2] J=5	4-9	6 4 4/4	4.0-9.0	+0.0035	0.0007	1	173,929.750			
Ne I (³ / ₂) nf 2[7/2] J=4	4-8	5 3 3/3	4.0-8.0	-0.0015	0.0006	1	173,929.750			
Ne I (³ / ₂) ng 2[11/2]° J=6	5-7	3 1 1/1	5.0-7.0	+0.0012	0.0000	1	173,929.750			
Ne I (¹ / ₂) ns 2[1/2]° J=1	3-20	18 16 16/16	1.7-18.7	+1.3027	0.0266	1	174,710.170			

Ne I ($\frac{1}{2}$) nd 2[3/2] $^{\circ}$ J=1 3-20 18 16 16/16 3.0-20.0 +0.0137 0.0120 1 174,710.170

Ne I ($\frac{1}{2}$) nd 2[5/2] $^{\circ}$ J=3 3-10 8 6 6/6 3.0-10.0 +0.0178 0.0064 1 174,710.170

Ne I ($\frac{1}{2}$) ns 2[1/2] $^{\circ}$ J=0 3-10 8 6 6/6 1.7-8.7 +1.3207 0.0265 1 174,710.170

Si I 3p·nd $^1D^{\circ}$ J=2 3-5 3 1 1/1 2.4-4.5 +0.5412 0.0461 1 65,747.760

Si I 3p·nd $^3F^{\circ}$ J=2 3-5 3 1 1/1 2.6-4.6 +0.3718 0.0179 1 65,747.760

Si I 3p·nd $^3F^{\circ}$ J=3 3-5 3 1 1/1 2.6-4.7 +0.3546 0.0309 1 65,747.760

Si I 3p·nd $^3F^{\circ}$ J=4 3-5 3 1 1/1 2.6-4.7 +0.3188 0.0886 1 65,747.760

Si II 3s 2 np $^2P^{\circ}$ 3-10 8 6 6/6 1.8-9.1 +1.0223 0.2717 2 131,838.140

Si II 3s 2 nf $^2F^{\circ}$ 4-11 8 6 6/6 3.9-10.9 +0.0845 0.0370 2 131,838.140

Si II 3s 2 ng 2G 5-11 7 5 5/5 5.0-11.0 +0.0176 0.0044 2 131,838.140

Si II 3s 2 ns 2S 4-9 6 4 4/4 2.6-7.6 +1.3943 0.0468 2 131,838.140

Si II 3s 2 nd 2D 3-8 6 4 4/4 2.9-7.7 +0.2289 0.1599 2 131,838.140

Zn II ns 2S 4-9 6 4 4/4 1.7-6.8 +2.2084 0.0679 2 144,892.600

Zn II nd 2D 4-9 6 4 4/4 3.0-8.0 +0.9583 0.0228 2 144,892.600

Zn II np $^2P^{\circ}$ 4-8 5 3 3/3 2.1-6.1 +1.8386 0.0616 2 144,892.600

Zn II nf $^2F^{\circ}$ 4-8 5 3 3/3 4.0-8.0 +0.0172 0.0044 2 144,892.600

Zn II ng 2G 5-7 3 1 1/1 5.0-7.0 +0.0038 0.0008 2 144,892.600

Totals: 153 channels, 1105 cells, bracket 1105/1105.

C.3 Flagged channels Channels whose δ spread exceeds 0.25 or whose V departs from $4v/3$ by more than 5% are flagged. Every flag in this collection has an identified cause; none is unexplained.

specie	channel	δ	cause
s		spread	
Al I	3s 2 nd 2D	0.630	compilers relabelled this series; y 2D removed as a perturber
Al II	3snd 1D	0.390	local perturber
Al II	3snf $^3F^{\circ}$	0.409	δ reverses sign between n = 6 and 7 — local perturber
Li I	np $2P^{\circ}$	0.678	39 cells to n = 42; quotation coarsens above n = 33
Si II	3s 2 np $^2P^{\circ}$	0.272	perturber

C.4 Reproducibility

Every computation in this book is reproducible from the level tables above and the rules of Chapters 13 and 14. No fitted parameter enters at any point; the only inputs are measured or calculated energy levels, their stated uncertainties, and the ionisation limits listed in C.2.

Appendix D — Margins, and what was recomputed

Serving PART IV — THE RECORD. What the conclusions can bear, and what was inherited rather than recomputed.

A referee should treat two classes of claim differently, and this appendix separates them.

D.1 Margins on the load-bearing conclusions How far is each result from the threshold that would overturn it?

conclusion	statistic	margin	breaks down at
the periodic table is not closed	$E = 36$	exact arithmetic	—
closure $\Leftrightarrow \mathcal{R}(X) = X$	200/200 agreements	exact	—
$\mathbb{D} \geq \dim q - \dim p$	420 random maps	exact — rank-nullity	—
$V > 2$ for any monotone sequence	Proposition 14.1	exact, proved	—
Rydberg floor $V \geq 32/11$	2.909091 at $v = 2$	exact	—
E3 admissibility criterion	2,513/2,513, nine functions	exact	—
$\chi_{\wedge}$ total	30,000 ambient points	exact	—
order recovery on a tree	20/20, two cap settings	—	a cyclic constraint graph
the bracket holds	1,442/1,442	—	one reordering
perturbation bounds	tightest 1.40 cm^{-1}	deductive	a 1.4 cm^{-1} shift at that cell
isoelectronic interpolation	1.3% median	13× better than at an end	the sequence end
$V = 4v/3$ on real channels	0.06%–4.4% median	—	coarse quotation
void-free fraction stable	27.7–30.1%	2.4 points over a 100× range	—
filled-d core raises δ	+0.35 at low ℓ	—	$\ell \geq 3$, where cores converge

Three rows deserve a sceptic's attention first.

The bracket's 1,442/1,442 breaks on a single reordering. §25.3 states the condition exactly — $|\Delta T| > 2Z^2R/v^3$ — and §18.4 argues that the collection never enters the regime where it can occur. That argument, not the count, is what a critic should attack.

The perturbation bounds are as strong as the failure condition and no stronger. If the condition is wrong, 1,061 bounds go with it.

The isoelectronic interpolation is quoted at 1.3% and used at the end of a sequence, where the measured cost is 2.1% for a hold-out and **3.5% for a genuine fifth member**. The Sc VI route of §18.6 uses the last figure, not the first.

D.2 What was computed in this revision Fourteen results rest on computations performed for this book, each reproducible from Appendix C's tables and the rules of Part III:

$E(X)$ for seven indices, the periodic table among them · closure $\iff \mathcal{R}(X) = X$ on 200 constructions · $\mathbb{D} \geq \dim q - \dim p$ on 420 random nonlinear maps · the E3 criterion on 2,513 tests · χ^2 totality on 30,000 uniformly sampled ambient points · order recovery by tree propagation, 20/20 · the 1,105 cells added here, across all 153 channels · the failure condition and its inversion to 1,061 perturbation bounds · the isoelectronic hold-out and fifth-member tests · the C₆ cost-law check against Singer et al. · the antiprotonic-helium two-route check at 0.09σ · the census by antisymmetric-state enumeration · Proposition 14.1 re-derived · the three falsification tests of §25.6 — the book's own, distinct from §24.3.4.2's test of an external prediction

D.3 What was inherited and not recomputed Six results come from earlier work and were not re-run:

the rule-ablation costs of §17.2 · the Sr I node coverages and fine-structure figures of §17.2 · the Ti I channel inventory · the multi-target failure counts of §11.3 · the Sc VI prediction's input defects · the earlier verification's per-species medians.

And fifteen figures at 140–164 dpi, produced by an earlier plotting pipeline whose default was figure.dpi = 160: f01, f02, f03, f04, f05, f06, f07, f08, f09, f10, f11, f12, f13, f14, f18. Measured and recorded, not regenerated — the source data of the sessions that produced them is not in this one. B.2.17.2.3 gives the distribution.

These rest on methods described in Appendix C but were performed before this revision. Chapter 21 records two hundred and ninety-seven errors found in work of exactly that kind — which is reason to re-run them rather than to trust them.

The Sc VI defects are the most consequential of the six, because the book's single prediction rests on them, and §18.6's arithmetic was independently reproduced here — 735,091 against 735,092 — while its inputs were not.

One inconsistency was found in the inherited material and is recorded. The exclusion discussion of §17.2 states that Sc VI's "highest measured level is 5s at 696,400 cm⁻¹". **That figure and the stated 6(5s) = 0.9812 cannot both describe the 4S° channel:** from 696,400 the defect is 0.514, and the stated defect requires 5s at 648,096. **The 696,400 is Sc VI's highest measured level overall, not the 4S° 5s**, and the prose ran them together. The prediction is unaffected, having been computed from the defects.

D.4 Computational methods All computations in Python 3 with NumPy. Closure tests are either exhaustive over all pairs — stated where so — or on uniform random samples with the sample size reported. Box searches enumerate all 2^d corners of candidate base cells.

Ritz fits use ordinary least squares on δ against $1/n^2$, with σ from the standard prediction error including the leverage term $s \cdot \sqrt{1 + 1/k + (x^*_o - \bar{x})^2 / S^*_{xx}}$.

Quotation granularity q is read from the source table as the smallest nonzero increment appearing in it — 10⁻³ cm⁻¹ for most ASD entries, coarsening to 0.1 or 1 cm⁻¹ in the high- n regions of several channels — and enters v_V directly.

Weighted centroids over fine-structure components use statistical weights $2^j + 1$.

Where a channel's neighbours are not consecutive in n , the bracket is still evaluated but V is not, since $V = 4x/(h|p-1|)$ assumes unit step. That exclusion is applied silently in the totals and is the reason the V counts are smaller than the cell counts in Appendix C.

Appendix E — The mathematics, indexed

Serving PART IV — THE RECORD. The book's own mathematics as a closed index — thirty-two elements, sixteen fibres, $E = 0$.

The book's method, applied to the book's own mathematics. Every mathematical element is a cell; the coordinates are its kind, its language, its status, how it was verified, and whether a precedent exists. The index is then closed, and $E(X)$ counts the questions the structure admits and the book does not answer.

E.1 Why the first attempt failed, and what it established Five coordinates were assigned: kind, language, status, verification, precedent. The closure admitted hundreds of cells. The reason is a precondition the book had never stated:

$\mathcal{R}$ recovers MONOTONE bounds. It applies to ORDERED coordinates only.

Kind and language are categorical. Ordering definition < mechanism < formula < theorem is arbitrary, and a monotone bound on an arbitrary ranking admits everything above it. Λ 's own coordinates are all ordered — $n, \ell, k, q, e, f, g, 2S$ are counts — so the book never met the case.

The repair is the shape Λ already has: fibre over the categorical axes, and put the ordered coordinates inside.

E.2 The structure

coordinate	values	kind
kind	definition · mechanism · formula · theorem · law · method · measurement	fibre
language	order · combinatorics · analysis · complexity · physics · algebraic geometry	fibre
status	withdrawn < conjectured < measured < verified < proved	ordered
verification	cited < sampled < exhaustive	ordered
precedent	none found < found	ordered

Thirty-two elements occupy sixteen of the forty-two possible fibres.

E.3 The two constraints, found by closing $E(X) > 0$ in three fibres, and the excess cells shared a form. §12's diagnostic loop says a separating form in the excess means a missing constraint. Two were found:

verification = cited $\Rightarrow$ precedent = found. You cannot cite a precedent that does not exist.
 status $\geq$ measured $\Rightarrow$ verification $\geq$ sampled. If the book measured it, the book verified it.
 Cited can coexist only with conjectured or withdrawn — a claim nobody here checked.

Both have the form of Λ 's own bounds — a monotone bound of one coordinate by another — and neither was stated before the index forced it.

E.4 The two elements the closure demanded One cell survived both constraints: an analysis formula both proved and exhaustively verified. The book had $V = 4\nu/3$ (measured, exhaustive) and λ^2 as the Newton decrement (proved, cited) and nothing in between.

E.4.1 $\lambda^2 = (2/3)T$, proved and exhaustively verified For a Rydberg term $T(\nu) = Z^2R/\nu^2$:

$$\lambda^2 = (T')^2/T'' = (2Z^2R/\nu^3)^2 / (6Z^2R/\nu^4) = (2/3) \cdot Z^2R/\nu^2 = (2/3)T$$

Proved in one line. And now verified exhaustively: 12,000 evaluations over $Z = 1 \dots 30$ and $\nu = 1.01 \dots 60$, zero violations, worst relative error 7.1×10^{-16} — floating-point exactness across the whole domain.

The identity was stated in Chapter 16 and cited to Nesterov-Nemirovskii. It had never been checked here.

E.4.2 The pointwise-minimum reconstruction $\Lambda = \{x : x_i \leq \min_j \phi_{ij}(x_j) \text{ for every } i\}$ — verified at 976 = 976, exact.

$\mathcal{R}$ recovers forty-eight bounds and takes the smallest at each cell. No single bound on g is tightest anywhere near everywhere — q on 92.6% of cells, k on 39.9%, f on 35.1%, down to n on 11.2%. The binding constraint changes cell by cell, and the pointwise minimum is the operator's entire content.

E.4.3 And it is not the coupling F carries F has one $\min()$. The recovered system has one at every cell.

	in F — the defining system	under $\mathcal{R}$ — the recovered
		system
coordinates with >1 bound	one: $g \leq \min(q, 4f+2)$	all eight
bounds per coordinate	0 or 1, except g	seven
which bound binds	fixed by the construction	changes cell by cell
cells with two bounds tied at the minimum	—	976 of 976 — 100%

Every one of the seven is tightest somewhere. For g : q on 93% of cells, k on 40%, f on 35%, ℓ on 20%. For q : k on 100%, ℓ on 66%, n on 48%. No coordinate has a single binding partner.

F 's coupling is structural and occurs once — the junction of the caterpillar, where g meets both q and $4f+2$. $\mathcal{R}$'s coupling is pointwise and occurs everywhere.

Both give 976, and they are different mechanisms. §6.5 states the defining coupling and stops there, which is correct for the expression and incomplete for the operator. The tie at every cell is the stronger fact: the recovered system is degenerate at every point, and no cell is governed by one bound alone.

Order-theoretic, exhaustive over all 976 cells, and no precedent found for the statement in this form.

E.4.4 Two perspectives of one definition, and one of them carries more F and $\mathcal{R}$ are the same definition seen twice. F says what Λ contains; $\mathcal{R}$ says what Λ requires. Both give 976 cells and neither is derived from the other's coupling.

$\mathcal{R}$ is idempotent. $\Lambda \rightarrow$ fifty-six recovered bounds $\rightarrow$ generates 976 $\rightarrow$ the same fifty-six bounds, and again the same. One application takes any presentation to a canonical form; further applications change nothing.

The defining system is not that fixed point. F carries seven bounds; the normal form has sixteen that constrain, of fifty-six recovered. Both generate the same set.

E.4.4.1 And three of the four directions are functions

$\Lambda \rightarrow$ the recovered bounds	unique — this is $\mathcal{R}$
the recovered bounds $\rightarrow \Lambda$	unique — generate
the defining bounds $\rightarrow \Lambda$	unique — this is Chapter 2
$\Lambda \rightarrow$ the defining bounds	NOT a function

Many presentations give the same Λ . So $\mathcal{R}$ forgets which one was used, and the forgetting is not recoverable.

A presentation and its normal form are two perspectives of one definition, and the map between them runs one way.

E.4.4.2 Which is what the single expression contributes F carries information Λ does not: the order in which the bounds were imposed.

The caterpillar of Figure 6.1 — a path of seven nodes with one pendant at k — is a fact about the presentation, not about the set. §6.7 proves that a single closed expression exists because the constraint

graph is a tree, and a tree is an ordering of impositions. Λ alone supplies no such order; §7.9 counts 1,113,045,672 maximal chains, and the construction picked one.

That is why Chapter 6 and Chapter 10 are separate chapters. Chapter 6 needs the presentation. Chapter 10 recovers only what survives forgetting it.

E.4.4.3 A refinement of P21 P21 says a definition contained in the lattice is true in every language it can be written in. §6.1 gives Λ in six, and the languages agree.

They do not carry equal amounts. The generating-function perspective carries a nesting order the set-theoretic perspective does not, and the order is real — it determines whether a closed form exists at all.

Every perspective on a contained definition is TRUE. They are not INTERCHANGEABLE. A perspective may carry structure the object does not, and a perspective that forgets it is still correct about the object.

That is not a weakening of P21. It is the reason §6.7's proof has a hypothesis — the tree — which is a property of how Λ was written and not of what Λ is.

E.5 The closed index With both constraints and both elements: $E = 0$ in all sixteen fibres, and **the theorem · order row carries the correction of §E.5.4** — five elements were printed here for years against a stated total of thirty-two, and the five that fit no row are A.4, A.5, A.11 and Chapter 11's two theorems.

fibre	cells	E
definition · order	1	0
definition · combinatorics	1	0
mechanism · order	1	0
formula · combinatorics	2	0
formula · analysis	3	0
formula · physics	1	0
theorem · order	10	0
theorem · combinatorics	3	0
theorem · analysis	2	0
law · complexity	1	0
method · order · combinatorics · analysis	1 each	0
measurement · combinatorics	2	0
measurement · physics · algebraic geometry	1 each	0

Thirty-two elements over sixteen fibres. No combination of status, verification and precedent is admitted by the structure and missing from the book, in any kind and any language.

The column sums to thirty-two, which is the figure this appendix has always stated and which its own table contradicted until §E.5.4 enumerated the elements and found the five it had no room for. Registers 219 to 222 hold the three positions the count was argued into before it was counted; register 269 records that the table itself stayed wrong for four entries after the correction was written a page below it.




Figure 25.1. The mathematics of this book as a cell set. Rows are kinds, columns are languages, and each number is how many elements occupy that fibre — thirty-two elements in sixteen of forty-two possible fibres. Within every occupied fibre the three ordered coordinates close: $E = 0$, once two constraints are stated. The dots are fibres nothing occupies, and each is a question the structure poses and the book answers by saying so.

E.5.1 Rechecked, and one thing the appendix does not enable Every $\mathcal{R}$ claim above was recomputed independently from Λ 's construction: fifty-six recovered bounds, sixteen of them constraining, the tightest bound on g being q on 92.6% of cells, k on 39.9%, f on 35.1% and l on 20.4%, the tightest on q being k on 100%, l on 66.2% and n on 47.7%, and every one of the 976 cells carrying at least two bounds tied at the minimum on some coordinate. All reproduce exactly.

What could not be rechecked was $E = 0$ itself, because the appendix printed fibre totals and never named its elements. F.4.1 had already ruled on that against the book — *an assertion is not an enumeration* — and the ruling was not applied here. It is applied now, as far as the book's own text permits, and the attempt found two things before it found anything else.

The stated count and the printed table disagree, and §E.5.2 settles it against the table. §E.5's fibre table lists $1 + 1 + 1 + 2 + 3 + 1 + 5 + 3 + 2 + 1 + (1 \text{ each} \times 3) + 2 + (1 \text{ each} \times 2) = \text{twenty-seven elements}$ across sixteen fibres, and the appendix says thirty-two in two places. The table is complete in fibres — it prints sixteen and the appendix states sixteen — so on the face of it the per-fibre counts exhaust the population and twenty-seven is the count. The enumeration of §E.5.2 overturns that: five proved results fit no fibre, and five is the gap. **The table is short by five rows; thirty-two is right.** Nothing outside the book was needed to see it, only the addition nobody performed. Registers 219, 220 and 221 carry all three positions, in the order they were held.

Seventeen are enumerable, and they close. Appendix A names its mathematical elements — A.1 through A.17, ten in the ledger and seven written out — and the ledger's own "where" column fixes two of the five coordinates for each: *in full* gives status = proved and verification = exhaustive, *computed* and *forced by* give verified, and anything cited gives cited. Indexed on that pair the seventeen occupy three distinct cells and **$E = 0$** : nothing the two coordinates admit is absent. That is a real closure over a stated population, which is what F.4.1 asked for.

And the fibre table names nine of the rest by itself. A fibre containing one element, where the book contains exactly one mathematical object of that kind and that language, determines its occupant. Nine do:

fibre	the element it can only be
definition · combinatorics	$E(X) = \mathcal{R}(X) - X $, §1.1
mechanism · order	$\mathcal{R}$, the closure operator, §11.2
formula · physics	the Rydberg term $T(v) = Z^2 R/v^2$
law · complexity	the reorderability law, §23.3

method · order	order recovery by tree propagation, Ch. 10
method · combinatorics	the collection procedure, Ch. 17
method · analysis	the bracket, Ch. 15
measurement · physics	the 1,442-cell bracket test, Ch. 18
measurement · algebraic geometry	the 540 Kreuzer-Skarke proposals, §24.3.4

And the remaining eighteen resolve from the book too, by three filters that need no memory.

Appendix E predates §7.11, so the tower's proofs A.14 to A.17 cannot be among its counts. Two apparent candidates are duplicate labels for elements already listed — Theorem 9.1 is A.2 by its own printed tag, Theorem 10.1 is A.6, "adjunction never repairs". And a candidate whose language belongs to another fibre is not a candidate: $T(v)$ sits in formula · physics and vacates formula · analysis. Applied together the three filters leave every fibre with exactly the population it counts.

E.5.2 The twenty-seven, enumerated

fibre	element	status · verification · precedent
definition · order	a closed index, §11.1	proved · exhaustive · found
definition · combinatorics	$E(X) = \mathcal{R}(X) - X $, §1.1	proved · exhaustive · none found
mechanism · order	$\mathcal{R}$, the closure operator, §11.2	proved · exhaustive · none found
formula · combinatorics	$F(z)$, the single expression, §6	proved · exhaustive · found
formula · combinatorics	the rank polynomial of Λ , §3.3	verified · exhaustive · found
formula · analysis	$\lambda^2 = (2/3)T$, §E.4.1	proved · exhaustive · found
formula · analysis	$V = 4v/3$, §16	measured · exhaustive · none found
formula · analysis	the bracket width, §16	verified · exhaustive · none found
formula · physics	the Rydberg term $T(v) = Z^2 R/v^2$, §15	verified · exhaustive · found
theorem · order	A.1 Λ is closed	proved · exhaustive · none found
theorem · order	A.2 $X \text{ closed} \Leftrightarrow \mathcal{R}(X) = X$ (Theorem 9.1)	proved · exhaustive · found (§22.11)
theorem · order	A.3 projections of closed sets are closed	proved · exhaustive · none found
theorem · order	A.6 adjunction never repairs (Theorem 10.1)	proved · exhaustive · none found
theorem · order	A.7 "may precede" is necessary	proved · exhaustive · none found
theorem · combinatorics	A.8 the total-order criterion at $d = 2$	proved · exhaustive · found (§23.3.3)
theorem · combinatorics	A.9 μ_Λ in closed form	proved · exhaustive · found
theorem · combinatorics	A.10 the tree factorisation of the void count	proved · exhaustive · none found

theorem · analysis	A.12 the bracket is limit-free	proved · exhaustive · none found
theorem · analysis	A.13 the failure condition, and its inversion	proved · exhaustive · none found
law · complexity	the reorderability law, §23.3	proved · exhaustive · none found
method · order	order recovery by tree propagation, Ch. 10	verified · exhaustive · none found
method · combinatorics	the collection procedure, Ch. 17	verified · sampled · none found
method · analysis	the bracket, Ch. 15	verified · exhaustive · none found
measurement · combinatorics	$ \Lambda = 976$ with $E(\Lambda) = 0$, §2.3	measured · exhaustive · none found
measurement · combinatorics	1,113,045,672 maximal chains, §7.9	measured · exhaustive · none found
measurement · physics	the 1,442-cell bracket test, Ch. 18	measured · exhaustive · none found
measurement · algebraic geometry	the 540 Kreuzer-Skarke proposals, §24.3.4	measured · sampled · found

E = 0 in all sixteen fibres, enumerated — and 4 unfibred, which §E.5.5 explains and §25.1.1 flags. Both are printed so neither can be read for the other. Every fibre carries exactly the count §E.5 prints, no element violates either constraint of §E.3, and no combination of status, verification and precedent is admitted by any fibre and absent from it. The closure this appendix asserted is now the closure it exhibits.

Precedent is read as a record of the search, not of the world — *none found* means this work found none, which is what Chapter 22 says is the only direction a search can settle. On that reading the coordinate is determinate for every element, and P20's prohibition is respected rather than evaded.

E.5.3 What makes something an element, and the five that test it Five results proved in Appendix A fit no printed count: A.4 $\mathbb{D} \geq \dim q - \dim p$, A.5 χ_Λ is total, A.11 v is inadmissible as an axis, and Chapter 11's two theorems on definability and on the σ -algebra of predicates. Whether they are elements cannot be argued without saying what an element is, and this appendix never said. It says now.

An ELEMENT of this index is a mathematical statement that (i) the book asserts on its own authority rather than importing; (ii) the book labels — an Appendix A number, a theorem number, or a name it uses in more than one place; (iii) is one of the seven kinds, taken at the granularity of the result and not of the numbers it yields; and (iv) stands at the time of writing, since a statement already withdrawn belongs to the register and not to the index.

Clause (iii) is the boundary between this appendix and Appendix G, and it did not exist to be drawn until G did: a result is one element here and many occurrences there. $|\Lambda| = 976$ is a single measurement in E.5.2 and appears in G's count fibre wherever it is printed. Clause (iv) is why Theorem 7.1 is absent — withdrawn at register 199, and the status coordinate's *withdrawn* value is for statements withdrawn after being indexed, not for statements never indexed.

Applied literally the criteria return thirty-two. Thirteen labelled proofs predating §7.11 (A.1 to A.13, since A.14 to A.17 are the tower's and post-date this appendix), two numbered theorems that are not duplicate labels (Theorem 9.1 carries the tag A.2 and Theorem 10.1 is A.6), and seventeen named but unlabelled elements — the definitions, the mechanism, the formulas, the law, the three methods and the four measurements enumerated in §E.5.2. Thirteen plus two plus seventeen is thirty-two, and twenty-seven plus the five in question is thirty-two.

So the five are elements, and §E.5's table is short by five rows, not the count by five elements. Two independent routes reach it — the criteria applied forward, the gap counted backward — which is Chapter 11's own standard for a defended claim: two disjoint routes, not one done twice.

E.5.4 The five, placed, and the table corrected All five are theorem · order, and each placement is read off its own proof rather than its subject. A.4 is a statement about order dimension. A.5 proves totality from chains, comparisons and a finite intersection of decidable conditions — no limit and no real-valued function appears in it. A.11 derives inadmissibility from A.6's join-preservation condition. Theorem 11.1 is max, min and a conjunction over coordinates. **Theorem 11.2 is the one that had to be argued:** its σ -algebra is a measure-theoretic object, but *analysis* in this book means the real-valued Rydberg calculus — λ^2 , $V = 4\pi/3$, the bracket's limit-freedom are what occupy its analysis fibres — and 11.2's proof is upper bounds and least upper bounds. The σ -algebra is notation for generated information, not a result in analysis. It is order.

A.5 also corrects its own verification: §11.5 proves it in one line and checks it on **30,000 uniformly sampled ambient points**, so it is proved · sampled, not proved · exhaustive. The distinction matters because E.3's second constraint is what would have caught the error had it gone the other way.

correction to §E.5	from	to
theorem · order	5	10
elements, total	27	32

With the five placed, theorem · order carries ten elements over two distinct cells — proved · exhaustive · none found and proved · exhaustive · found — and **E = 0 there as in every other fibre**. Neither constraint of §E.3 is violated. The count the appendix stated twice is right, the table it printed was short, and both are now reconciled by enumeration rather than by argument.

Q item L is closed. Register 210, 219, 220, 221 and 222 carry the four positions this number was held in and the work that settled it.

E.5.5 The promotion tested, and what the test exposed about E Part II is titled THE LAW and its contents are quantified over every index, not over Λ ; this appendix classes all of them *theorem* and reserves *law* for one element. The reclassification was run rather than argued, under two readings of the word.

reading of "law"	theorem	law	fibres	E
as printed	15	1	16	0
A — holds of every index	4	12	18	0
B — holds of every index AND forbids or governs	10	6	17	0

Reading A empties the kind it was meant to refine: four theorems survive out of fifteen, which is a *reductio* on the criterion rather than a result about the book. Reading B is balanced and matches Chapter 13's own title — *what the law forbids* — so the law is the part that prohibits.

And all three close at E = 0, which is the finding. The index's own measure cannot adjudicate the index's own classification. Pressing on why exposes something the appendix should have said when it first adopted fibring:

fibration of the same thirty-two elements	fibres	E
kind × language, as printed	16	0
kind only	7	2
language only	6	3
none — one fibre	1	4

E falls as the fibration is refined, and reaches zero by refinement alone. Fibre finely enough and every fibre holds one cell, which is a box, and $E = 0$ says nothing. So $E = 0$ is not a claim about a set; it is a claim about a set **relative to a stated fibration**, and the coarser the fibration the stronger the claim. §E.1 adopted fibring as the repair for a closure that admitted hundreds of cells and did not ask what the repair cost. It cost this.

Λ is unaffected, and the contrast is the point. Λ 's coordinates are all ordered, so $E(\Lambda) = 0$ is computed at the coarsest fibration there is — a single fibre, 976 cells, no categorical axis to hide behind. That is the strong form of the claim. This appendix's $E = 0$, at sixteen fibres over thirty-two elements, is the weak form, and the book has been printing both under one name.

Which settles the promotion. Reading B adds a fibre and reading A adds two, so both make $E = 0$ cheaper than it already is. The classification stays as printed, not because it is right but because it is the coarsest of the three and therefore the strongest available statement. If a reason to promote is found, it will have to be a reason about the mathematics and not about the index, and it will make this appendix's closure claim weaker in exchange. Registers 232 and 233.

E.5.6 The re-closure this appendix owes, performed, and it fails §E.1 says $E = 0$ expires every time the mathematics grows. It has grown four times since §E.5.2 was enumerated, and this is the first re-closure actually run rather than promised.

new element	fibre	coordinates
A.18 the triangle region is join-closed	theorem · order	proved · exhaustive · none found
$\mathcal{R}$ is a closure operator (§25.4.1)	theorem · order	proved · exhaustive · found
A.19.0 the alphabet, $\Sigma(A - 1) = 17$	theorem · combinatorics	proved · exhaustive · found
Λ as 976 words in $\{0,1\}^{17}$ (§6.1.1)	formula · combinatorics	proved · exhaustive · found

Thirty-two elements become thirty-six. The first run of this re-closure gave **$E = 1$** : theorem · order admitted *proved · sampled · none found* and contained nothing there, the corner opened jointly by $\mathcal{R}$ -as-closure-operator arriving with *sampled* verification and A.18 arriving with *none found*.

§12's loop offers two resolutions and the right one was available. The sampling was not necessary: extensivity, monotonicity and idempotence of $\mathcal{R}$ each have a two-line proof, now written at §25.4.1, and three thousand random subsets were corroboration standing in for an argument nobody had attempted. With $\mathcal{R}$ recorded as **proved · exhaustive · found**, theorem · order returns to two cells and **$E = 0$ across all sixteen fibres at thirty-six elements.**

The entry is kept because the sequence is the point: **the re-closure found a defect that was not in the index but in the evidence for one of its elements**, and the repair was to prove a thing that had been sampled. Register 282.

E.6 What this does and does not establish It does not prove the mathematics correct. Audits 2 and 3 do that — 20 of 20 on Λ 's structure over all 475,800 pairs, and 24 of 24 on the stated equations.

What it establishes is that nothing is MISSING that the structure admits. By P23 an object is complete when all its definitions are identified and no question can be had about them. The first condition is now checkable and checked. The second remains what P23 says it is: a standing claim.

And the method is the book's own. A regex over fifty-five thousand words returned three hundred and seventy fragments of punctuation. An index returned two missing constraints and two missing elements.

Appendix F — Q, indexed and closed

Serving PART IV — THE RECORD. The book's own ignorance as a closed index — eight items, $E(Q) = 0$ fibred by domain, and see §25.1.1 flag 2 for what a fibred zero is worth.

Amended by the tower: eight items closed at $E(Q) = 0$ fibred, four entered by §7.11.7 with re-closure owed. Twelve stand. Appendix E indexes the book's mathematics. This one indexes what the book does not know. Eight items, four ordered coordinates, one categorical fibre, and $E(Q) = 0$ — reached by adding one item the list was missing and one constraint the coordinates had not stated.

F.1 The coordinates

coordinate values	kind
domain physical · bibliographic · mathematical · computational	fibre
blocks nothing < a novelty assessment < one stated claim < a result the book relies on	ordered
obstacle retrievable < buildable < requires an object that does not exist	ordered
cost hours < days < unbounded	ordered
depends the book does not depend on it < partly < yes	ordered

Ordered coordinates only, fibred by the categorical one — the precondition §E.1 states, and the shape Λ itself has.

P19 requires $Q(X) = \emptyset$. The set has been worked through, and this is where it stands.

F.1.1 Recomputed at fifteen items, and closed **F.1.1** Four entered by §7.11.7, of which the self-duality count and the four-cap check closed on recomputation; two more entered by the repairs of register 211 and Appendix G — item L and item M. **L, stated precisely at §E.5.1's close:** of Appendix E's twenty-seven elements, nine are named by the fibre table and eighteen sit in seven fibres holding forty-three candidates between them, so L asks first *which named candidate occupies each of the eighteen slots* — a disambiguation, not a recall. Then, for all twenty-seven, it asks the three coordinates the fibre does not give: status, verification and precedent. Seventeen have status and verification determinable from Appendix A's ledger wording. **Precedent is the coordinate the book records nowhere per element** — three statements about precedent appear in the whole text — so it is the residue of the residue, and the only part of L that is recall rather than reading. **M** closes the numbers index. At thirteen items $E(Q)$ was 2: the mathematical fibre admitted novelty · retrievable · days · no and one claim · buildable · days · no, and Q contained neither.

§12's diagnostic loop offers exactly two resolutions for an admitted-and-absent cell, and both were run. **The exclusion route is unavailable.** To empty the first cell the coordinates would need novelty $\wedge$ retrievable $\Rightarrow$ cost $\leq$ hours, which items B and C already violate; to empty the second they would need buildable $\Rightarrow$ blocks $\leq$ novelty, which item M violates by itself. No monotone constraint excludes either cell without contradicting Q's own contents — which is what it means for a cell to be genuinely admitted rather than an artefact of a missing bound.

The occupation route succeeds, and both occupants were already written in this book. Item N is whether anyone has lattice-ordered the seniority scheme — §22.9 names it, enters it in $E(\text{search})$, and never carried it to Q. Item O is the retest Chapter 10's tree-propagation guarantee needs on Λ_9' , whose f-g-2S' cycle §7.11.1 says makes the guarantee lapse and *must be retested rather than inherited* — stated as an obligation and never indexed as one. With both entered, **fifteen items and $E(Q) = 0$ in every fibre, and 1 unfibred.**

That is item H's mechanism a third and fourth time: the questions the closure names are not new questions, they are obligations the book states in prose and its indices never receive. The leak runs one way — prose to index — and nothing in the book currently checks it. Bibliographic 3, computational 3, mathematical 3, physical 2 distinct cells. The items entered were: the spectroscopic v-check per channel

against measured terms — a lookup against re-staged NIST terms, not a conceptual question, and NIST carries seniority labels for d-shell terms, so the tenth axis is checkable per channel where the ninth alone was not; the 8-versus-112 resolution of §3.3; cap-stability of the join-closure half-result; and the earlier work's 2J_c containment claim. Their coordinates are assigned in the session record; E(Q) at twelve items has not been recomputed, and the eight-item closure below is stated for the eight it covers. Recorded here so the count cannot silently stand at eight.

F.1.2 The open set as this session leaves it — thirteen items, listed Items E, J and L closed during the work, and the self-duality count with them. Three entered late and were described as *entered in Q* in the sections that named them without ever reaching this appendix — register 239's failure, repeated three times. They are entered here.

item	question	blocks	obstacle	cost
A	the Sc VI level, measured	one claim	nonexistent	unbounded
B	Paschen & Götze §III, read	nothing	retrievable	days
C	the nine named literatures, entered	nothing	retrievable	unbounded
D	the Kreuzer-Skarke Hodge pairs, read	one claim	retrievable	days
F	the reorderability decision problem, searched	nothing	buildable	days
G	the complexity literature, entered	nothing	retrievable	hours
H	whether Sc VI has already been measured	one claim	retrievable	days
I	the spectroscopic v-check per channel	one claim	retrievable	days
K	the earlier 2J _c containment claim, tested against its own	nothing	retrievable	hours
M	Appendix G closed — E(G) = 33 resolved	one claim	buildable	days
N	whether anyone has lattice-ordered the seniority scheme	nothing	retrievable	days
O	tree propagation retested on Λ_9 's f-g-2S' cycle	one claim	buildable	days
P	the 1,442 split into independent and Ritz-derived tests	one claim	retrievable	hours

Every *blocks* value of *novelty* fell to *nothing* under §22.2.1 — six of the thirteen. **E(Q) = 0 in all four domain fibres and 1 unfibred**, both printed per §25.1.1 flag 2. Register 256. F.2 Closed Do the perturbation bounds tighten at higher order? Yes, 1,585-fold — median 25.96 cm⁻¹ at order 1, 0.0164 at order 6. §16.10

Is there an order maximising the bound, given σ ? No interior optimum. Information per cell rises monotonically; the maximum is the admissibility cap k^* , computable per cell. §16.11

Does the cell count change at matched order? Yes: 518 at order 1, 34 at order 6 — with containment at 100% for every order once the premise is checked. §16.10

What is V for the convex bracket on measured levels? It grows with order — 157 to 3.2×10^6 on a matched node set. What holds at every order is *containment*: 560 tests, zero failures. **The floor of Proposition 15.2 belongs to the classical bracket alone.** §16.10

Do all accelerators pay $1/|p-1|$? Yes, at order m : $A^m(x^p) = (-1)^m x^p/(p-1)^m$, proved for $m = 1 \dots 4$. Richardson is the exception — a different kernel, no $(p-1)$ at all. §19.6

Does the claim set close, and under what measure? Yes — a Gröbner basis of eight, against which every relation derived since reduces to zero. $E(\text{claim set}) = 0$. §22.8

Why is $F(-1) = 2$? Because l and f each have two values, so $F_{\text{box}}(-1) = 0$, and the tree's couplings remove the vanishing factors. §6.8.1

Does the amplification have a general form? $A = N[\varphi(\Lambda \cup \{y\})] - N[\varphi(\Lambda)] - 1$, exactly, and to 7% as fibre $\times (0.972 - 0.154 \cdot \text{descendants})$. §11.7

F.3 Open, and only these $A = \text{Sc VI } 3s^2 3p^3 ({}^4S^\circ) 6s \ 3S^\circ_1$. Committed prediction, unmeasured. [735,860, 737,380] $\text{cm}^{-1} \cdot 13.5615\text{--}13.5895 \text{ nm} \cdot 91.338 \text{ eV}$. The lower edge tests monotonicity, the upper tests convexity. Needs a spectrometer.

B — Paschen & Götze §III, Wie findet man eine Serie und ihre Grenze?, and Runge, Encyklopädie 1926, pp. 783–820. Both unread. Springer Book Archives, public domain since 2017, scanned with OCR and not exposed.

C — nine unentered literatures. Named in §22.7, with the vocabulary: excess width, Newton decrement, partial identification, zero-error capacity.

D — verifying 540 Kreuzer–Skarke predictions. §24.3.4 computes $E(X)$ on the $\chi = \pm 6$ slice specified in print and gets 540 predicted Hodge pairs, none falsified, two confirmed by name. [§7.11.6 supplies the calibration this claim owes: an $E(X)$ proposal can be vetoed by structure the index does not carry, and these 540 are proposals of that kind.] What remains is the lookup: the 30,108 pairs are in a gzipped file this session could not read.

E — the physical meaning of the target-spin axis. Answered in §7.11.1 by exact vector coupling: $2S' \leq g$ is a true bound, zero violations, and not tight in the two ways §7.11.2 names. The residue, which is a lookup and not a question: whether the measured target terms of the 153 channels saturate $\min(g, 4f+2-g)$. Q gains four further items from the tower session, listed at §7.11.7, and $E(Q)$ at the new count is owed a recomputation before it is claimed.

F — the reorderability decision problem, and it is not this book's problem. §23.3 establishes the law and the procedure. Λ 's own case is answered: its structure is constructed in Chapter 2, its embedding is tight — height $17 = \Sigma(|A_i| - 1)$ — its join-irreducible poset has $|J(\Lambda)| = 17$ and width $7 \leq 8$, and by Dilworth and Larson the chain decomposition of $J(\Lambda)$ is the axis system. Nothing about Λ is open.

What is open is the general question: given an arbitrary cell set with no construction, decide reorderability in time polynomial in $|X|$ and d . §23.3.5 states the law, the procedure, and the residue — realisability.

F.4 The eight, classified by their first empty cell B.2.16's grid asks, for each item, what the FIRST unmet requirement is.

item	first empty cell	the obstacle
A Sc VI	the level is measured	an instrument
B Paschen & Götze	the scan is text-searchable	a digitisation artefact
C nine literatures	each is entered	labour
D Kreuzer–Skarke	540 predictions looked up	a file format
E target spin	the selection rule is confirmed	a spectroscopic fact

F reorderability	realisability of the hard relations	a missing object
G complexity literature	the papers are read	a library
H Sc VI, already measured?	a search is run	an afternoon

G — the sublattice complexity literature, unread. Ryter & Schmid 1994 and Adams–Dwinger–Schmid 1996. §23.3's residue may be settled there.

H — whether Sc VI ($4S^0$)6s has already been measured. §15.2's exclusion says the level "exists and is simply unmeasured," and §18.6.5 records that nothing was fetched for the bracket. Neither is a search. The book asserts the absence of a measurement without recording an attempt to find one, and a published value would settle item A in an afternoon.

This item was not on the list. It was produced by closing the list — §F.5 gives the computation.

F.4.1 Q, closed §25.2 asserted $|Q| = 7$. An assertion is not an enumeration, so the seven were indexed — coordinates blocks, obstacle, cost, depends, fibred by domain — and closed.

domain	items	E(fibre)
bibliographic	B, C, G	0
computational	D	0
mathematical	F	0
physical	A, E	2

Two cells admitted and absent, both physical. One is blocks one claim · retrievable · days — and that is item H, which the list did not contain.

The other is buildable · unbounded, and it is excluded by a constraint the coordinate system had not stated:

obstacle = buildable $\implies$ cost $\leq$ days. If a thing can be built, its cost is bounded by the building. Unbounded cost belongs to instruments and to objects that do not exist.

With that bound stated and H added, $E(Q) = 0$ in every fibre.

F.4.2 And the column that matters All eight items have depends = no.

Not one open question blocks a result this book relies on. Q is a list of what would be added, not of what is missing.

Six are retrievable, one is buildable, one requires an object that does not exist. Two have unbounded cost — a spectrometer and a theorem. Three block a stated claim; none blocks a result.

$|Q| = 8$, and it is enumerated rather than asserted. Six are blocked by something outside the index — an instrument, a scan, two units of work, a measurement, a search. One is blocked by something that does not exist. And one is a library.

And the distinction that matters is not 'internal' against 'external'. It is retrievable against requires a new object — because the first kind closes with effort and the second does not close at all until someone builds something. Four of the six are retrievable.

F.5 Precedence, pushed as far as it goes Every component of §23.3 was found in print after being derived from scratch. The register carries the corrections; this is the ledger.

what was derived what already had it

the recovered bound ϕ_{ij} Rival 1973/74 — sublattices are the complement of a closed family of irreducible intervals, and Lemma 3.5(ii) of Siggers states ϕ_{ij} exactly

"subdirect" — every Larson 1998; named, with the subdirect/tight distinction the book needed value used

tightness $\Leftrightarrow$ acyclic Siggers 2016, Lemma 5.1(iii) constraint digraph

sublattices ↔ preorder		Siggers 2016, Cor. 4.2 — verified here on four boxes, offset by the empty	
extensions		sublattice	
the axis system as a Birkhoff 1937 · Dilworth 1950 · Larson 1998 chain decomposition of J(Λ)			
the step law, 3·2^(d−2)	Rival 1973 — \	K ≤	L for a
		(3	maximal
		/2	sublattice
		)\	

the subcube-removal Chen-Koh-Tan — maximal sublattices are $L - [a, b]$ construction

the large sublattices of catalogued; the maximal ones in finite topology, Sharp, Stephen $2^{\lfloor r \rfloor}$

the unbounded Tucker 1972 obstruction families

C1P and the linear Booth & Lueker 1976 · Fulkerson & Gross 1965 algorithm

the constraint-language Schaefer 1978 dichotomy

Nine components, nine precedents. The subcube-removal construction is Rival's interval removal; the "step law" is his $3/2$ bound sharpened to the Boolean case, which was already catalogued elsewhere.

F.6 And the complexity question itself may be settled Ryter & Schmid, Deciding Frattini is NP-complete, Order 11 (1994). Adams, Dwinger & Schmid, Maximal sublattices of finite distributive lattices, Algebra Universalis 36 (1996) 488–504.

So decision problems about sublattices of finite distributive lattices have a complexity literature, and at least one of them is settled NP-complete.

This book has not read it. §23.3's residue may already be closed there, and until that is checked the claim "open" is a claim about this author's reading, not about the state of knowledge.

That is recorded as item G rather than left implicit.

F.7 What precedence costs the book Nothing in §23.3 is withdrawn. The derivations agree with the literature wherever both exist, and the agreement is itself a check — Siggers' bijection was verified against an independent enumeration, and Rival's $3/2$ bound contains the measured $3/4$.

What is withdrawn is the implicit claim of novelty. The section states results; it does not claim to state them first, and §F.5 is where the difference is written down.

B.2.10 says: enumerate targets before searching. It should have said, and now does: SEARCH BEFORE DERIVING. Nine components and a day, against nine citations.

F.8 What closed, and what it cost Two items closed during this work. The target-spin axis was taken on the grid's advice — its first empty cell was reachable by labour alone — and Λ_9 passed six independent tests. §23.3's characterisation closed at $d = 2$ and its law at every d .

The cost is in Chapter 21. Sixty-odd corrections come from §23.3 alone.

An index that records its own rules must also record when it failed to consult them. This chapter is where that is done.

Appendix G — The numbers, indexed

Serving PART IV — THE RECORD. The book's own printed numbers as a cell set — built, indexed and **open by thirty-three cells**. $E(G) = 33$ across six fibres, and §25.1.1 flag 2 governs how a fibred figure should be read.

Five things in this book were indexed before this appendix existed: its mathematics (E), its ignorance (F), its errors (Chapter 21), its data provenance (C) and its terms (the Index). Its numbers were not among them, and of the five only four were closed — E was asserted, as §E.5.1 now records. What is not indexed cannot be closed, and what is not closed carries defects that no audit is shaped to find. Register 210 was one such defect and 219 is a second, both in Appendix E and both found by trying to recompute rather than by reading. Two instances are too few to estimate a rate, which is why §G.3.1 measures the population instead.

G.1 The population 995 distinct claim-bearing numbers, 2,433 occurrences: 1,166 inside tables and 1,267 in prose. 1,715 of the occurrences — 70% — stand in a paragraph carrying no method, code or section cue. That is not a count of uncheckable numbers. Most are recomputable from their context. It is a count of numbers the book cannot currently sort into recheckable and not, which is the defect.

G.2 The coordinates

coordinate values	kind
quantity count · density · defect · bound · measurement · rate	fibre
definition asserted < stated < stated with inputs	ordered
inputs named < enumerated < printed	ordered
verification asserted < recomputed < recomputed under variation	ordered

Ordered coordinates fibred by the categorical one, per §E.1's precondition — the only shape $\mathcal{R}$ accepts.

G.3 The index, built Every claim-bearing occurrence is assigned its four coordinates mechanically from the paragraph that carries it — the quantity's kind from its units and vocabulary, the definition from whether a method and an input set are both named, the inputs from whether they are named, enumerated or printed, and the verification from whether the paragraph records a recomputation and whether that recomputation varied the caps. The rules are stated so the assignment can be rerun; they are coarse, and being coarse and rerunnable is the point, since a judgment made once per number would not be checkable at all.

fibre	occurrences	distinct cells	admitted	E
count	1,680	14	22	8
defect	291	15	24	9
measurement	204	12	20	8
bound	183	9	13	4
rate	56	9	13	4
density	39	5	5	0
all	2,453	—	—	33

$E(G) = 33$, across the six fibres of G.2. Thirty-three combinations of definition, inputs and verification are admitted by the structure of what the book already prints and occupied by nothing it prints. Each is a number the book could state to the standard it already meets elsewhere and does not — and by §12's loop each resolves either by a number that belongs there or by a constraint the coordinates are missing. Neither has been done. Unlike Q, this index is open and is not claimed otherwise.

G.3.1 And the distribution is the finding Of 2,453 indexed occurrences, **2,316 carry verification = asserted** — no recomputation recorded anywhere near them — against 115 recomputed and 22 recomputed under variation. On the definition axis, 1,998 are asserted, 145 stated, and 310 stated with

their inputs. So 94% of the numbers this book prints are, by its own coordinates, at the weakest verification the coordinate admits. That is not a claim that 94% are wrong. It is a measurement of how much of the book's arithmetic stands on the reader's trust rather than on the book's own standard, and it is the first time the quantity has had a number.

G.3.2 The tripwire The press runs the extraction on every render and reports the count before any reading — a D1-type check in Chapter 11's sense: mechanical, prior, and indifferent to what the author believes. It reports and does not fail the build, because a failing build on 94% would stop the book from printing at all. The threshold is a decision the author has not made, and recording that it is undecided is better than choosing it silently.

G.4 What closing it would establish, and what it costs E.6 claims nothing is missing that the structure admits, about the book's mathematics. Closing G extends that from the mathematics to the arithmetic, and makes §25.2's S3 — that a reader given the claims and the evidence could reconstruct the structure — true in the strong sense rather than the prose one. The cost is F.4-shaped: buildable, days to unbounded, depends = no. Entered in Q as item M, enumerated rather than asserted away.

G.4.1 Two ratios the book had never measured about itself Appendix G indexes the numbers this book prints and had not indexed the two that describe its own shape. They are named here because a quantity without a name cannot be tracked, and both are measurements the method demands of anything else it examines.

The withdrawal ratio. Words of register against words of argument — Chapter 21 against Chapters 1 to 25 without it. **19,504 : 38,817 = 0.502**. One word withdrawn for every two asserted.

The method ratio. Words of Appendix B — the principles, the twenty protocols, the twenty prime audits — against the whole book but for itself. **11,811 : 75,459 = 0.157**. Against the other six appendices alone it is 0.877: the machinery is now very nearly the size of all the mathematics it audits.

What each is for. §25.1.1's fourth referee flag says the register is longer than the results and the book has no answer to the fair version of the question. The withdrawal ratio is that question made numeric: at 0.502 the register is half the argument, and the flag asked not whether that is large but **whether it is falling**. It is not tracked over time and cannot be, because entries carry no date; §25.1.2's split of thirty-four defects to eight self-inflicted is one session's measurement and there is no second point. **A ratio with one observation is a number, not a trend**, and that is the honest state of referee flag 4.

The method ratio has no referee behind it and may matter more. Appendix B has grown faster in these sessions than anything it governs. Whether a method that is 88% the weight of its own subject is thorough or top-heavy is a judgement in §B.2.20's sense and not a computation — but it is now a judgement with a figure attached, which it was not. Register 296.

G.4.2 The ideal withdrawal ratio is 1 : 1, and the book cannot yet say where it stands §G.4.1 measures the withdrawal ratio in words and leaves the number uninterpreted. Words are a proxy. The quantity is a ratio of **claims**, and on that basis the ideal has a derivation rather than a preference.

A claim's survival carries information, and 1 : 1 is where it carries the most. If withdrawals and survivors stand in ratio r , the chance a given claim survives is $1/(1+r)$, and the news that it did carries $-p \log_2 p - (1-p) \log_2 (1-p)$ bits. That is maximal at $p = \frac{1}{2}$, which is $r = 1$. **Below 1 : 1 survival is cheap news** — almost everything survives, so surviving says little. **Above it, withdrawal is cheap** — the record is mostly failure and the process is spending more than it keeps. At parity, the fact that a claim stands is worth exactly one bit, and that is the most a register can ever tell a reader about any single claim in it.

Measured, the book does not have one answer. It has a bracket. 190 claim-numbers stand in the register. What stands against them depends entirely on what a surviving claim is, and the book has never said:

surviving claims taken as	count	ratio	p(survive)	bits
------------------------------	-------	-------	------------	------

Appendix E's mathematical elements	36	5.28 : 1	0.159	0.633
those plus the Index's terms	87	2.18 : 1	0.314	0.898
every claim-bearing number in Appendix G	2,831	0.07 : 1	0.937	0.339

The three readings span a factor of seventy-five, and 1 : 1 lies inside them. That is a bracket in this book's own sense: the quantity is bounded by measured neighbours and undetermined between them, so the prediction is **consistent with the evidence and not established by it** — which is precisely what §15 says a bracket means and precisely the standing §18.6 has.

What would settle it is a definition, not a measurement. §E.5.3 closed the oldest item in Q by saying what an element is; the same move is available here and has not been made. Until *claim* is defined as *element* was, the withdrawal ratio has no value, only a range — and the register, which exists to make the book's failures countable, cannot count its own denominator. Entered in Q. Register 297.

G.5 And it needs its own falsification An index that cannot fail is not a check. The test: sample n printed numbers, attempt recomputation from the book alone, report the rate. This session is the first such sample, $n = 8$ — seven quantities recomputed from the construction and Appendix E's $E = 0$ — and the observed failure is one in eight, Appendix E, whose inputs the book does not print. A rate from eight is not a rate; it is a method demonstrated once, and G.3.1 gives the population it should be drawn from.

END MATTER

Index

This index satisfies the law of this book.

Construction Terms are ordered by specificity: $s \sqsubseteq t$ when s is a specialisation of t . The index is closed iff

$$s \sqsubseteq t \Rightarrow \text{loc}(s) \subseteq \text{loc}(t)$$

every location of a specialisation being also a location of its generalisation. That is a down-set condition in §11.4 form, so **a closed index is a sublattice of (terms × locations)**.

Verified: 51 terms, 129 entries, 40 specialisation relations, **0 violations — E(index) = 0.**

What closure buys. The term list is recoverable by projection (S1); the specificity order is recoverable by tree propagation, a term hierarchy being a tree (S2); $\text{loc}(\cdot)$ is recoverable as the union over specialisations (S3); and every (term, location) pair is present or absent — **no see also limbo, no partial entries** (D3).

An index with E > 0 requires the reader to already know where to look — which is the periodic table's failure at the scale of a back-matter page.

Entries bracket §17.2 · §17.2 · §17.2 · §17.2 · §17.2 · §17.2 · §25.3 · §25.1 · §25.3 · §16.4 · App A

bracket failure §25.3 · §16.4 bracket width §25.1 interiority §17.2 · §17.2 limit-free §25.3 · App A the four rules §17.2 · §17.2 · §17.2 · §17.2 · §17.2 · §17.2 · interiority §17.2 · §17.2

closure §1.1 · §2.3 · §11.1 · §11.2 · §11.1 · §11.2 · §11.3

closed index §1.1 · §2.3 · §11.1 closure defect §1.1 · §11.2 closure operator §11.2 $\mathcal{R}(X) = X$ §11.1 · §11.1 · §11.2 · §11.3

cost of a guarantee §25.5 · §25.1 · §17.2 · §25.3 · §18.6

$V = 4v/3$ §25.1 · §25.3 the floor §17.2 · §25.3 the pole §25.5 · §25.3 v_V §18.6

extension §25.1 · §17.2 · §25.3 · §16.4

addable cells §25.1 adjoinable axes §17.2 imposable constraints §25.3 · §16.4

prediction §18.6 · §18.6 · §17.6 · App E

E(X) as budget §18.6 *Sc VI* §18.6 · App E *second route* §17.6

quantum defect §17.2 · §25.3 · §18.6 · §17.13

isoelectronic law §17.2 · §17.13 penetration §25.3 · §18.6

retrieval §25.1 · §17.2 · §25.5

retrieval redundancy §25.1 · §25.5 route set §25.1 · §17.2

self-defence §11.3 · §11.4 · §11.5 · §17.5 · §17.6 · §21.8 · §25.3 · App D

totality §11.5 · §17.5 · §21.8 · App D · process index §17.5 · §21.8 · App D two routes §11.4 · §17.6 $\mathbb{D}_{\text{def}}$ §11.4 · §25.3 $\mathbb{D}_{\text{phys}}$ §11.3

self-reference §11.1 · §11.2 · §11.3 · §11.4

alphabet recovery §11.1 *bound recovery* §11.2 *order recovery* §11.3 · §11.4

the collection §18.6 · §17.11 · §17.12 · §17.13 · App C

decline modes §17.11 *isoelectronic pairs* §18.6 · §17.13 *provenance* App C *the census* §17.12

the index Index

index, self-referencing Index

the record Ch 18 · App E

margins App E *recomputed* App E *withdrawals* Ch 18

index, self-referencing *this page*

The last entry is a cell of the index referring to the index. It is not a joke; it is the fixed point made visible, and its presence is what makes $\mathcal{R}(\text{index}) = \text{index}$ rather than merely $\text{index} \subseteq \mathcal{R}(\text{index})$.

References

Every work cited in this book, with what it is cited for. Entries marked [unread] were named but not obtained; §25.2 lists them as open.

R.1 Spectroscopic data Kaufman, V. & Martin, W. C. (1991). *J. Phys. Chem. Ref. Data* 20,

775. — Al I, Al II levels. Kimura, K., Makino, K., Yamada, H. & Yoshizumi, Y. (2024). — characterising the integer points of 2-decomposable polyhedra by closedness under operations. **$E(X) = 0$ is 2-decomposability**; the class is defined by closure under operations including min and max, which is §9's condition. §22.12 U1.

Kramida, A. E. & Martin, W. C. (1997). *J. Phys. Chem. Ref. Data* 26,

1185. — Be I levels. NIST Atomic Spectra Database. — Ar II, Be II, Bi I, C II, Ca II, Cd II, Ga I, He I, He II, Hg II, K II, Li I, Li II, Mg II, N II, Na II, Ne I ($^3/2$), Ne I ($^1/2$), Si I, Si II, Zn II.

Sansonetti, J. E. (2008). *J. Phys. Chem. Ref. Data* **37**, 1659. — Na I.

Sansonetti, J. E. (2008). *J. Phys. Chem. Ref. Data* **37**, 7. — K I.

Sugar, J. & Corliss, C. (1985). *J. Phys. Chem. Ref. Data* **14**, Suppl. 2. — Sc VI and iron-group limits.

Sugar, J. & Musgrove, A. (1990). *J. Phys. Chem. Ref. Data* **19**, 527; (1995) **24**, 1803. — further ionisation limits.

R.2 Exotic systems and collective quantities Hori, M. et al. (2011). *Nature* 475, 484; and (2006) *Phys. Rev. Lett.* 96, 243401. — antiprotonic helium, both accessed via CODATA 2010, Table XII (arXiv:1203.5425) and arXiv:1304.4330.

Korobov, V. I. (2008). *Phys. Rev. A* **77**, 042506. — antiprotonic helium transition energies.

Singer, K., Stanojevic, J., Weidemüller, M. & Côté, R. (2005). *J. Phys. B* **38**, S295. — long-range Rydberg interaction coefficients.

R.3 Series structure and quantum defects Racah, G. (1942). *Phys. Rev.* 62, 438–462. — jK coupling; the scheme §7.11 uses for Rydberg channels where core splitting dominates exchange. §22.9.

Racah, G. (1943). *Phys. Rev.* 63, 367–382. — seniority; the tenth axis of the tower, and the exact sets that fix its density. §§7.11.1, 22.9.

Rydberg, J. R. (1890). The term formula $T = Z^2R/v^2$. — the object of Part III.

Ritz, W. (1903). The quantum-defect expansion $\delta = \delta_0 + \delta_z/n^2$. — used for every two-point solve.

Edlén, B. (1960). *Handbuch der Physik* 27, 80. — δ falls with l ; S and P penetrate; footnote 78 is the 1960 precedent for $\mathbb{D} = 0$ when parameters equal observations (§11.4.1).

Edlén, B. (1964). *Encyclopedia of Physics*. — examined at one remove (§22.5).

Paschen, F. & Götze, R. (1922). *Seriengesetze der Linienspektren*. Berlin: Springer. — **§III, Wie findet man eine Serie und ihre Grenze?**, [unread]: Springer Book Archives, scanned with OCR and not exposed.

Runge, C. (1926). *Die Serienesetze in den Spektren der Elemente*, in Sommerfeld (ed.), *Encyklopädie der Mathematischen Wissenschaften*, 783–820. — [unread], found in the closing search; a second survey of the same subject.

R.4 Bounds, acceleration and the cost of rigour Boole, G.; Nörlund, N. E.; Steffensen, J. F. — the ratio of first to second central differences, in the finite-difference literature (§16.1).

Aitken, A. C. (1926). *Proc. R. Soc. Edinburgh* **46**, 289. — Δ^2 acceleration; **$\lambda^2 = y^2/y''$ is its correction term** (§16.8.1).

Seki Kōwa (c. 1680). — the same acceleration, three centuries earlier.

Richardson, L. F. (1911). — extrapolation with no $(p-1)$ denominator; the exception in §19.6.

Shanks, D. (1955). *J. Math. Phys.* **34**, 1. — iterated Δ^2 ; pays $(p-1)^{-m}$.

Moore, R. E. (1966). *Interval Analysis*. Prentice-Hall. — **the cost-of-rigour framing, stated plainly** (§16.8.4).

IEEE Std 1788-2015. — interval arithmetic; requires enclosures with bounded excess width.

Manski, C. F. (1989, 2003). *Partial Identification of Probability Distributions*. — **the law of decreasing credibility**, the same trade in econometrics.

Shannon, C. E. (1956). *IRE Trans. Inf. Theory* **2**, 8. — zero-error capacity: the same trade in bits.

Körner, J. & Orlitsky, A. (1998). *IEEE Trans. Inf. Theory* **44**,

2207. — zero-error information theory, surveyed. Nesterov, Y. & Nemirovskii, A. (1994). *Interior-Point Polynomial Algorithms in Convex Programming*. SIAM. — the Newton decrement and self-concordance, which λ^2 and §16.8.2 rediscover.

R.5 Order, lattices and closure Birkhoff, G. (1937). *Duke Math. J.* **3**, 443. — the representation theorem: a finite distributive lattice is the order ideals of its poset of join-irreducibles. §7.9's maximal-chain count rests on it, and §F.3's chain decomposition of $J(\wedge)$.

Dilworth, R. P. (1950). *Ann. Math.* **51**, 161. — chain decomposition, and the width bound §23.3 uses. (1960) *Combinatorial Analysis*, 85–90 — the lattice of maximum antichains.

Rival, I. (1973). Maximal sublattices of finite distributive lattices. *Proc. Amer. Math. Soc.* **37**, 417–420. (1974) *ibid.* **II**, **44**, 263–268. — every sublattice of a finite distributive lattice is the complement of a closed family of IRREDUCIBLE INTERVALS. That family is the recovered bound ϕ_{ij} of Chapter 10, and the $3/2$ bound on maximal sublattices contains §23.3's measured $3/4$.

Chen, C. C., Koh, K. M. & Tan, S. K. (1975). *Algebra Universalis* **5**, 88–97. — the maximal proper sublattices are exactly $L - [a, b]$ with $[a, b] \cap J = \{a\}$ and $[a, b] \cap M = \{b\}$. §23.3's subcube-removal construction is this.

Larson, R. (1998). Embeddings of finite distributive lattices into products of chains. *Semigroup Forum* **56**, 70–77. — tight embeddings correspond to chain decompositions of $J(L)$; "subdirect" is the recovered-alphabet condition, named.

Siggers, M. (2016). On the representations of finite distributive lattices (arXiv:1412.0011). — sublattices of a product of chains correspond one-to-one with spanned preorder extensions of C^∞ (verified here on four boxes); tight $\iff$ the digraph is acyclic; and the explicit warning that subdirect does not imply tight, which §23.3 walked into.

Adams, M. E., Dwinger, P. & Schmid, J. (1996). Maximal sublattices of finite distributive lattices. *Algebra Universalis* **36**, 488–504. — [unread], §F.6.

Ryter, Ch. & Schmid, J. (1994). Deciding Frattini is NP-complete. *Order* **11**, 257–279. — [unread]. §23.3's residue may be settled here.

Tucker, A. (1972). A structure theorem for the consecutive 1's property. *J. Combin. Theory B* **12**, 153–162. — the unbounded forbidden submatrices; §23.3's obstruction route closes on them.

Baker, K. A. & Pixley, A. F. (1975). *Math. Z.* **143**, 165–174. — polynomial interpolation and the Chinese remainder theorem in universal algebra; an algebra with a majority term has every subalgebra of a product determined by its binary projections. **A.2 is this theorem specialised to products of chains**; see §22.11.

Bergman, G. M. — the Double-projection Theorem: a sublattice of a finite product of lattices is determined by its two-fold projections. The lattice form of Baker-Pixley, the median $m(x, y, z) = (x \wedge y) \vee (x \wedge z) \vee (y \wedge z)$ being the majority term. §22.11.

Booth, K. S. & Lueker, G. S. (1976). *J. Comput. Syst. Sci.* **13**, 335–379. — PQ-trees; consecutive-ones in linear time, which is why unbounded obstructions do not preclude a polynomial procedure.

Dunz, — *Seriengesetze der Linienspektren* (Leipzig, 1911). — earliest complete compilation of the series laws, chronologically closest to Ritz; §22.5's route set and Q item B. Retrieved as a scan, unread: the volume is in Q, not in evidence.

Fulkerson, D. R. & Gross, O. A. (1965). *Pacific J. Math.* 15, 835–855. — interval matrices, the foundation of the above.

Schaefer, T. J. (1978). The complexity of satisfiability problems. *STOC*, 216–226. — the boolean CSP dichotomy. §23.3's arity law is this theorem applied to a constraint language read off a cell set.

Sperner, E. (1928). *Math. Z.* **27**, 544. — the antichain bound; §3.4's largest level.

Galois connections / formal concept analysis. — $E(X)$ as a closure defect; the general setting for $\mathcal{R}$.

Stahl, J. & Wille, R. (1984). — reorderability and scale-measurability of a context; §23.3's original citation, withdrawn and reinstated.

Yannakakis, M. (1982). *SIAM J. Alg. Disc. Meth.* 3, 351. — the complexity of order-dimension problems. Also withdrawn and reinstated.

Janet, C. (1928). — the left-step periodic table; $E(\text{Janet}) = 0$.

Sharp, H. and Stephen, C. — maximal sublattices of the Boolean lattice studied in the setting of finite topologies; cited via the catalogue of large sublattices of 2^r . [unread]

R.6 Cited in the transfer tests Poincaré, H. (1890). — non-integrability of the three-body problem; §24.1 states that nothing here touches it.

Gödel, K. (1931). — incompleteness; P20's bound on what a language can close.

Lagrange, J.-L. (1772) and Euler, L. (1767). — the five equilibrium points; §24.1.1's bracket test.

Hill, G. W. (1878). — the Hill sphere; **the gravitational analogue of a capacity bound** (§16.14.1).

Mardling, R. A. & Aarseth, S. J. (2001). *MNRAS* **321**, 398. — the hierarchical-triple stability criterion; monotone and convex, $V \approx 100$.

Batyrev, V. V. (1994). *J. Alg. Geom.* 3, 493 (arXiv:alg-geom/9310003). — dual polyhedra and mirror symmetry; the theorem behind §24.3.1.

Kreuzer, M. & Skarke, H. (2000). *Adv. Theor. Math. Phys.* 4, 1209 (arXiv:hep-th/0002240). — all 473,800,776 reflexive 4-polytopes; 30,108 distinct Hodge pairs. The 30,108-pair file is published and was not read in this work (§25.2).

Huang, W. & Taylor, W. (2019). *JHEP* (arXiv:1809.05160). — the counts 248,305 / 248,305 / 495,515 used in §24.3.1.

Candelas, P., de la Ossa, X. & Rodriguez-Villegas, F. (2007). *Triadophilia* (arXiv:0706.3134). — the thinly populated tip of the Hodge plot.

R.7 What is not cited, and why No source is cited for the lattice itself. Λ is constructed in Chapter 2 from the four quantum numbers and the exclusion principle, both standard, and every structural result about it in this book was computed rather than looked up. §22 states what was searched for and not found.

Physical constants — $R = 109,737.31568 \text{ cm}^{-1}$ — are used without citation, as a metre is.